

PRINCETON  
UNIVERSITY



# HEP analyses with gradients

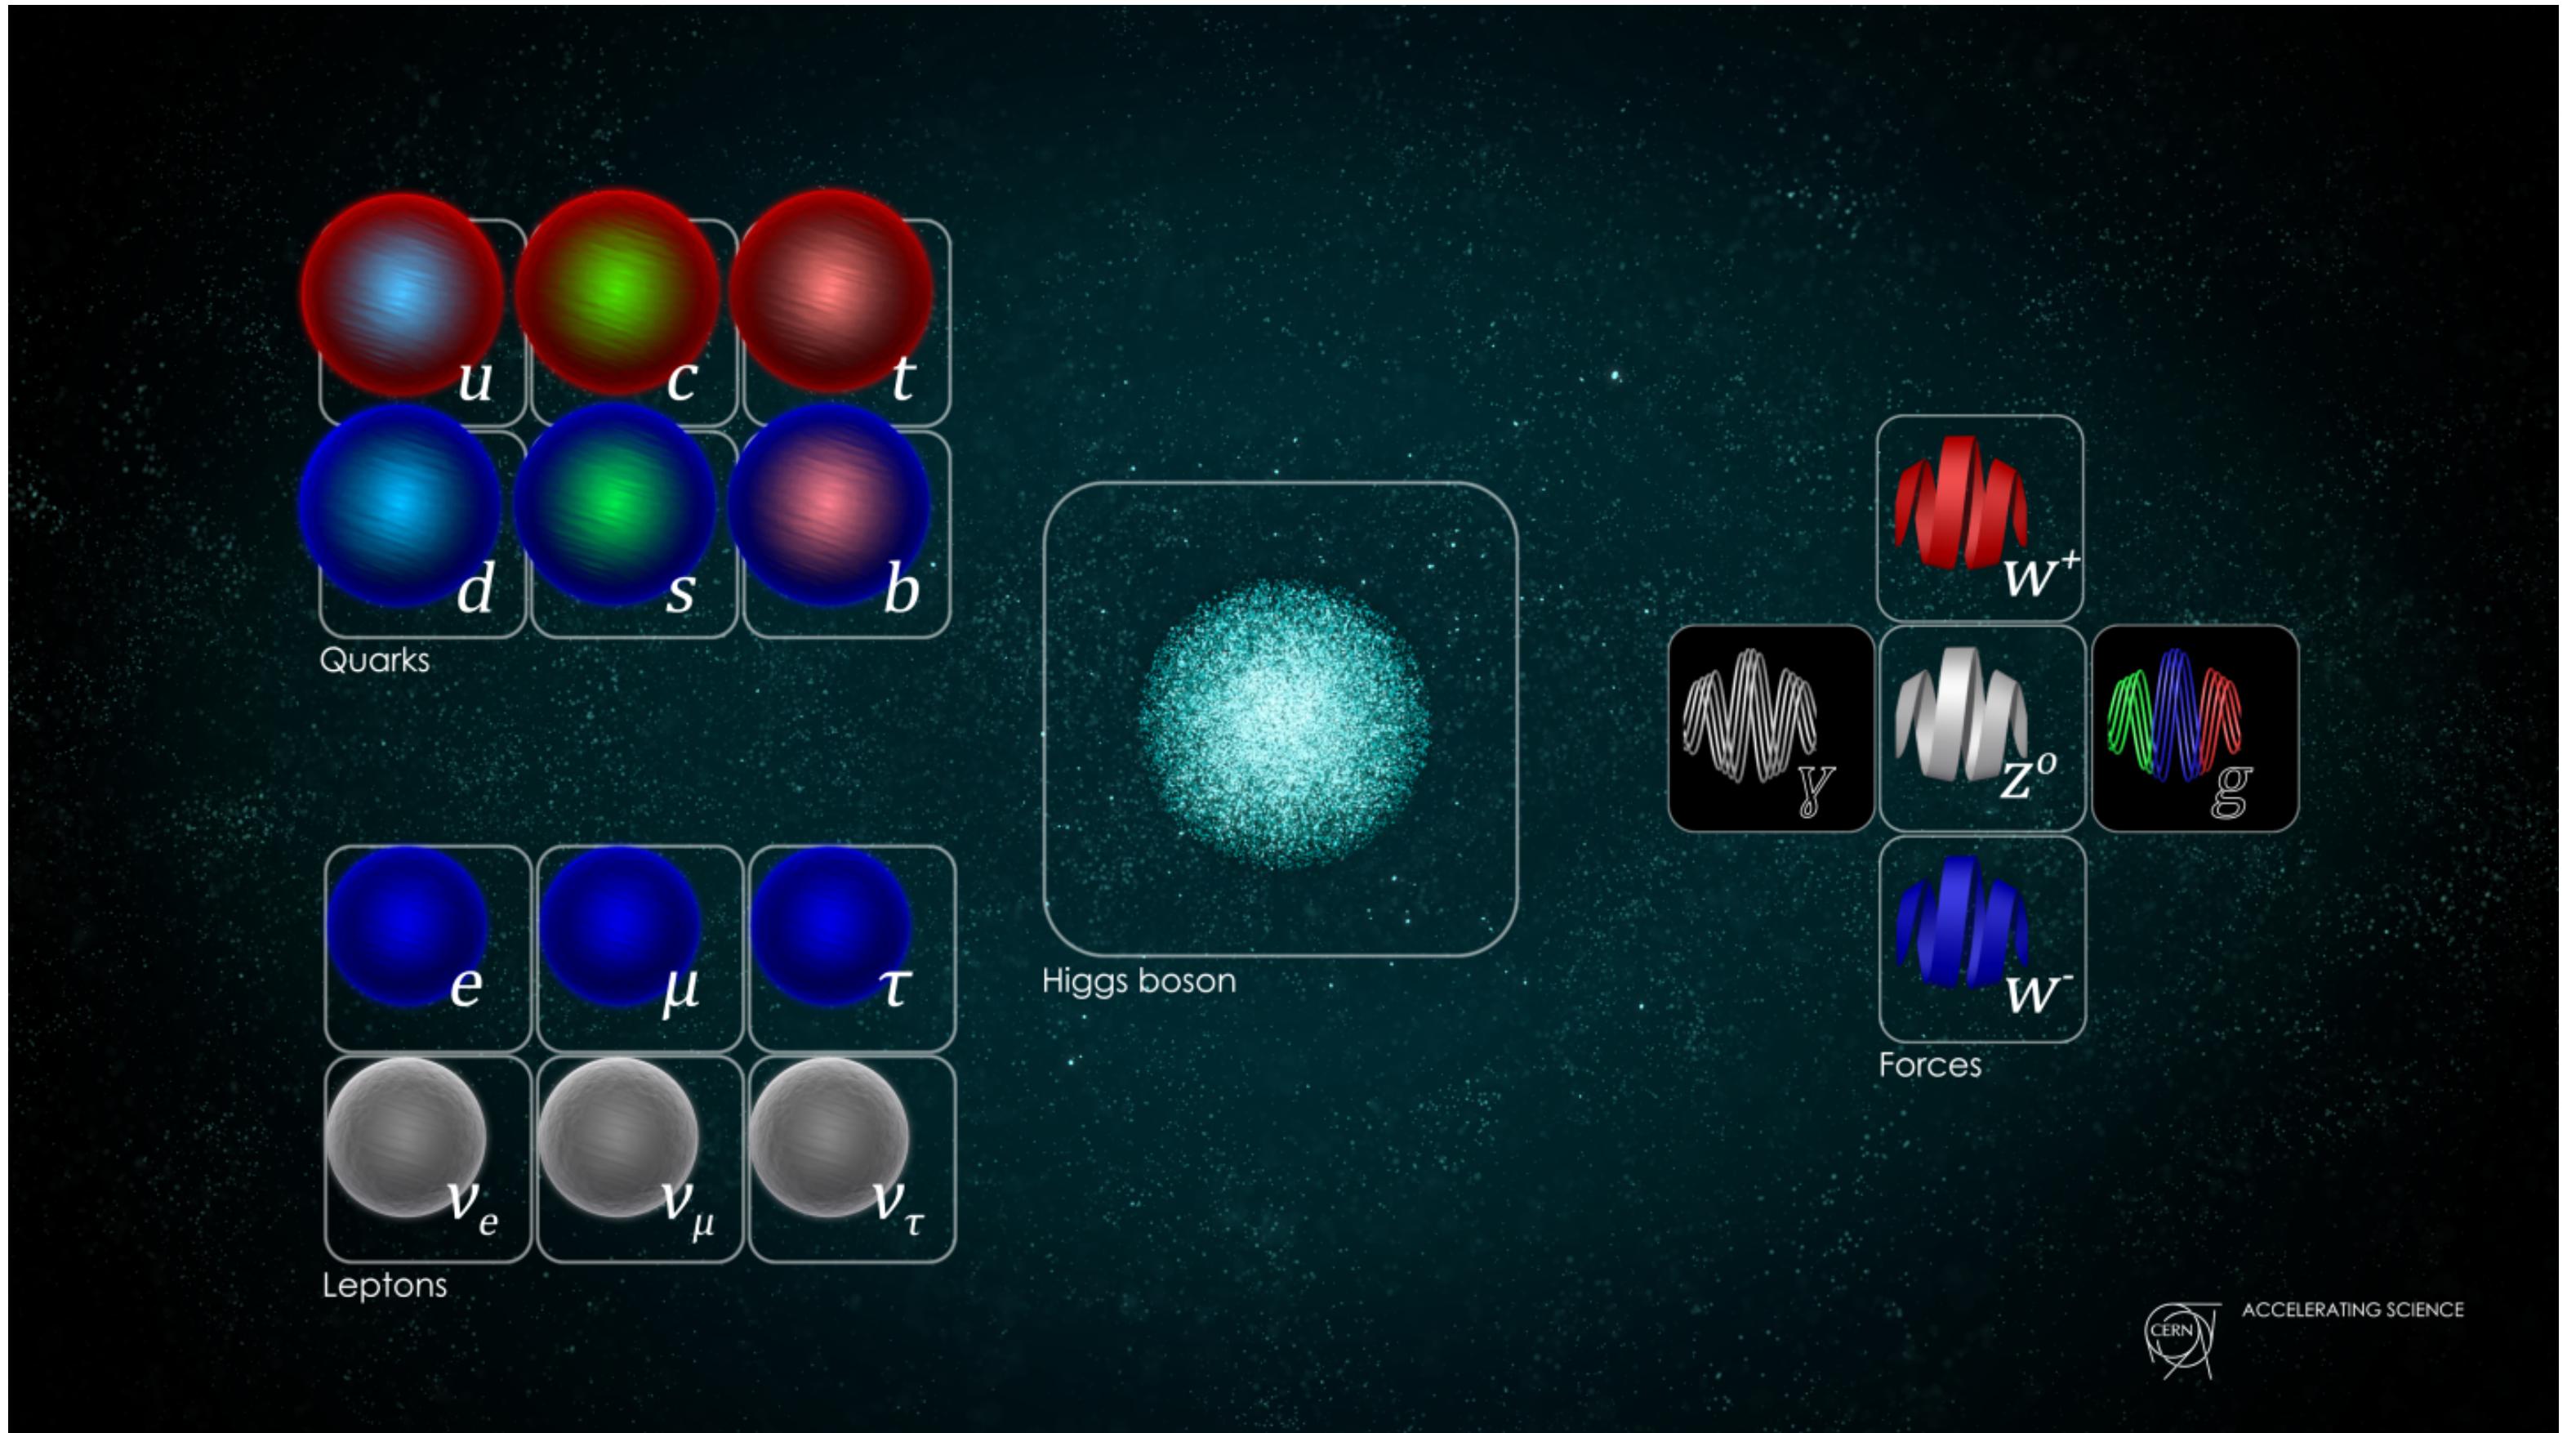
Differentiating a CMS-like analysis with the scikit-HEP ecosystem

What are HEP measurements

# What is HEP work about?

## The Standard Model (SM)

Our best description of the universe at the quantum scale is the *Standard Model (SM)*

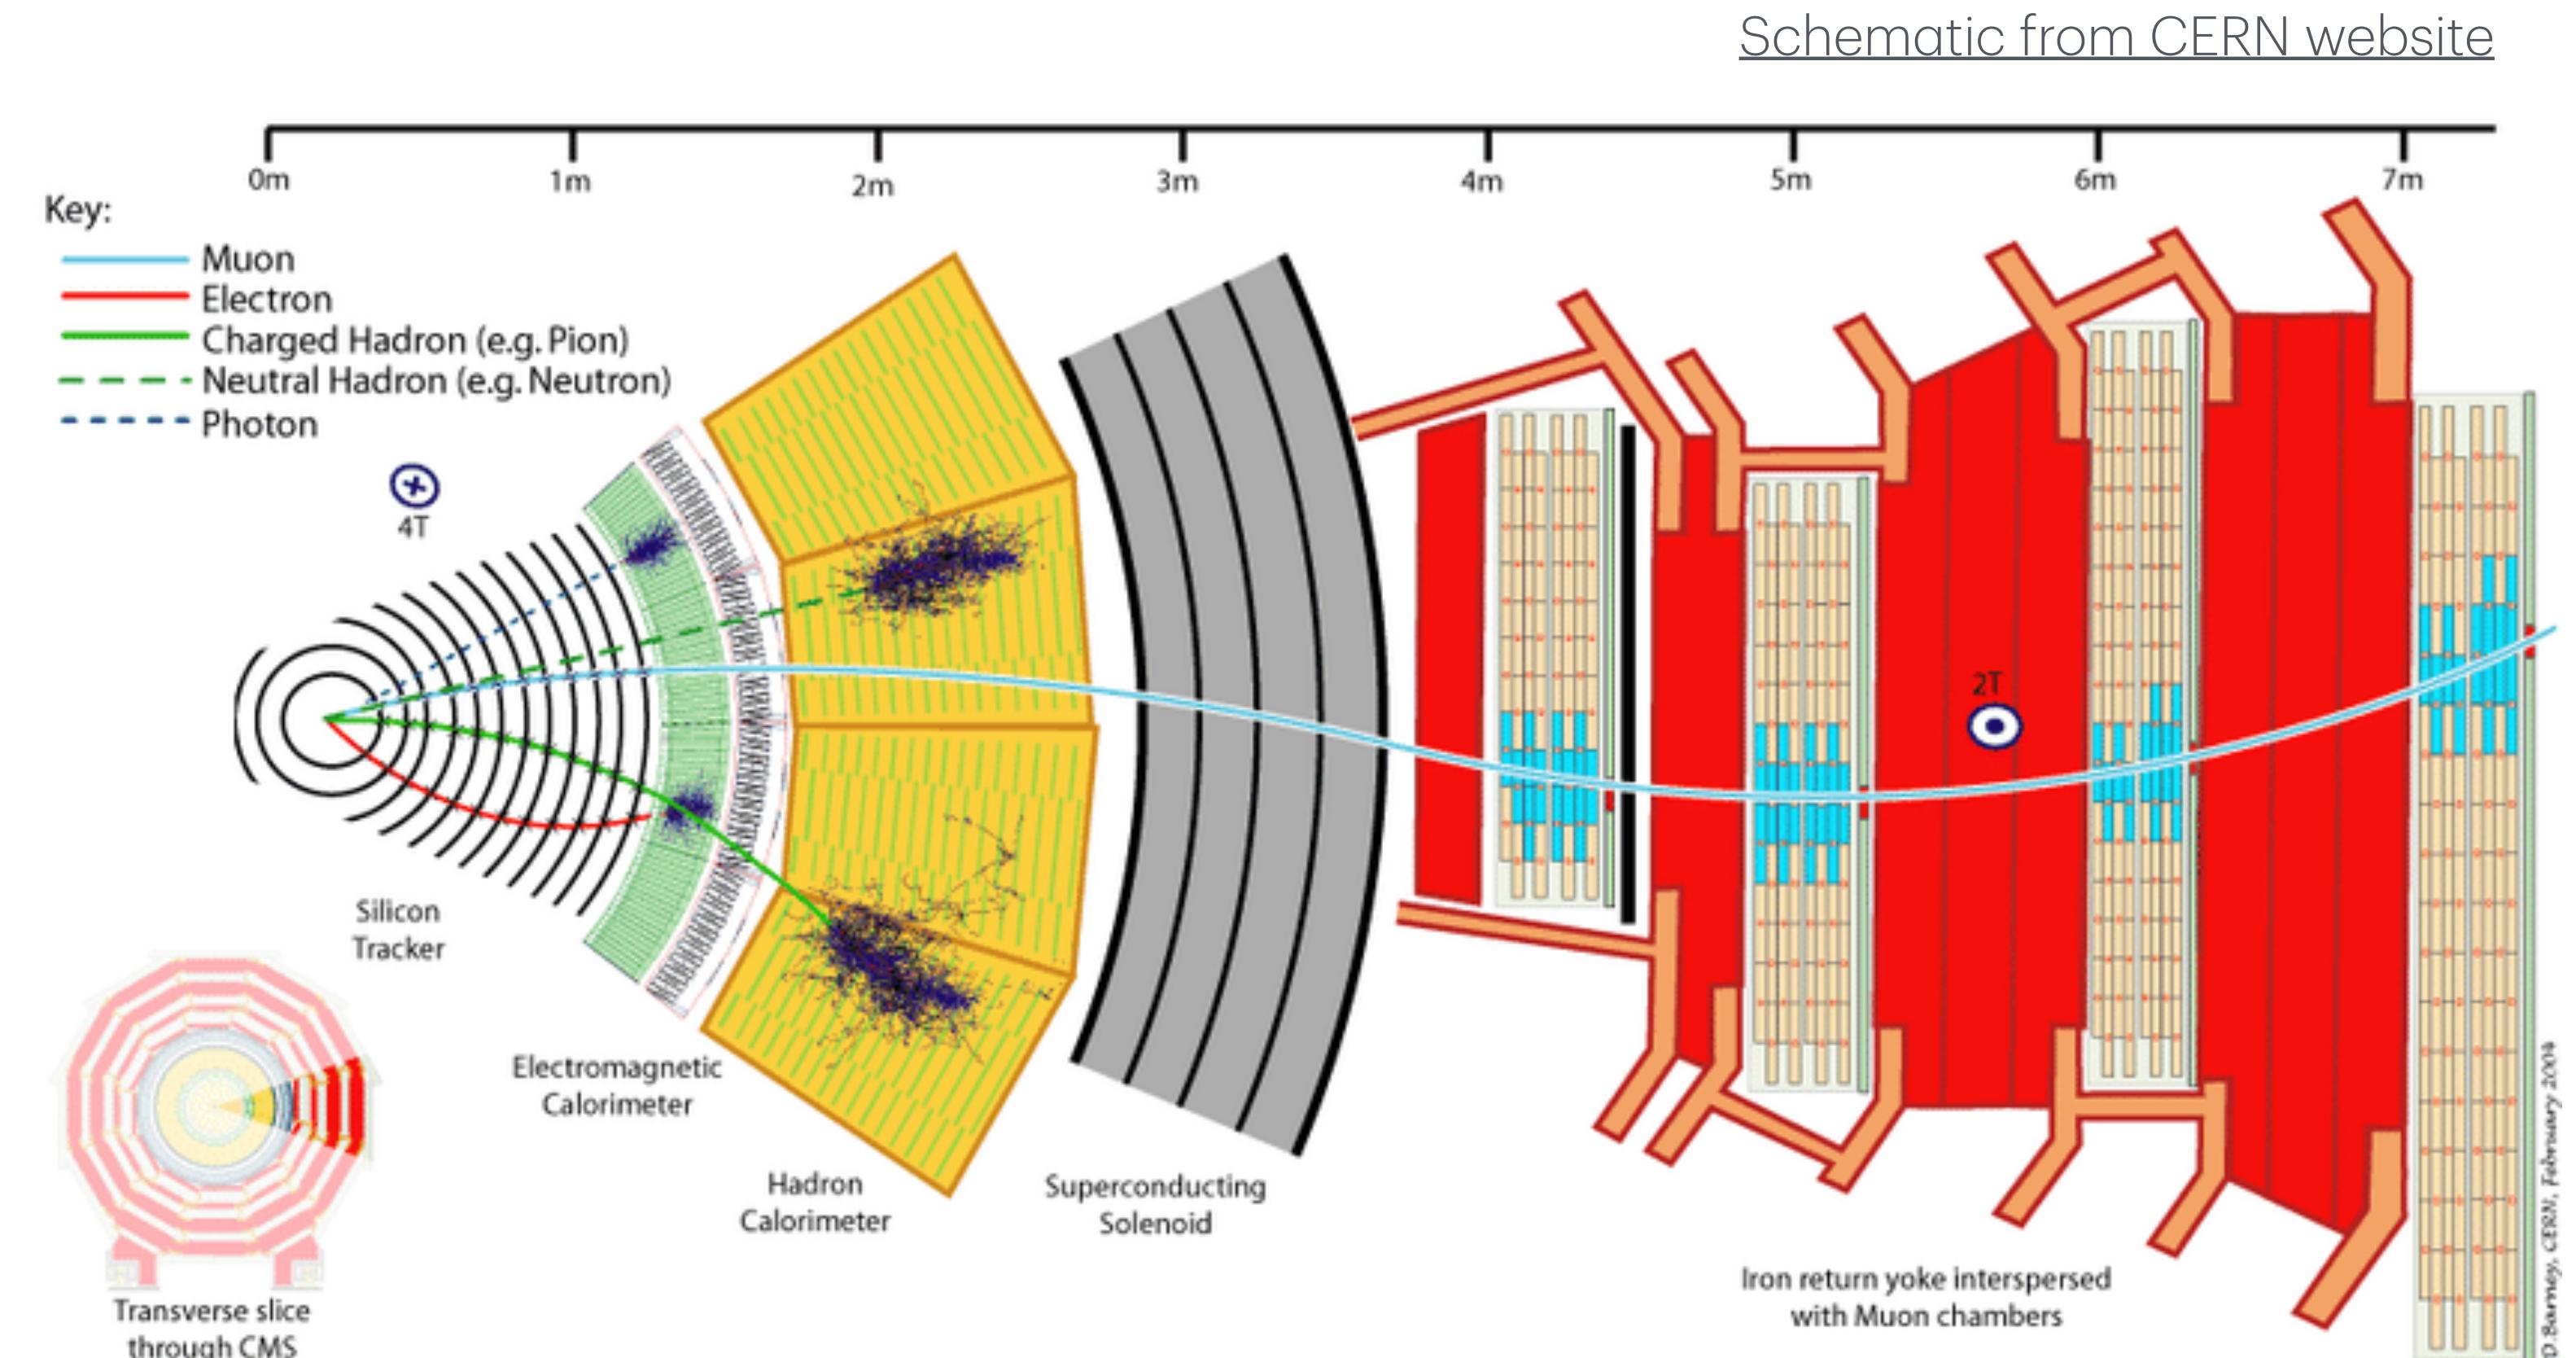


[Image from CERN website](#)

# What is HEP work about?

Testing the SM (and other theories)

We build big expensive experiments (e.g. CMS) to test theoretical predictions



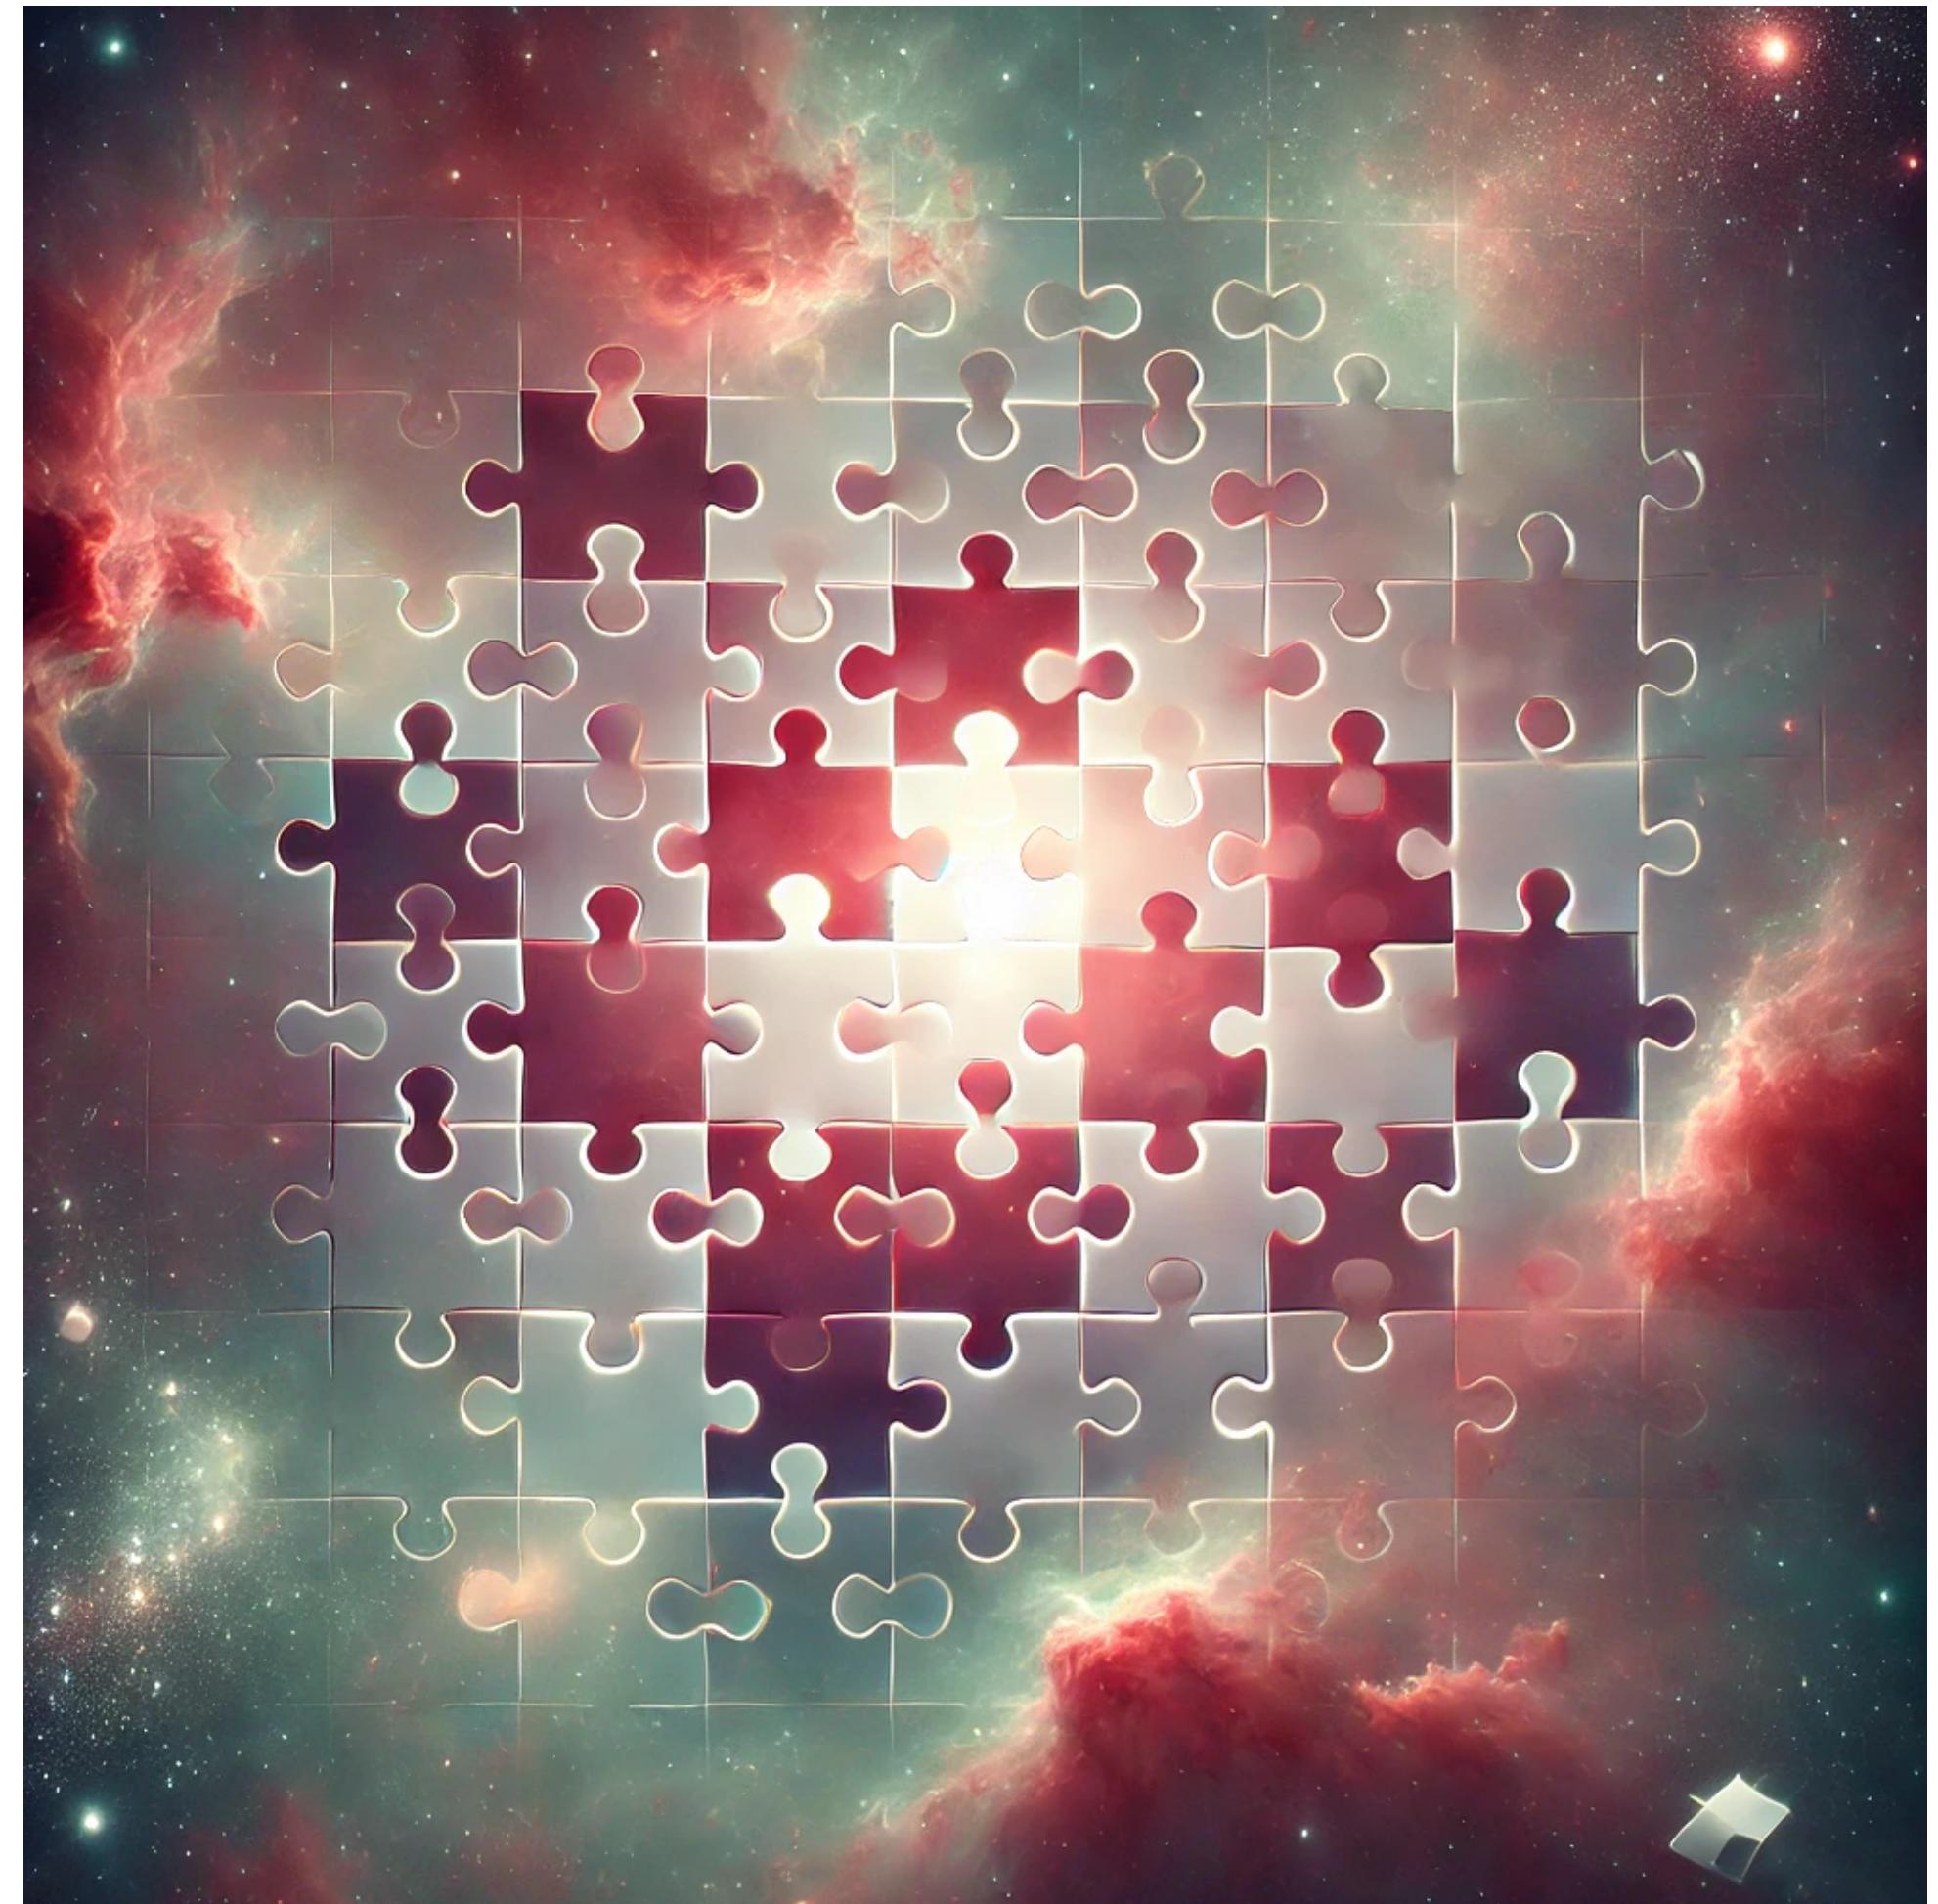
# What is HEP work about?

What we know so far

---

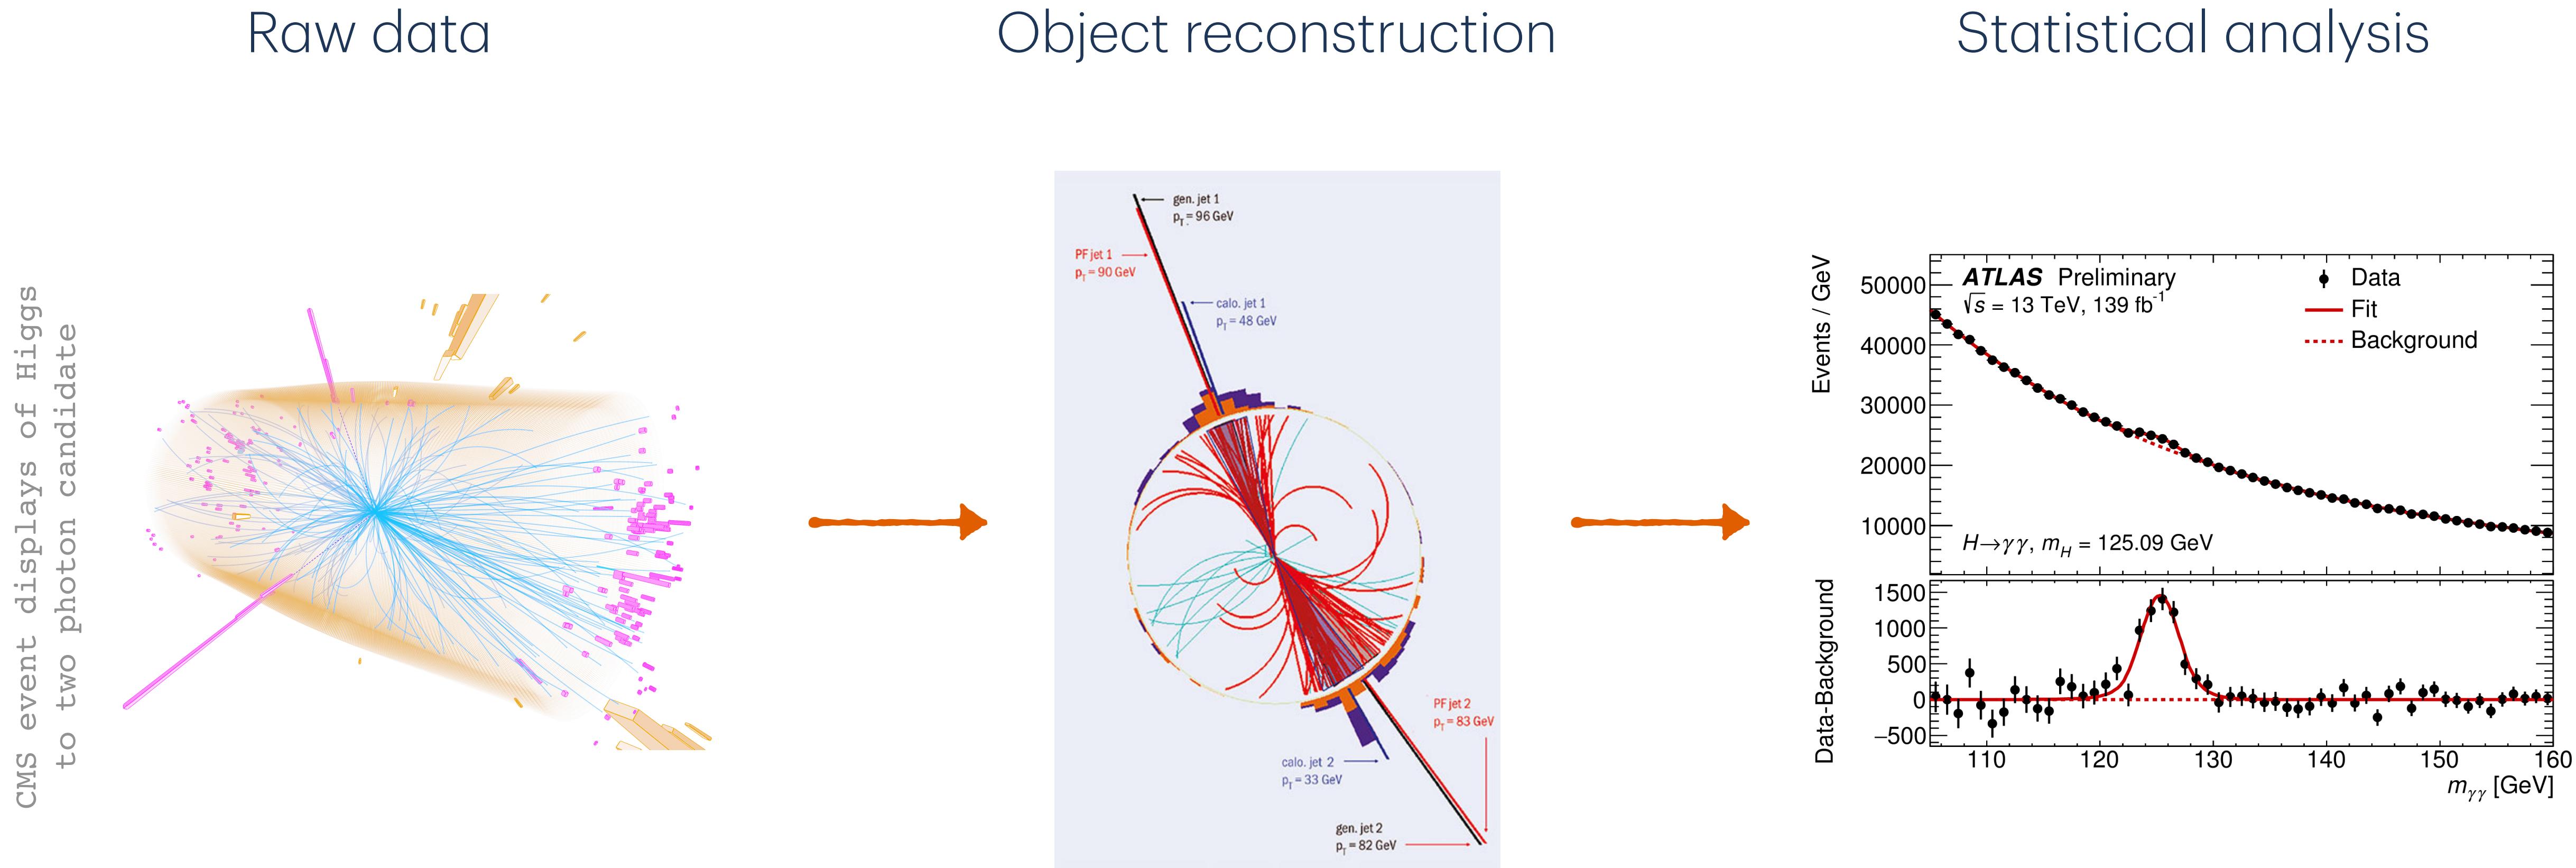
Two main headlines:

1. Almost all SM predictions have been confirmed to high precision
2. The SM is unable to explain some fundamental observations in nature — must look Beyond SM (BSM)



# What is a HEP measurement?

## An overview of the pipeline

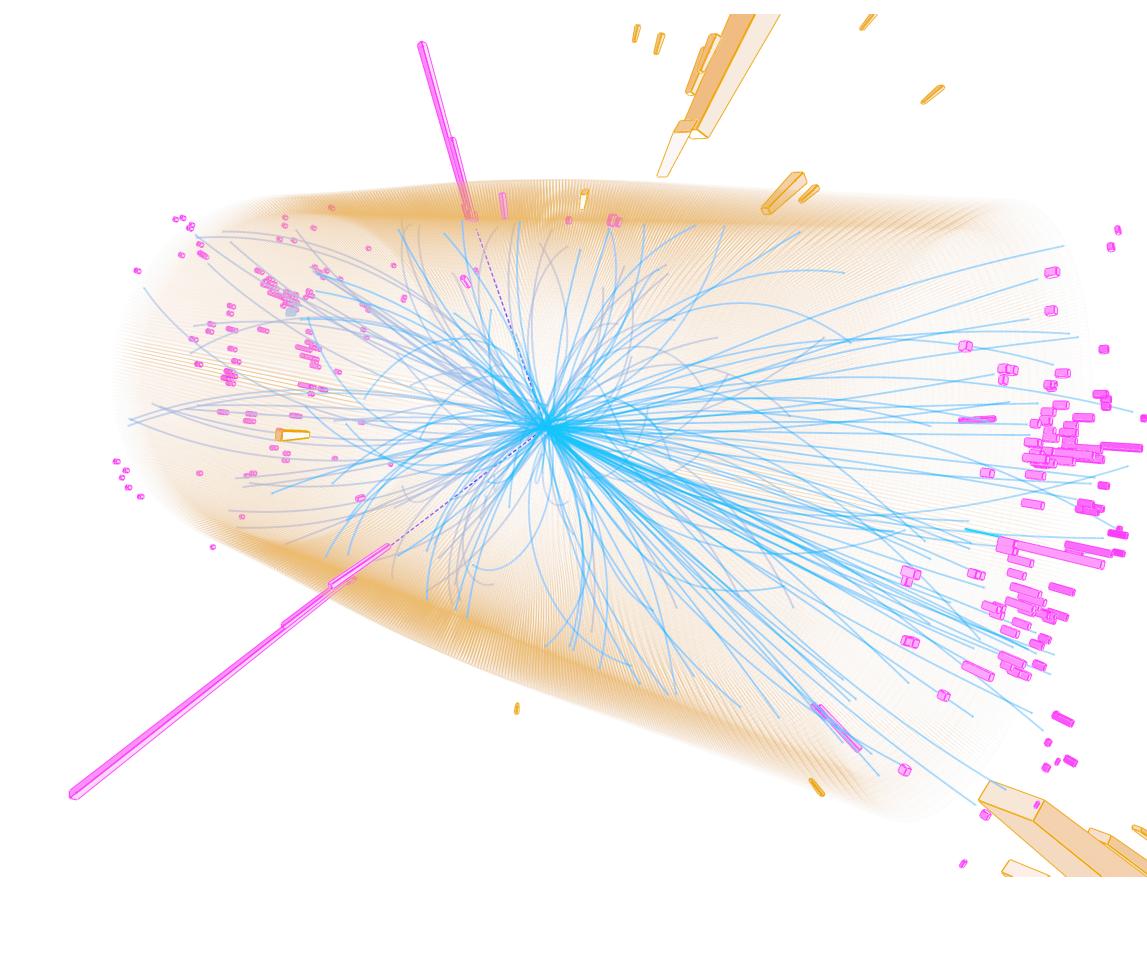


# What is a HEP measurement?

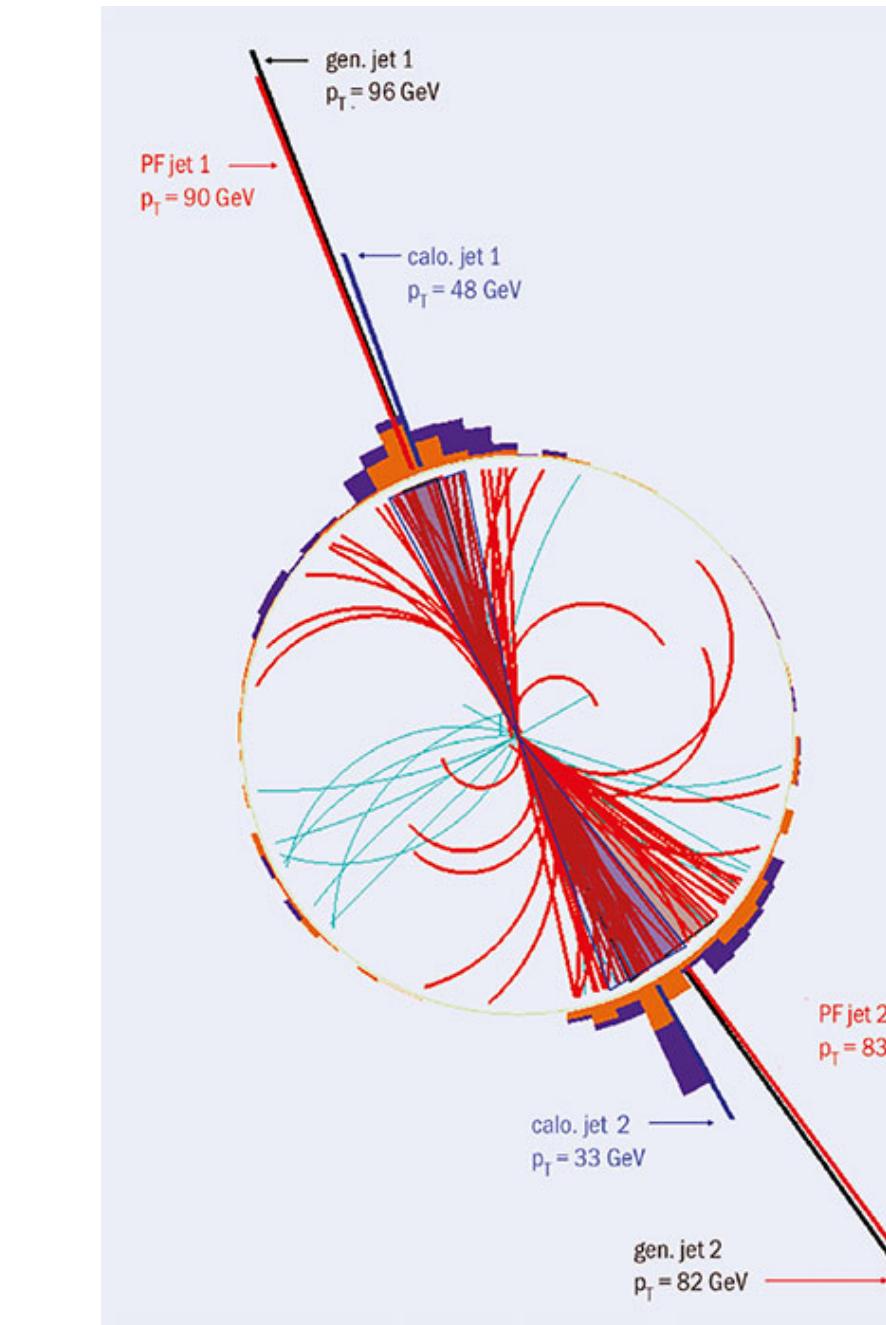
## An overview of the pipeline

Raw data

CMS event displays of Higgs  
to two photon candidate

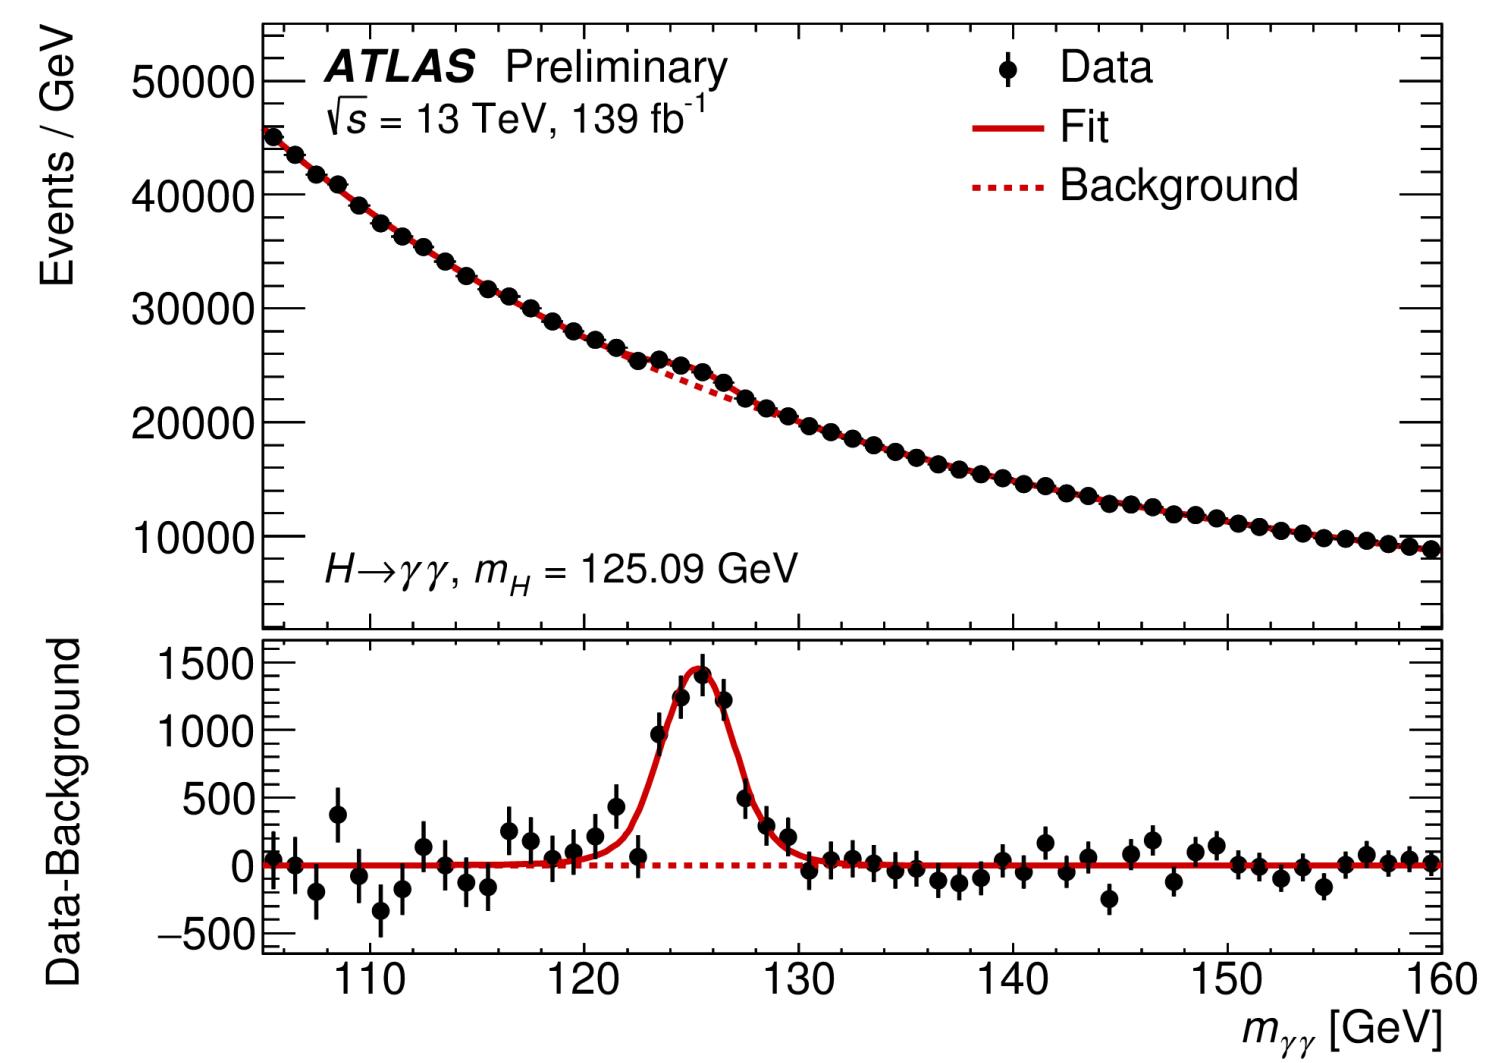


Object reconstruction



Statistical analysis

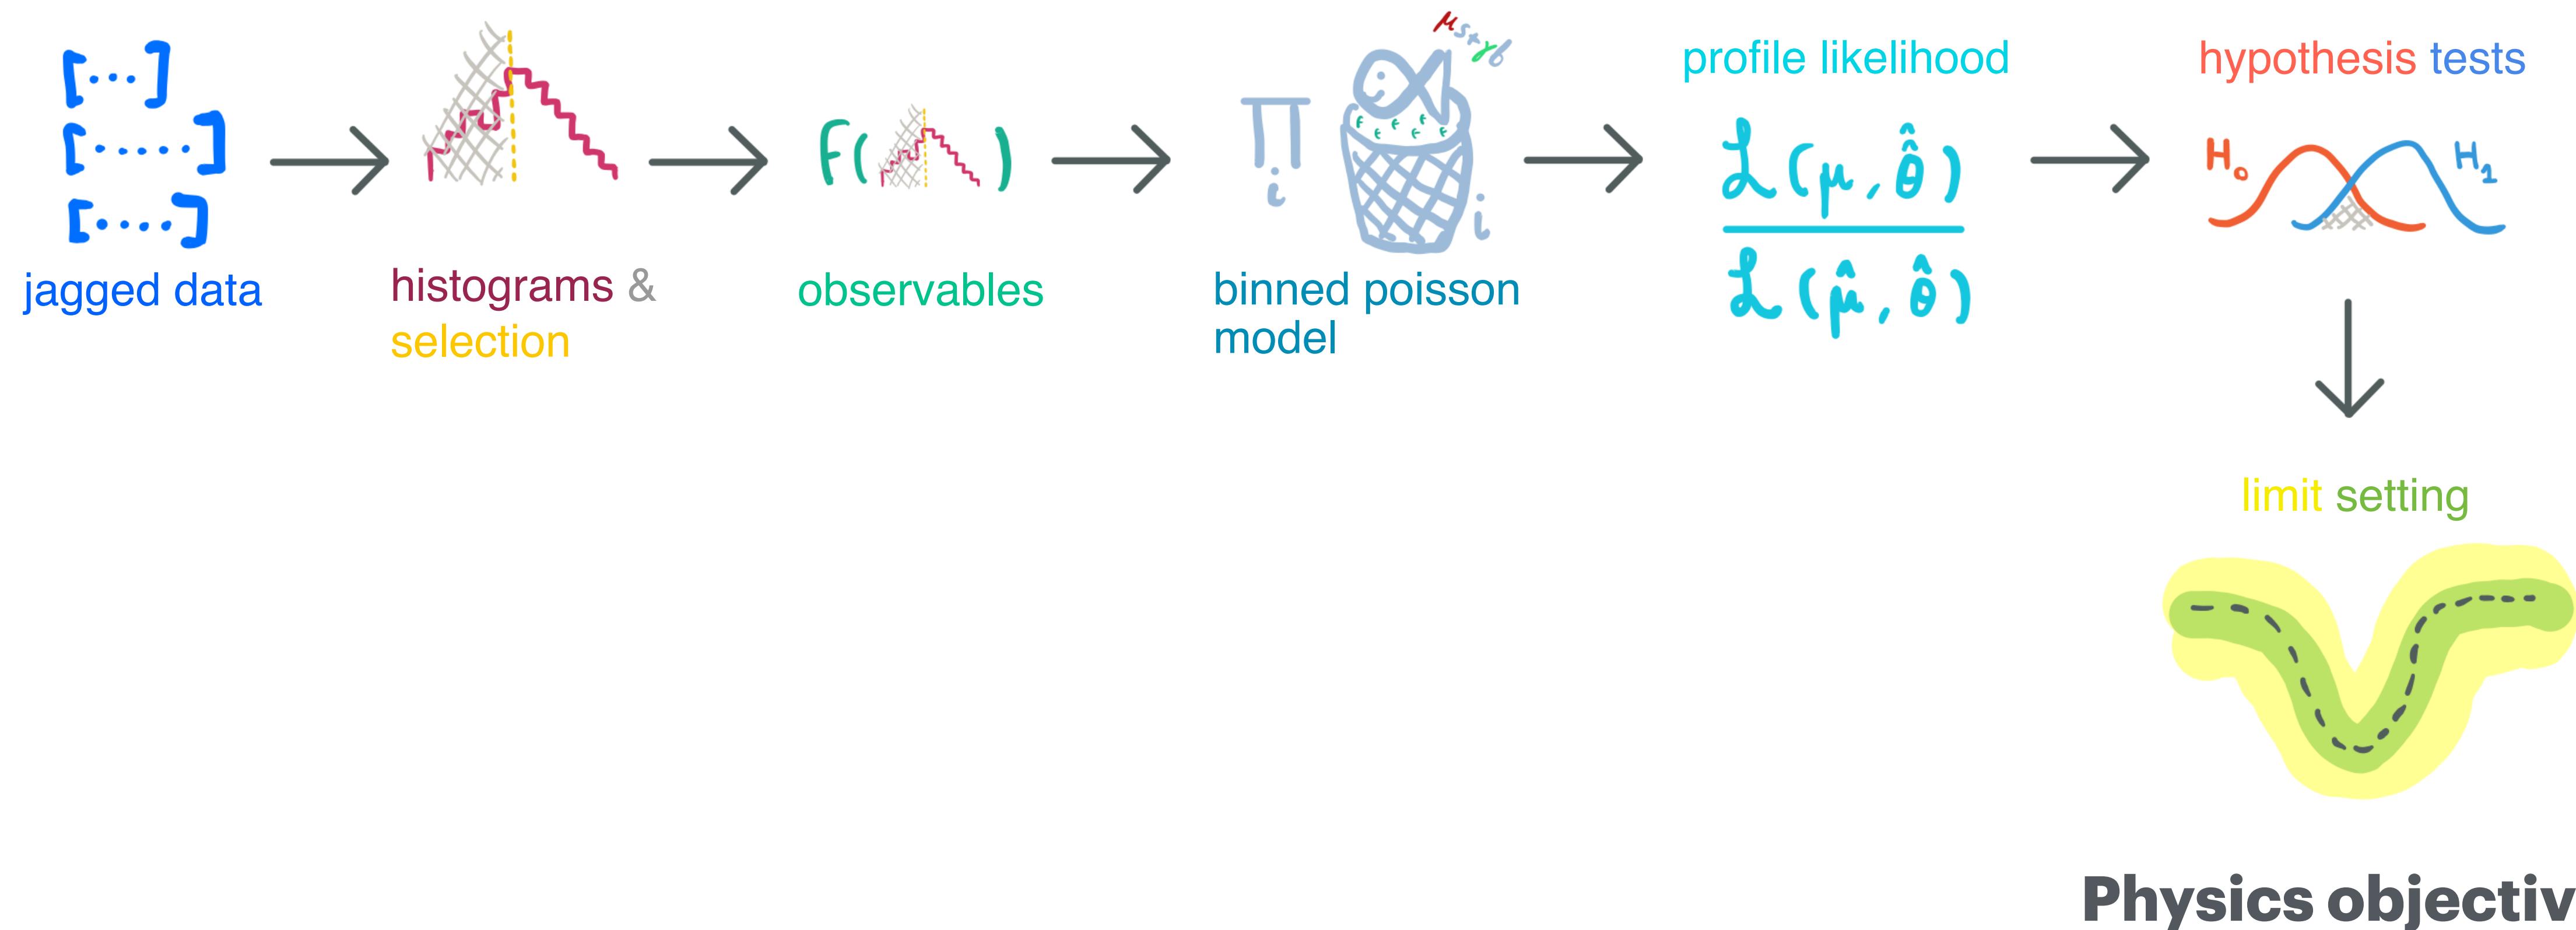
Focus of today



# A HEP statistical analysis

From reconstructed particles to a physics objective

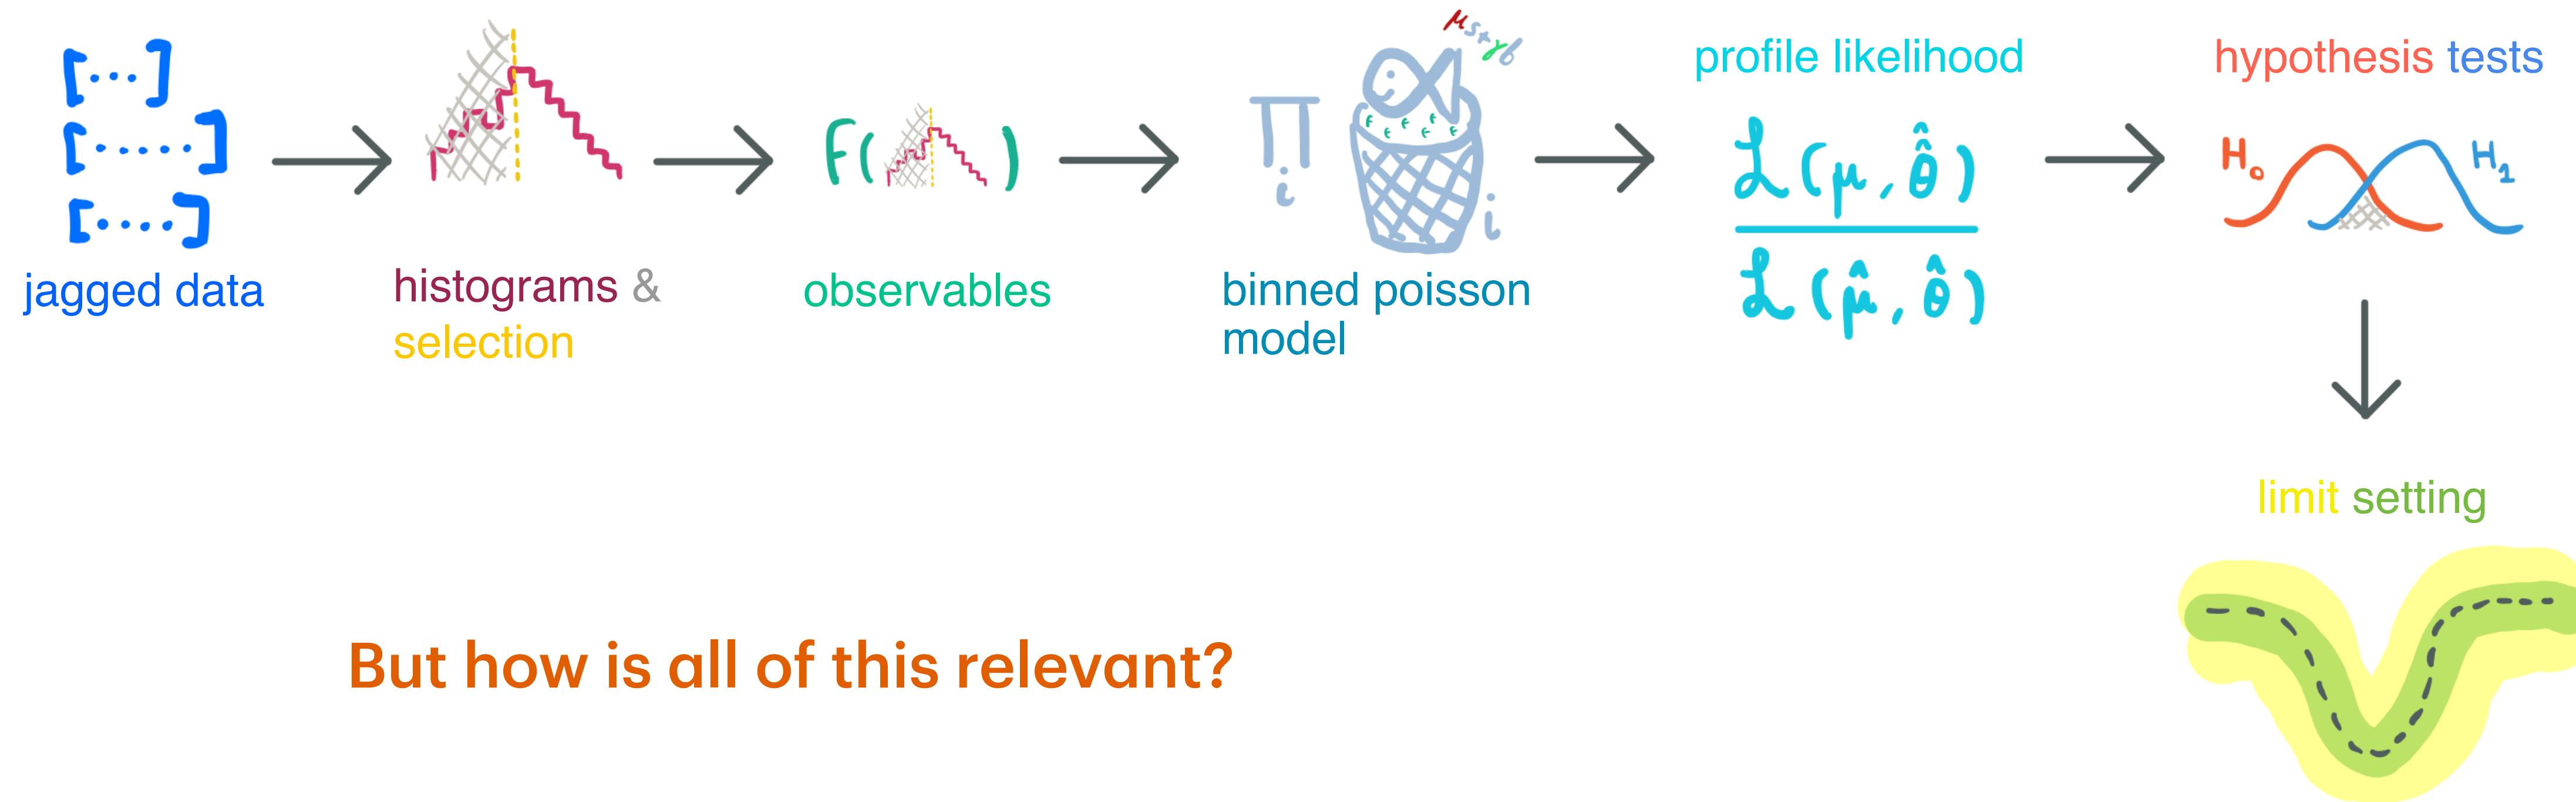
From [Nathan Simpson's talk](#)



# A HEP statistical analysis

From reconstructed particles to a physics objective

From [Nathan Simpson's talk](#)

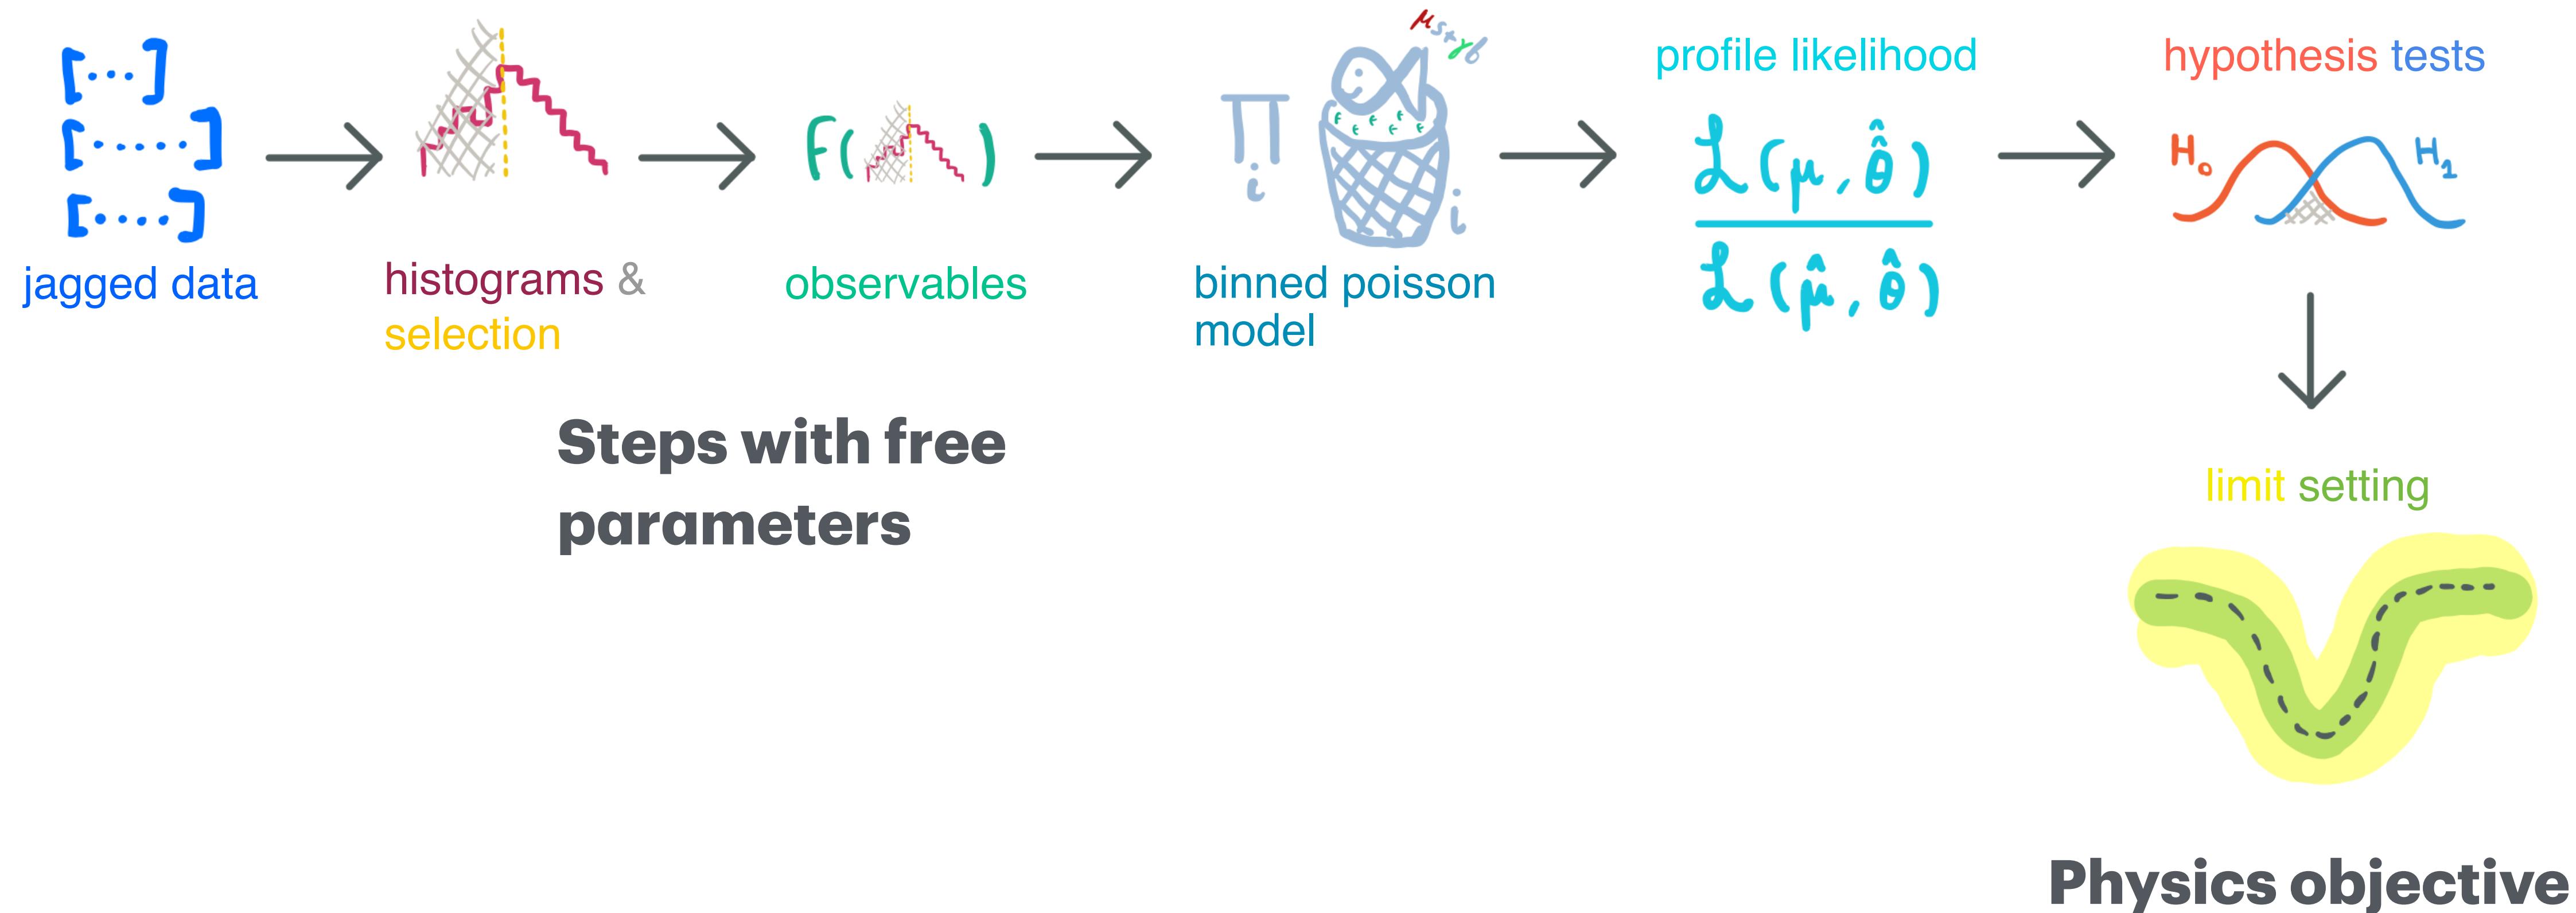


But how is all of this relevant?

# A HEP statistical analysis

From reconstructed particles to a physics objective

From [Nathan Simpson's talk](#)



# A HEP statistical analysis

From reconstructed particles to a physics objective

From [Nathan Simpson's talk](#)



# The need for HEP gradients

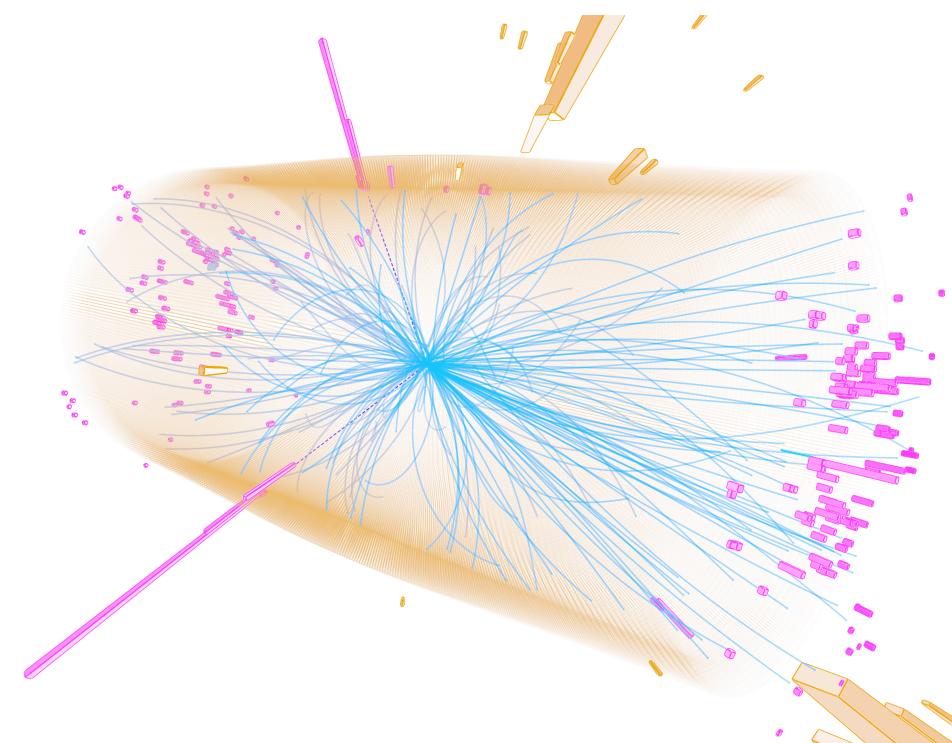
# Why trace gradients?

Towards end-to-end learning?

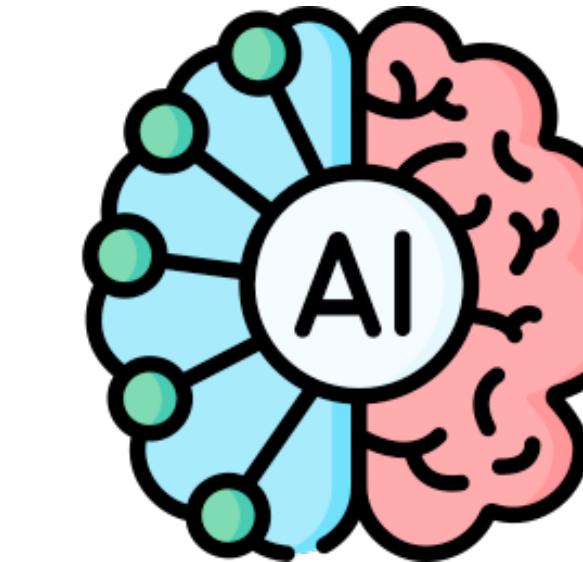
AI/ML entering more and more steps in the HEP pipeline →  
learnable end-to-end analyses?

CMS event displays of Higgs  
to two photon candidate

Raw data



Black box



Physics objective



Higgs found!

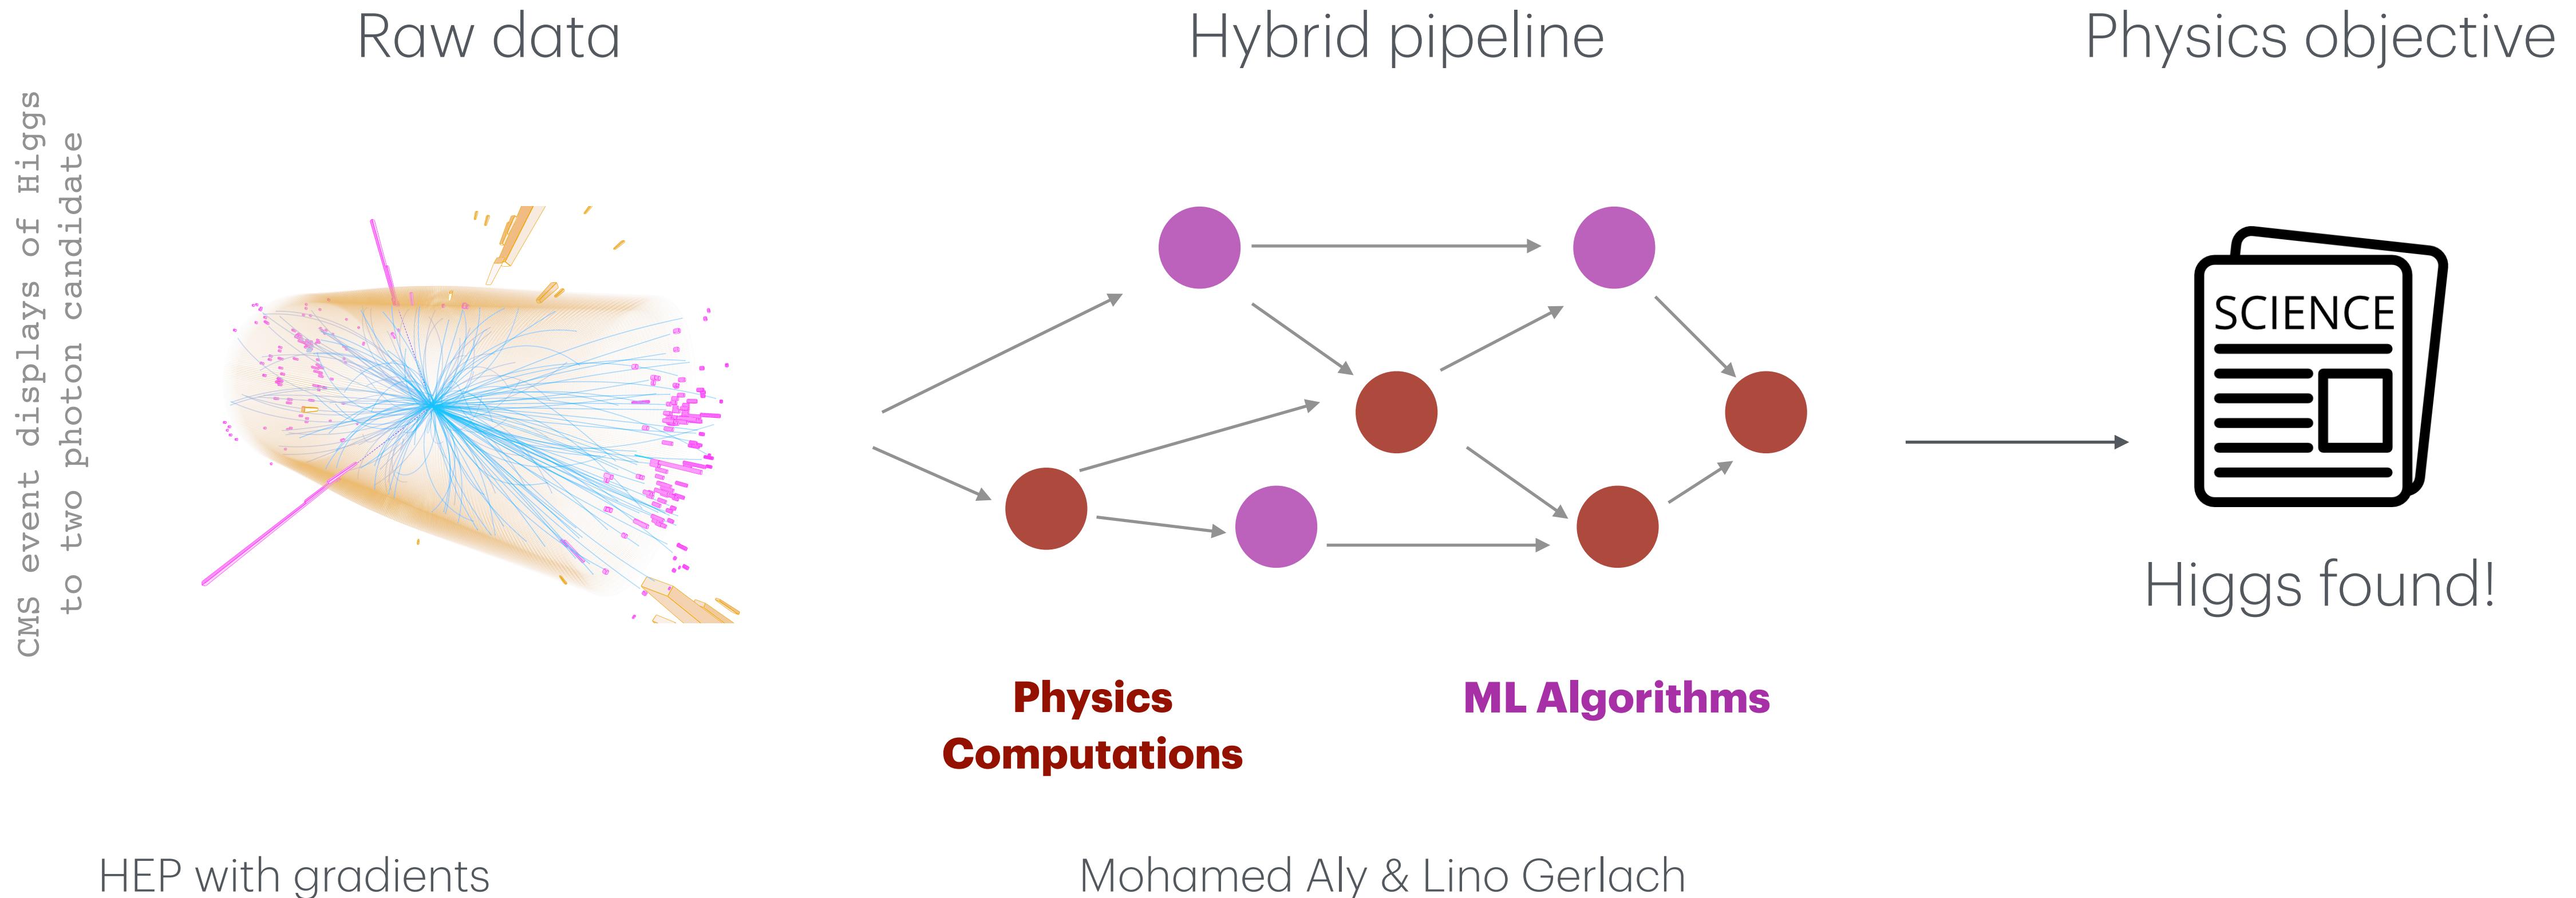
Do we trust this?

# Why trace gradients?

Inserting physics knowledge

Inspired by Lukas Heinrich's talk

More realistically: a physics-informed hybrid approach →  
knowledge inserted via data flows between nodes

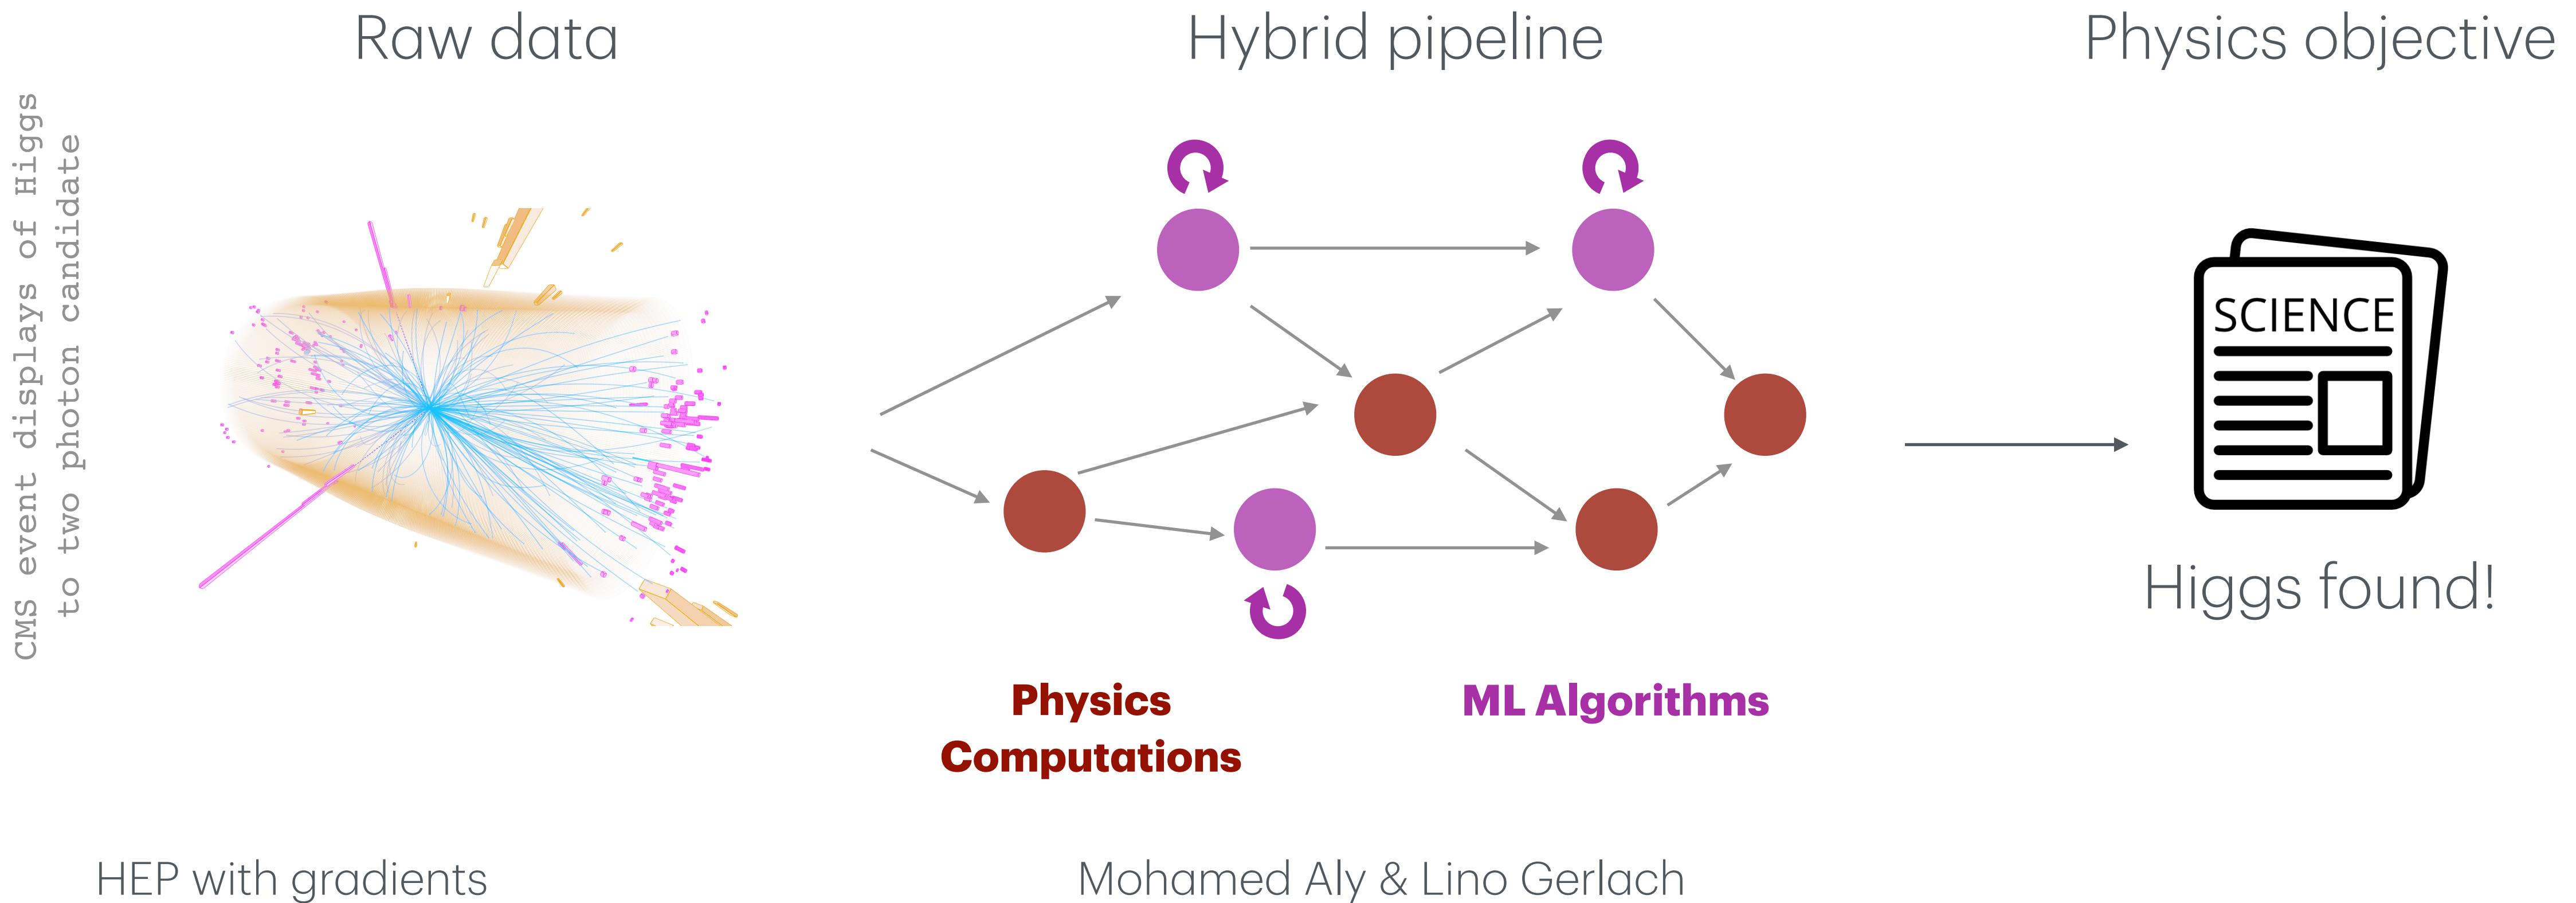


# Why trace gradients?

Holistic optimisation of hybrid pipeline

Inspired by Lukas Heinrich's talk

Each ML step is trained independently to optimise an objective

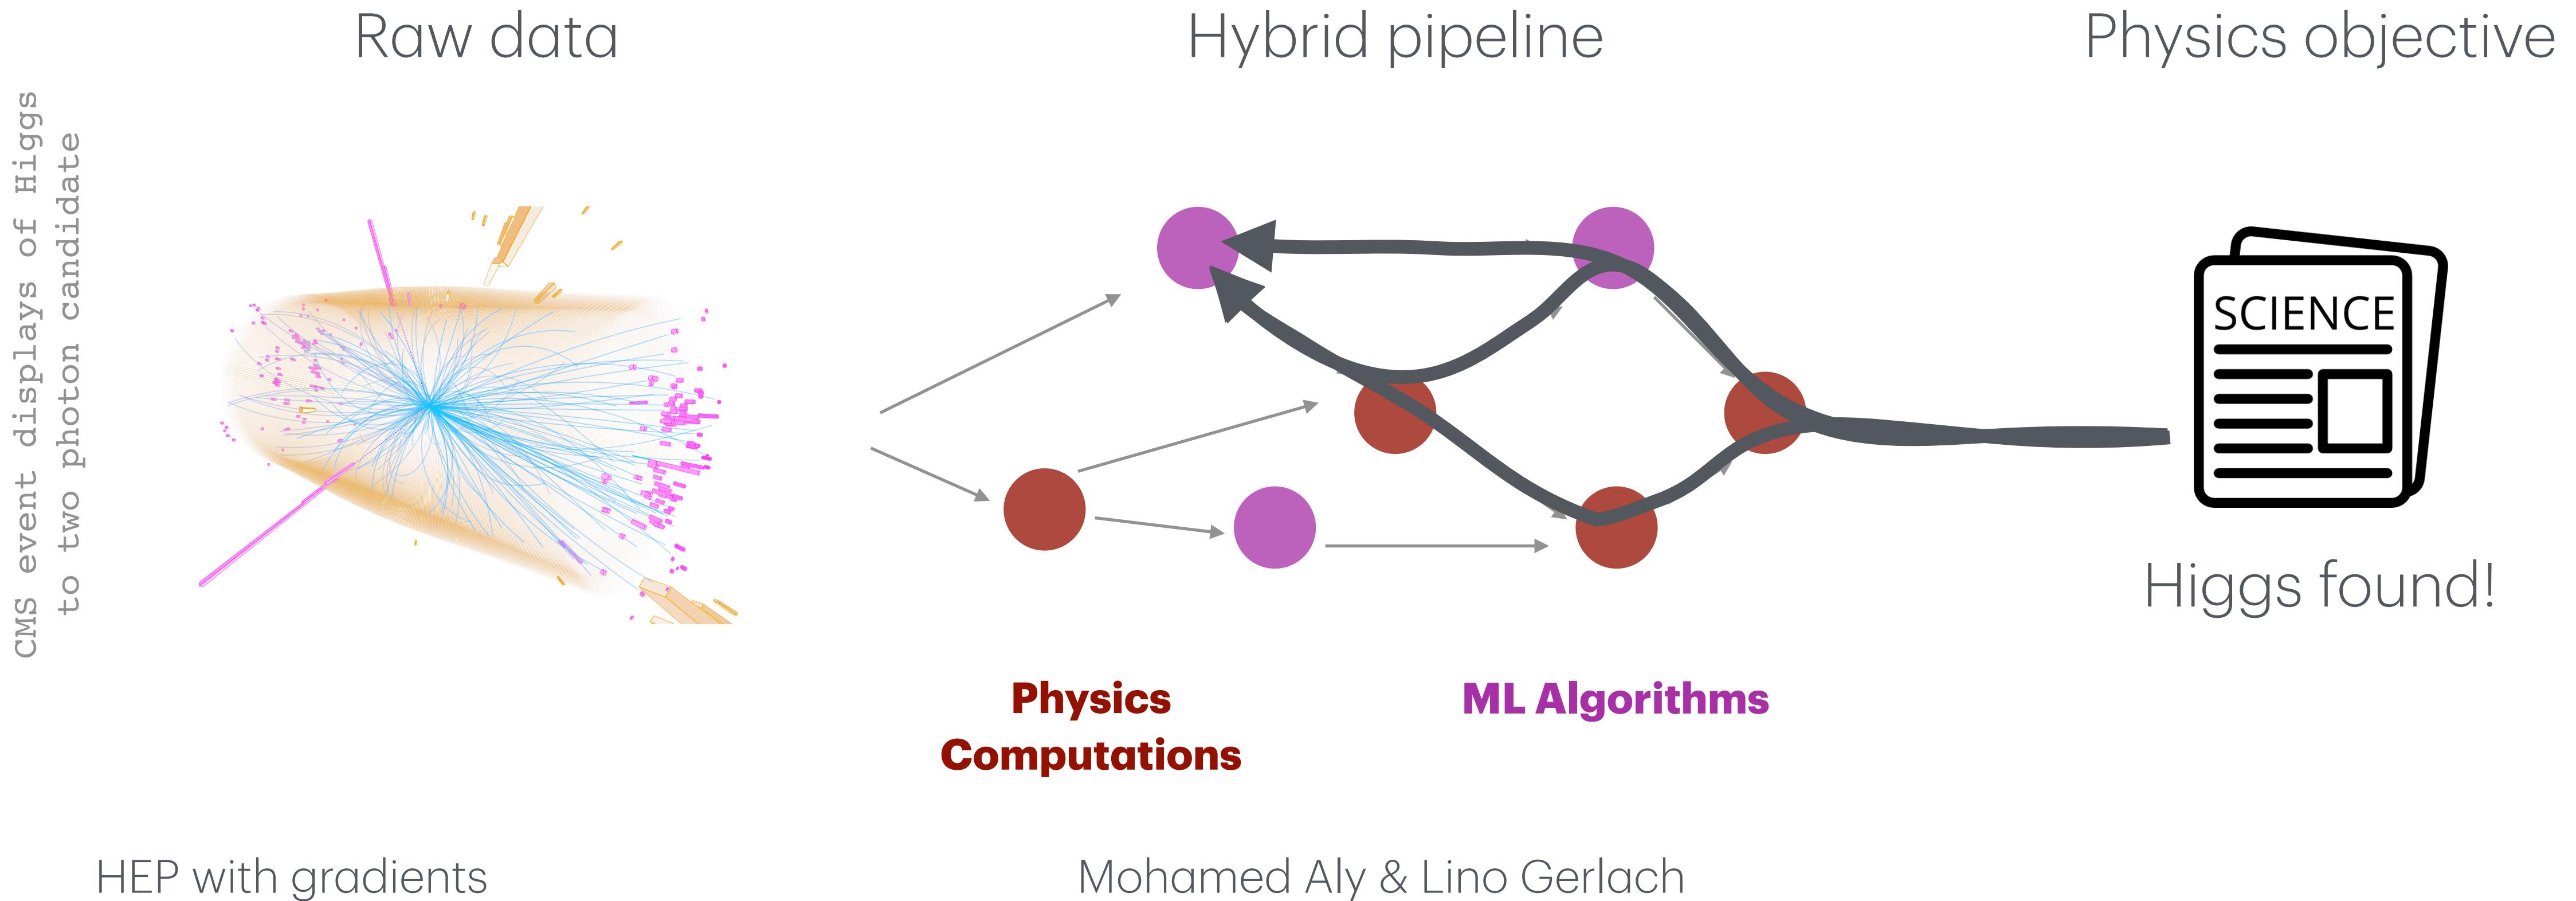


# Why trace gradients?

## Holistic optimisation of hybrid pipeline

Inspired by Lukas Heinrich's talk

Ideally, we want to optimise everything for our physics objective  
— i.e. back propagate through the pipeline

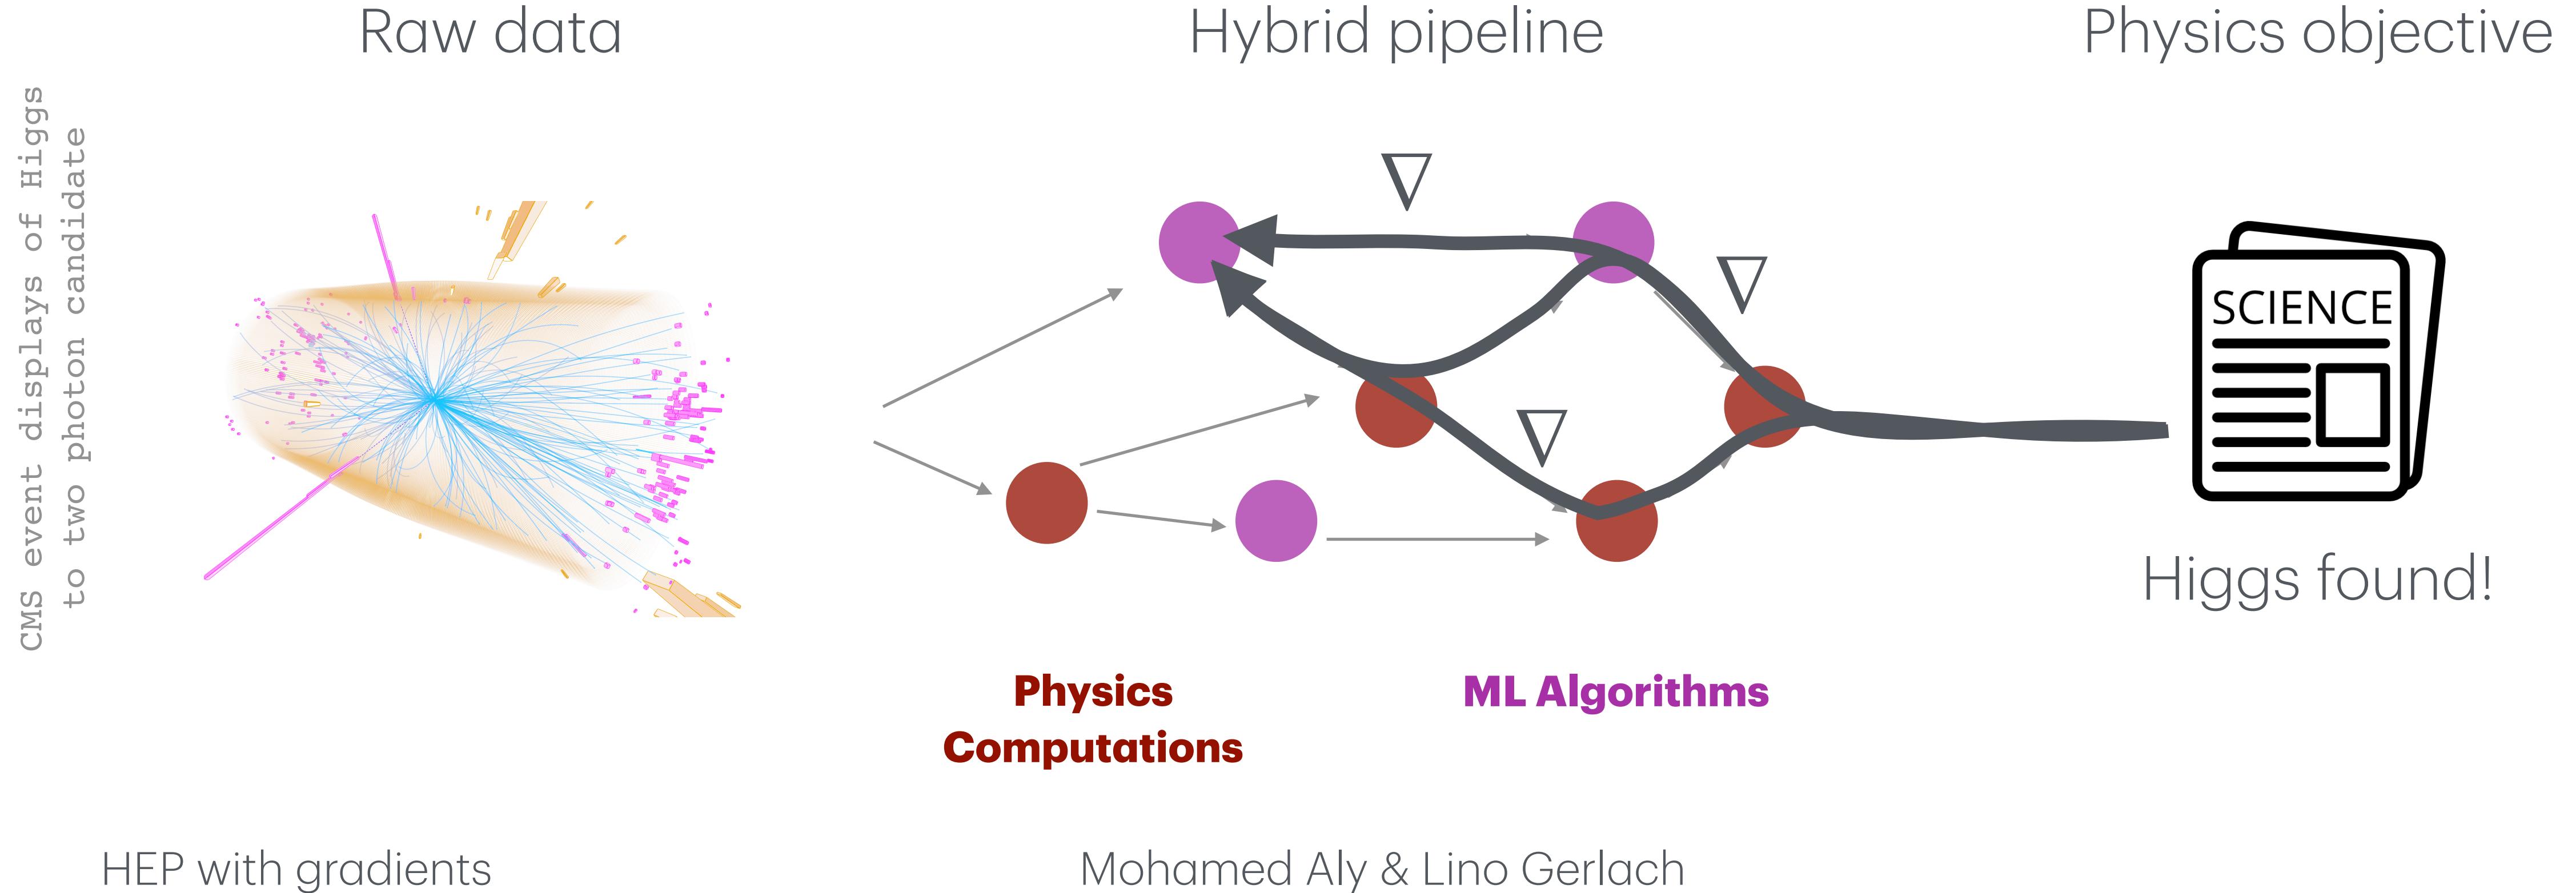


# Why trace gradients?

Holistic optimisation of hybrid pipeline

Inspired by Lukas Heinrich's talk

Need to track gradients for holistic optimisations!

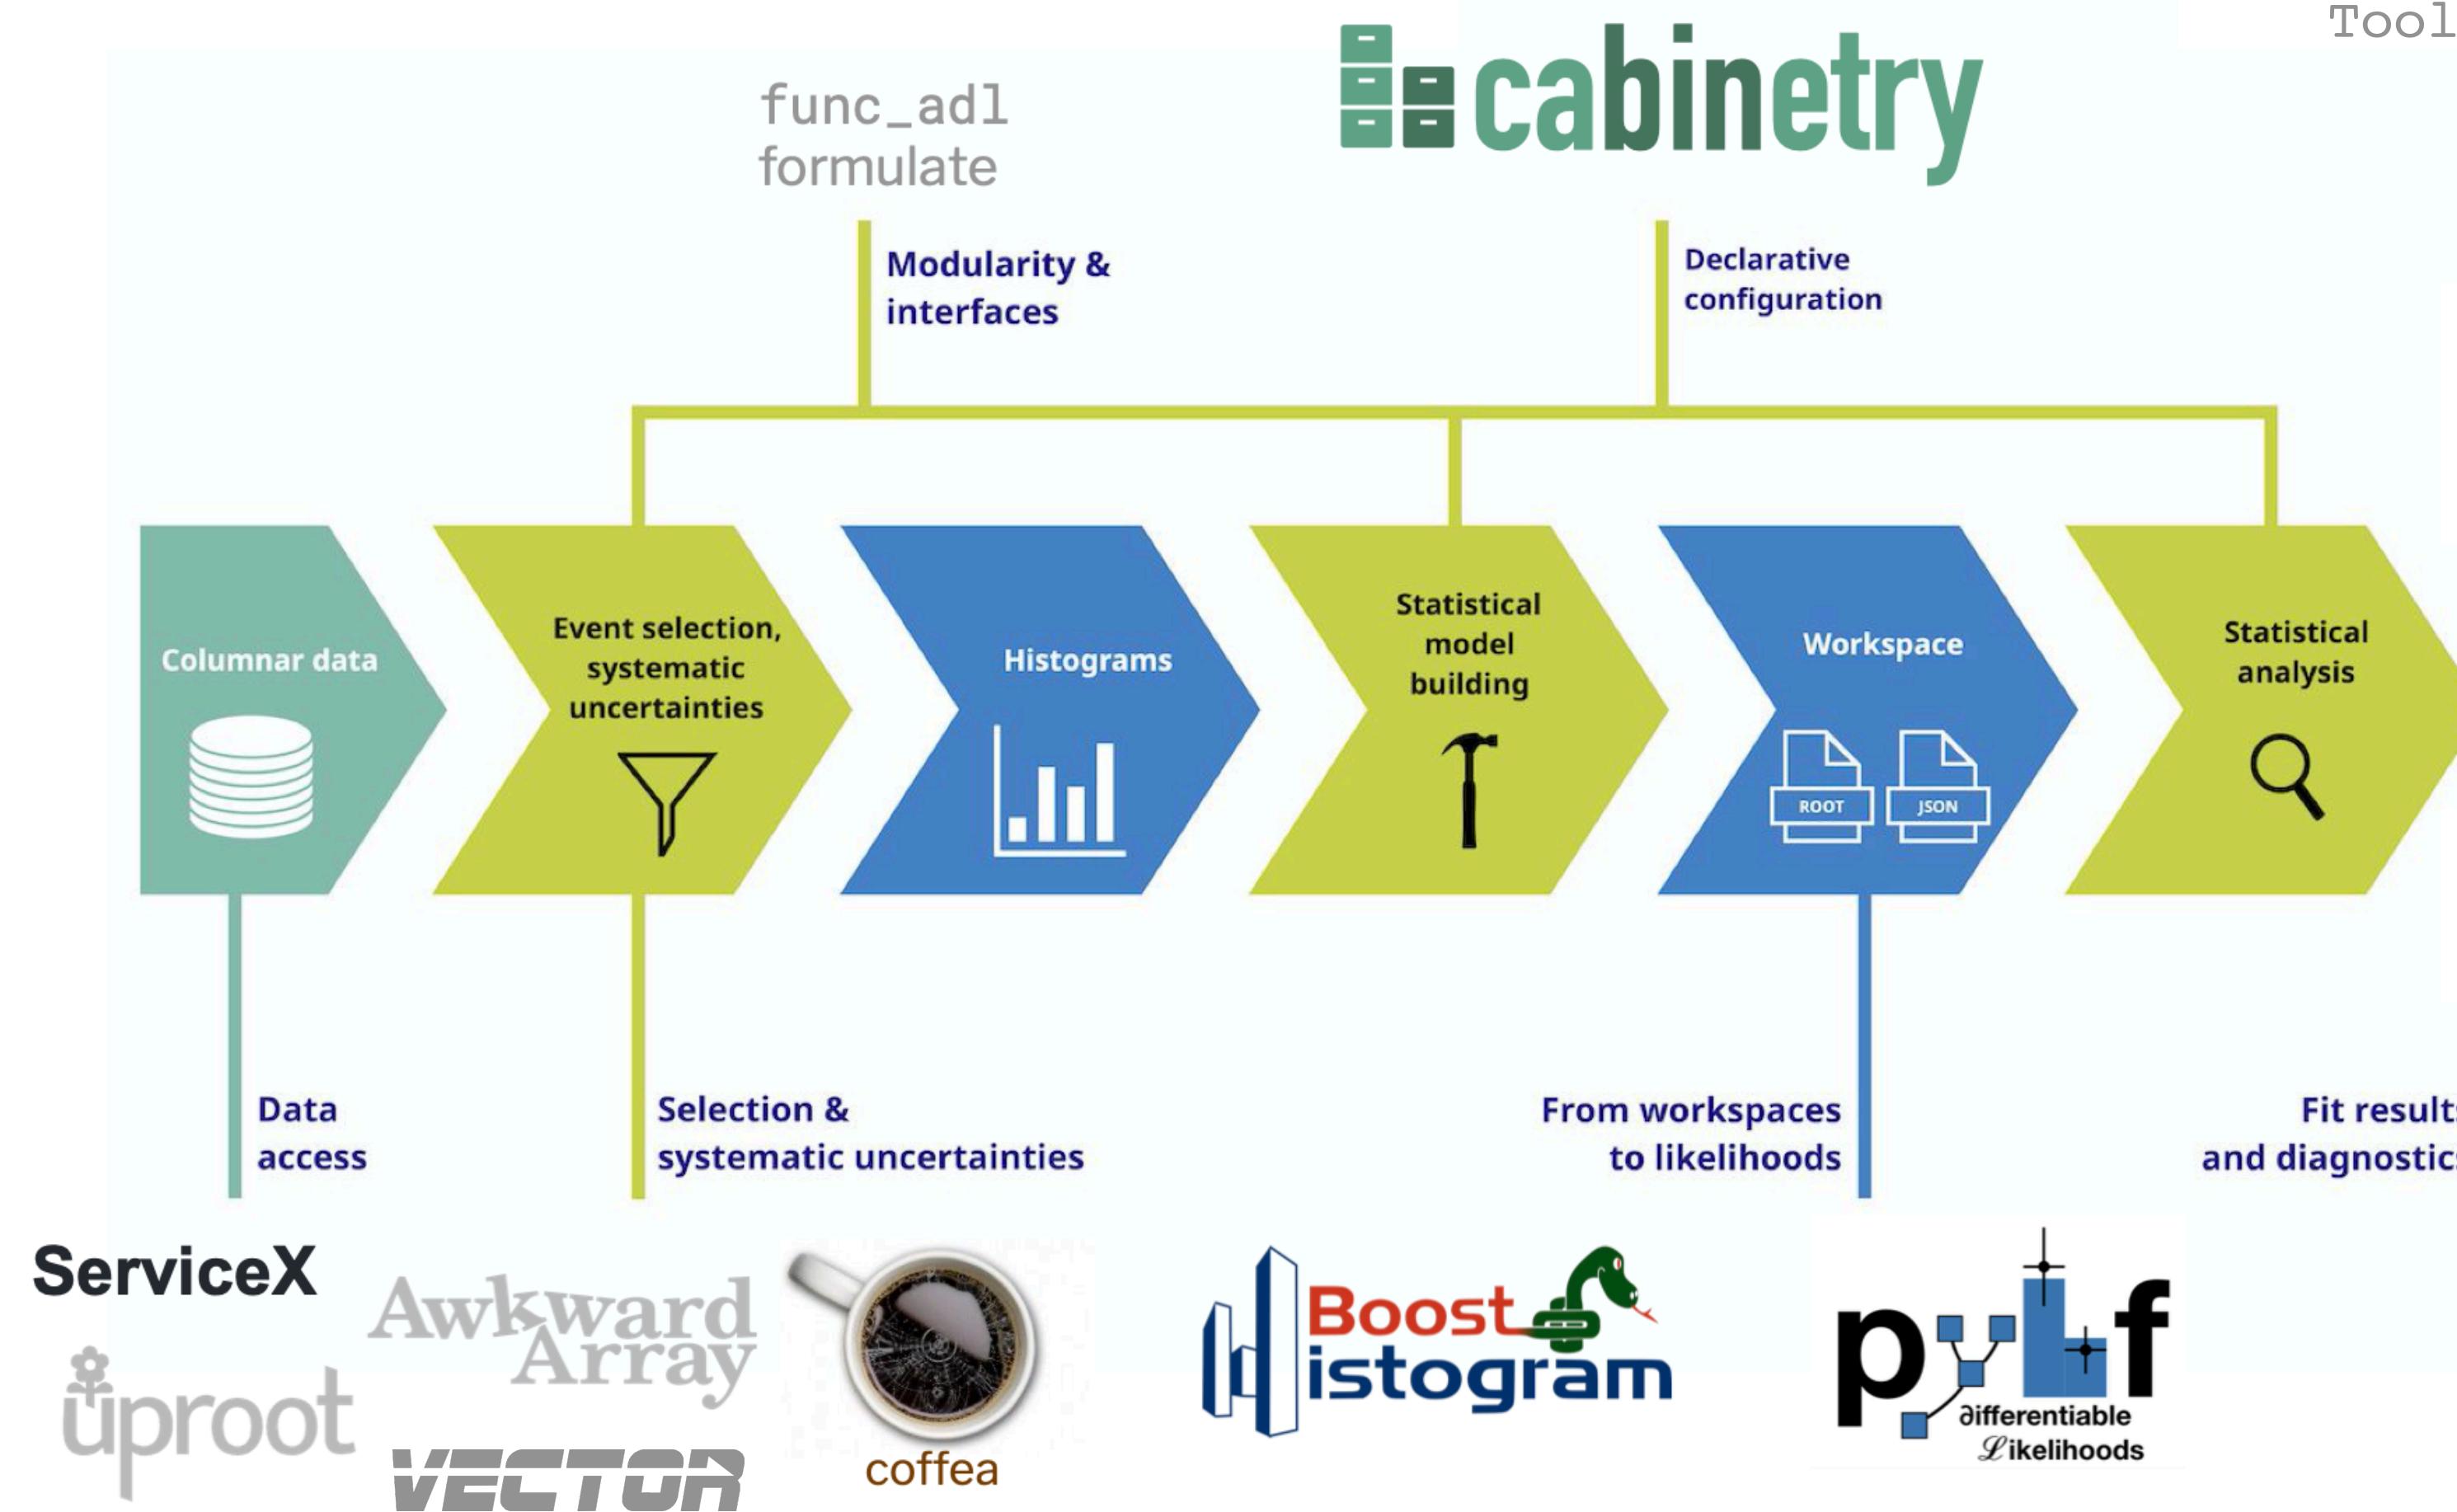


The path to HEP gradients

# The scikit-HEP Ecosystem

Open-source pythonic data analysis tools

Modified from [IRIS-HEP website](#)

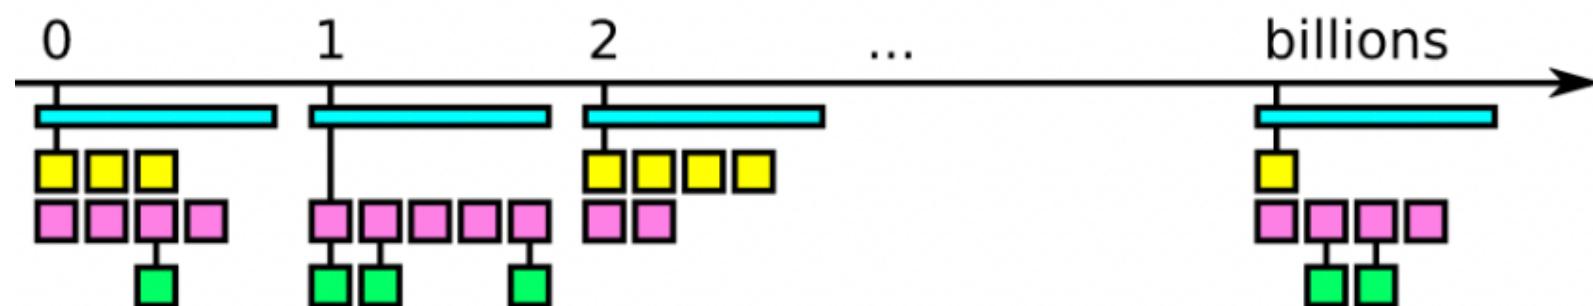


# The Scikit-HEP Ecosystem

Particularly important

## Awkward Array

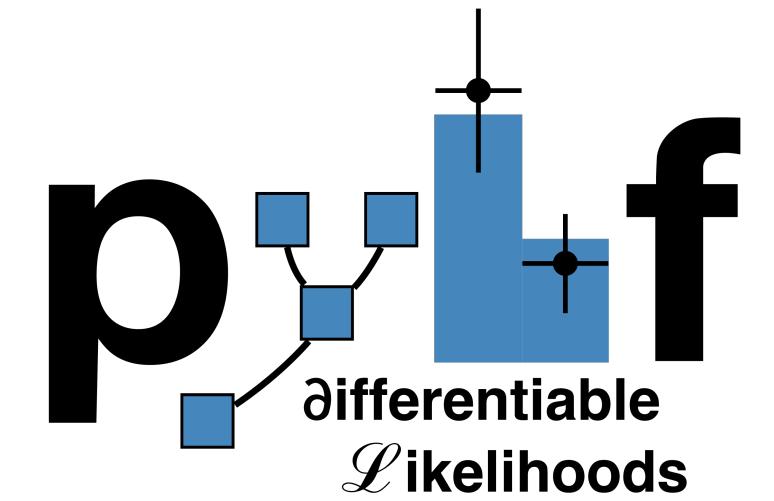
Ragged arrays with  
numpy-like idioms



HEP with gradients

## vector

Support for Lorentz-  
vectors (4-vectors)



Constructing  
differentiable  
likelihoods

Mohamed Aly & Lino Gerlach

# The Scikit-HEP Ecosystem

Particularly important

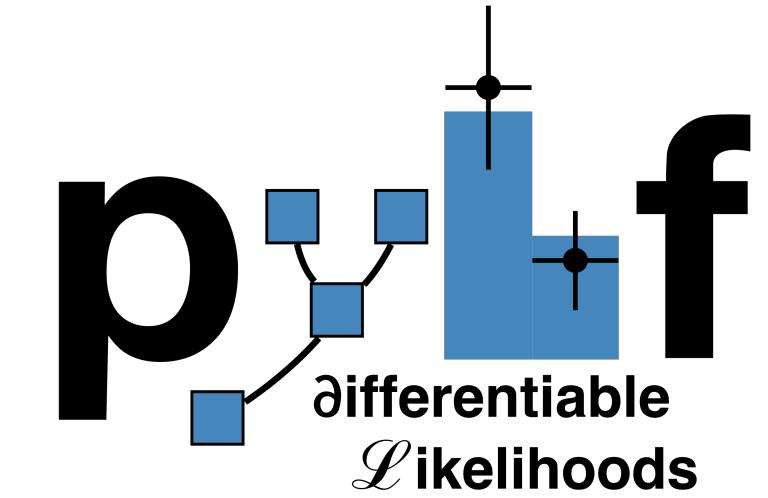
---



Ragged arrays with  
numpy-like idioms



Support for Lorentz-  
vectors (4-vectors)



Constructing  
differentiable  
likelihoods

We need to trace gradients through all the tools...

# The Scikit-HEP Ecosystem

Particularly important

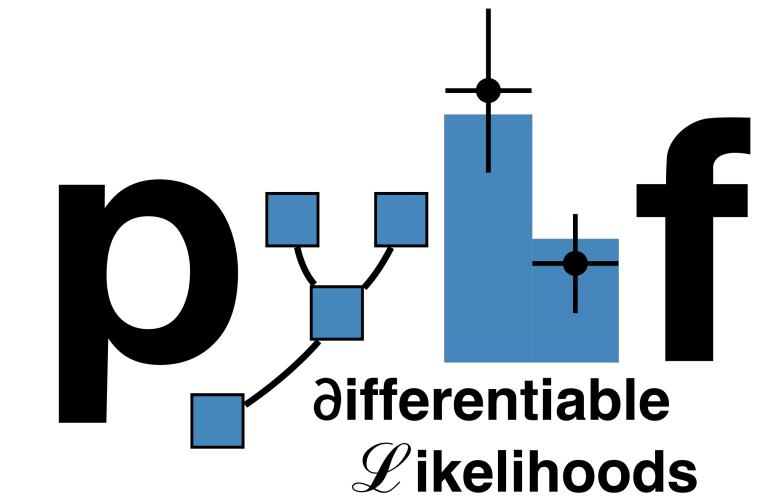
---

## Awkward Array

Ragged arrays with  
numpy-like idioms

## vector

Support for Lorentz-  
vectors (4-vectors)

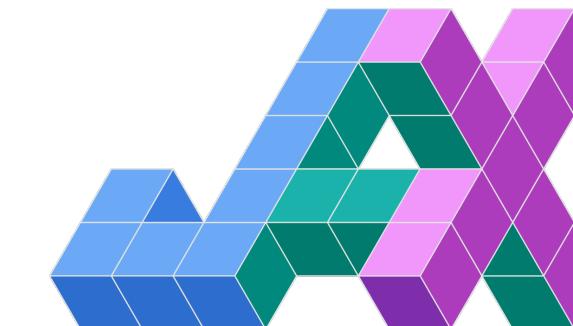


Constructing  
differentiable  
likelihoods

We need to trace gradients through all the tools...



HEP with gradients



Is this enough?

Mohamed Aly & Lino Gerlach

# The Scikit-HEP Ecosystem

Particularly important

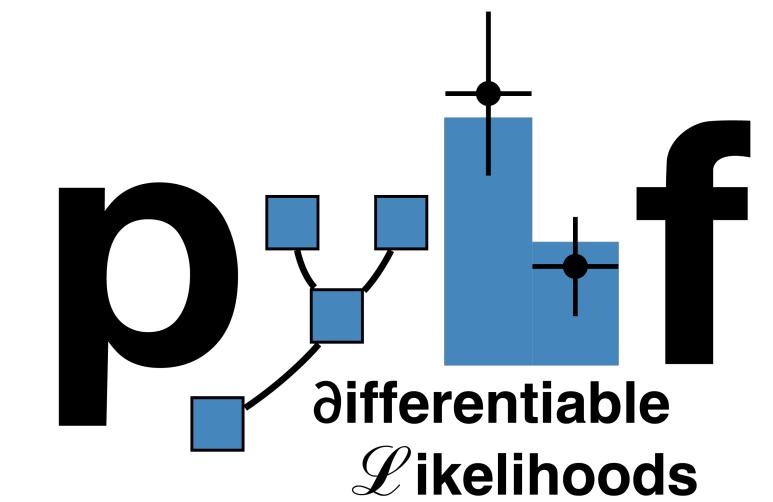
---



Ragged arrays with  
numpy-like idioms

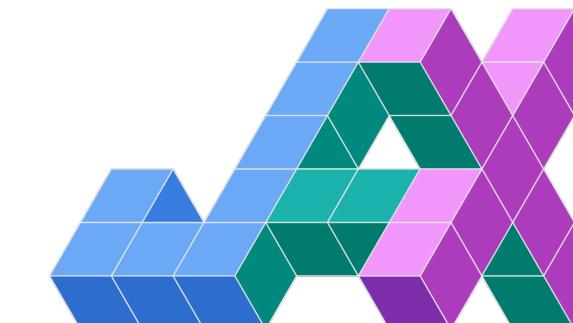


Support for Lorentz-  
vectors (4-vectors)



Constructing  
differentiable  
likelihoods

We need to trace gradients through all the tools...



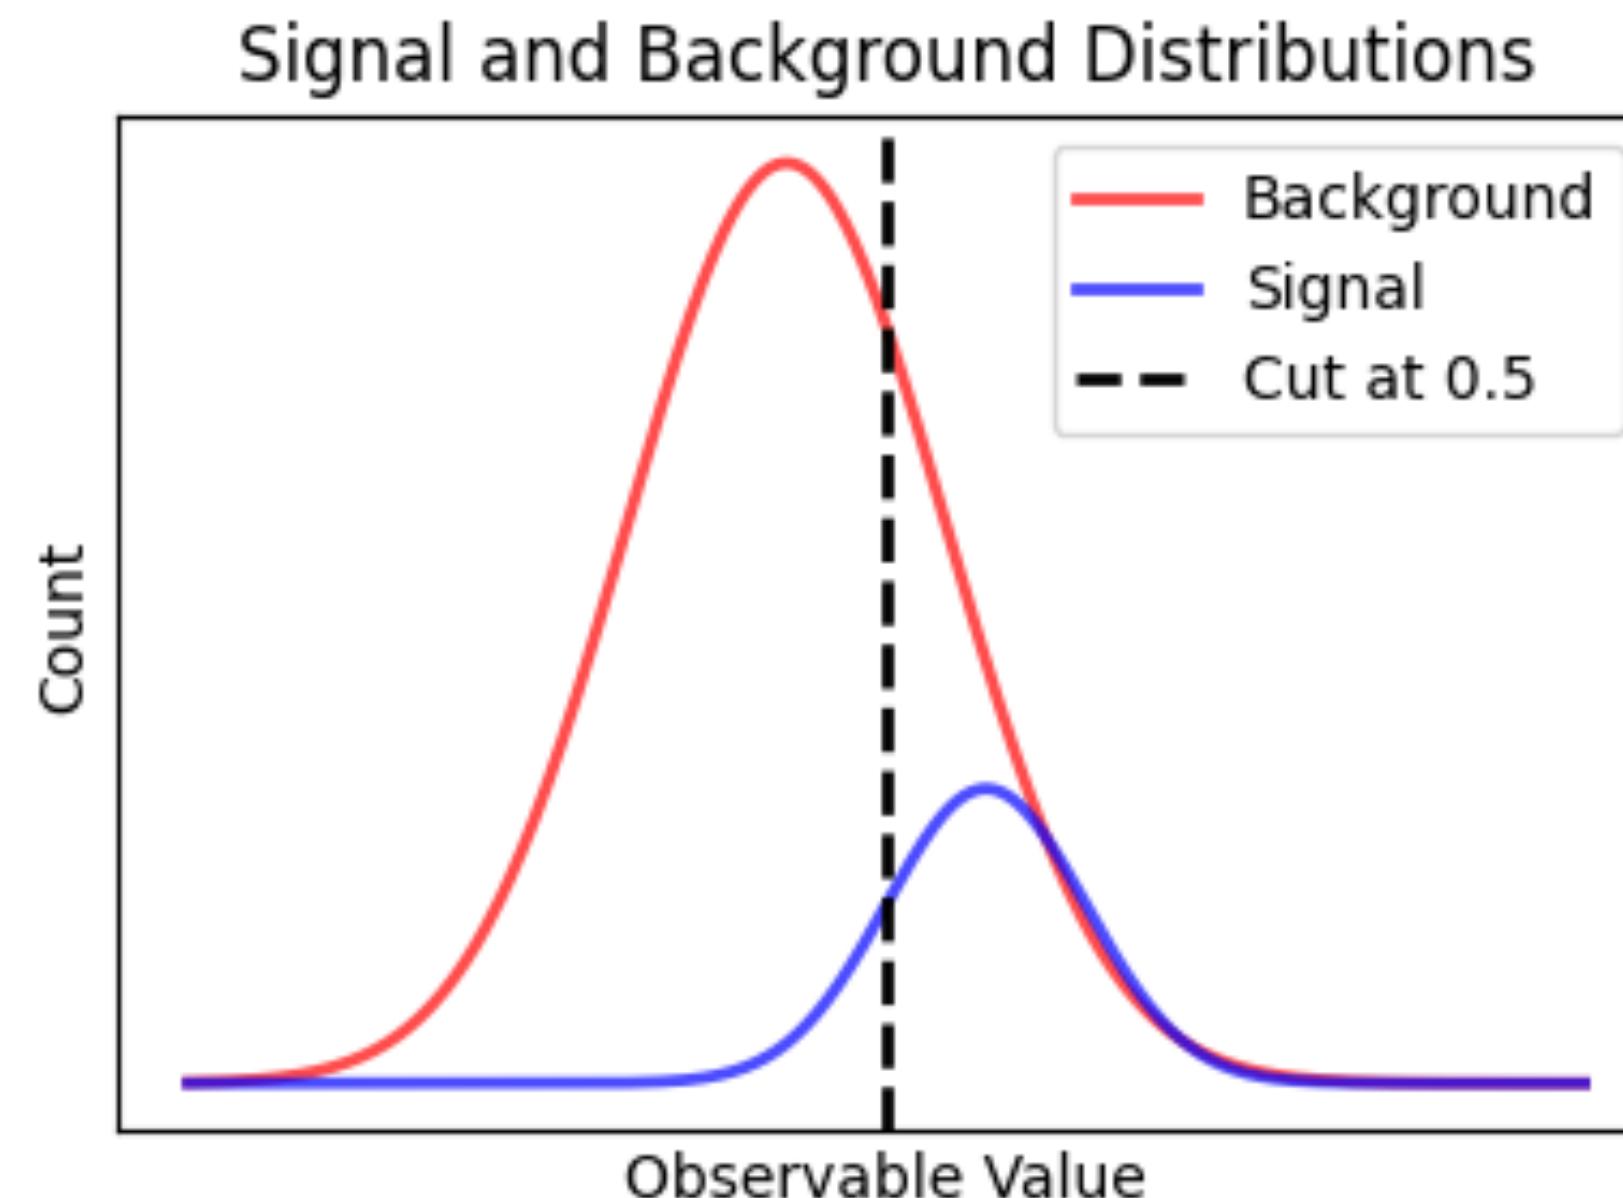
Is this enough?  
Not quite...

# Differentiable Relaxations

Making discrete primitives differentiable

Typical operations applied on HEP data are non-differentiable

e.g. applying a threshold, sorting

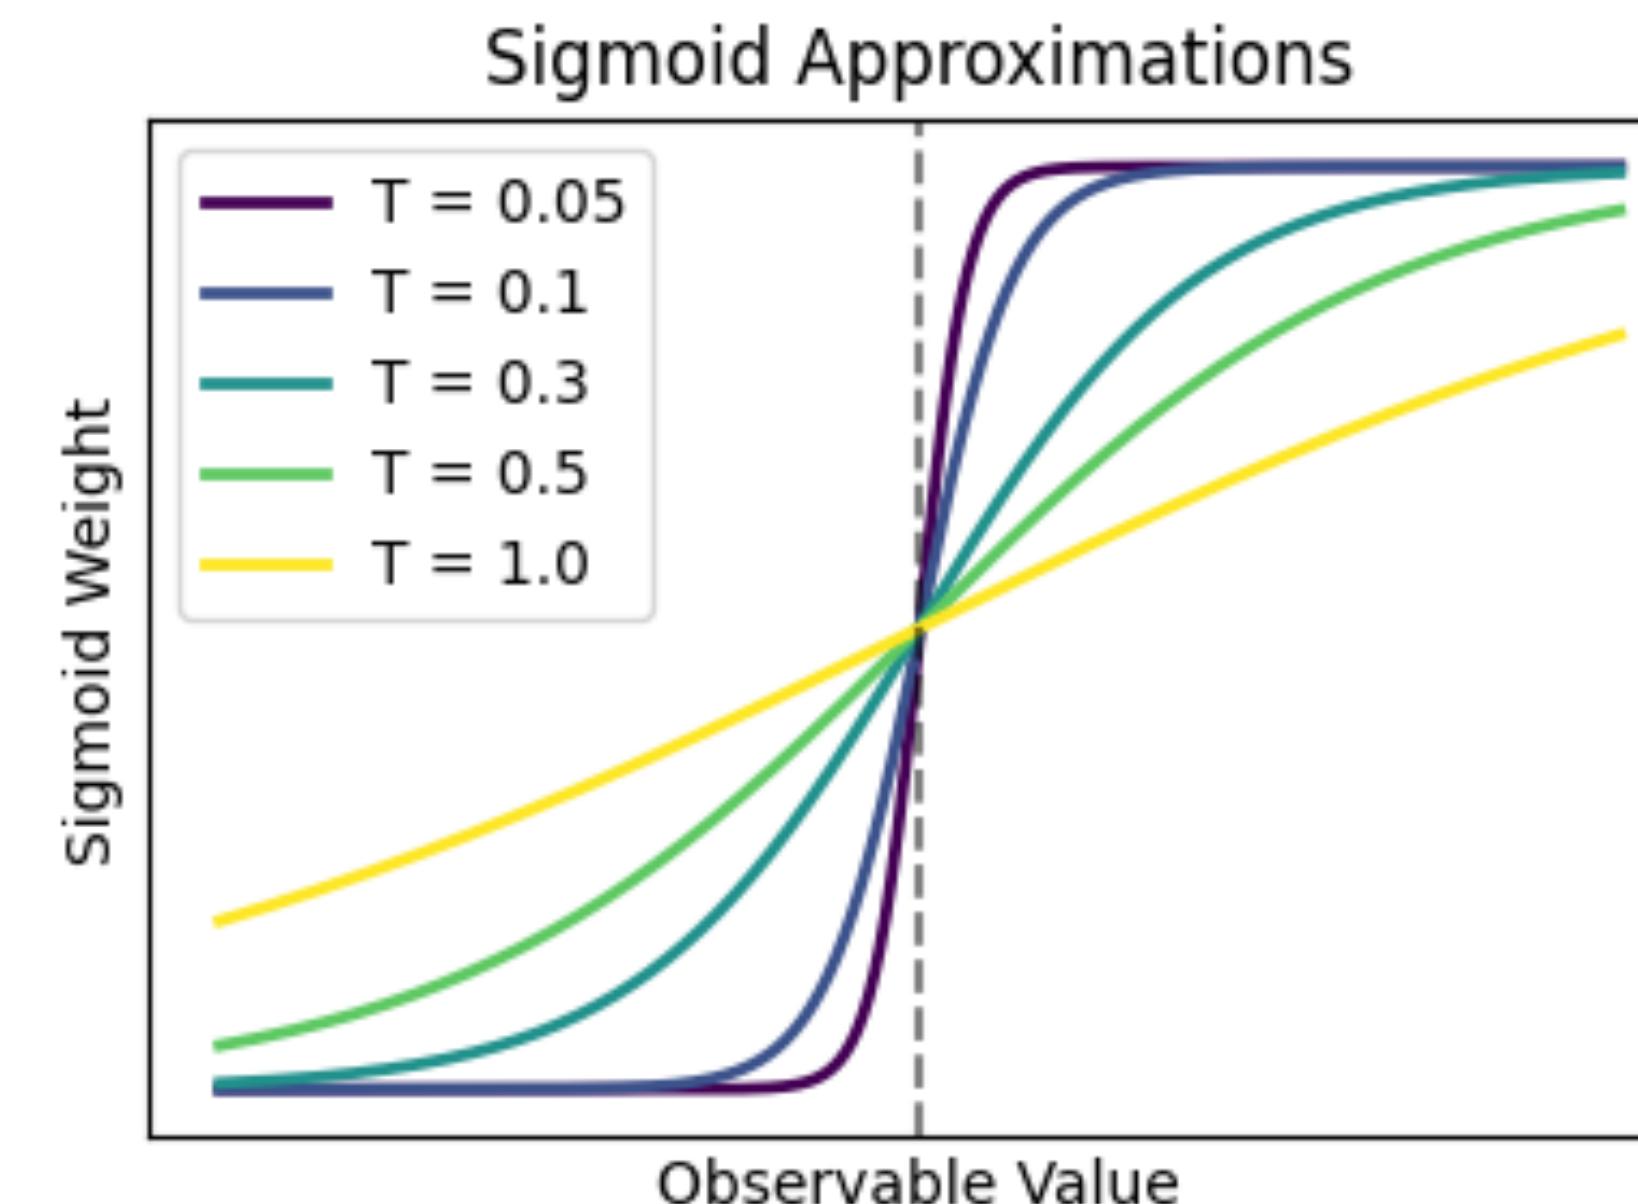


# Differentiable Relaxations

Making discrete primitives differentiable

We can approximate (smoothen) the operations with differentiable functions → differentiable relaxations

Soft sorting demo for a mass reconstruction problem

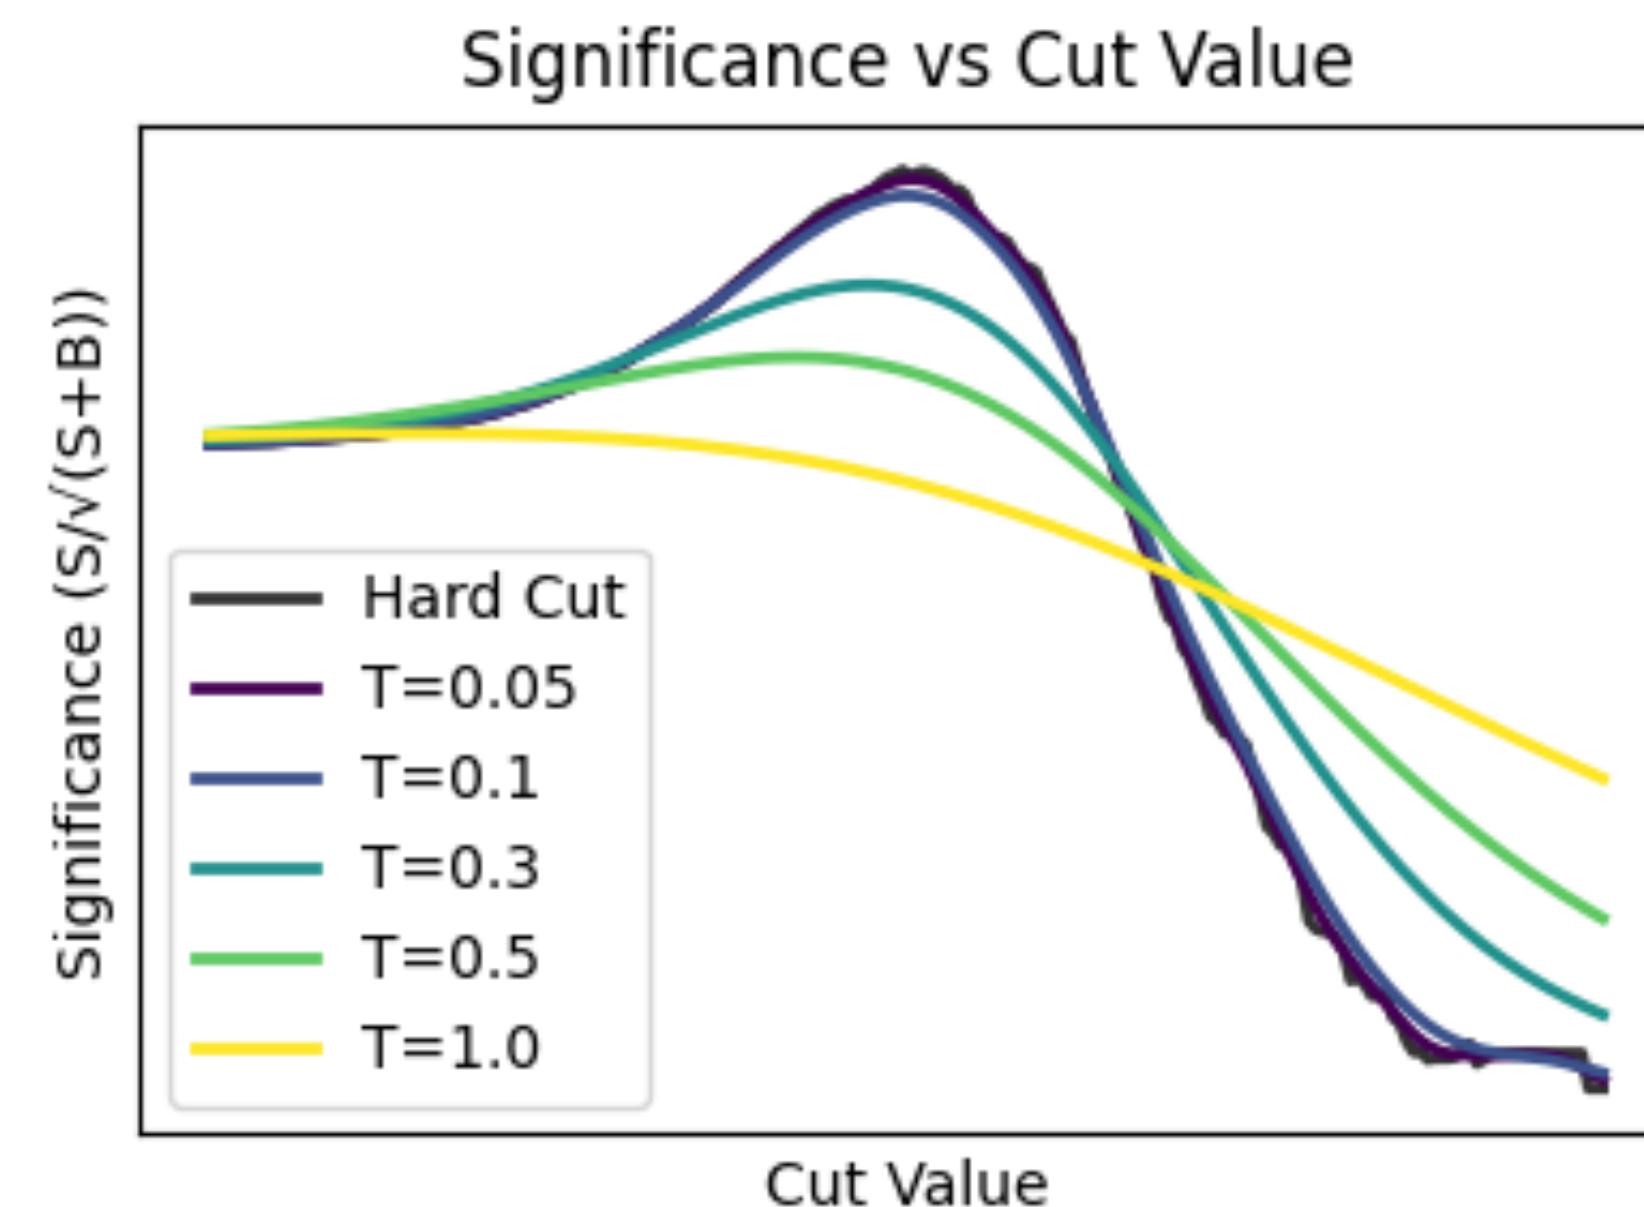


Smoothing  
controlled by the  
“Temperature” (T)

# Differentiable Relaxations

Making discrete primitives differentiable

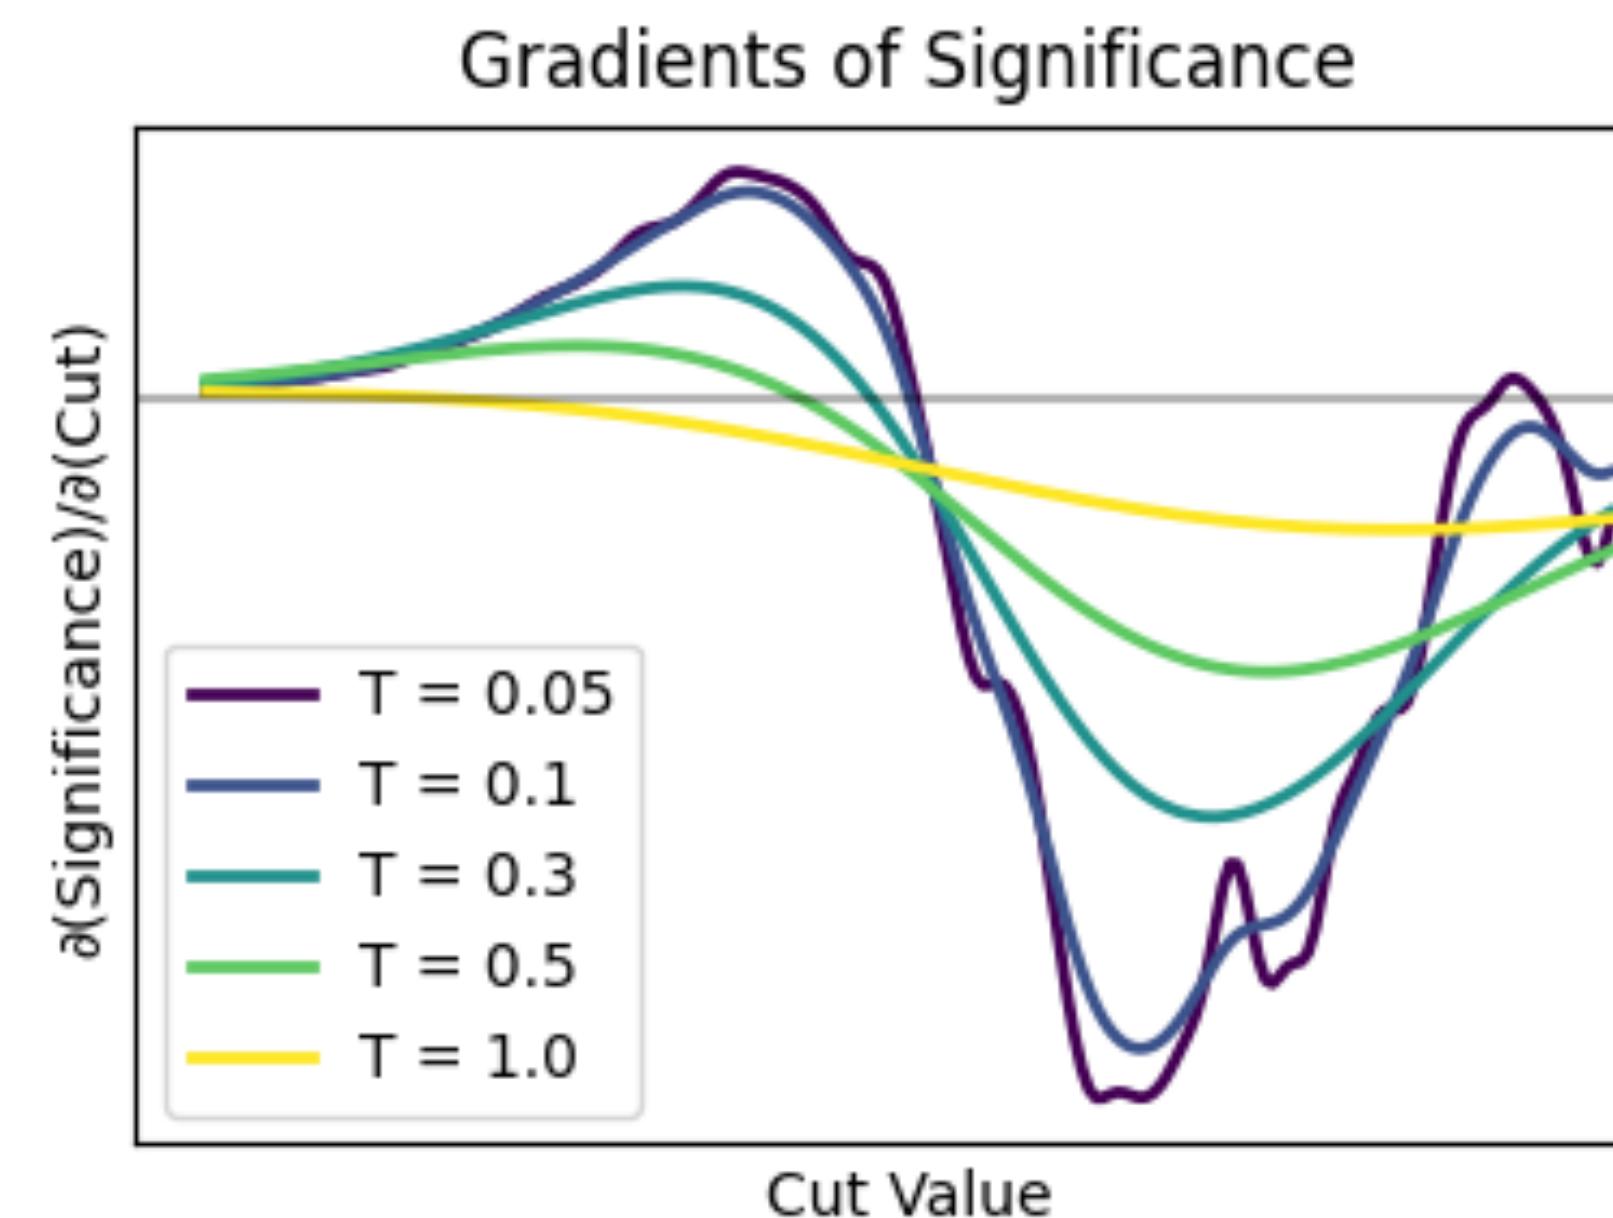
Optimal significance value is less robust against softening than its location (i.e. cut value)



# Differentiable Relaxations

Making discrete primitives differentiable

A trade-off: faithfulness to a hard threshold and stability of in gradient evaluation

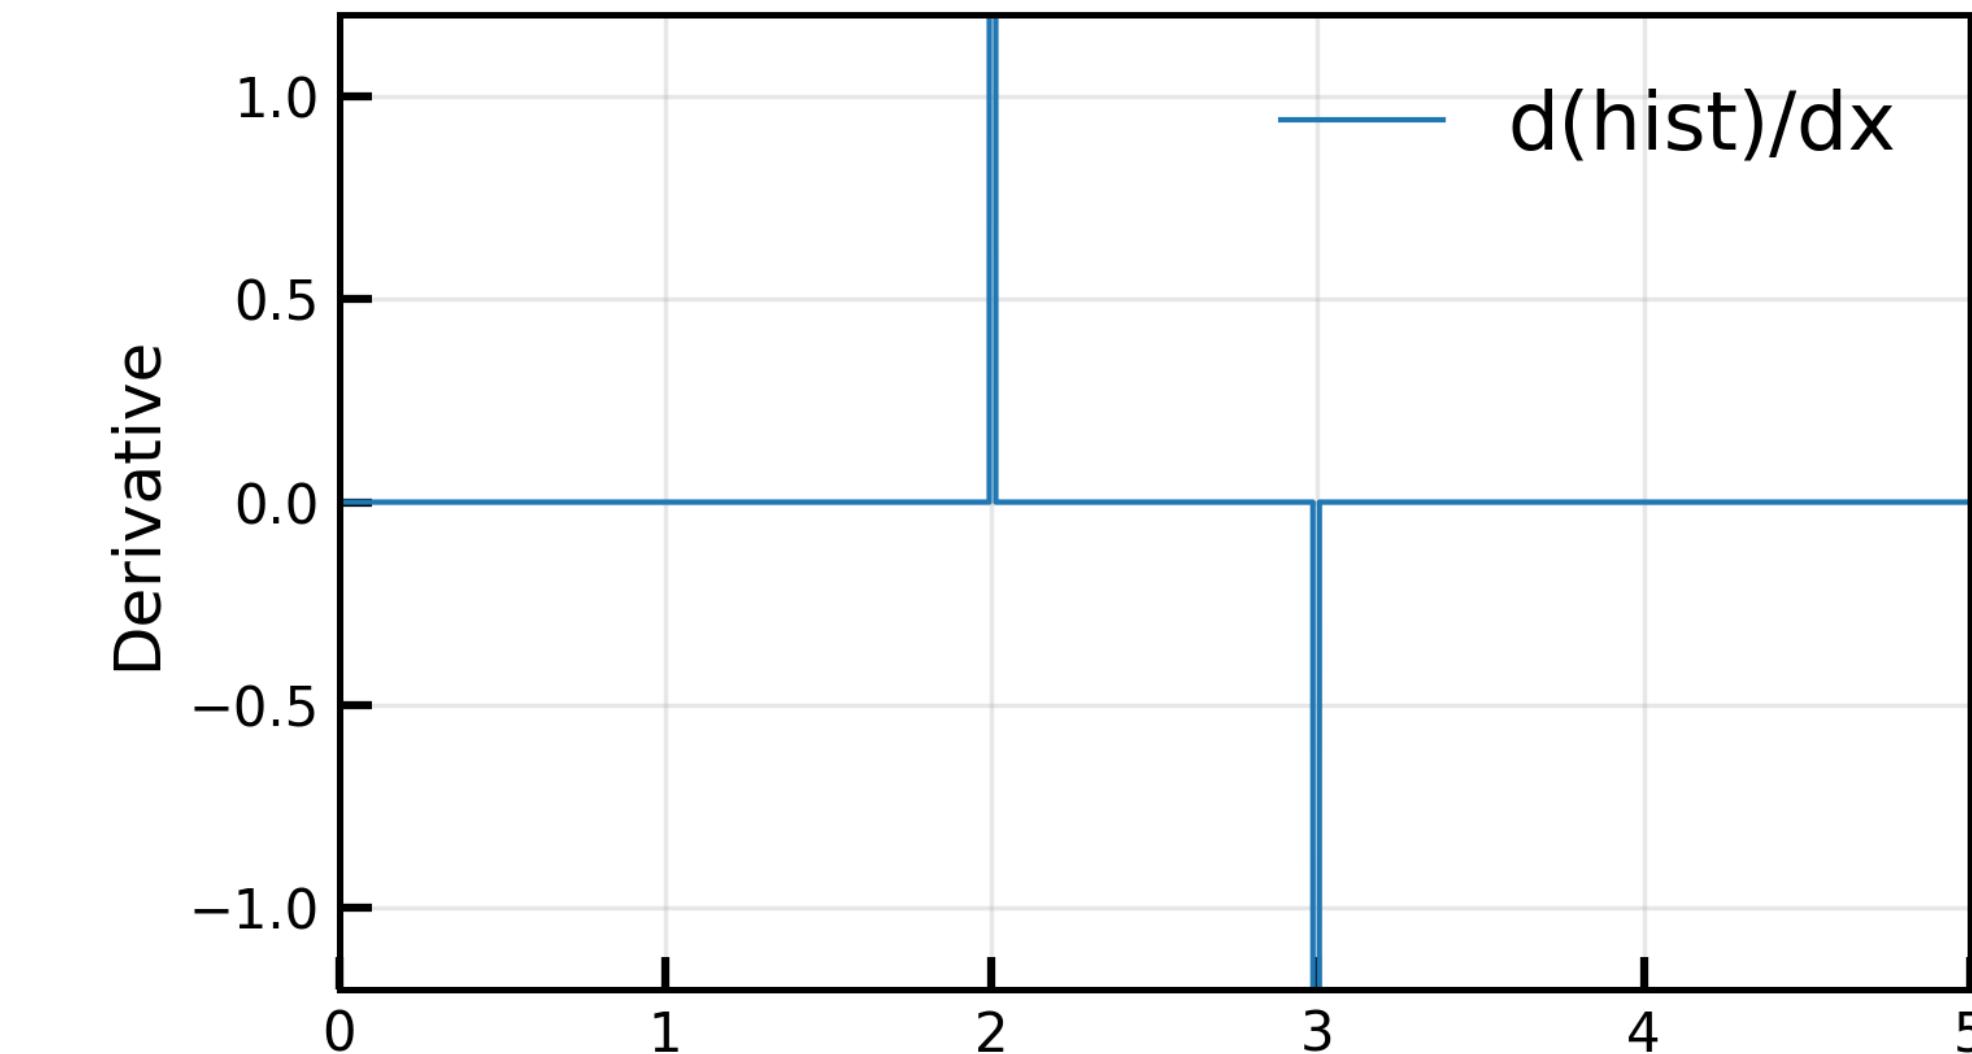
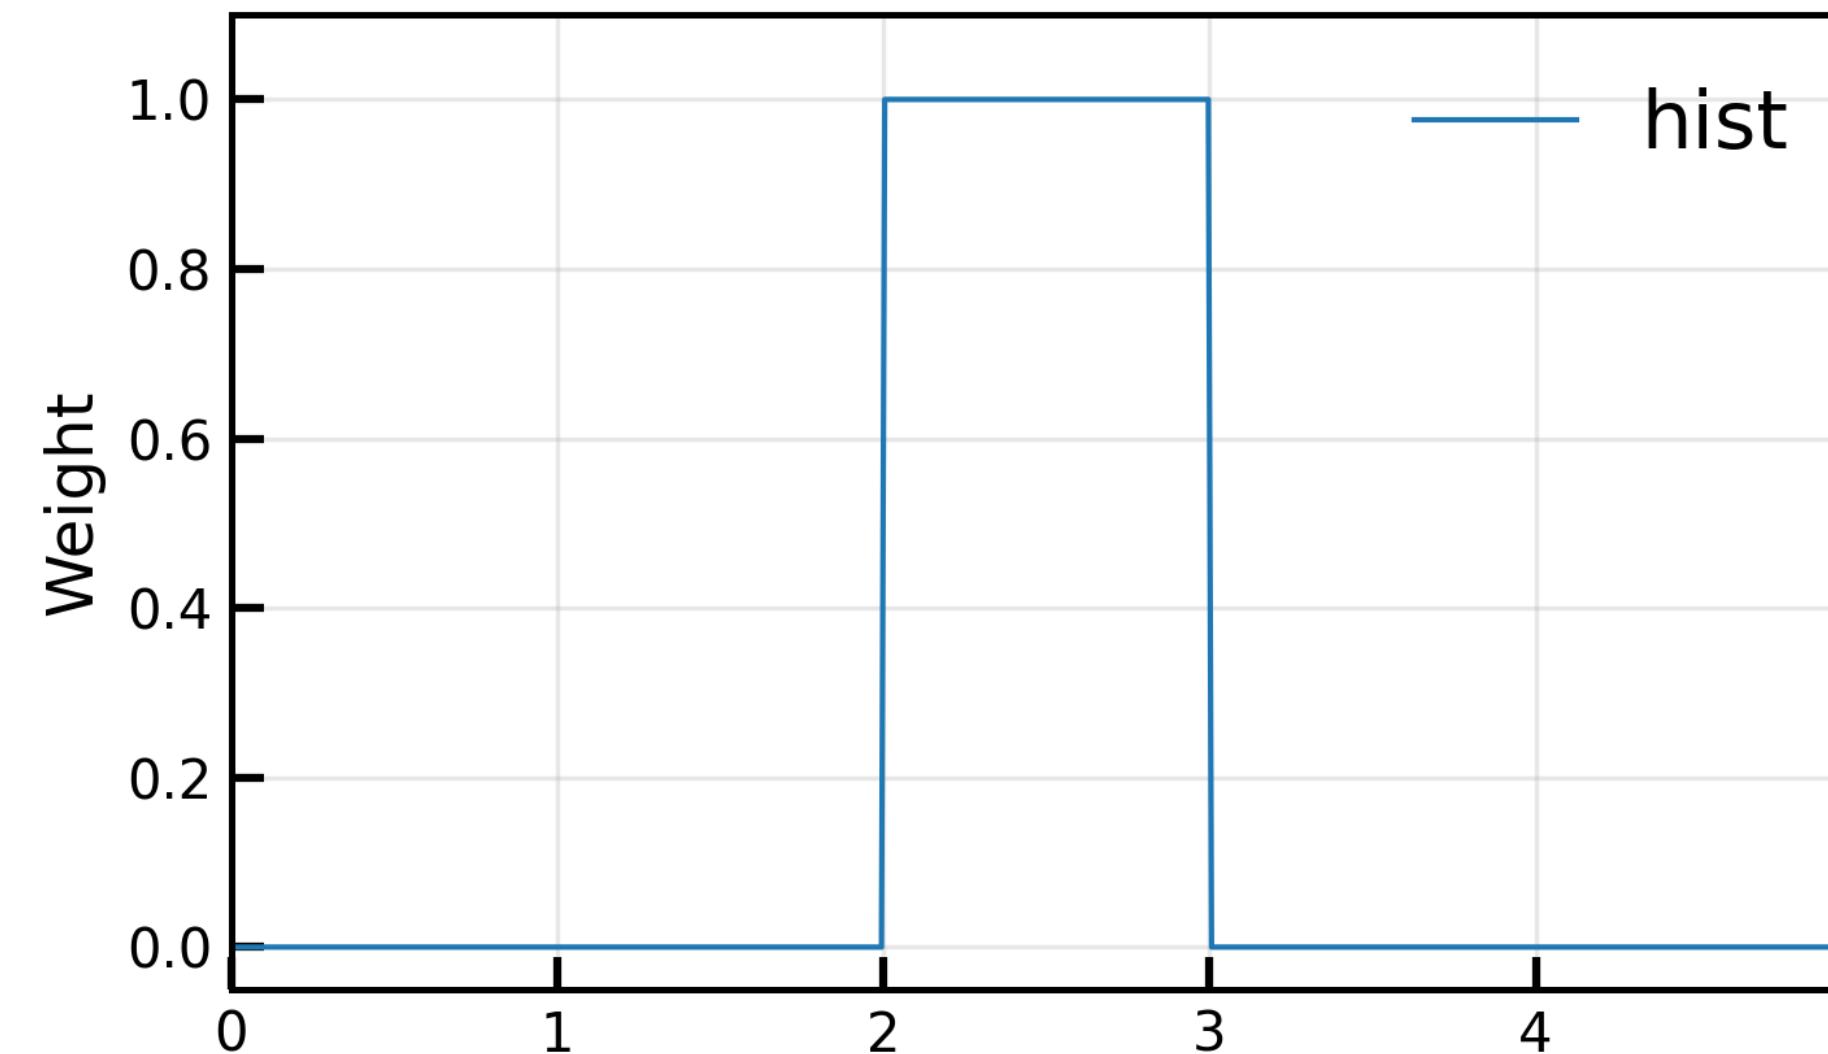


# Differentiable Relaxations

## Making histograms differentiable

Histograms are central (non-differentiable) data structures in HEP analyses

Vary data point: land in same bin, 0 gradient, land in another bin, infinite gradient

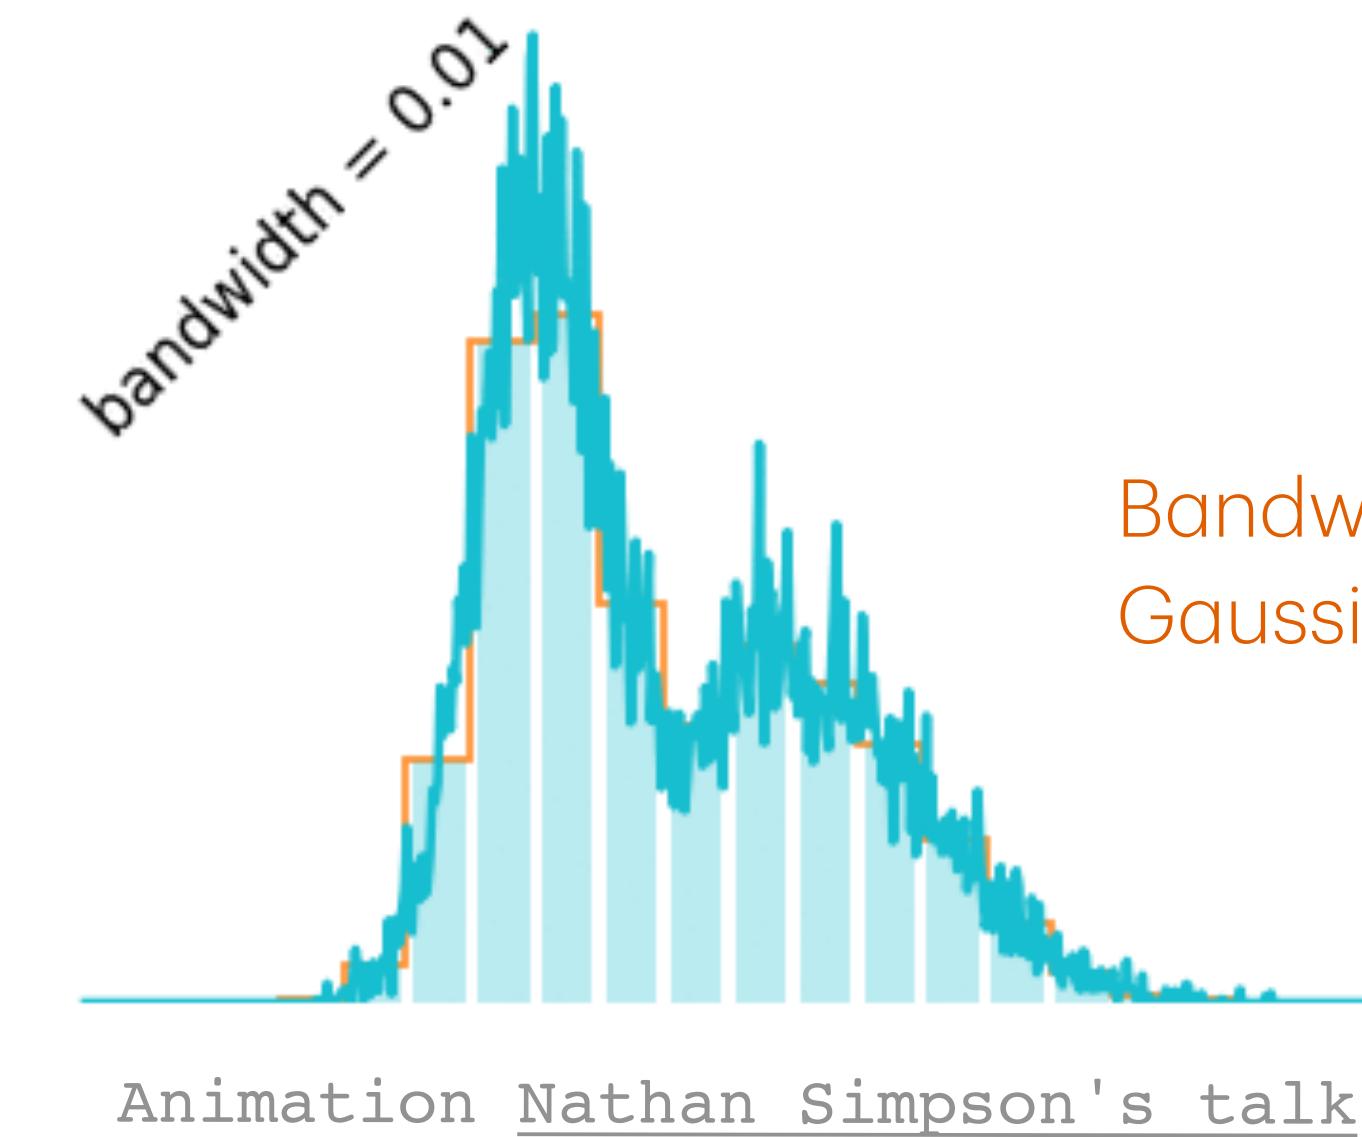
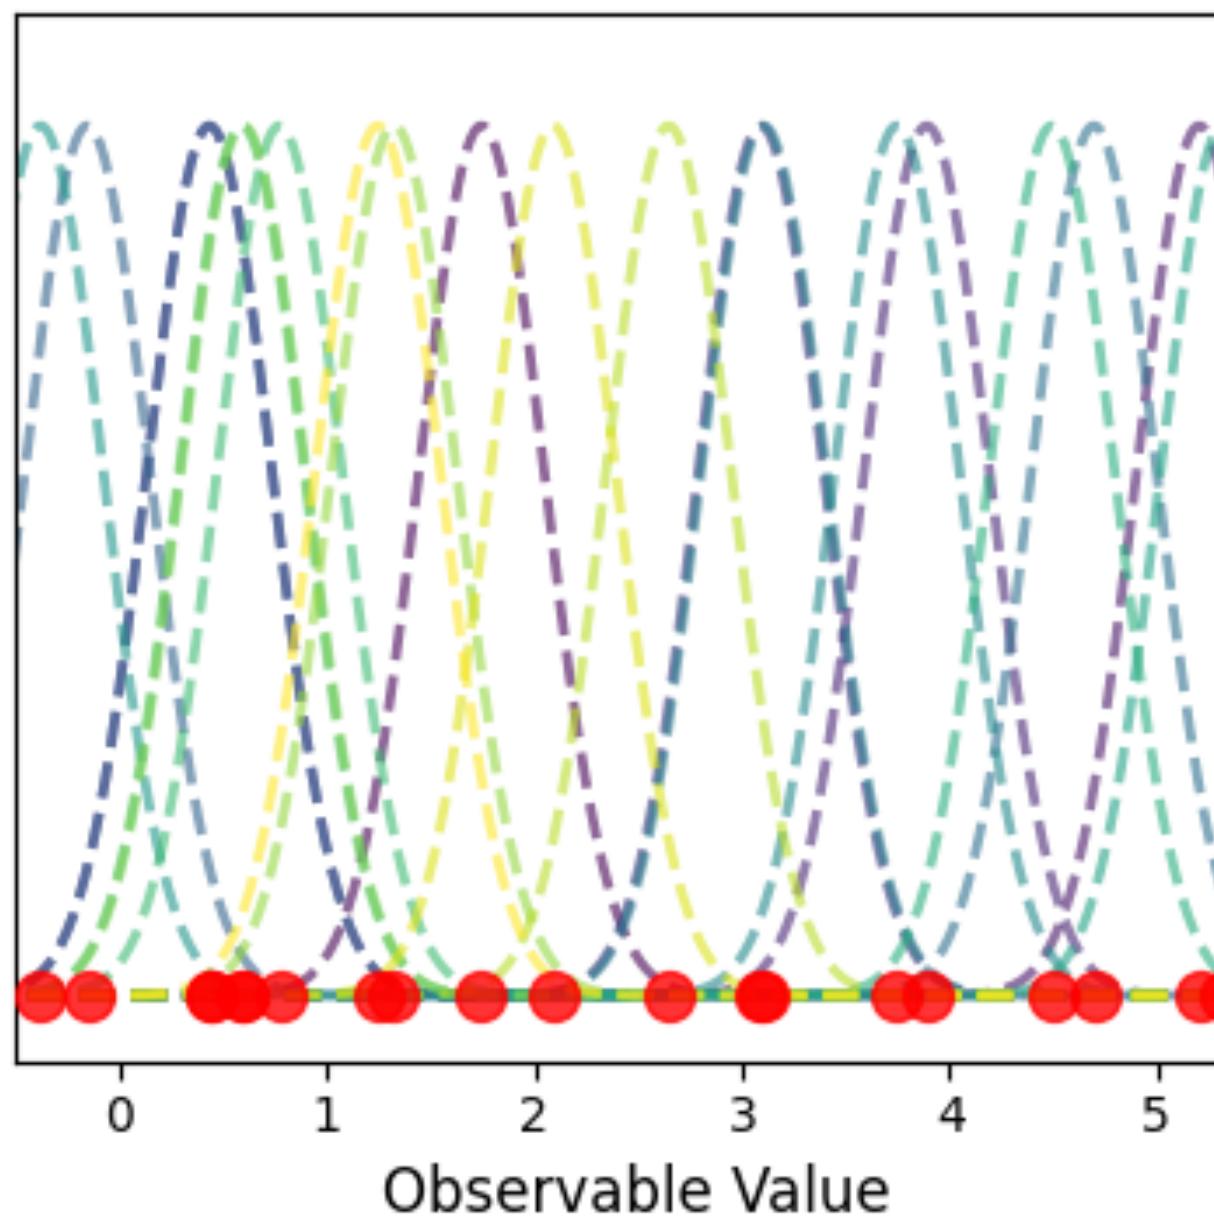


# Differentiable Relaxations

## Making histograms differentiable

Relaxation: binned kernel density estimate (bKDE)

Step 1: Define 1 Gaussian per-data point, and sum all Gaussians

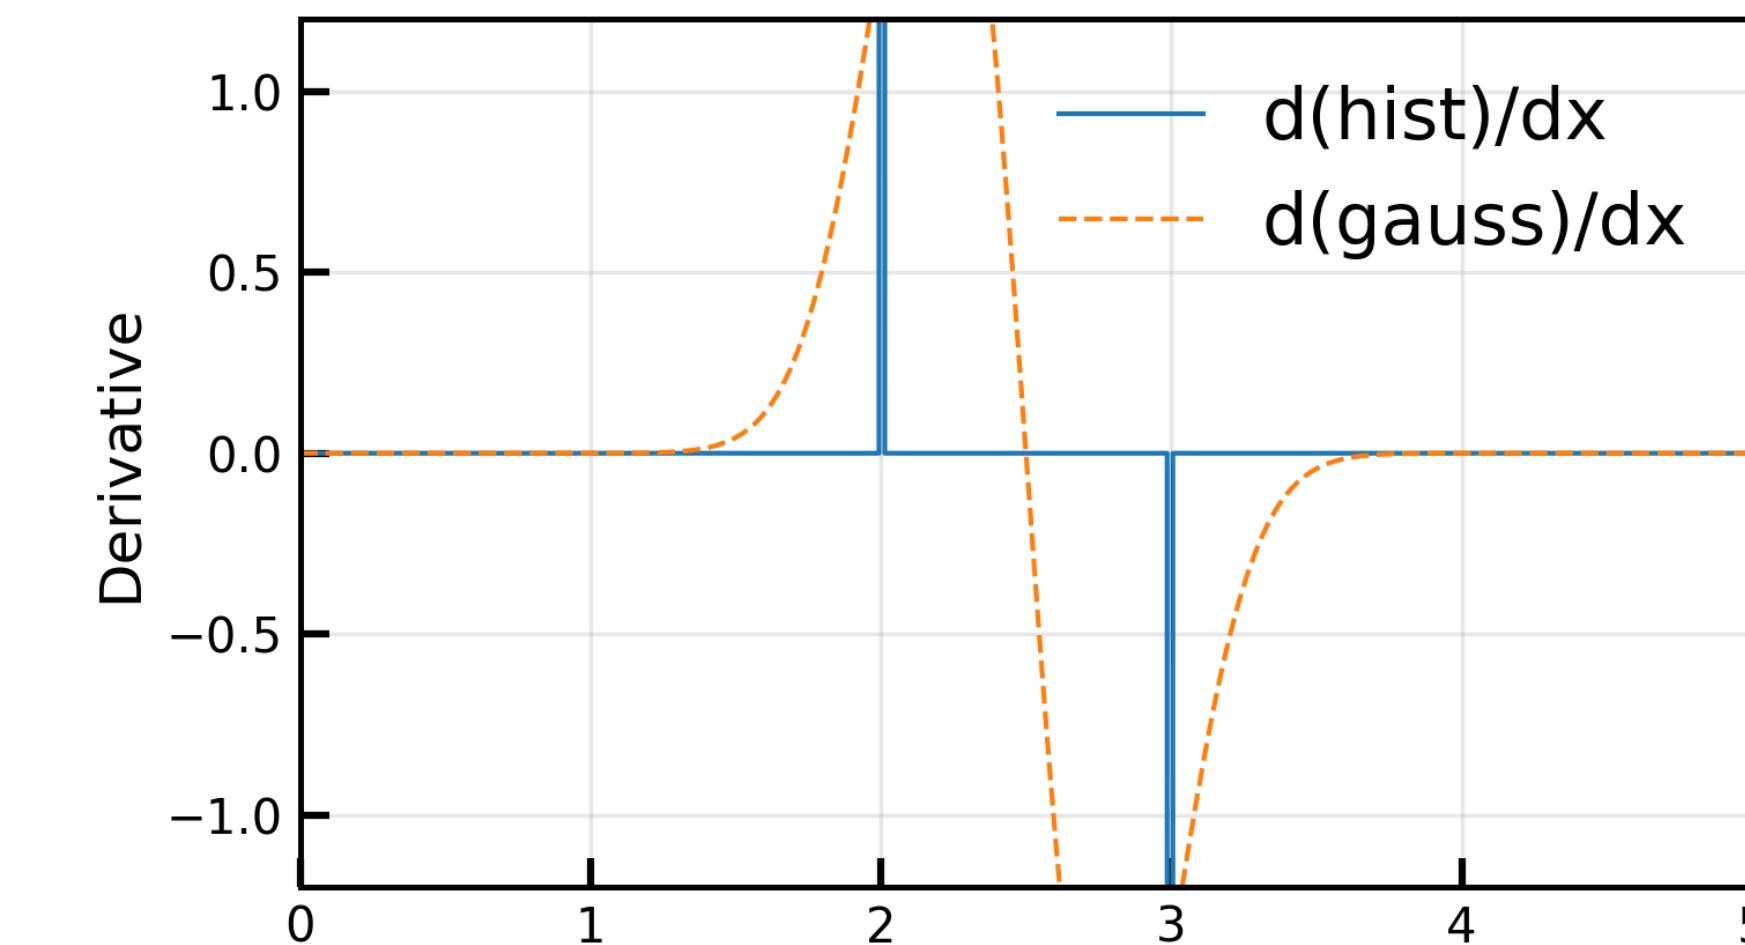
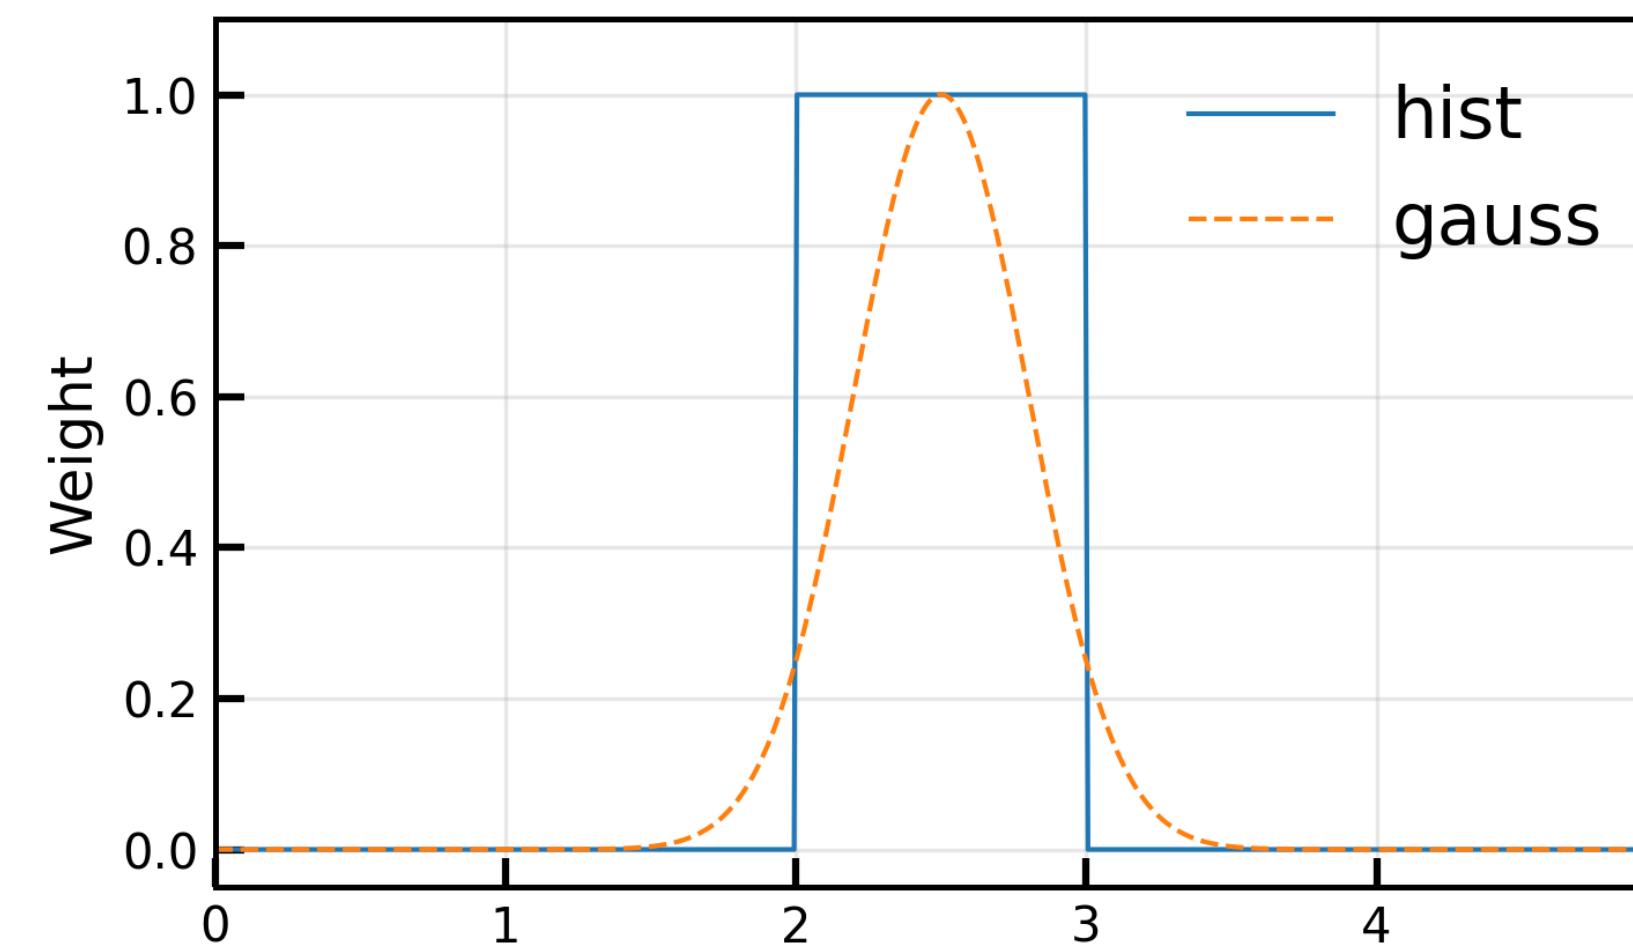


# Differentiable Relaxations

## Making histograms differentiable

Relaxation: binned kernel density estimate (bKDE)

Step 2: Integrate the Gaussian between bin-edges to get bin content



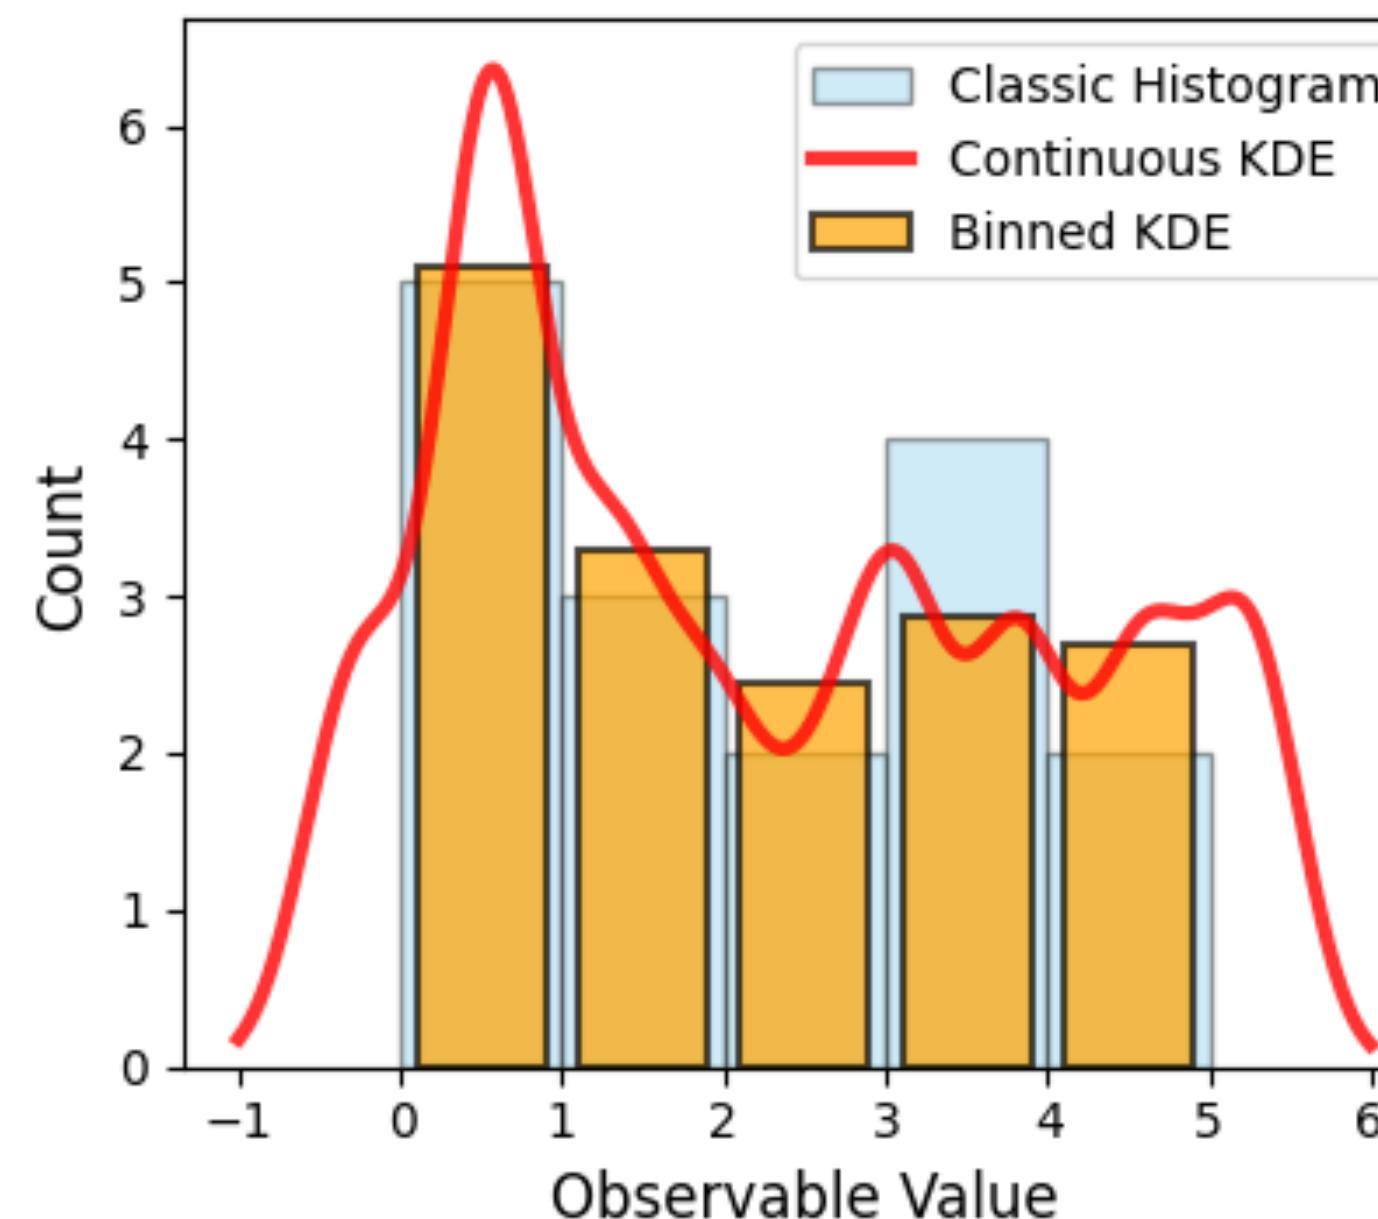
# Differentiable Relaxations

## Making histograms differentiable

Relaxation: binned kernel density estimate (bKDE)

[Great tutorial on this by Nathan Simpson](#)

Step 3: Differentiate each bin with respect to each data-point

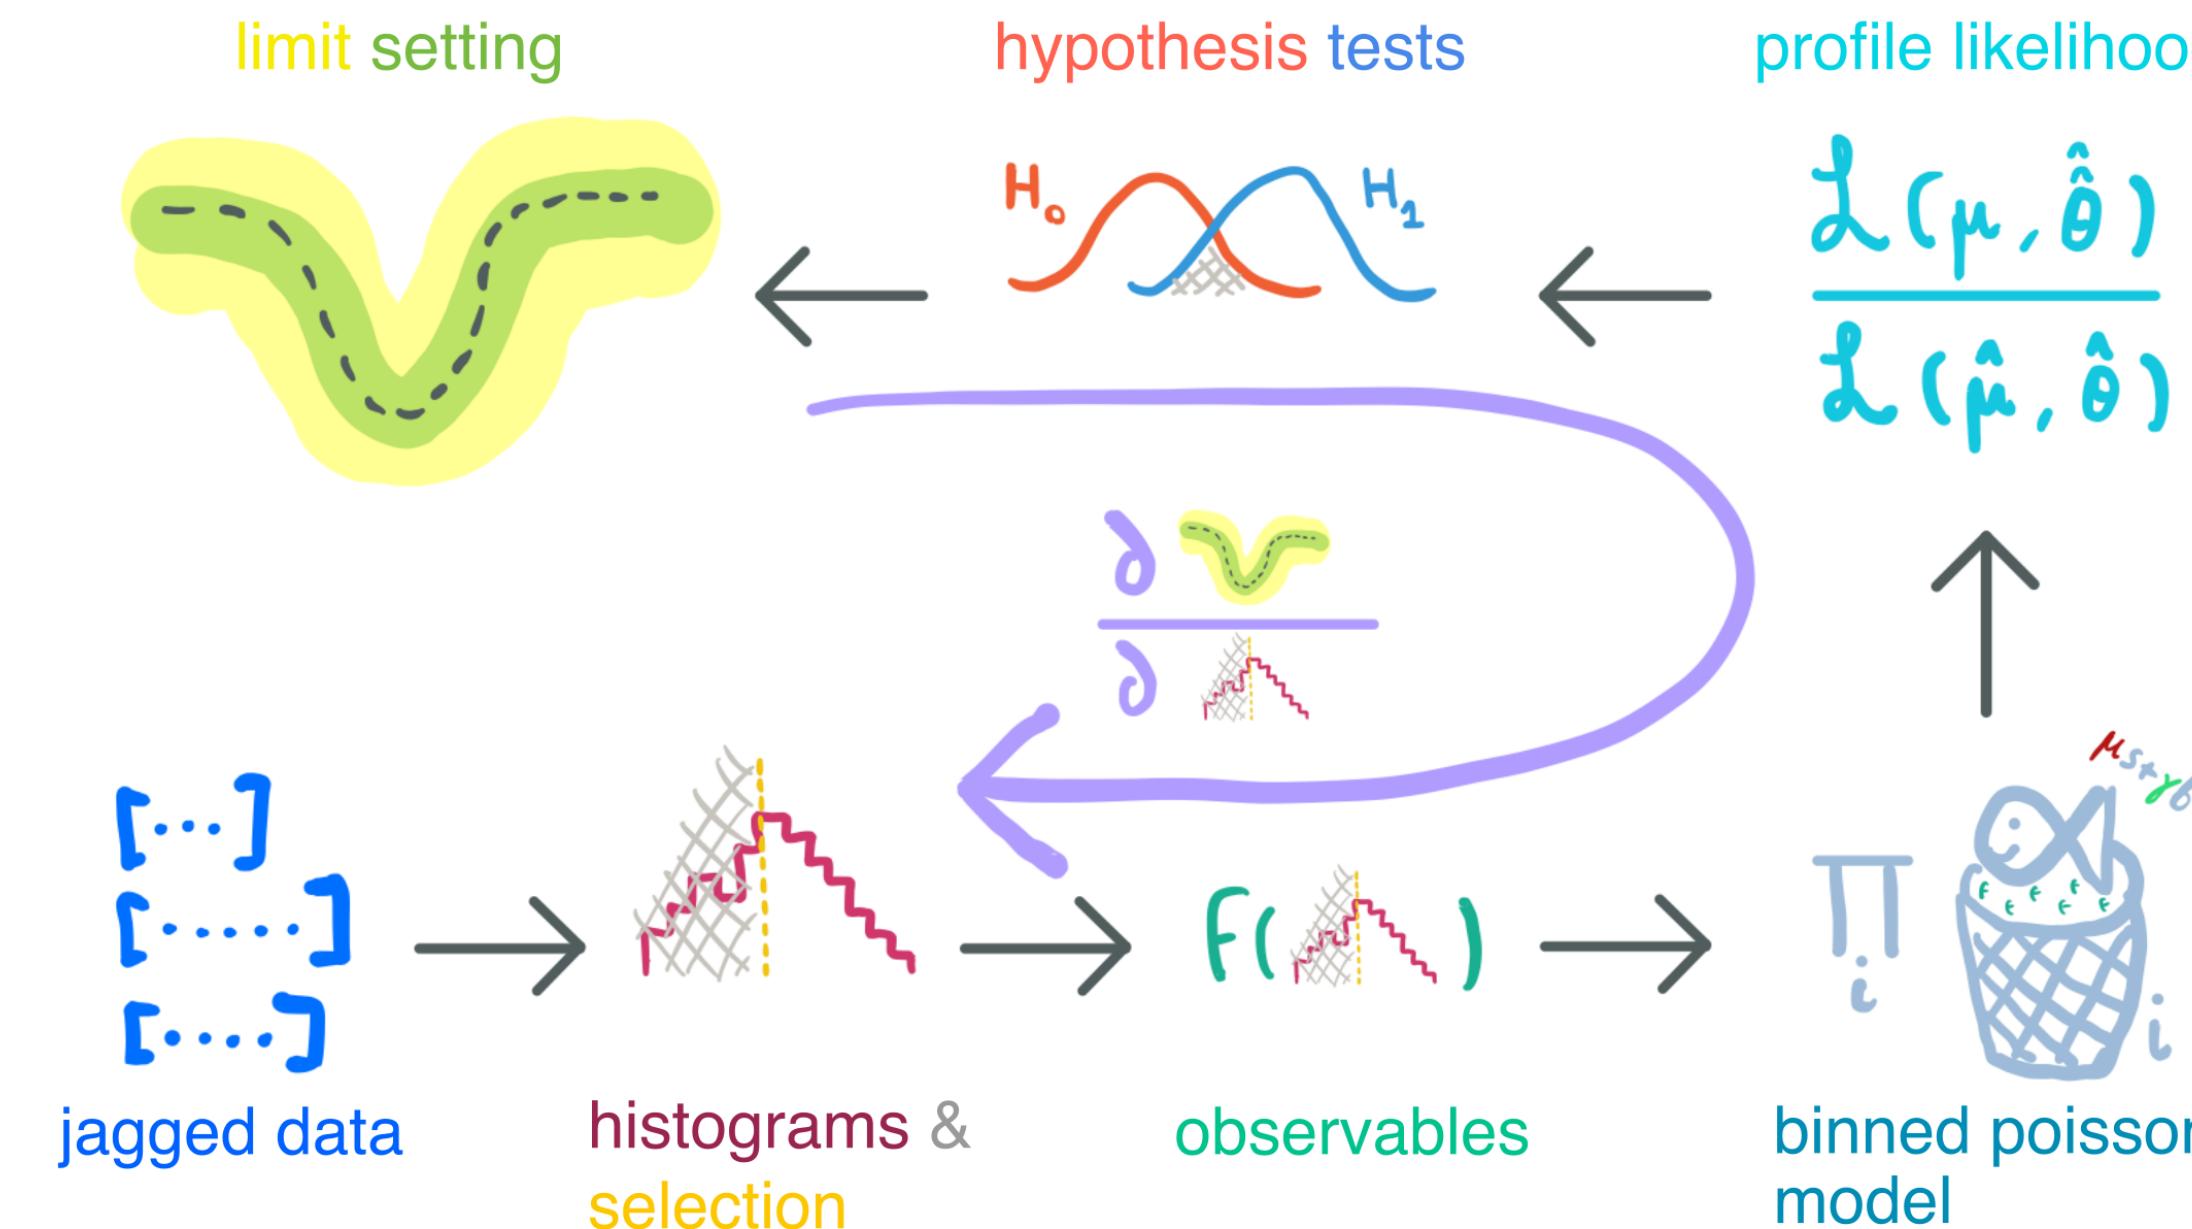


# Fixed-Point Differentiation

Optimising the analysis while fitting to data

From [Nathan Simpson's talk](#)

With the gradient of the significance with respect to all parameters →  
we can optimise all parameters to maximise significance



# Fixed-Point Differentiation

Optimising the analysis while fitting to data

Caveat: the significance requires performing fits

Significance is a function of:

$$\frac{\mathcal{L}(\mu, \hat{\theta})}{\mathcal{L}(\hat{\mu}, \hat{\theta})}$$

Run fit (nested minimiser) for each iteration in  
parameter optimisation?

Very expensive...

# Fixed-Point Differentiation

Optimising the analysis while fitting a model to data

Solution from Implicit function theorem: → fixed-point differentiation

$$F(x, y)$$



A well-behaved function with a  
ROOT  $F(x_0, y_0) = 0$

## Implicit function theorem:

There exists an implicit function  $\phi(x)$  such that  $F(x, \phi(x)) = 0$  for  $x \sim x_0$

with 
$$\frac{d\phi}{dx} = - \left( \frac{\partial F}{\partial y} \right)^{-1} \left( \frac{\partial F}{\partial x} \right)$$

# Fixed-Point Differentiation

Optimising the analysis while fitting a model to data

Significance involves likelihood minimisation to find the best-fit  
model parameters  $\theta^*(\alpha)$

$$\nabla_{\theta} \mathcal{L}(\alpha, \theta(\alpha)) = 0$$

Implicitly defines  $\theta(\alpha)$  and hence  $\frac{d\theta}{d\alpha}$

$$S(\alpha, \theta(\alpha)) \Rightarrow \frac{dS}{d\alpha} = \frac{\partial S}{\partial \alpha} + \frac{\partial S}{\partial \theta} \cdot \frac{d\theta}{d\alpha}$$

Plug in  
More math in the back-up

Can compute gradient of the objective under the condition that  
we are at the likelihood minimum! Optimization in single loop!

# Fixed-Point Differentiation

Optimising the analysis while fitting to data

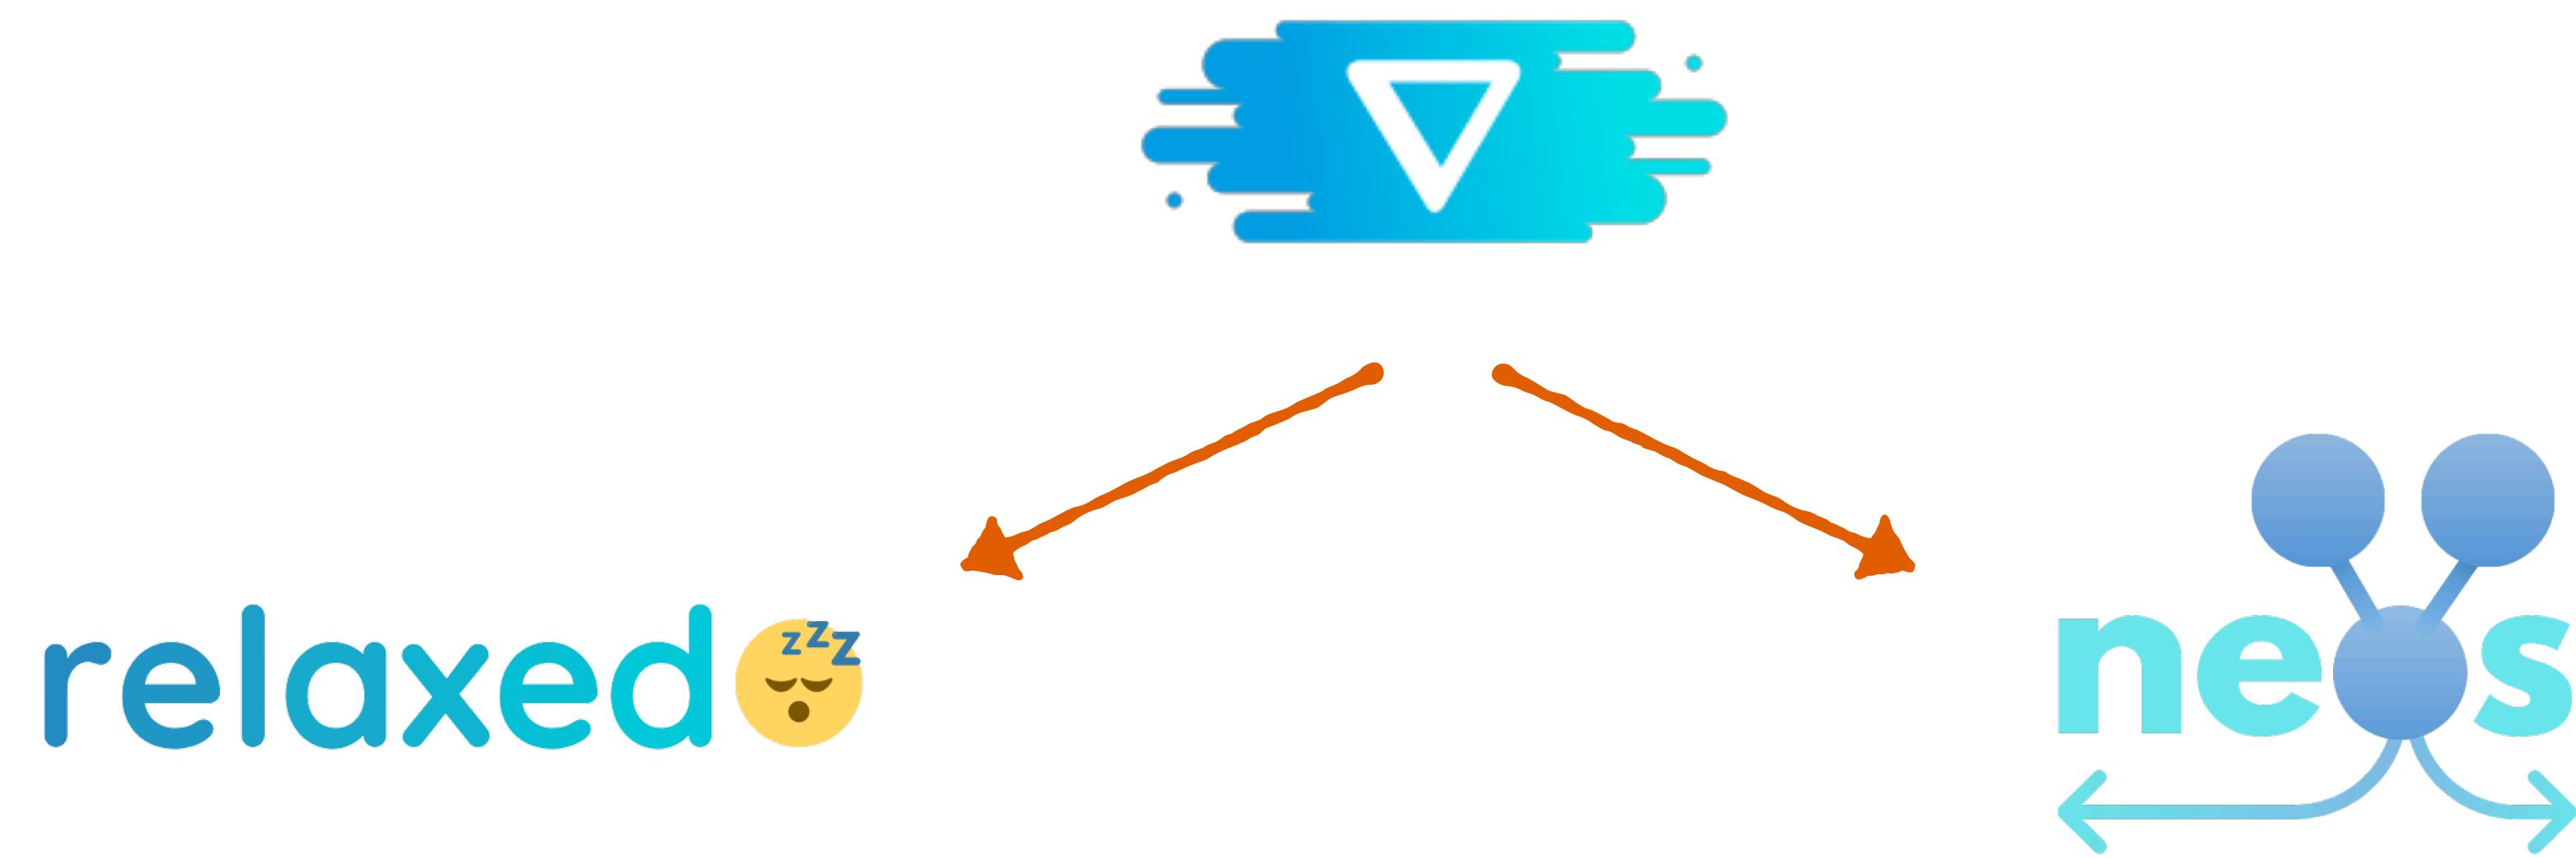
From [Nathan Simpson's talk](#)

The gradient of the objective is computed under the condition  
that we are at the likelihood minimum!

# Where we stand

State of end-to-end differentiable HEP analysis

gradhep: collection of HEP differentiation tools



Differentiable  
relaxations

End-to-end differentiable  
analysis cookbook with  
relaxed and PyHF

# Where we stand

## State of end-to-end differentiable HEP analysis

Proof-of-concept for a differentiable analysis with simulated toy  
data → [arXiv:2203.05570](https://arxiv.org/abs/2203.05570)

### **neos: End-to-End-Optimised Summary Statistics for High Energy Physics**

**Nathan Simpson<sup>1</sup> and Lukas Heinrich<sup>2</sup>**

<sup>1</sup>Lund University

<sup>2</sup>Technical University of Munich

E-mail: [n.s@cern.ch](mailto:n.s@cern.ch), [lukas.heinrich@cern.ch](mailto:lukas.heinrich@cern.ch)

“[...] a long term goal would be to scale neos in application to open data from the major LHC experiments, [...], allowing the probing of how the performance benefits scale to realistic problems with many sources of uncertainty.”

# Purpose of this work

What do we want to do

---

Goal #1

Scale concept of differentiable analysis to more realistic  
HEP data from major LHC experiment (CMS)

Goal #2

Review state + compatibility of existing tools and identify  
challenges to be addressed towards complete analysis

# Applying gradients with CMS open-data

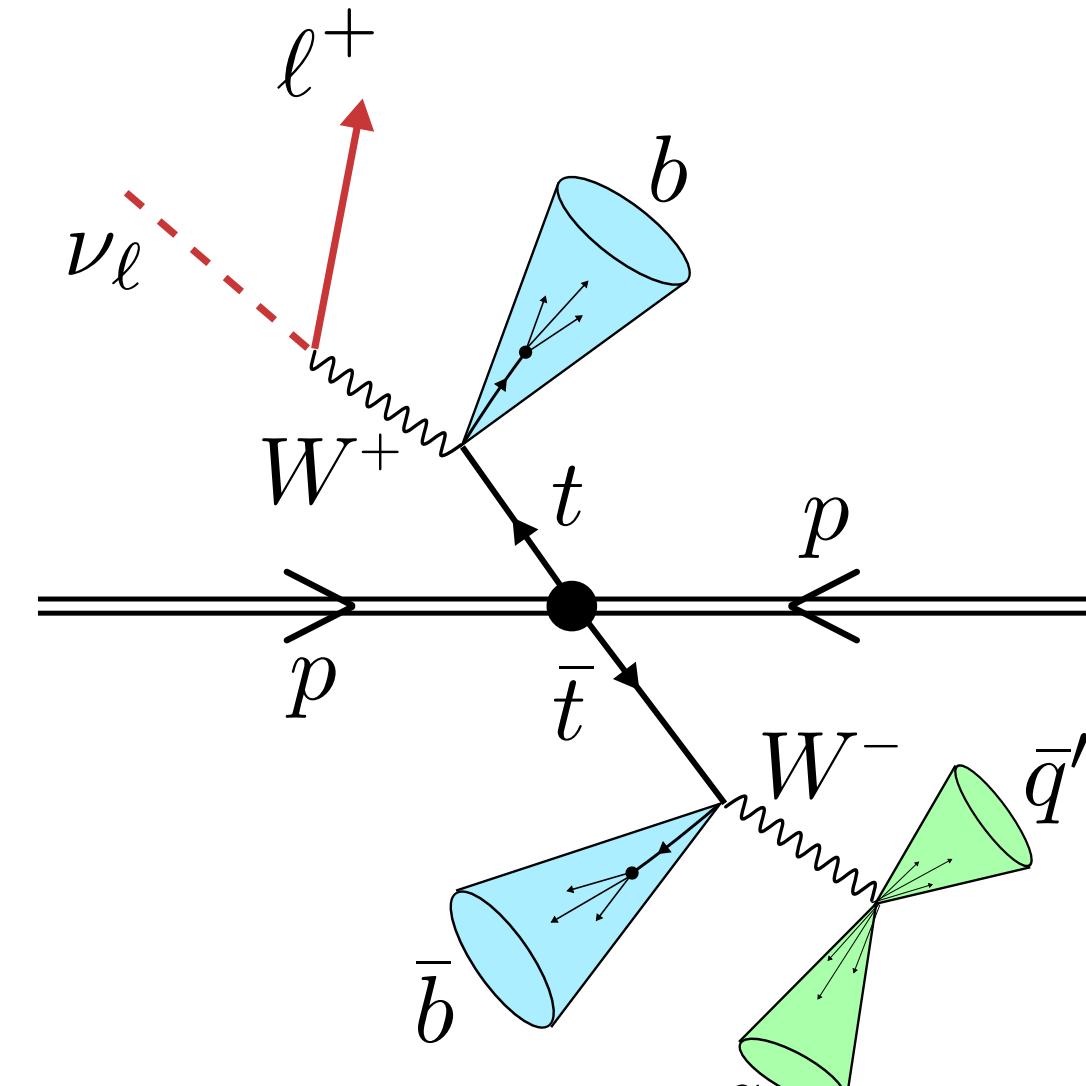


# Case study: a $Z' \rightarrow t\bar{t}$ analysis

Differentiating through CMS open data

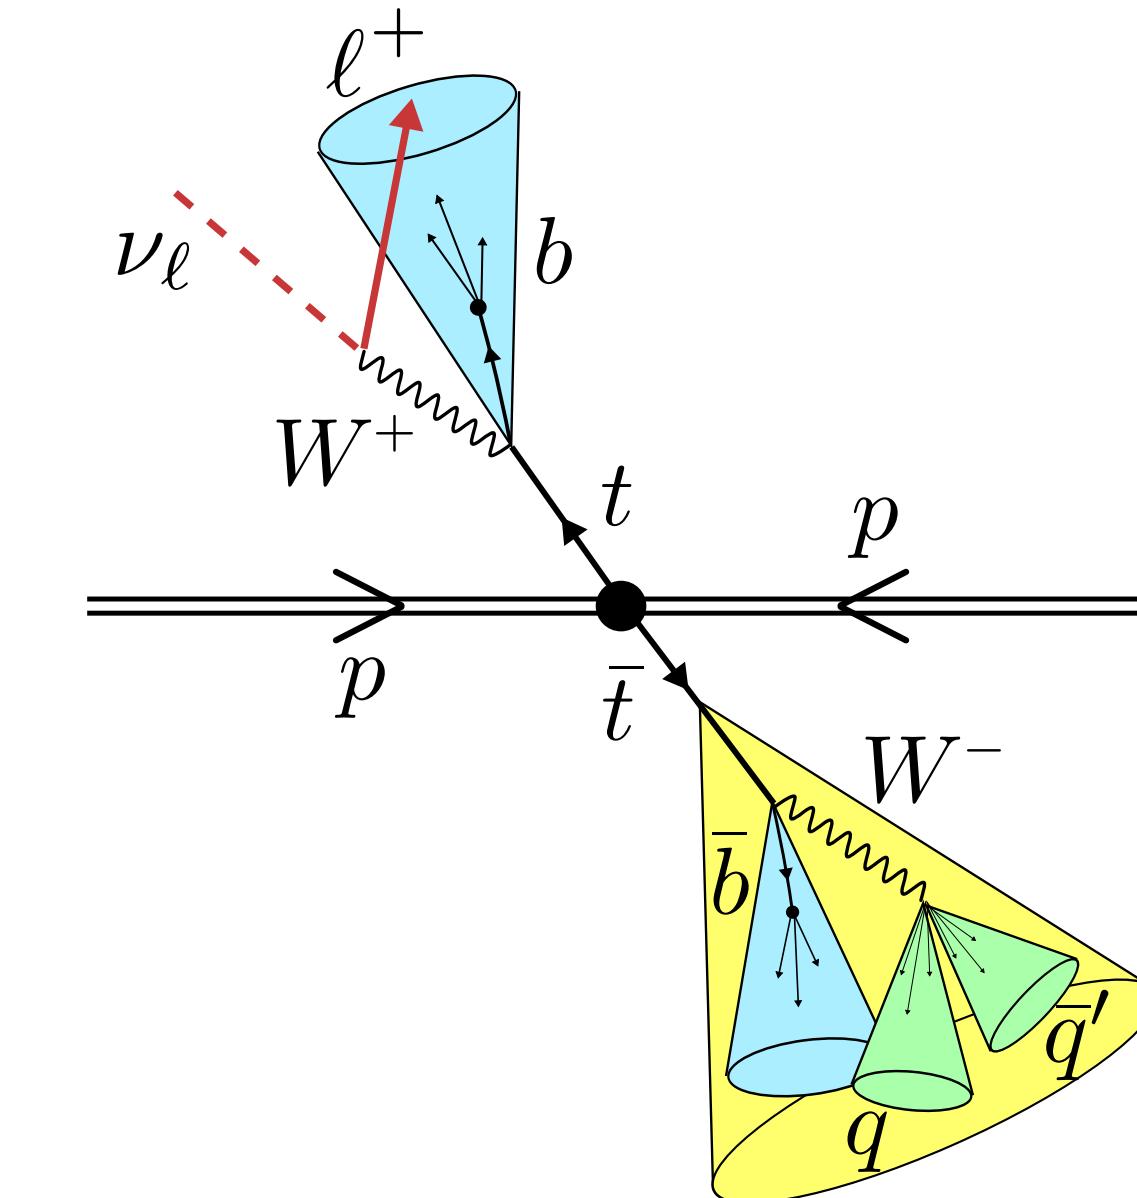
Diagram from [CMS open-data workshop](#)

- Searching for a hypothetical (heavy) BSM particle called  $Z'$  [[arXiv:1810.05905](#)]
- Produced in proton-proton collisions, decays into a pair of top quarks



(resolved)

Increasing  $Z'$  mass  
Collimated decay  
products



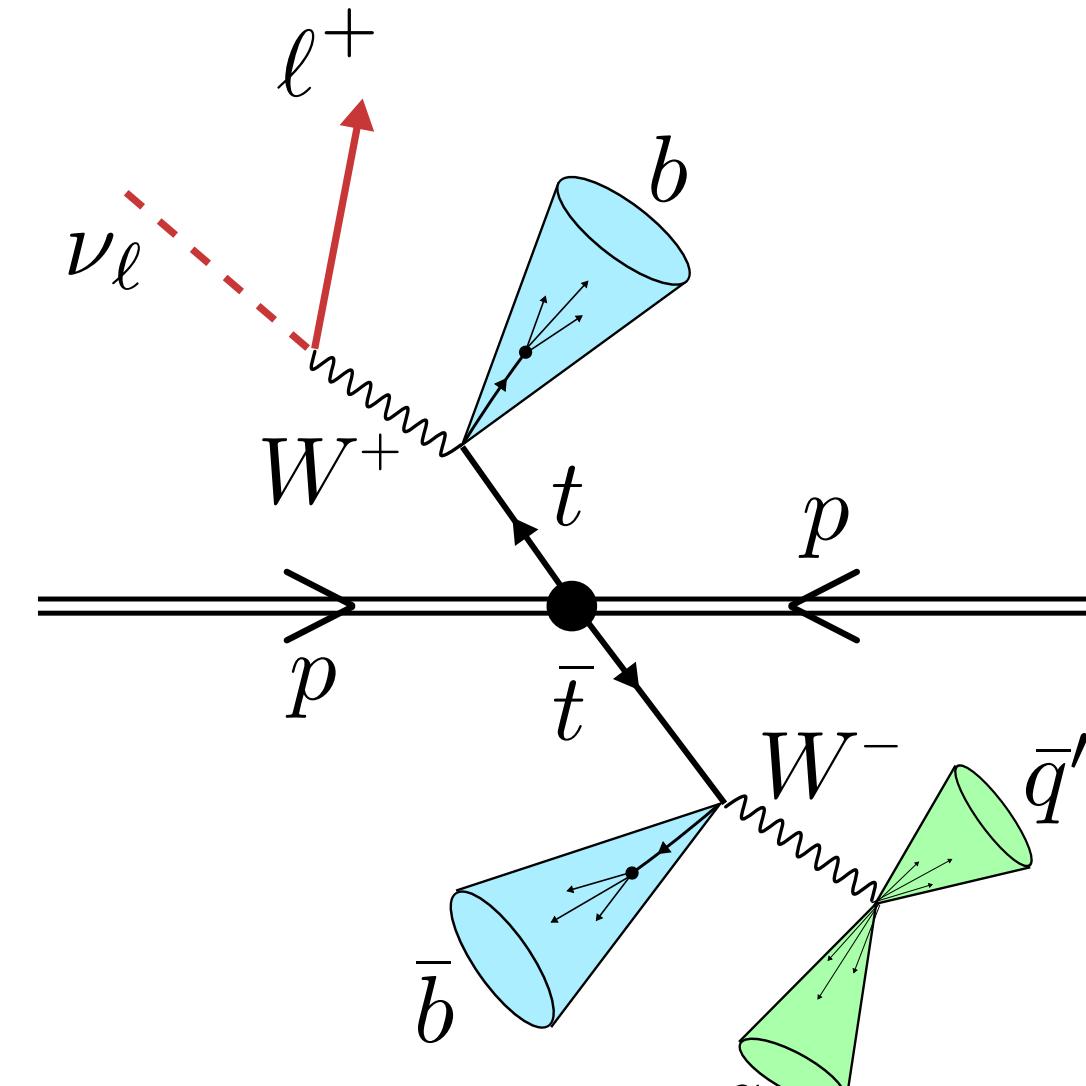
(boosted)

# Case study: a $Z' \rightarrow t\bar{t}$ analysis

Differentiating through CMS open data

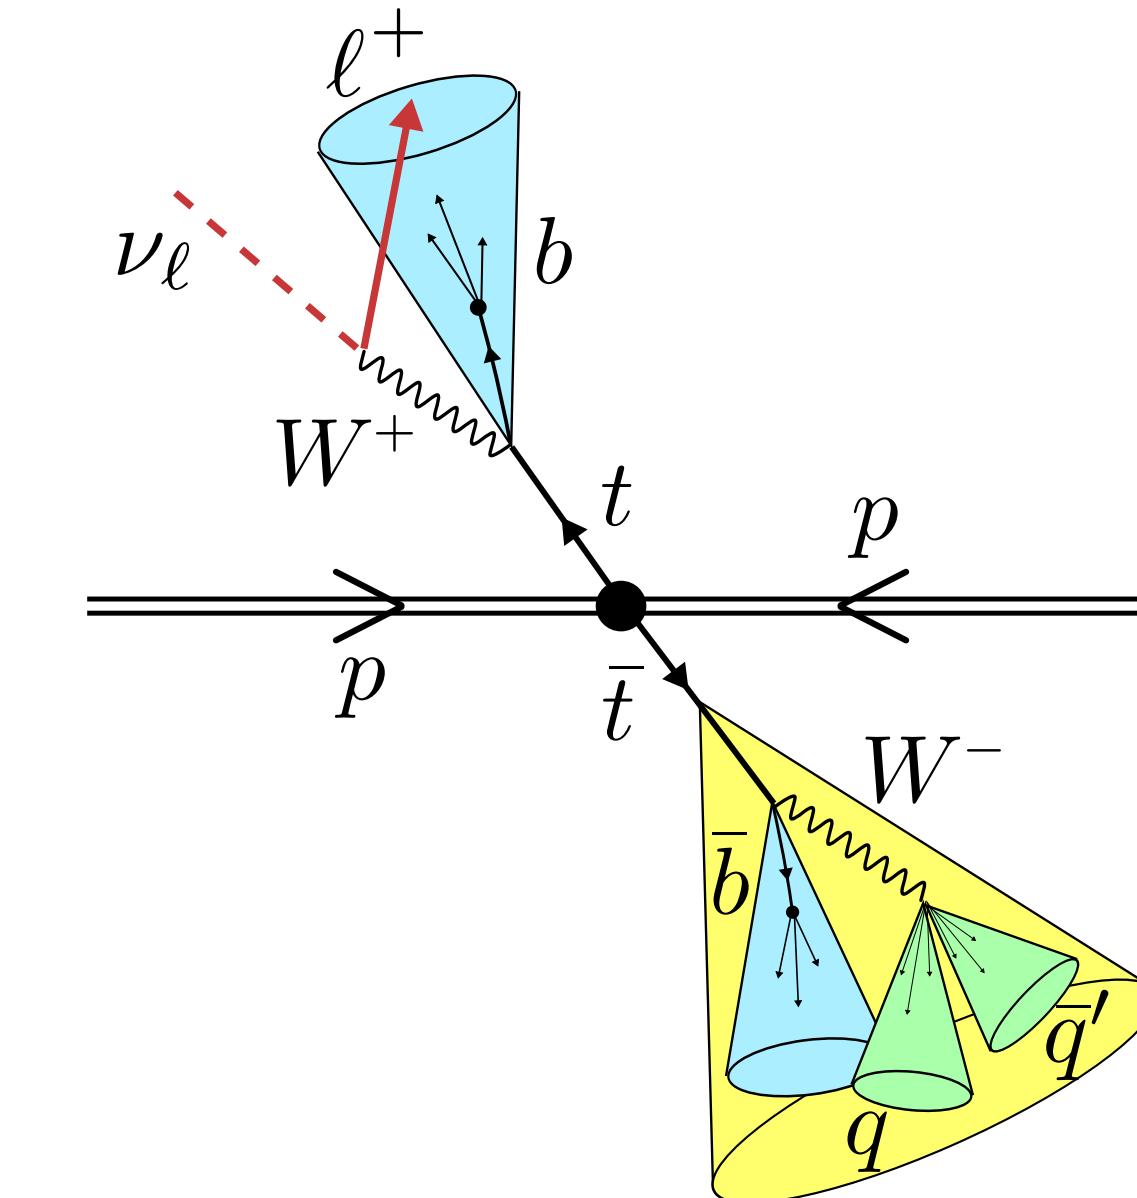
Diagram from [CMS open-data workshop](#)

- Searching for a hypothetical (heavy) BSM particle called  $Z'$  [[arXiv:1810.05905](#)]
- Produced in proton-proton collisions, decays into a pair of top quarks



(resolved)

Increasing  $Z'$  mass  
Collimated decay  
products

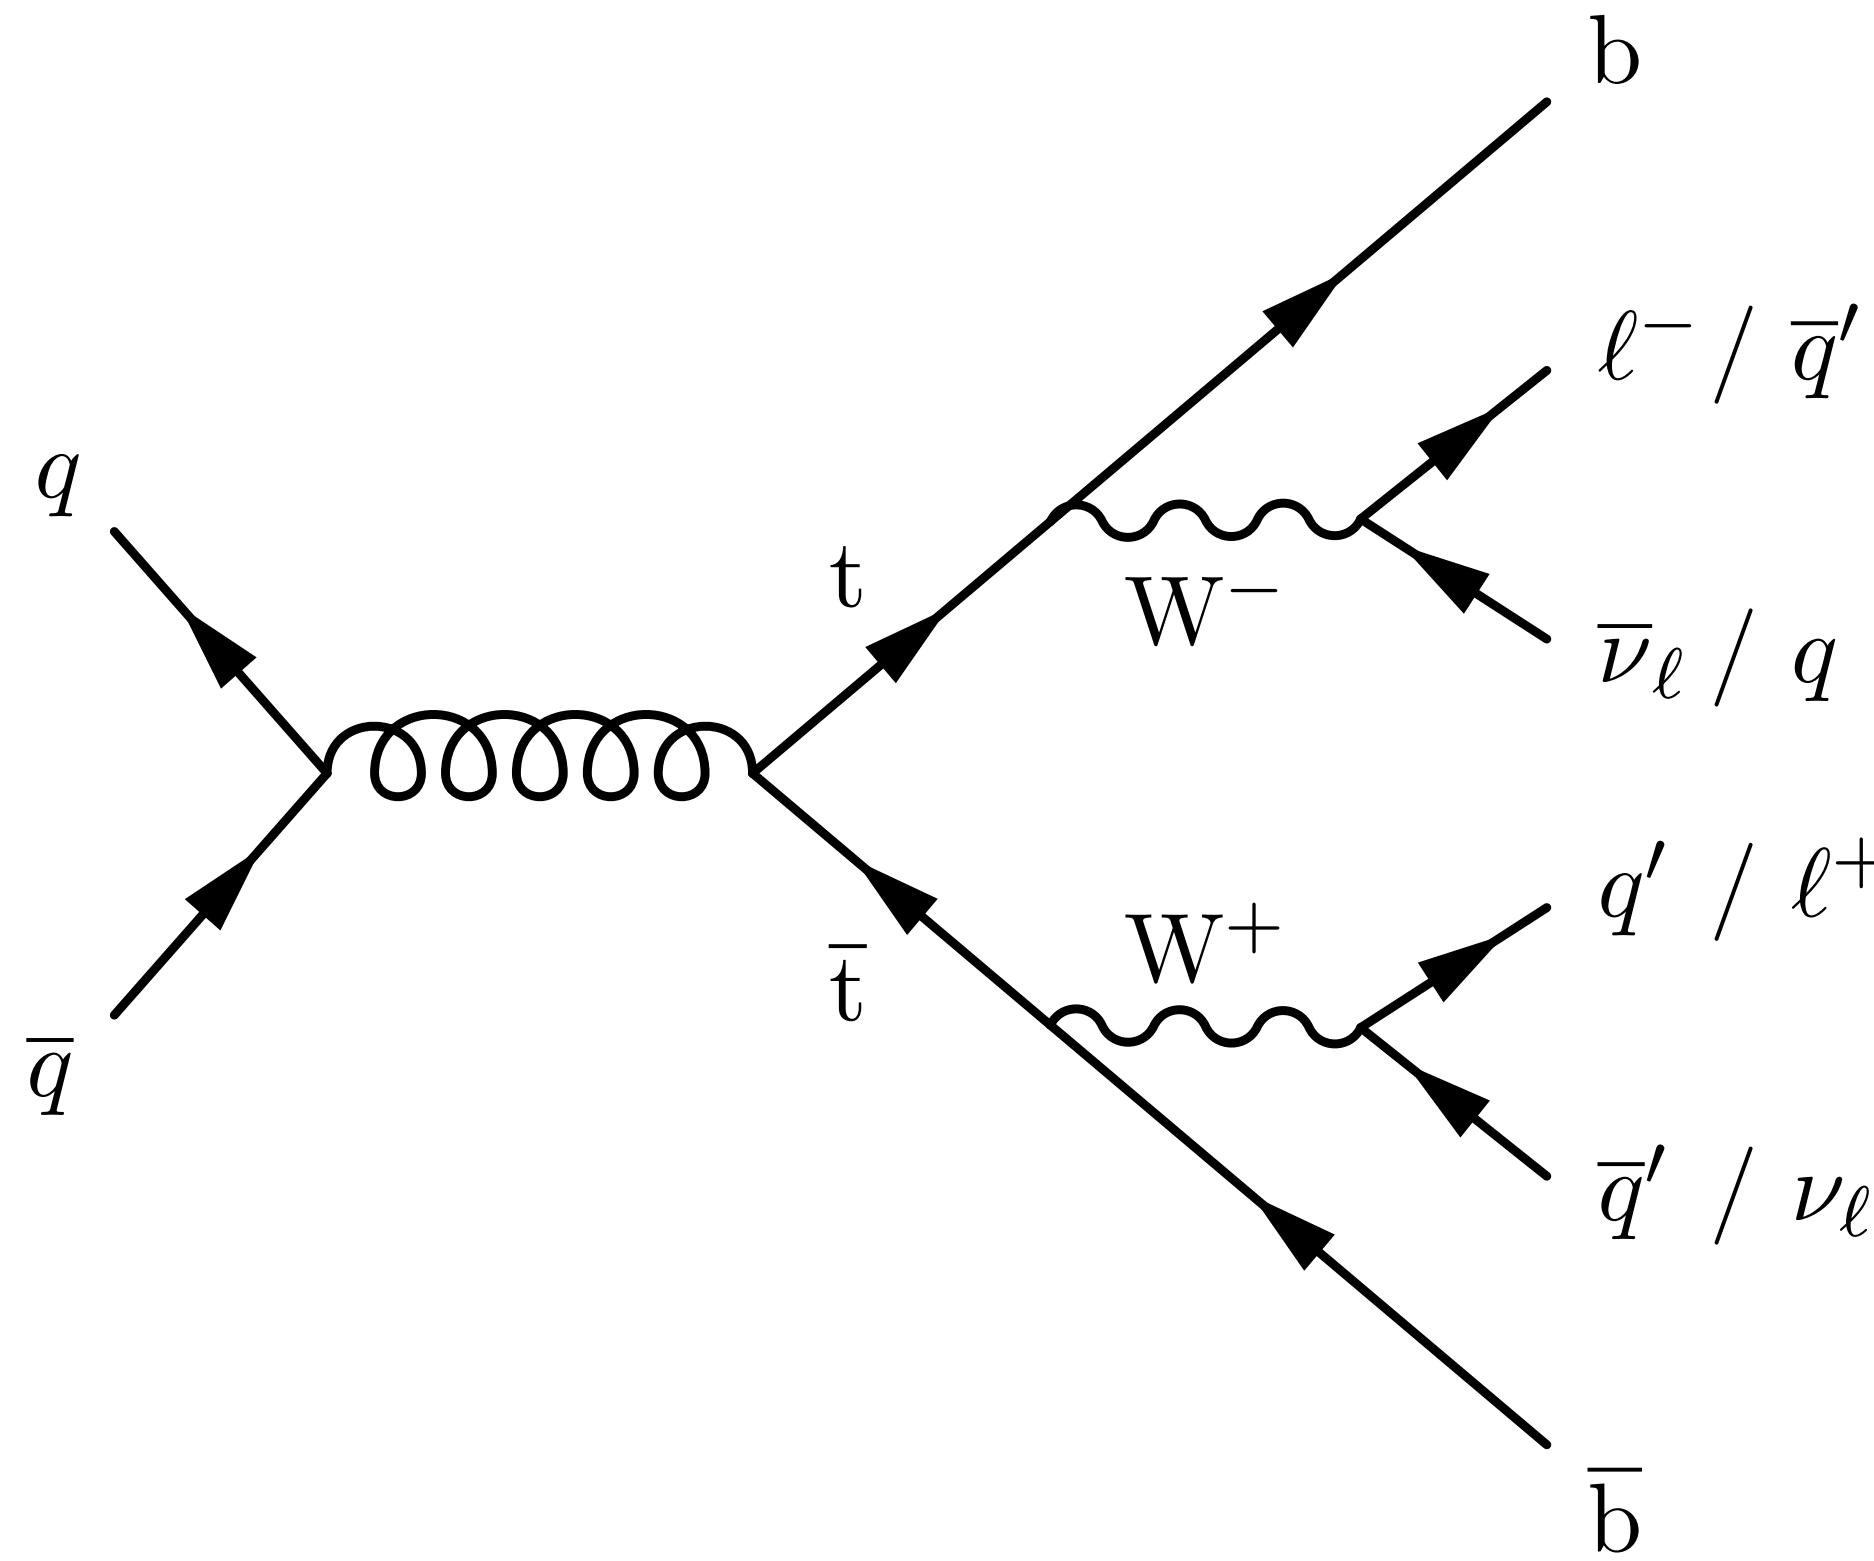


(boosted)

# The simplified analysis

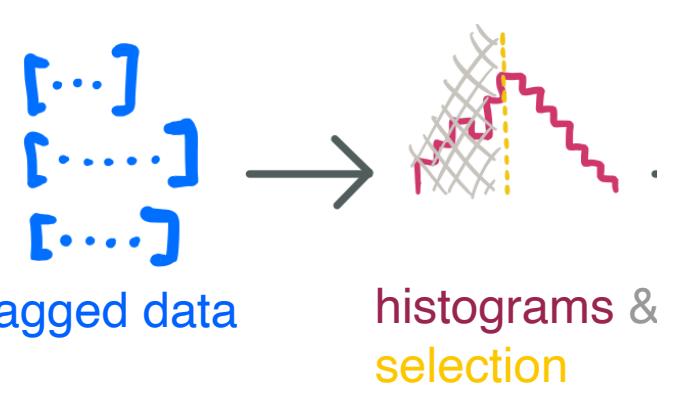
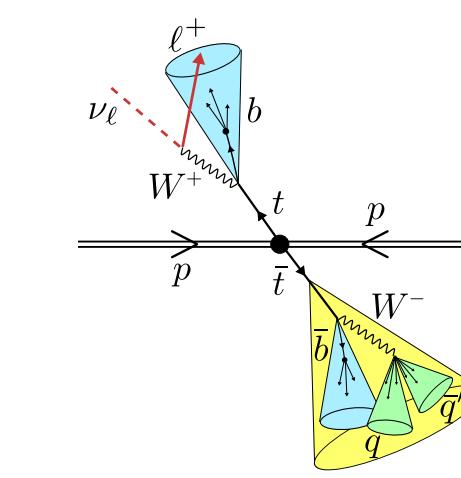
Based on CMS open-data workshop [2024]

Main background: the production of top-quark pairs from the SM



# The simplified analysis

Based on [CMS open-data workshop \[2024\]](#)



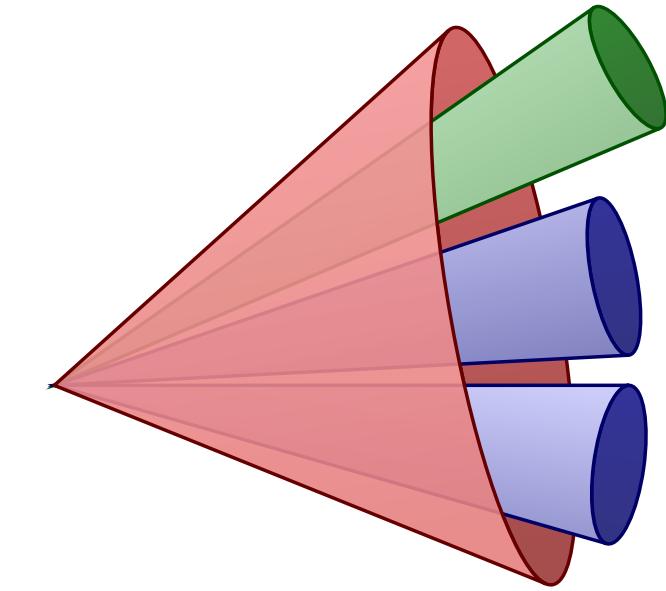
Selecting the signal and dumping the background



Exactly 1 muon



MET > 50 GeV



> 1 b-tagged jet

Proxy to b-quark

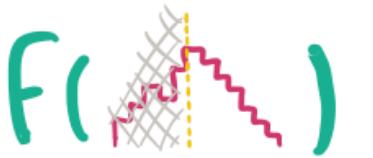
$$H_T = \text{MET} + \text{muon momentum} > 150 \text{ GeV}$$

Exactly 1 fat-jet

Proxy to top-quark

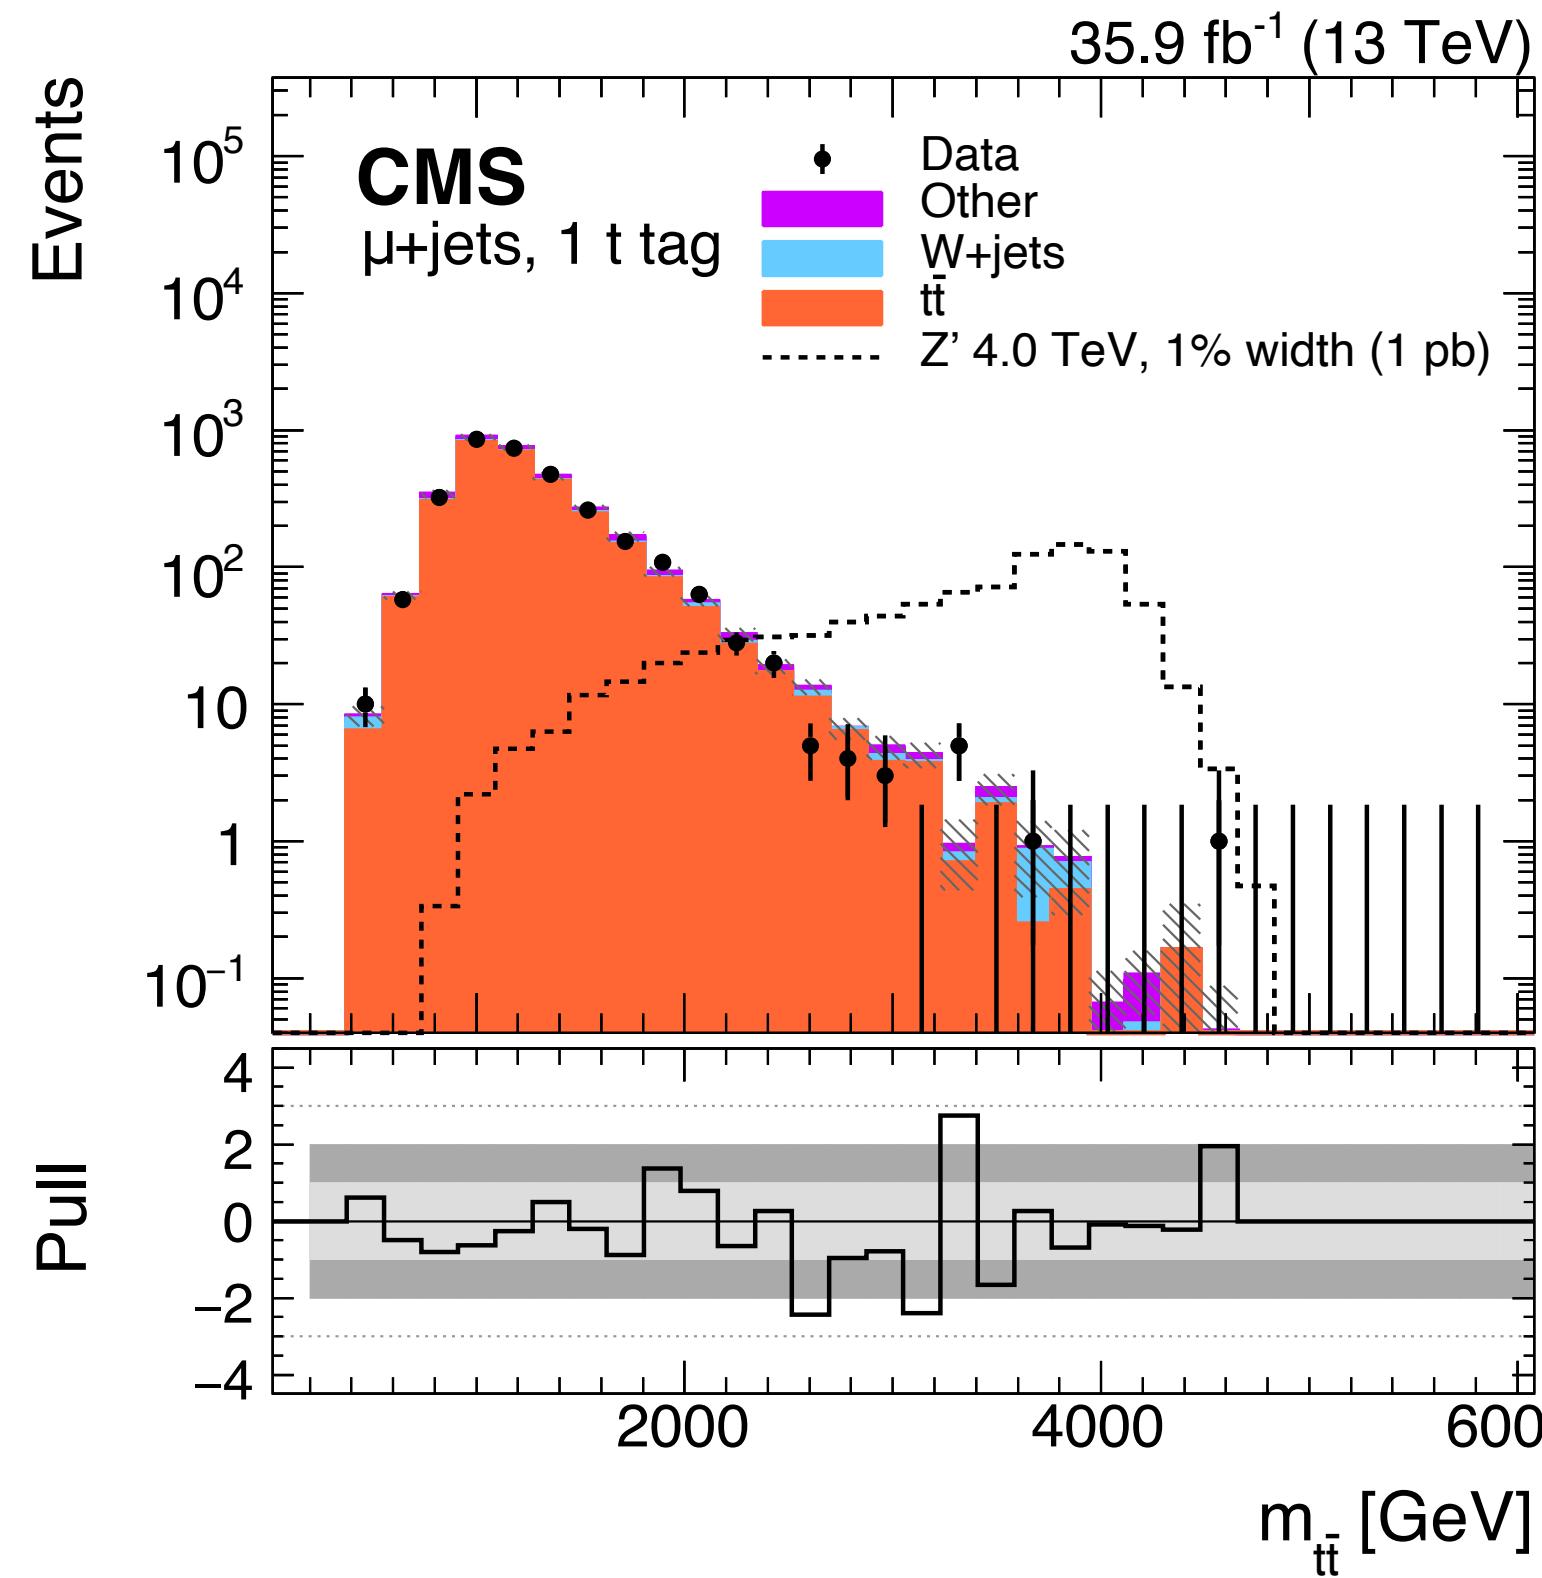
# The simplified analysis

Based on CMS open-data workshop [2024]



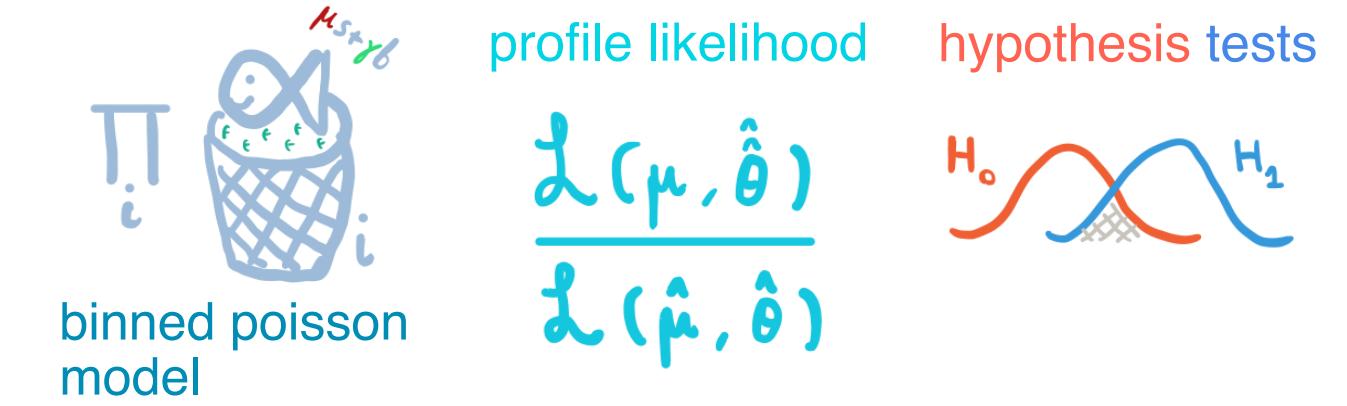
observables

Building a discriminating observable: mass of the  $t\bar{t}$  system



# The simplified analysis

Based on [CMS open-data workshop \[2024\]](#)



Perform a binned profile likelihood fit to extract signal significance [ $p$ -value]

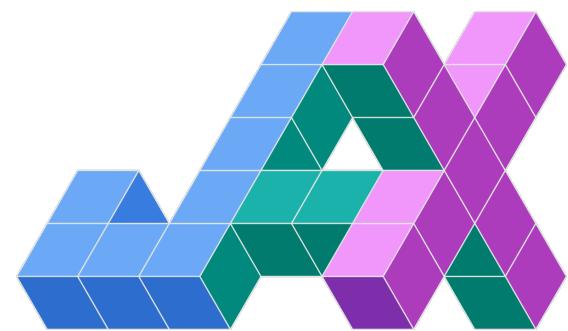
$$p(\vec{n}, \vec{a} | \vec{k}, \vec{\theta}) = \prod_{i \in N_{ch}} \text{Pois}(n_i | \nu_i(\vec{k}, \vec{\theta})) \cdot \prod_{j \in N_{\theta}} c_j(a_j | \theta_j)$$

Bin-by-bin Poisson

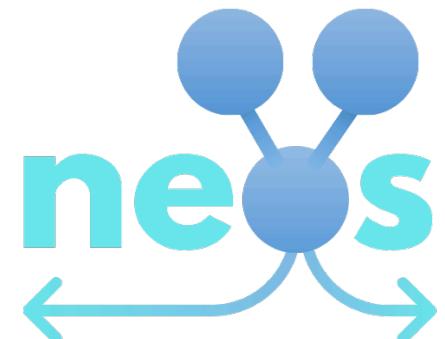
Priors of fit parameters

# Differentiating the analysis

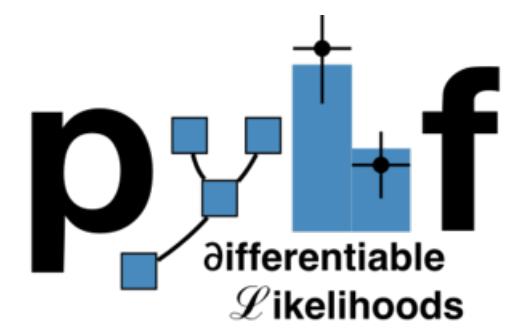
Tools we tried vs tools we used



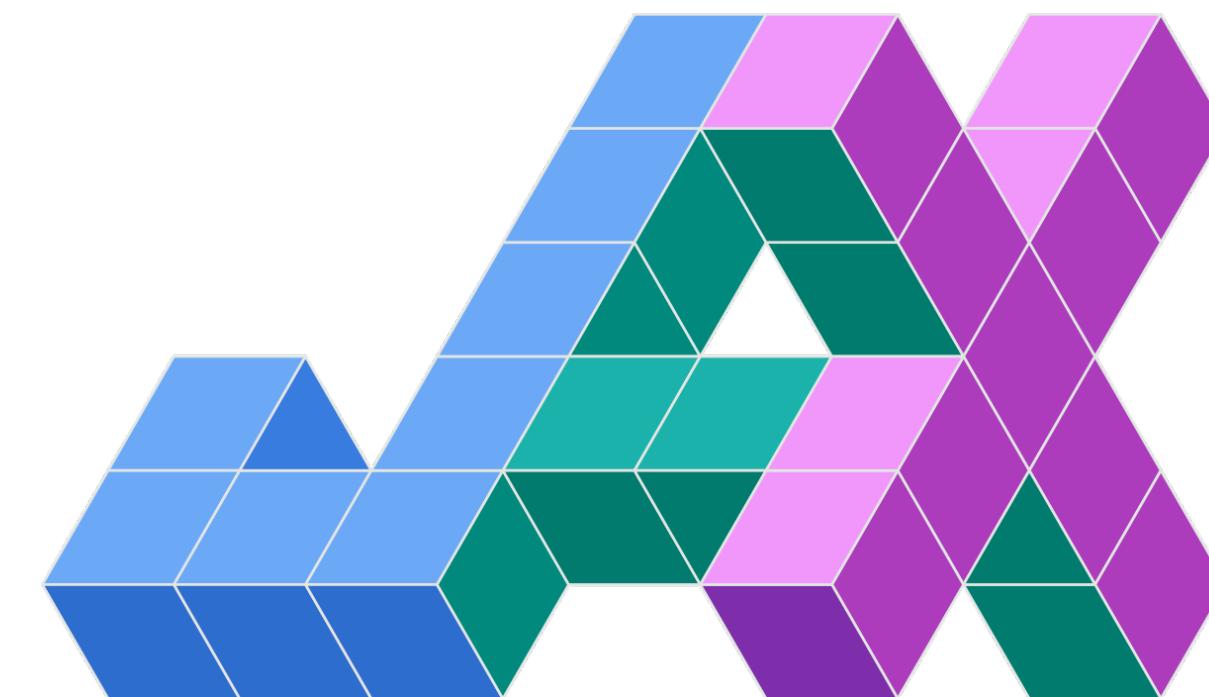
Awkward  
Array



relaxed 😴



eνerℳore



relaxed 😴

Awkward  
Array

VECTOR

# Differentiating the analysis

## Successes and compromises

Success #1: relaxing cuts and differentiable (KDE-based) histograms

```
# Initialise a histogram
histogram = analysis.init_histogram()
# Read the data as an awkward array
data = analysis.process_inputs(input_files)
# Define a cut on x > threshold
cut = data.x > threshold
# Filter the data
data_filtered = data[cut]
# Compute observable
observable = analysis.compute_observable(data_filtered)
# Extract histogram weights
event_weights = analysis.get_weights(data_filtered)
# Fill histogram
histogram.fill(observable, event_weights)
```



```
import jax
import jax.numpy as jnp
# Read the data as an awkward array
data = analysis.process_inputs(input_files)
# Define a relaxed version of the hard cut with temperature T
soft_weight = jax.nn.sigmoid((data.x - threshold)/T)
# Compute observable
observable = analysis.compute_observable(data_filtered)
# Extract histogram weights
event_weights = analysis.get_weights(data_filtered)
# Compute a KDE-based histogram
cdf = jax.scipy.stats.norm.cdf(
    binning,
    observable,
    scale=bandwidth,
)
weighted_cdf = cdf * soft_weight * weights
bin_weights = weighted_cdf[1:, :] - weighted_cdf[:-1, :]
diff_histogram = jnp.sum(bin_weights, axis=1)
```

# Differentiating the analysis

## Successes and compromises

---

### Compromise:

Cannot use jagged data with pure JAX → limited cut and observable implementations

# Differentiating the analysis

## Successes and compromises

---

### Compromise:

Cannot use jagged data with pure JAX → limited cut and observable implementations

Awkward arrays have a JAX backend  
but the functionality is still evolving  
[see [Saransh's talk](#)]

# Differentiating the analysis

## Successes and compromises

Success #2: writing a pure-JAX statistical model to pass to `relaxed`

```
# Extraa# 1) JIT-compiled Poisson log-PDF
@jax.jit
def poisson_ll(n, lam):
    return n * jnp.log(lam + 1e-12) - lam - jsp.special.gammaln(n + 1.0)

# 2) A model compatible with relaxed with manual scaling
class AllBkg(eqx.Module):
    obs: jnp.ndarray
    bkg: jnp.ndarray
    sig: jnp.ndarray

    def expected_data(self, pars):
        mu = pars["mu"]
        theta = pars["tt"]
        return [mu * self.sig + theta * self.bkg], []

    def logpdf(self, data, pars):
        (obs_list, _) = data
        exp_list, _ = self.expected_data(pars)
        return jnp.sum(poisson_ll(obs_list[0], exp_list[0]))
```

```
# 3) Dummy histograms
obs = jnp.array([12., 18., 25., 30.])
bkg = jnp.array([10., 10., 10., 10.])
sig = jnp.array([1., 2., 3., 4.])

# 4) Instantiate model and run hypotest in pure JAX
model = AllBkg(obs=obs, bkg=bkg, sig=sig)
data = ([obs], [])
p0, mle = relaxed.infer.hypotest(
    test_mu      = 0.0,
    data         = data,
    model        = model,
    init_pars   = {"mu": 1.0, "tt": 1.0},
    return_mle_pars = True,
    test_stat   = "q0",
)
```

`relaxed` handles implicit differentiation!

# Differentiating the analysis

## Successes and compromises

---

### Compromise:

Tedious to implement nuisance parameter modifiers from scratch

# Differentiating the analysis

## Successes and compromises

---

### Compromise:

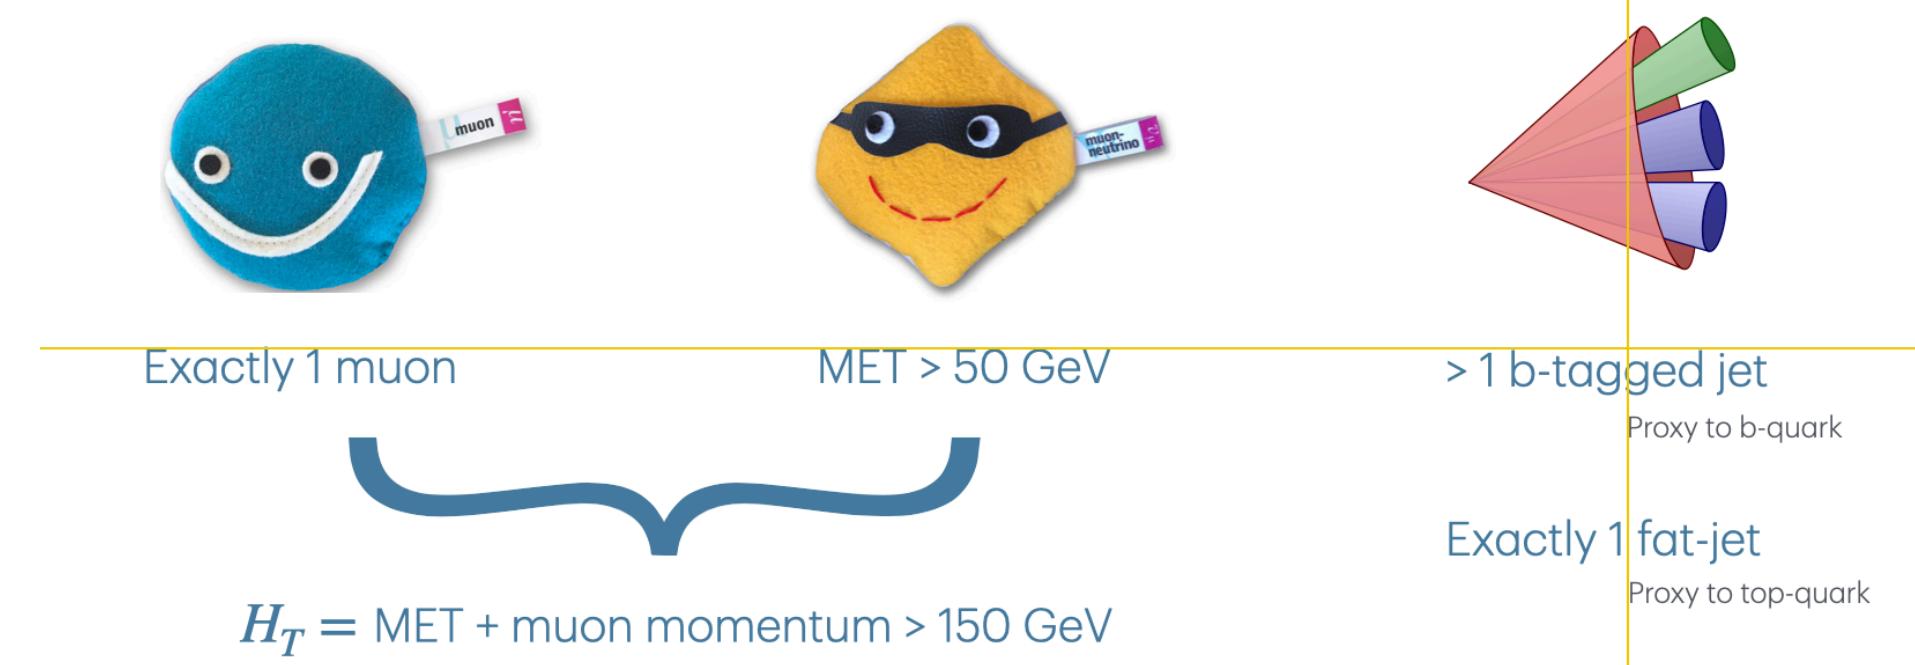
Tedious to implement nuisance parameter modifiers from scratch

hence we did not yet implement  
systematic uncertainties

# Optimisation

## Physics objectives and parameters to optimise

- Compute the derivative of the significance with respect to:
  - The MET cut threshold
  - The b-tagging score threshold used to count b-jets
  - The  $H_T$  cut threshold
- Optimise the signal significance via the  $p$ -value
  - Smaller  $p$ -value, higher significance



# Optimisation

## The gradient descent

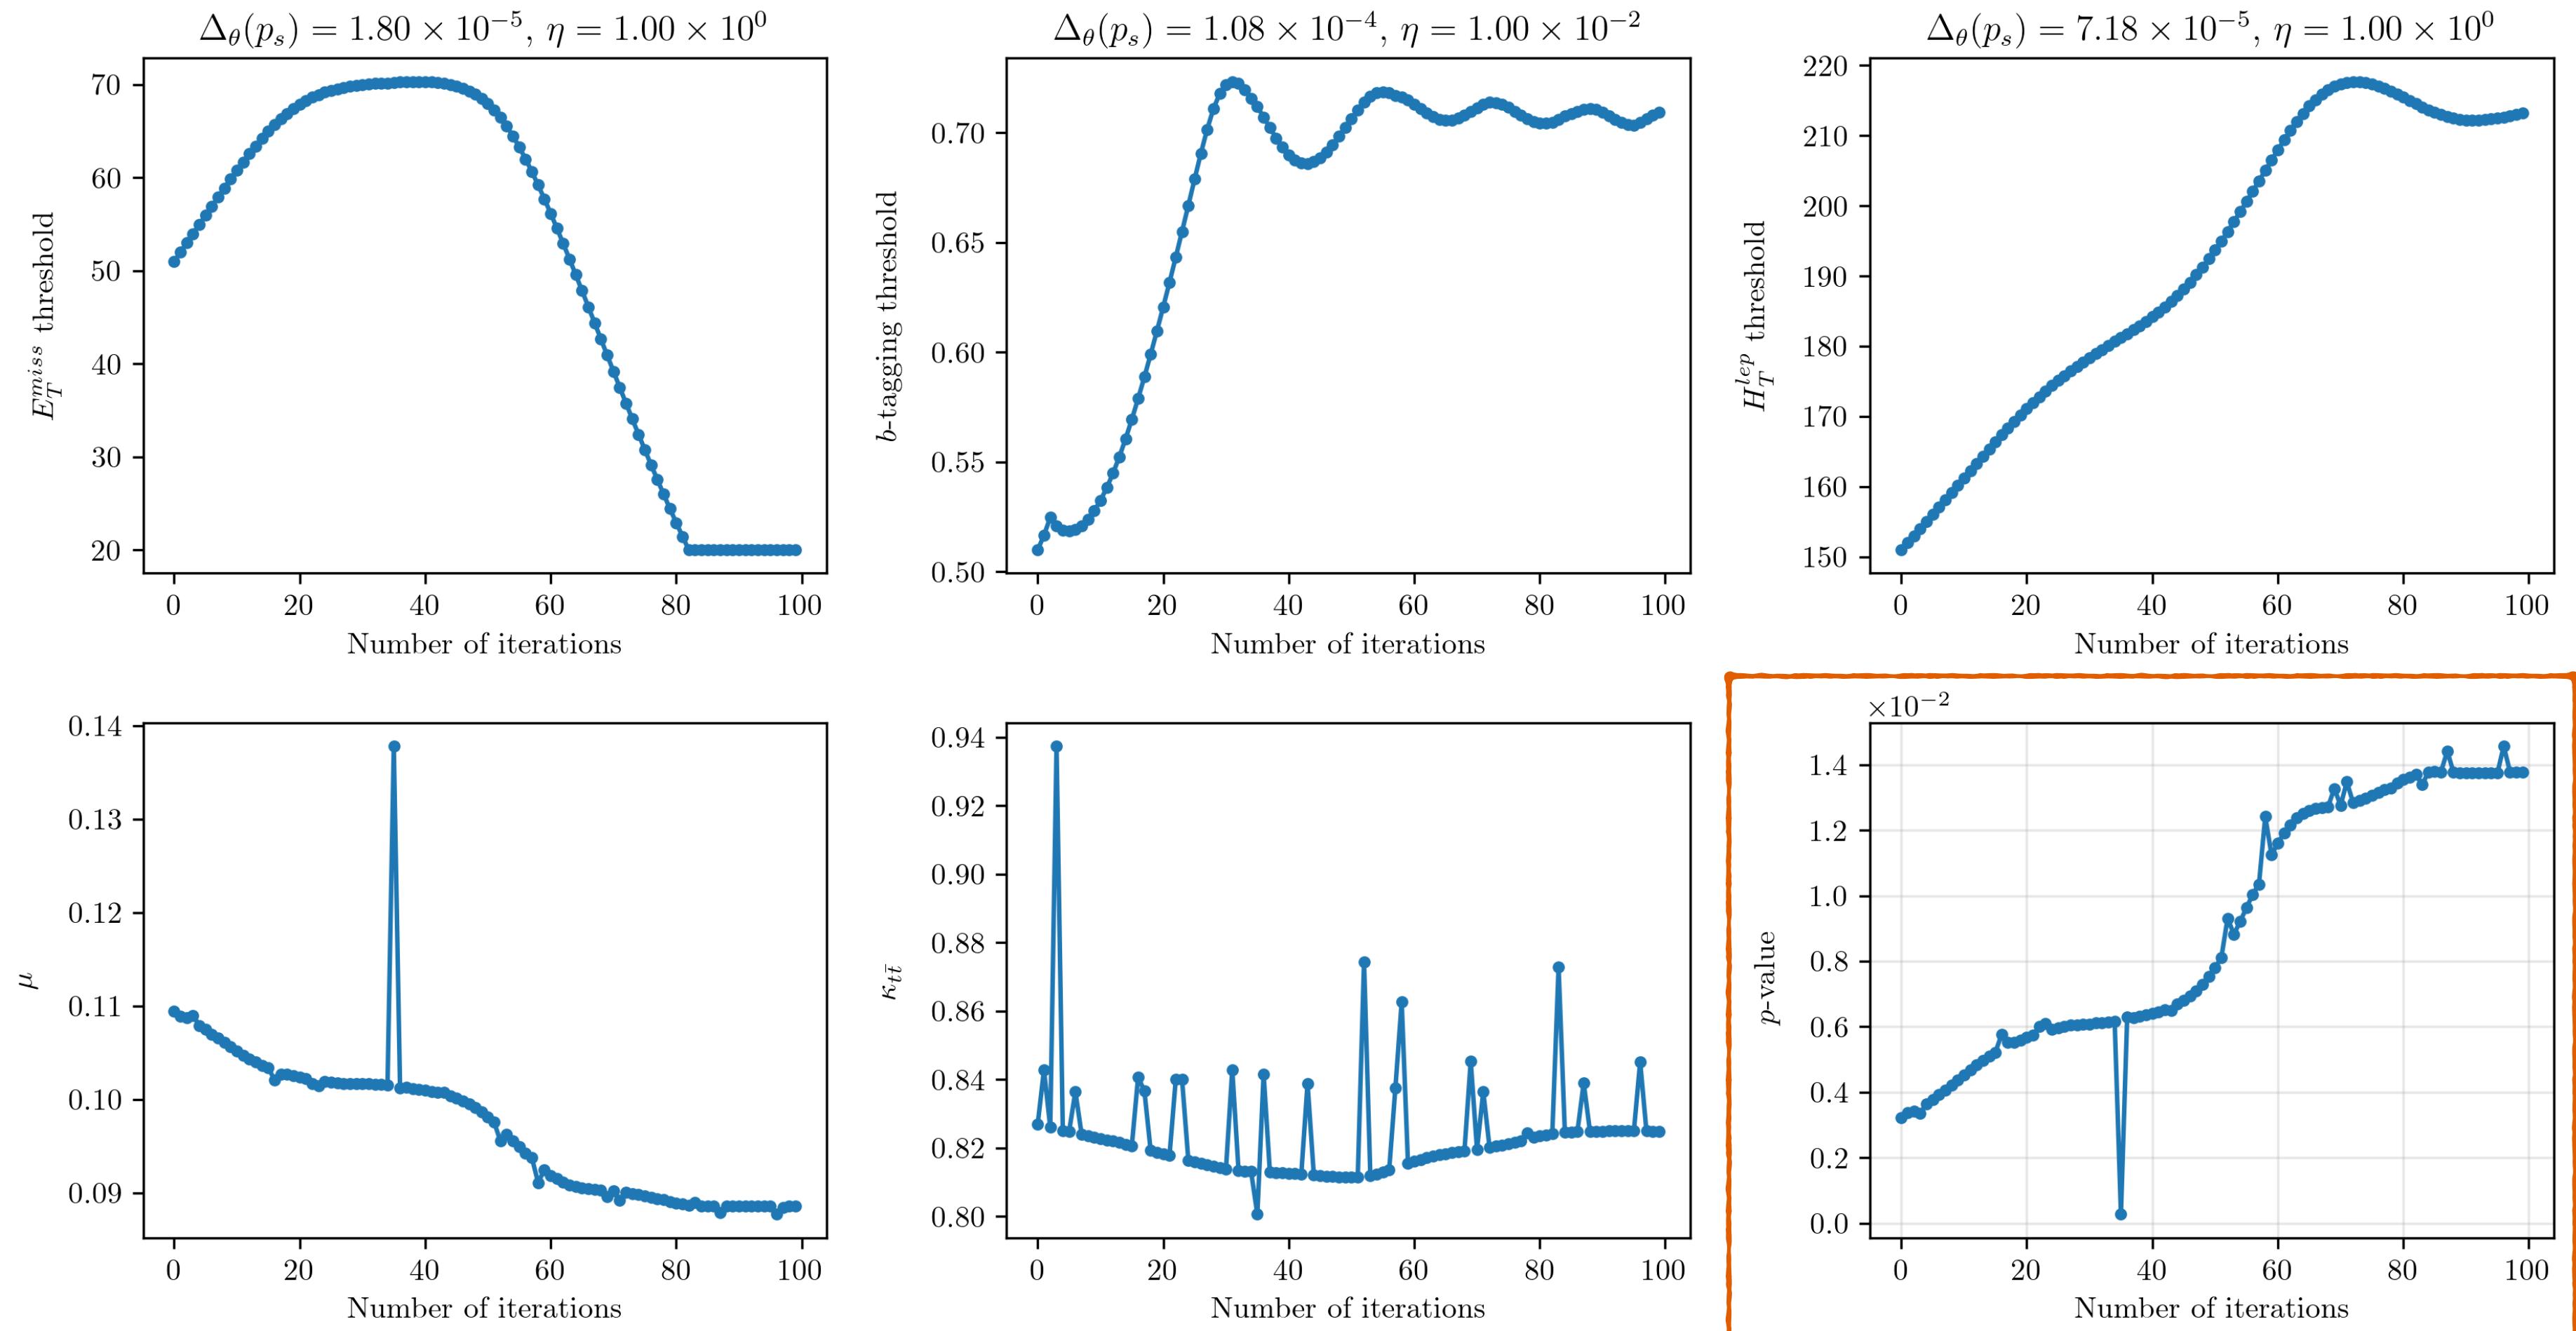
---

- `jaxopt.OptaxSolver` for back-propagation
- Implement parameter-wise learning rates and clipping functions
- Implicit differentiation of fit parameters in `relaxed`

# Optimisation

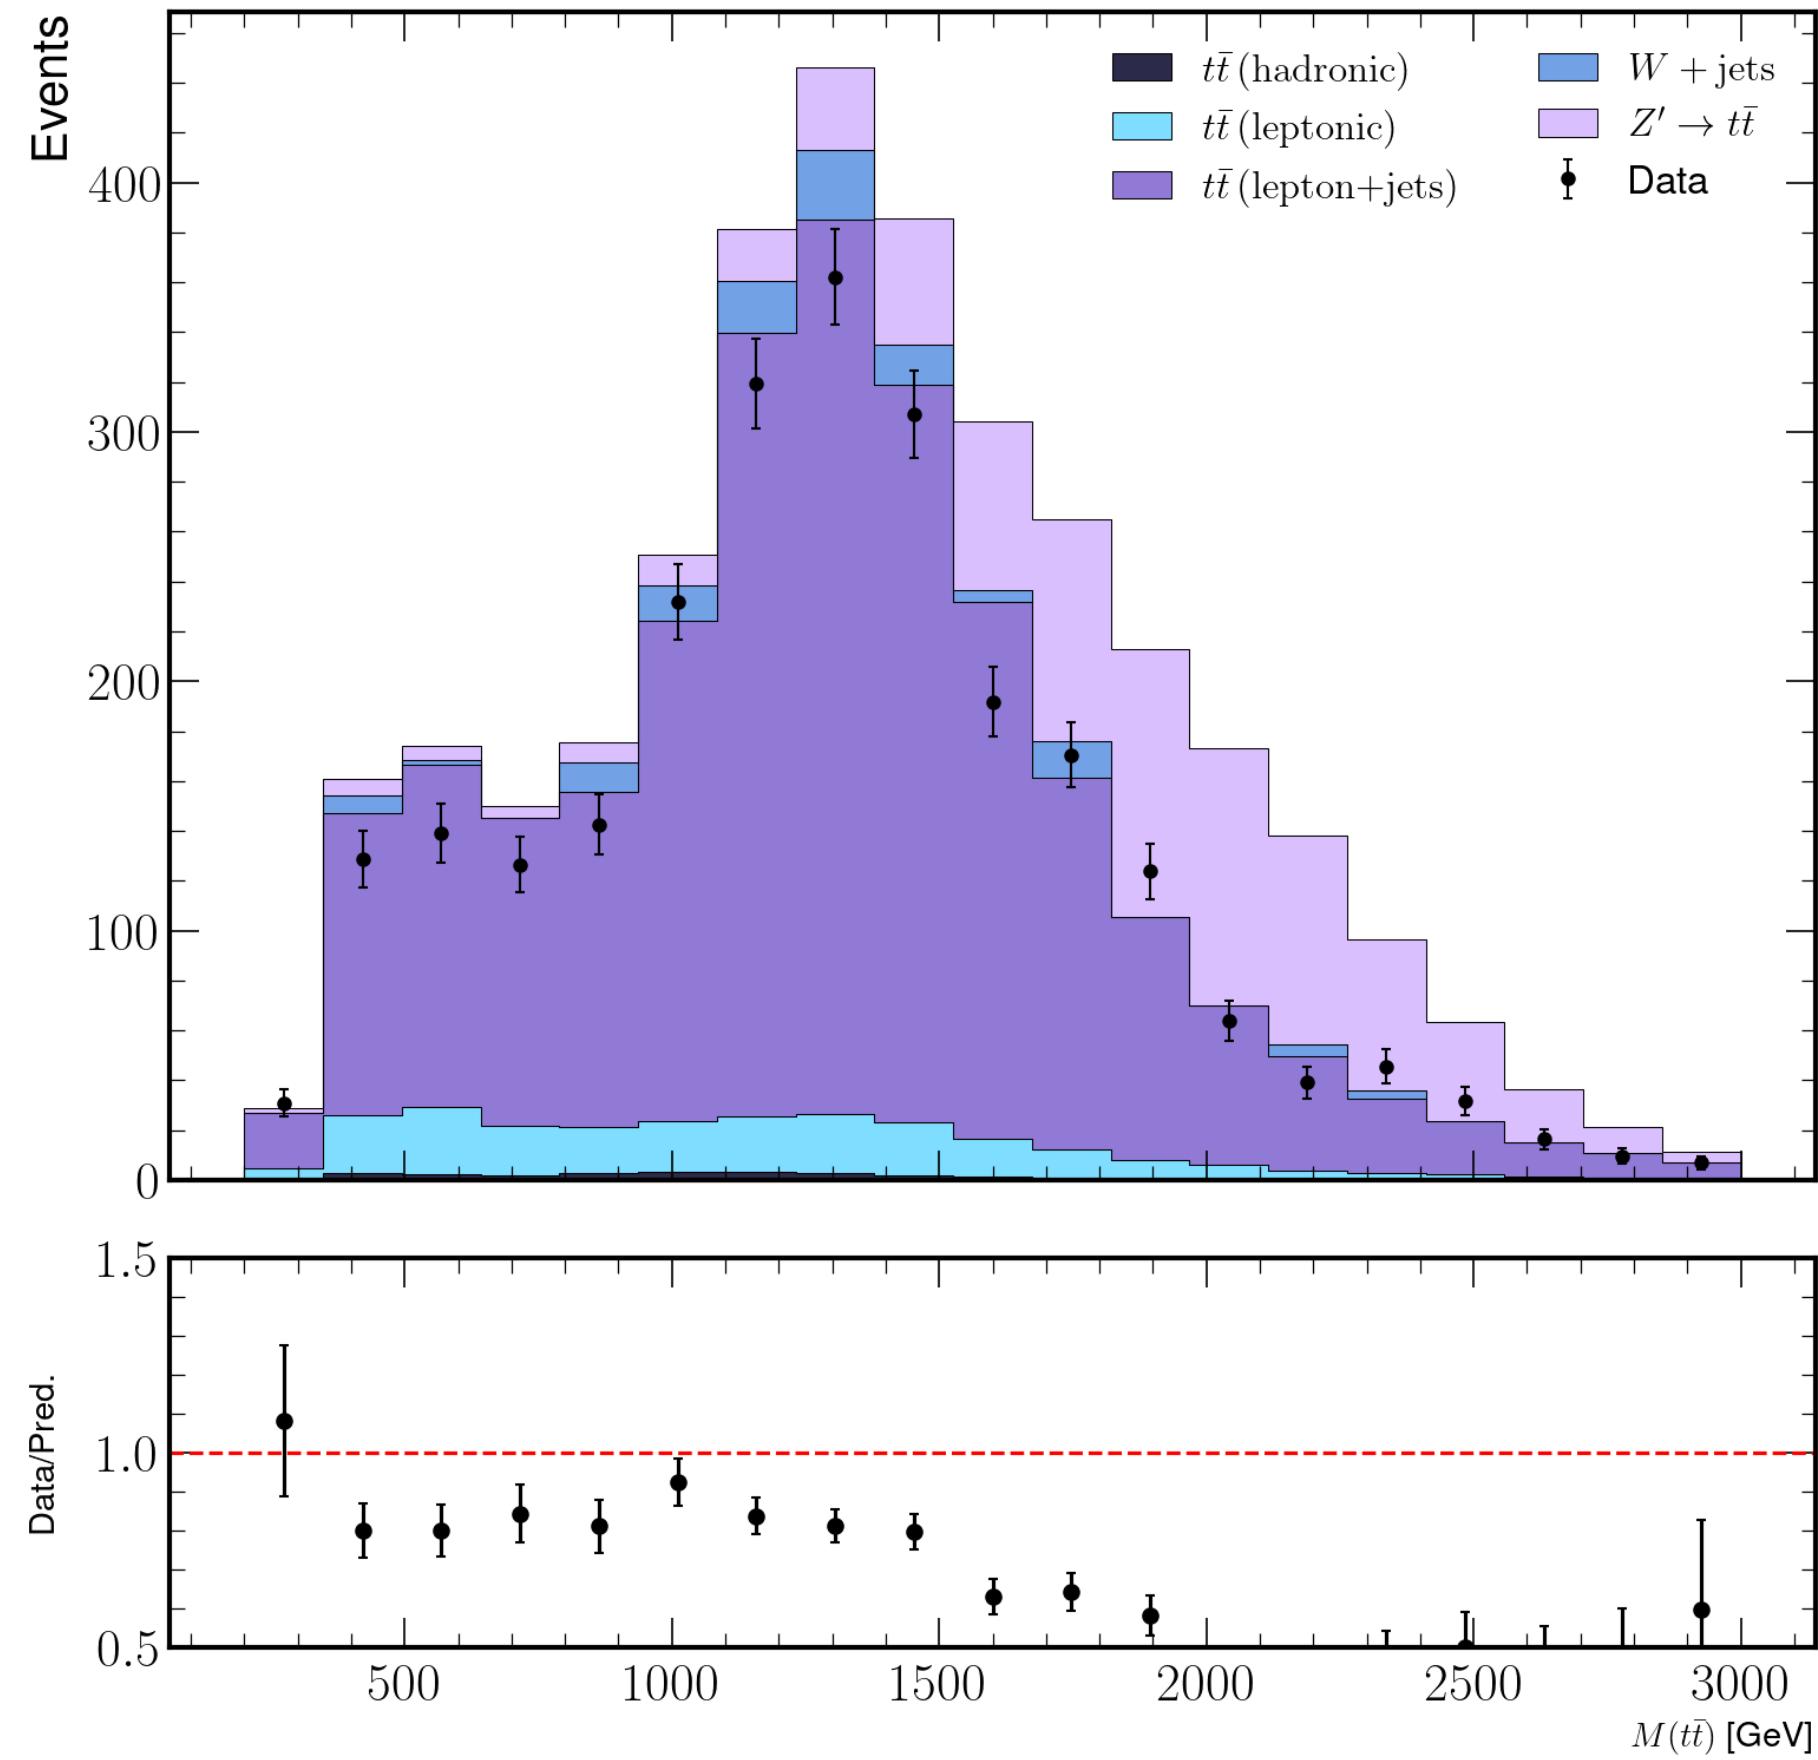
## The gradient descent

- 100 iterations in ~3 minutes without JIT
- JIT optimisations were difficult to implement
- 320% improvement in significance

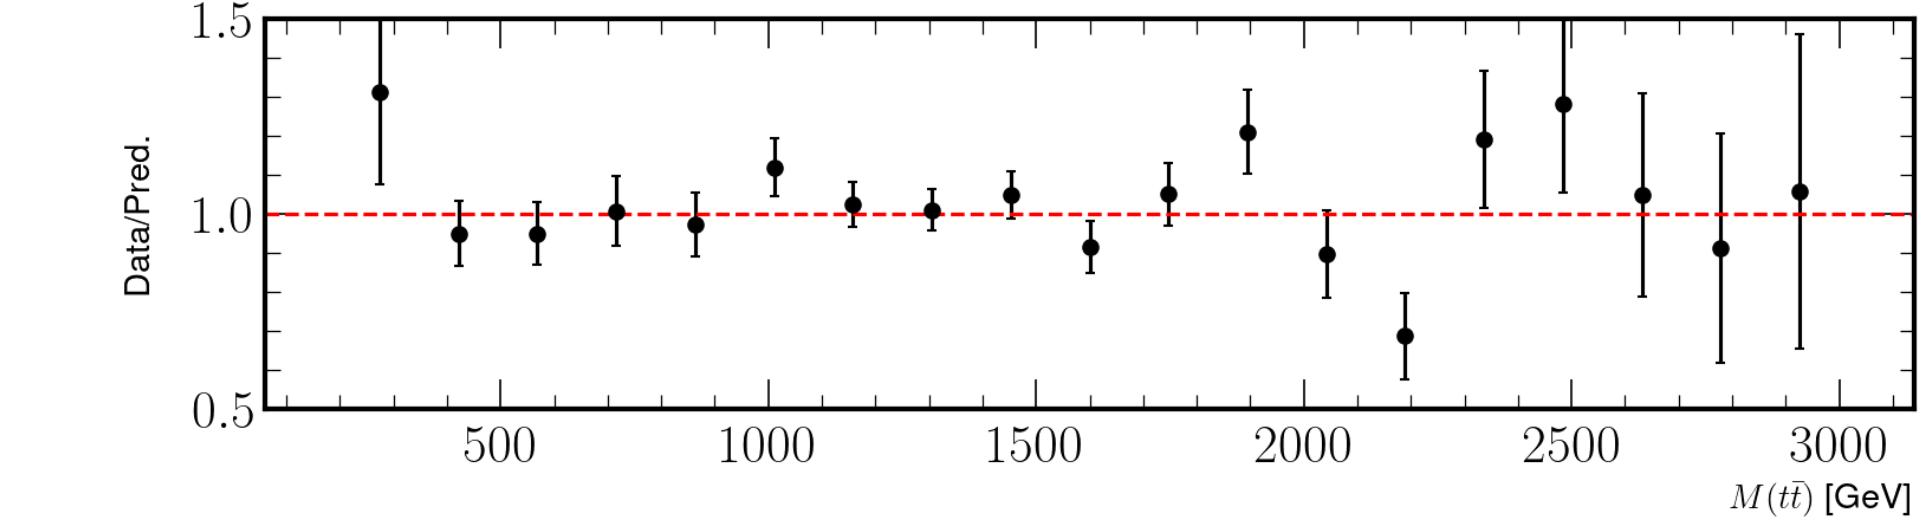
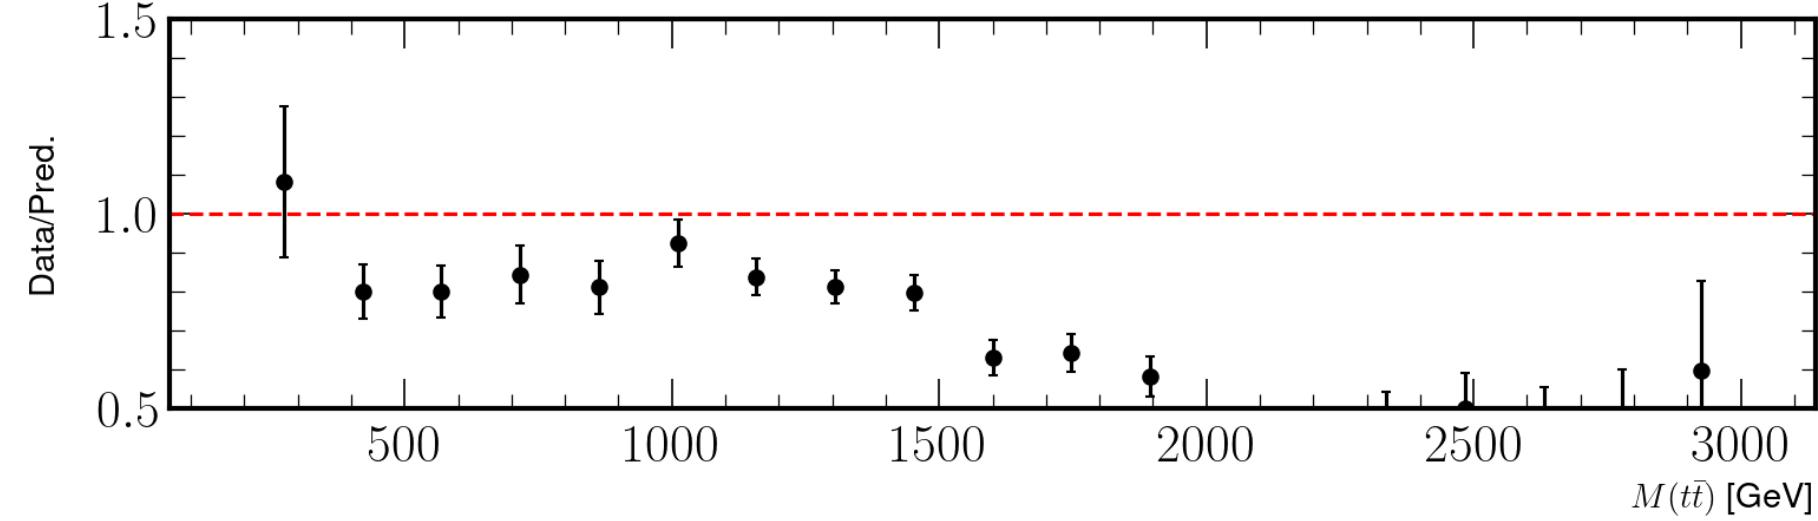
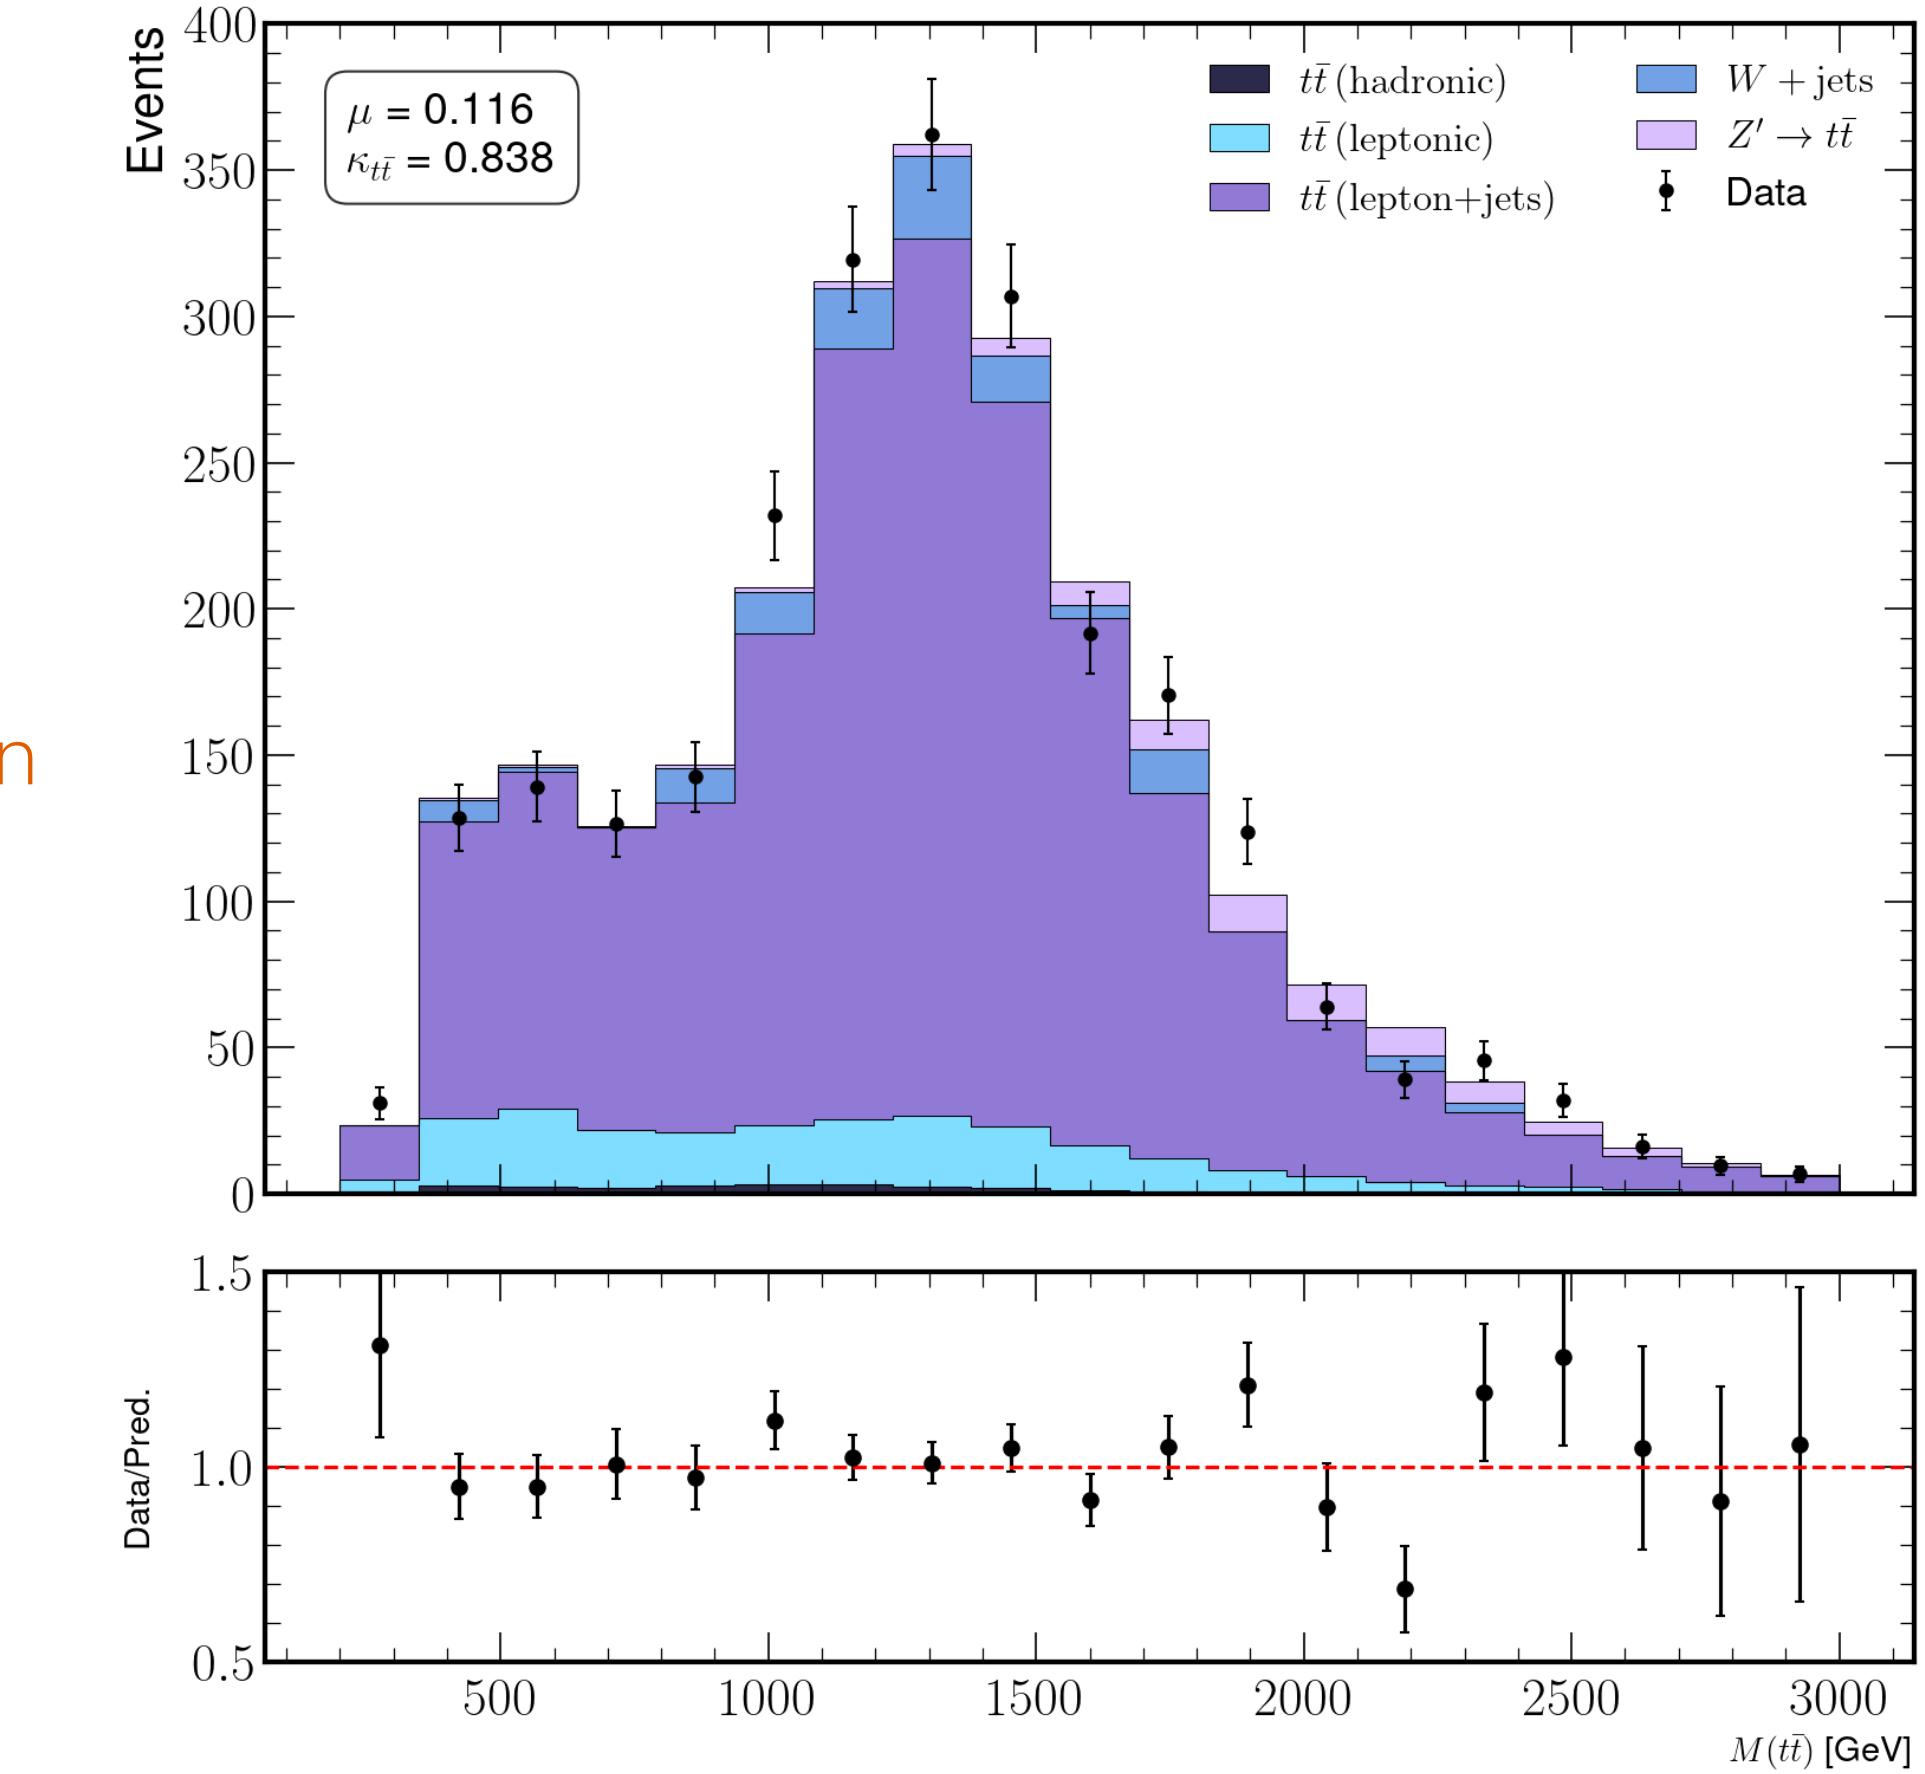


# Optimisation

## The fit results



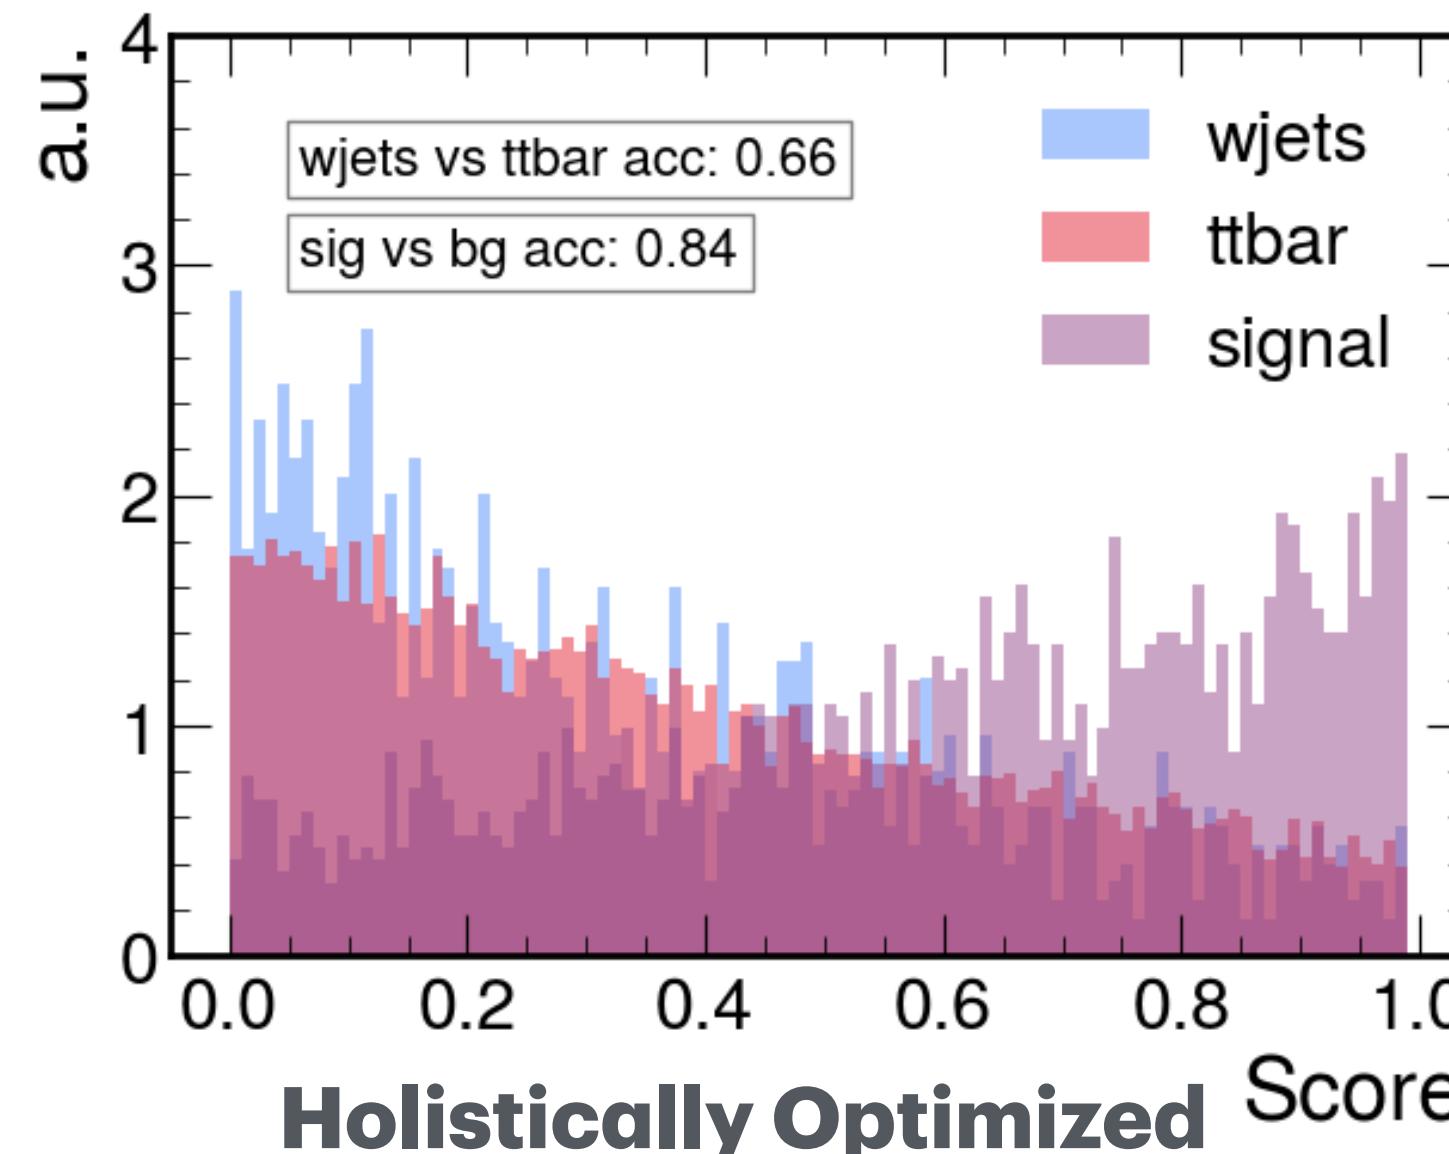
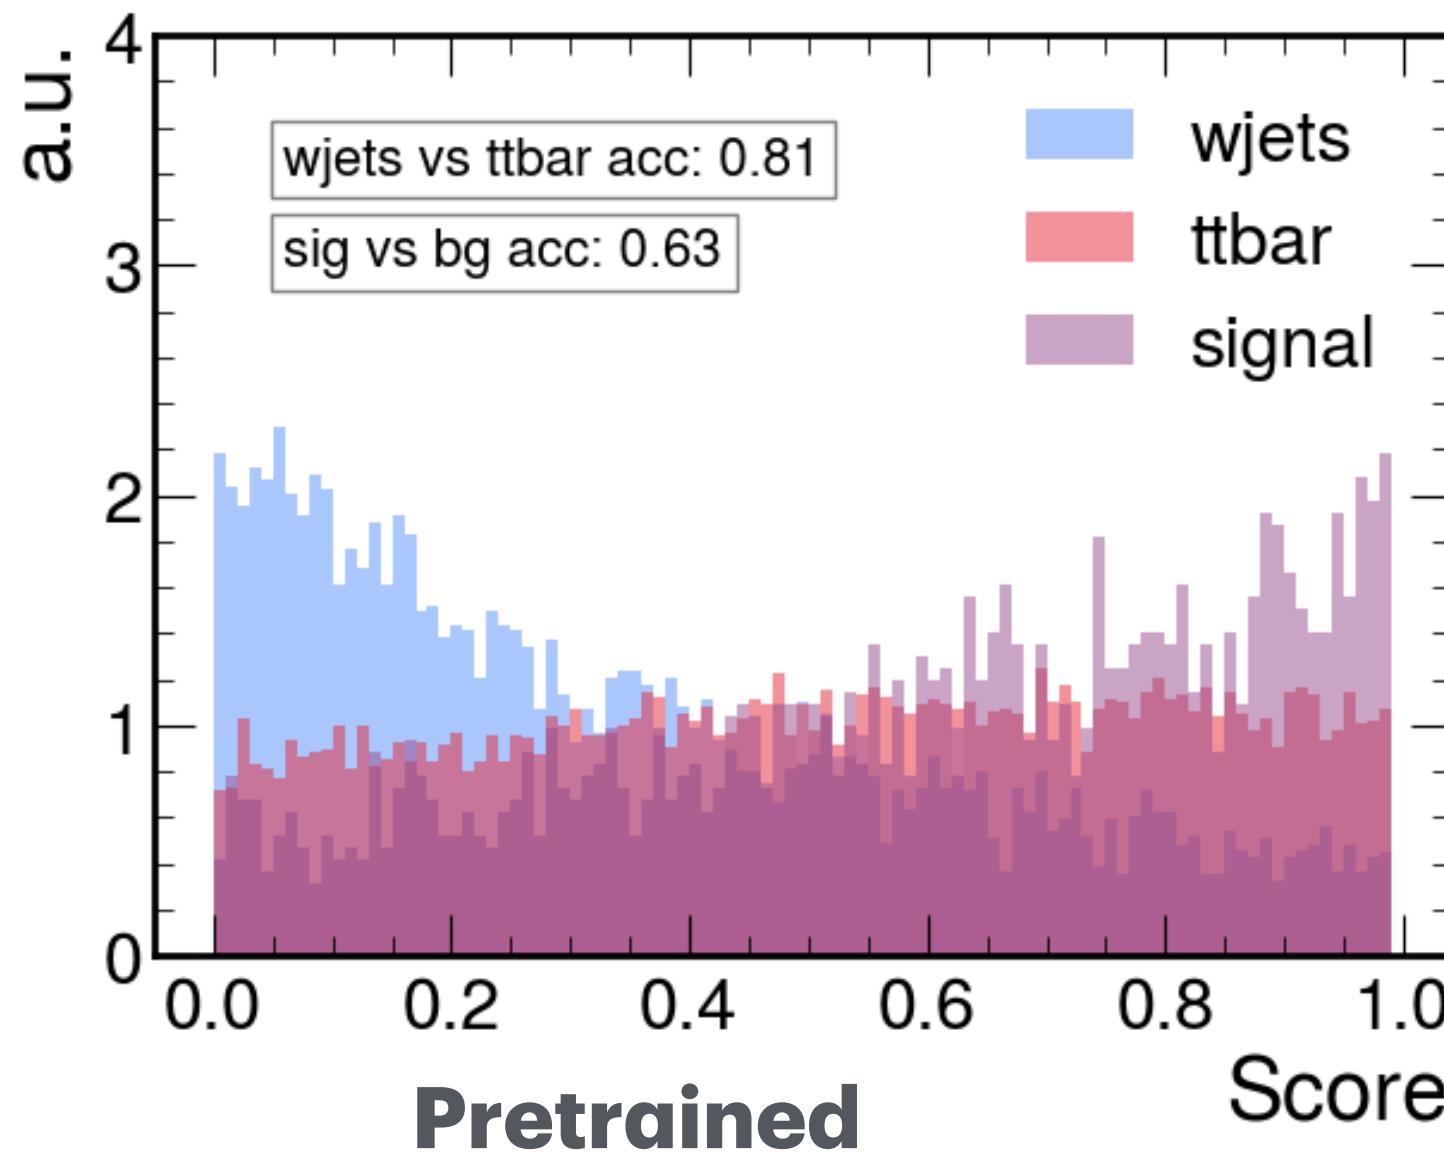
Holistic Optimisation



# MVA Classifier

## Holistically Optimizing ML steps

- Pre-train simple wjets / ttbar classifier from high-level kinematic variables
- Apply soft-cut on low value (0.05) to suppress wjets in signal region
  - Optimize NN's parameters as part of the entire pipeline (small learning rate)



- Re-training the wjets-vs-ttbar classifier in the holistic optimization turns it into a signal-vs-background classifier!
- No ground-breaking insight, but demonstration of feasibility to optimize O(1k) parameters!

# Conclusions

---

## Goal #1

Scale concept of differentiable analysis to more realistic HEP data from major LHC experiment (CMS)



## Goal #2

Review state + compatibility of existing tools and identify challenges to be addressed towards complete analysis



# Way forward

---

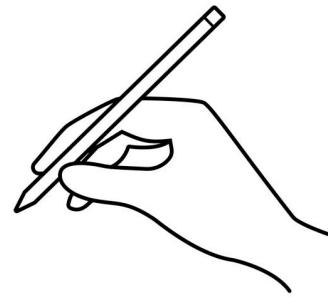
- Gradually make differentiable example more realistic
  - Also fit control regions, include systematic uncertainties
  - Optimize w.r.t more-and-more parameters
  - Include more realistic ML algorithms in the pipeline
- Eventually optimize (and publish) an actual physics measurement (w/ approval from collaboration)
- Document all missing features, bugs, and compatibility issues in GitHub issues
- Discuss with community how to proceed towards a differentiable Scikit-HEP ecosystem

---

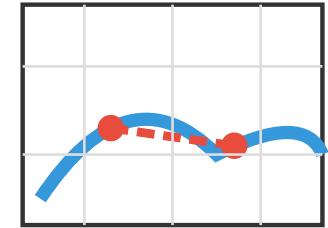
# Backup

# What is differentiable programming?

Four ways to calculate derivatives in a computer program



- Manual differentiation (calculate by hand and implement in code)



$\Delta y / \Delta x$

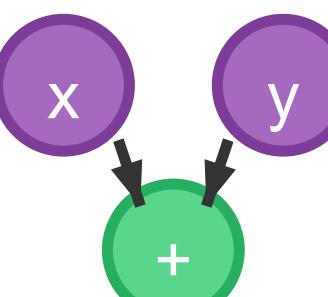
- Numerical differentiation (finite difference approximation)



$$x^3 \rightarrow 3x^2$$

- Symbolic differentiation (program calculates analytical derivative)

- Equations can become complex ("expression swell")



(val, grad)

- Automatic differentiation (store gradients, apply chain rule)

- Key idea behind backpropagation

# Fixed-Point Differentiation

Fitting the model and optimising the analysis - two loops?

- Physics objective  $S$  depends on fit result:
- Two optimisation loops? No, there is a trick!

$$S(\theta^*(\alpha), \alpha) \Rightarrow \frac{dS}{d\alpha} = \frac{\partial S}{\partial \alpha} + \frac{\partial S}{\partial \theta^*} \cdot \frac{d\theta^*}{d\alpha}$$

↑  
Best fit result

Implicit  
function  
theorem

$$f(x, y) = 0 \Rightarrow \exists \phi : f(x, \phi(x)) = 0 \text{ and } \frac{d\phi}{dx} = - \left( \frac{\partial f}{\partial y} \right)^{-1} \left( \frac{\partial f}{\partial x} \right)$$

Best fit  
Result:

$$\nabla_{\theta} \mathcal{L}(\theta^*(\alpha), \alpha) = 0 \Rightarrow \frac{d\theta^*}{d\alpha} = - \left( \frac{\partial^2 \mathcal{L}}{\partial \alpha \partial \theta} \right)^{-1} \left( \frac{\partial^2 \mathcal{L}}{\partial \theta^2} \right)$$

Combined  
Result

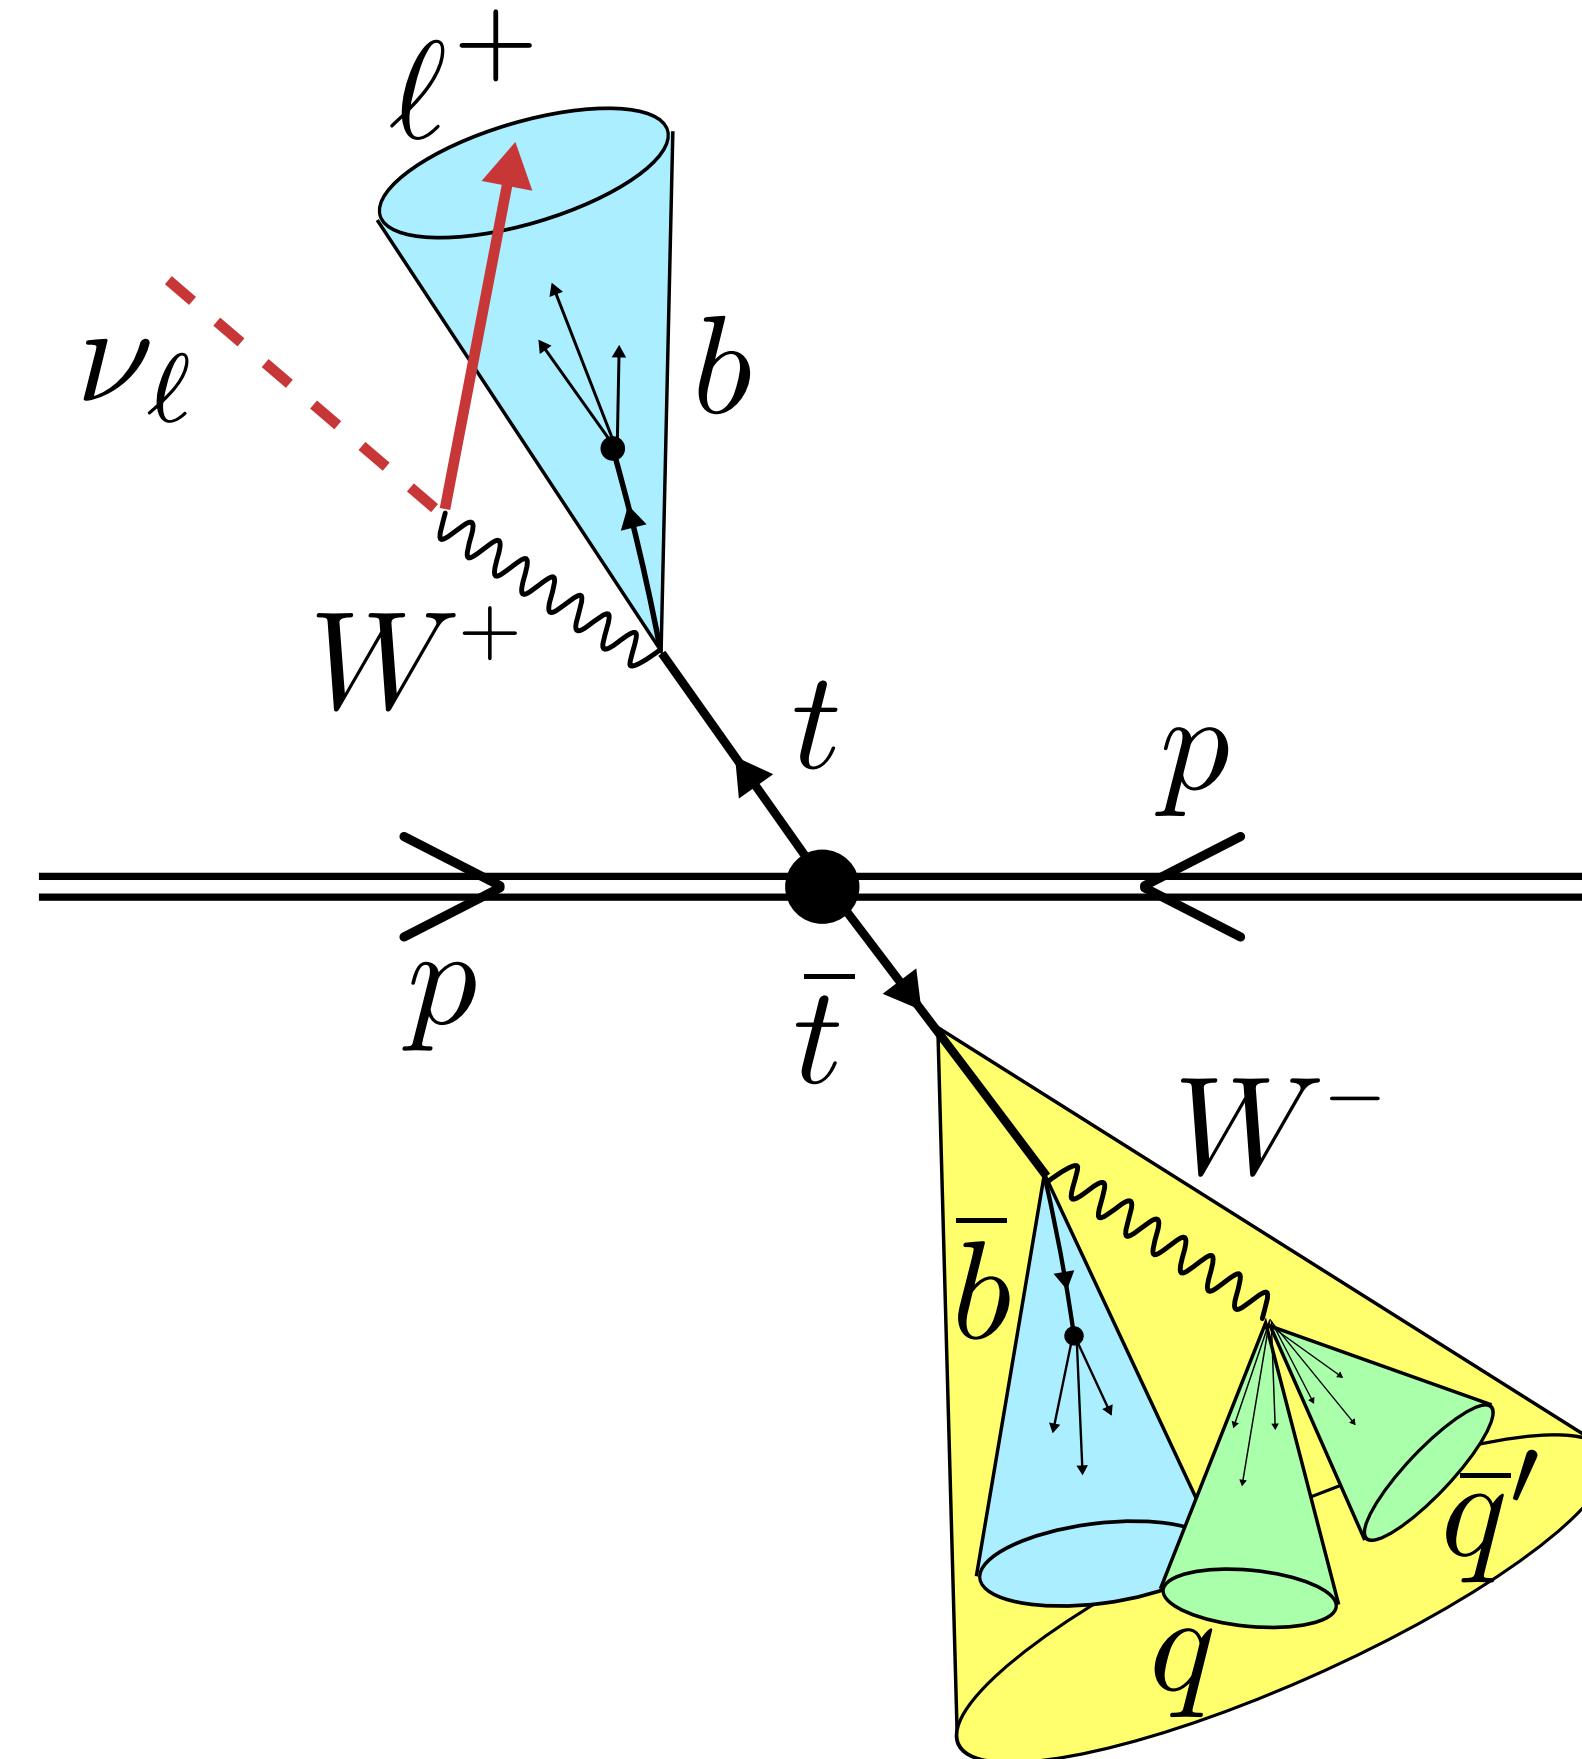
$$\frac{dS}{d\alpha} = - \frac{\partial S}{\partial \alpha} - \frac{\partial S}{\partial \theta^*} \cdot \left( \frac{\partial^2 \mathcal{L}}{\partial \theta^2} \right)^{-1} \cdot \left( \frac{\partial^2 \mathcal{L}}{\partial \theta \partial \alpha} \right)$$

- Gradient of Objective under condition that the likelihood is stationary
- Can run gradient descent along the best-fit manifold. Very neat!

# Dissecting the process

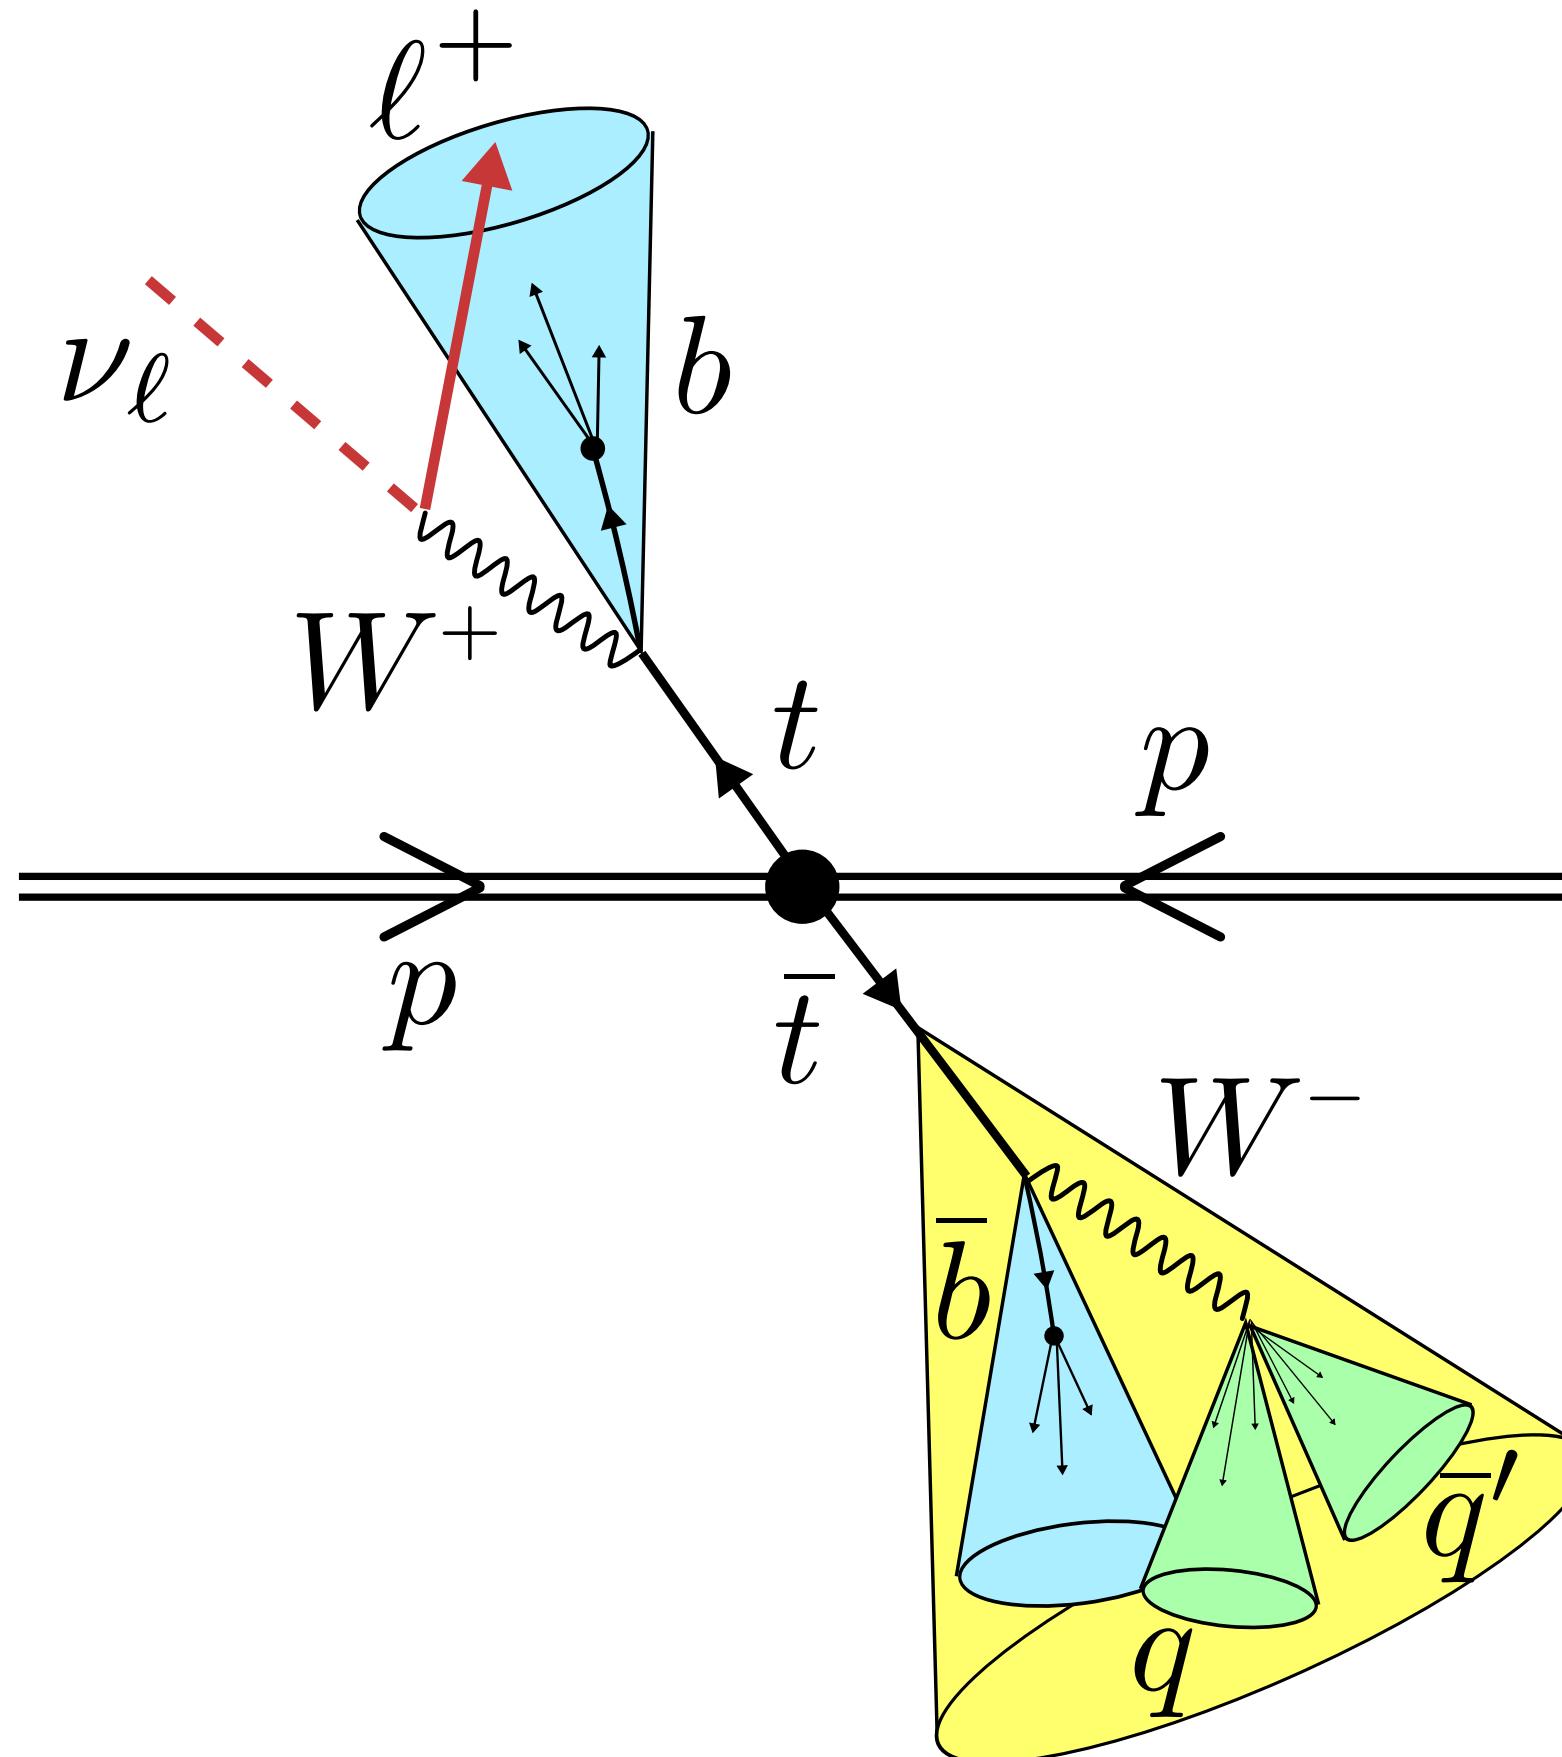
Getting the terminology right

Diagram from [CMS open-data workshop](#)



# Dissecting the process

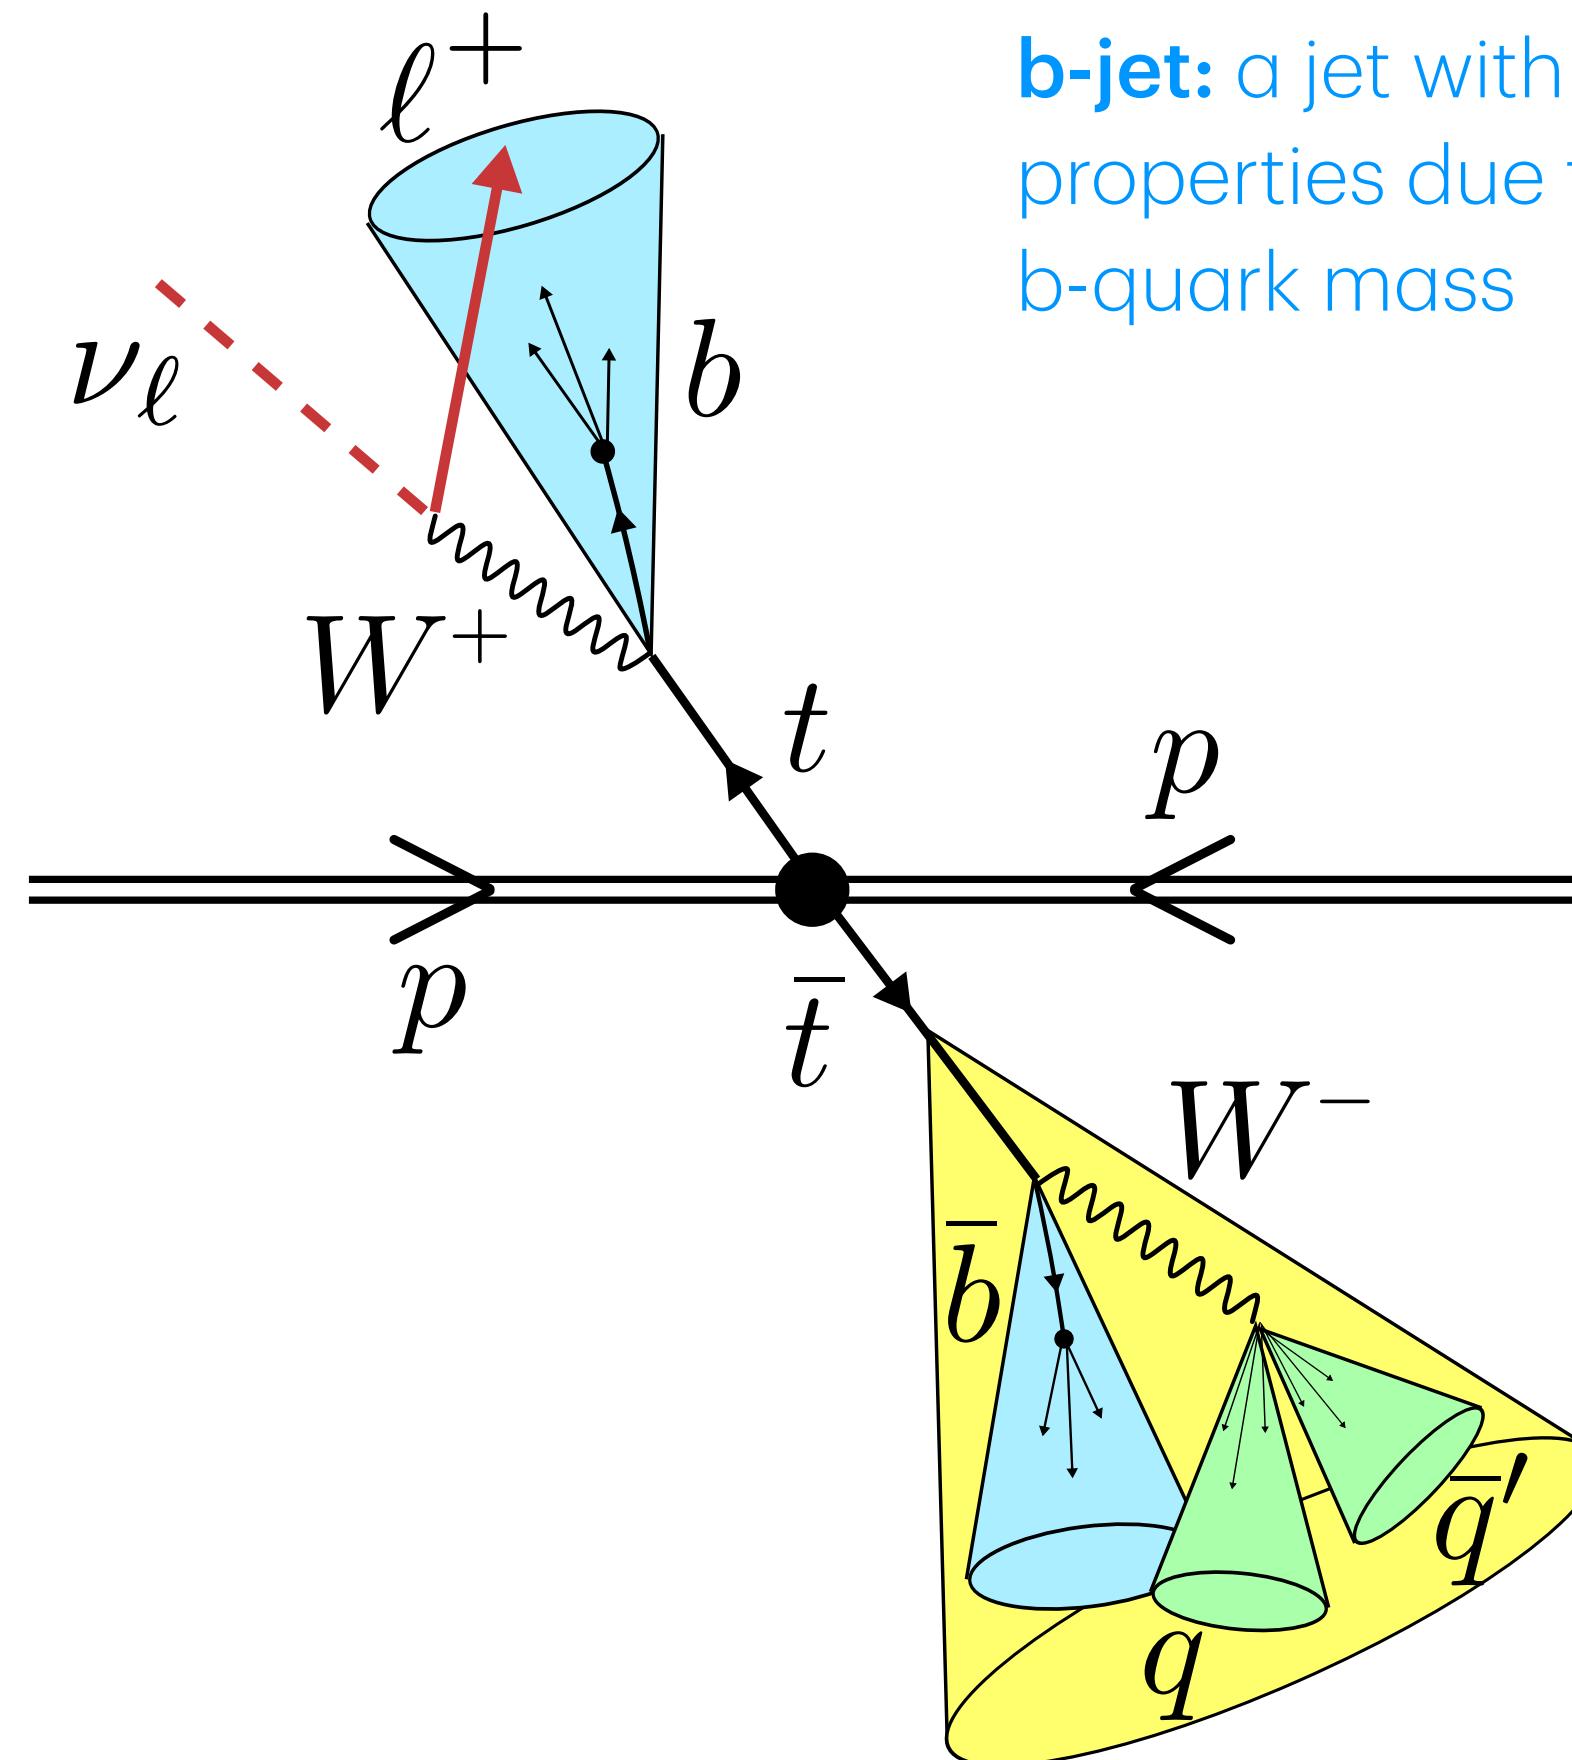
Getting the terminology right



**jets:** due to quarks  
“fragmenting” then  
“hadronising” into bound  
states (*hadrons*)

# Dissecting the process

## Getting the terminology right

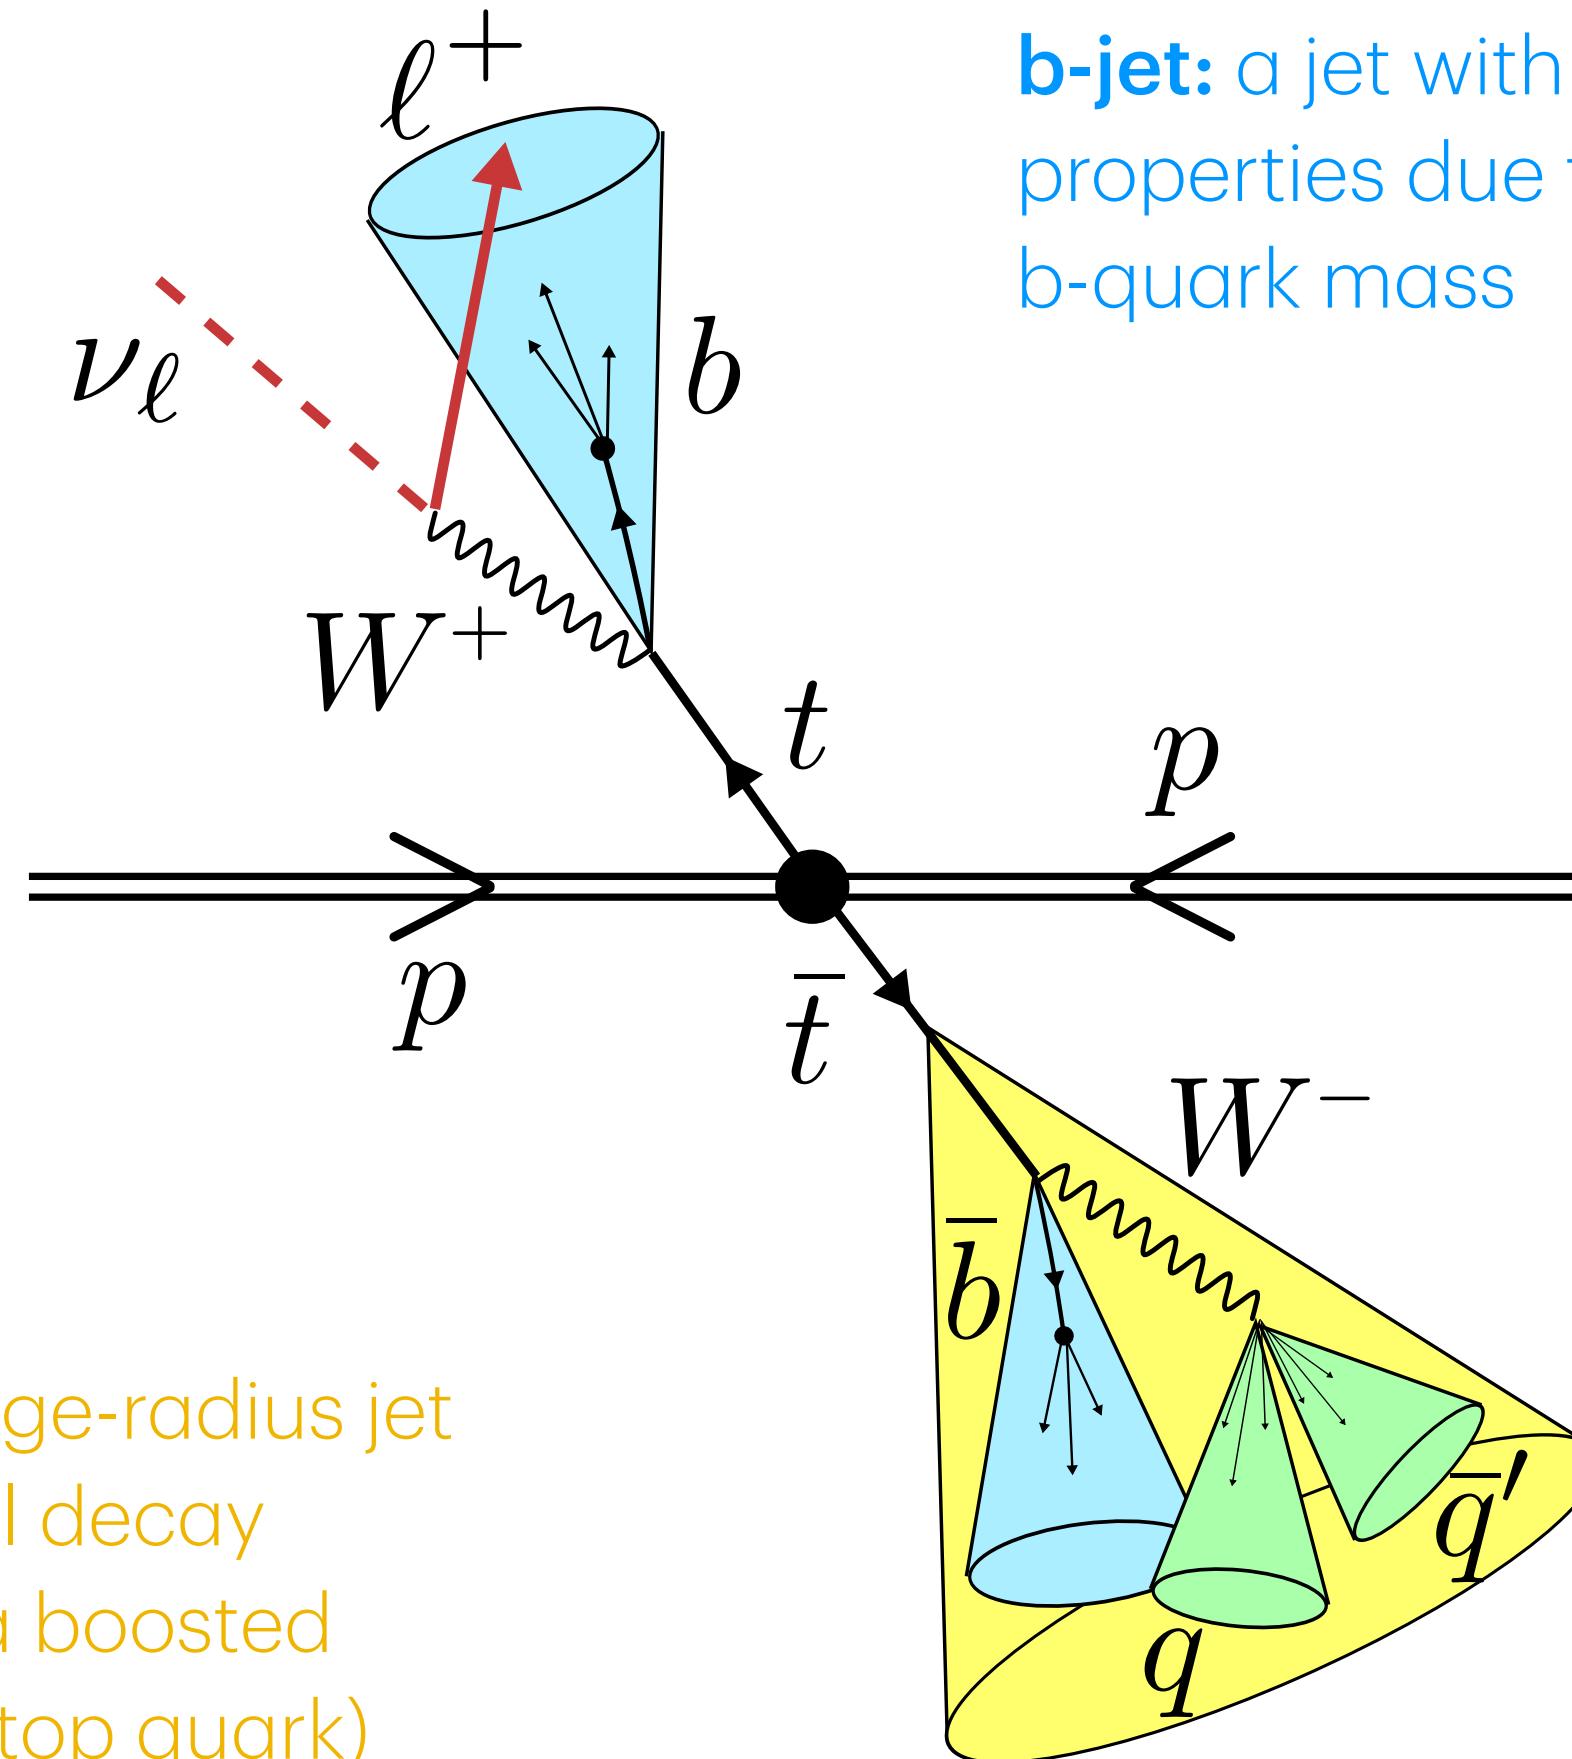


**b-jet:** a jet with special properties due to the heavy b-quark mass

**jets:** due to quarks "fragmenting" then "hadronising" into bound states (hadrons)

# Dissecting the process

## Getting the terminology right



**b-jet:** a jet with special properties due to the heavy b-quark mass

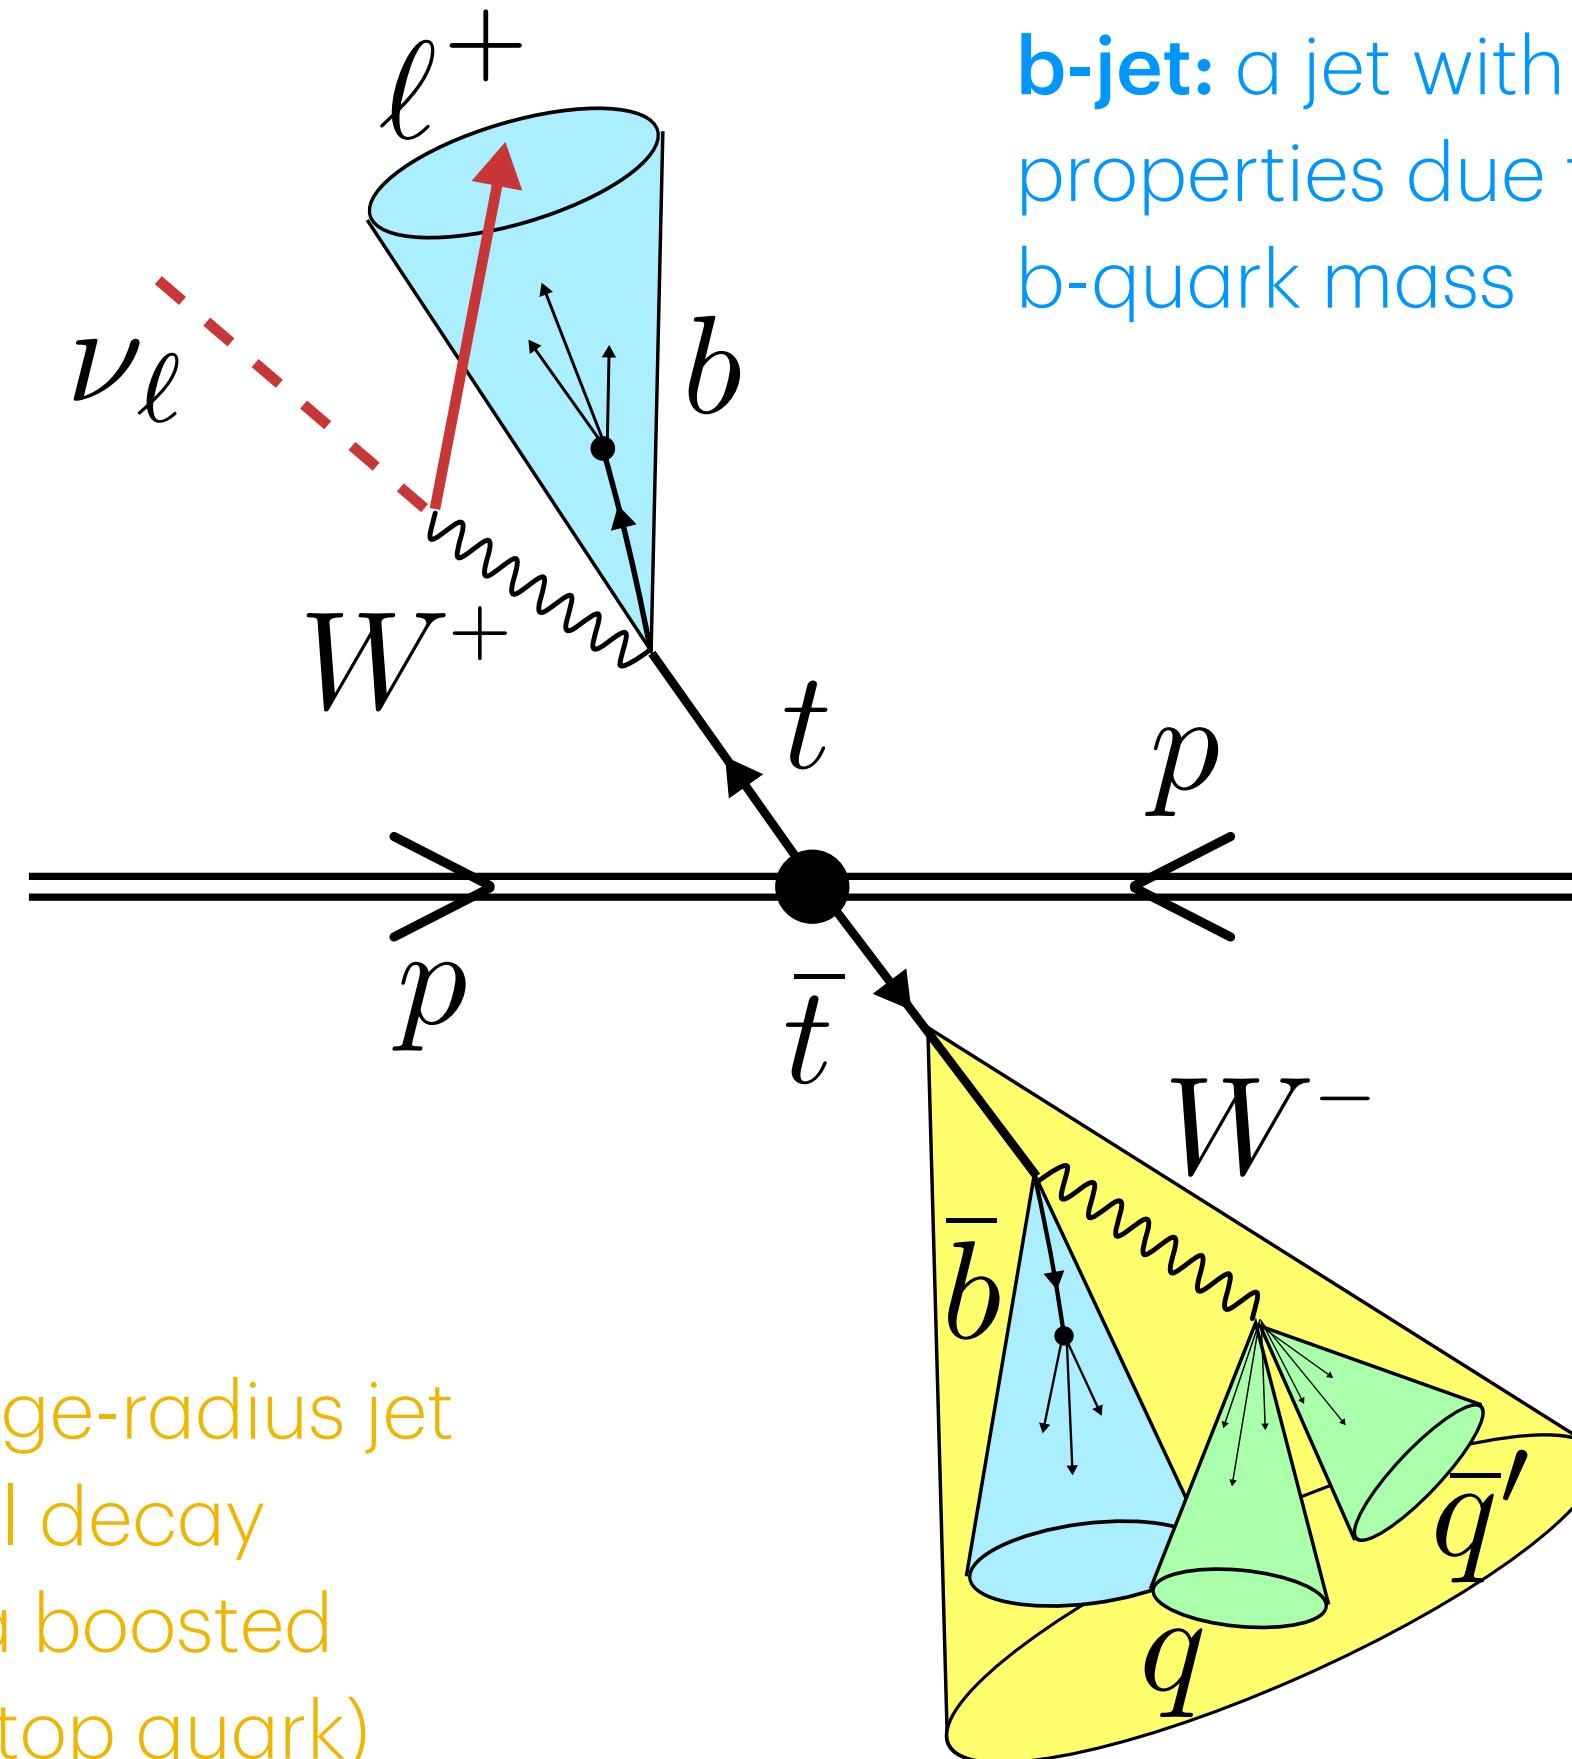
**fat-jets:** a large-radius jet containing all decay products of a boosted particle (e.g. top quark)

**jets:** due to quarks "fragmenting" then "hadronising" into bound states (hadrons)

# Dissecting the process

## Getting the terminology right

**neutrinos:** undetectable in CMS, their presence inferred via momentum conservation as missing energy (MET)



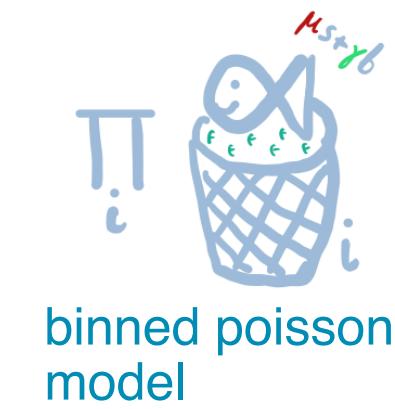
**b-jet:** a jet with special properties due to the heavy b-quark mass

**fat-jets:** a large-radius jet containing all decay products of a boosted particle (e.g. top quark)

**jets:** due to quarks "fragmenting" then "hadronising" into bound states (hadrons)

# The simplified analysis

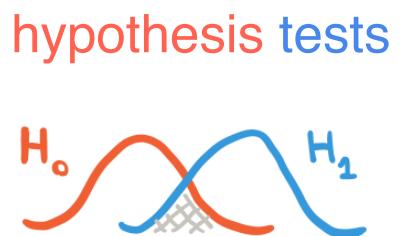
Based on [CMS open-data workshop \[2024\]](#)



$\prod_i$   
binned poisson  
model

$$\frac{\mathcal{L}(\mu, \hat{\theta})}{\mathcal{L}(\hat{\mu}, \hat{\theta})}$$

profile likelihood



hypothesis tests

Likelihood computed from observed data and parameterised model

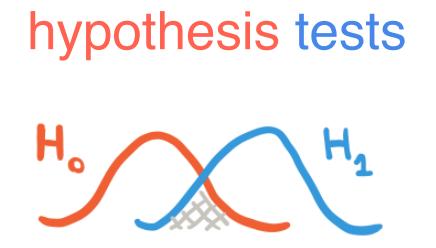
$$\prod_{i \in N_{ch}} \text{Pois}(n_i \mid \nu_i(\vec{k}, \vec{\theta}))$$

# The simplified analysis

Based on [CMS open-data workshop \[2024\]](#)



$$\frac{\mathcal{L}(\mu, \hat{\theta})}{\mathcal{L}(\hat{\mu}, \hat{\theta})}$$



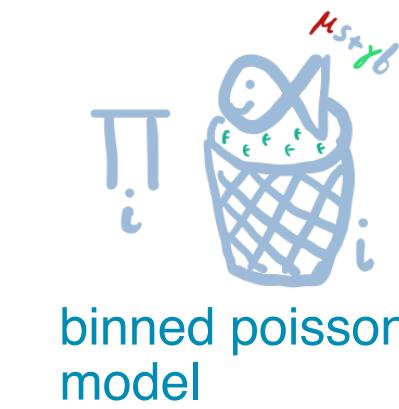
Likelihood computed from observed data and parameterised model

$$\prod_{i \in N_{ch}} \text{Pois}(n_i | \nu_i(\vec{k}, \vec{\theta}))$$

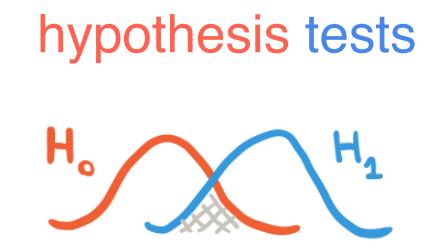
Observed data

# The simplified analysis

Based on [CMS open-data workshop \[2024\]](#)



$$\frac{\mathcal{L}(\mu, \hat{\theta})}{\mathcal{L}(\hat{\mu}, \hat{\theta})}$$



Likelihood computed from observed data and parameterised model

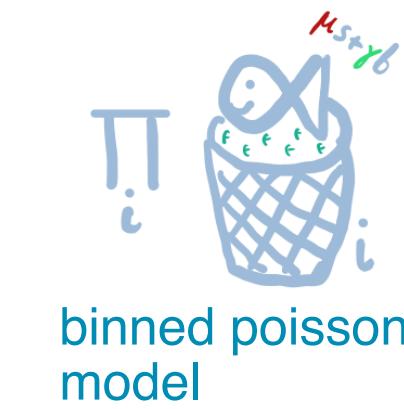
$$\prod_{i \in N_{ch}} \text{Pois}(n_i | \nu_i(\vec{k}, \vec{\theta}))$$

Observed data

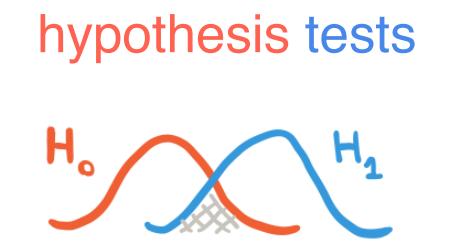
Prediction which depends  
on parameters of interest  $\vec{k}$   
and nuisance parameters  $\vec{\theta}$

# The simplified analysis

Based on [CMS open-data workshop \[2024\]](#)



$$\frac{\mathcal{L}(\mu, \hat{\theta})}{\mathcal{L}(\hat{\mu}, \hat{\theta})}$$



Likelihood computed from observed data and parameterised model

$$\prod_{i \in N_{ch}} \text{Pois}(n_i | \nu_i(\vec{k}, \vec{\theta}))$$

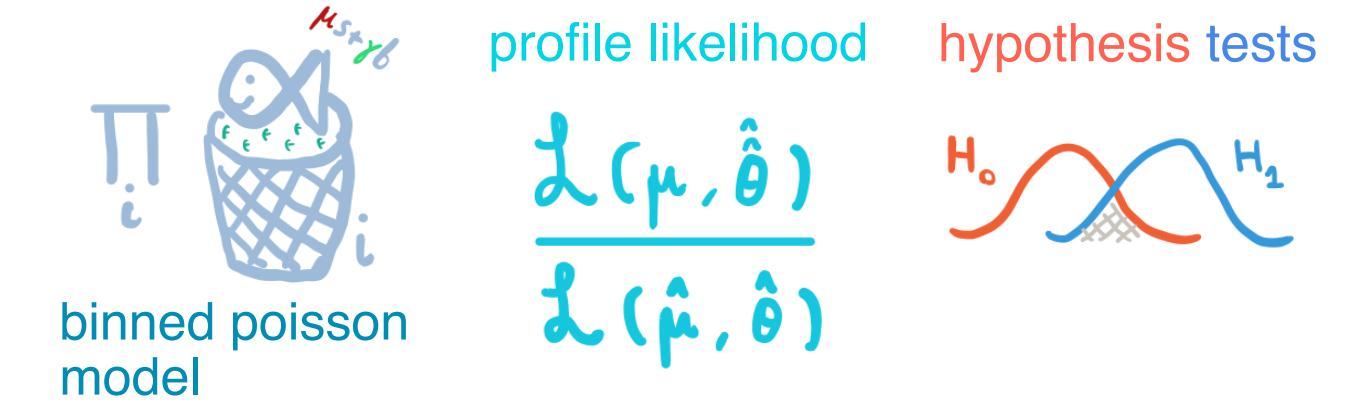
Observed data

Prediction which depends  
on parameters of interest  $\vec{k}$   
and nuisance parameters  $\vec{\theta}$

The signal prediction is  
parameterised via a  
signal strength  $\mu$

# The simplified analysis

Based on [CMS open-data workshop \[2024\]](#)



Perform a binned profile likelihood fit to extract signal significance [ $p$ -value]

$$p(\vec{n}, \vec{a} | \vec{k}, \vec{\theta}) = \prod_{i \in N_{ch}} \text{Pois}(n_i | \nu_i(\vec{k}, \vec{\theta})) \cdot \prod_{j \in N_{\theta}} c_j(a_j | \theta_j)$$

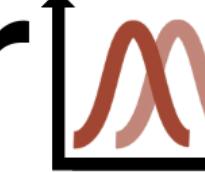


Constraint terms for  
systematic uncertainties  
from auxiliary measurement

# Differentiating the analysis

Tools that did not work very well together

---

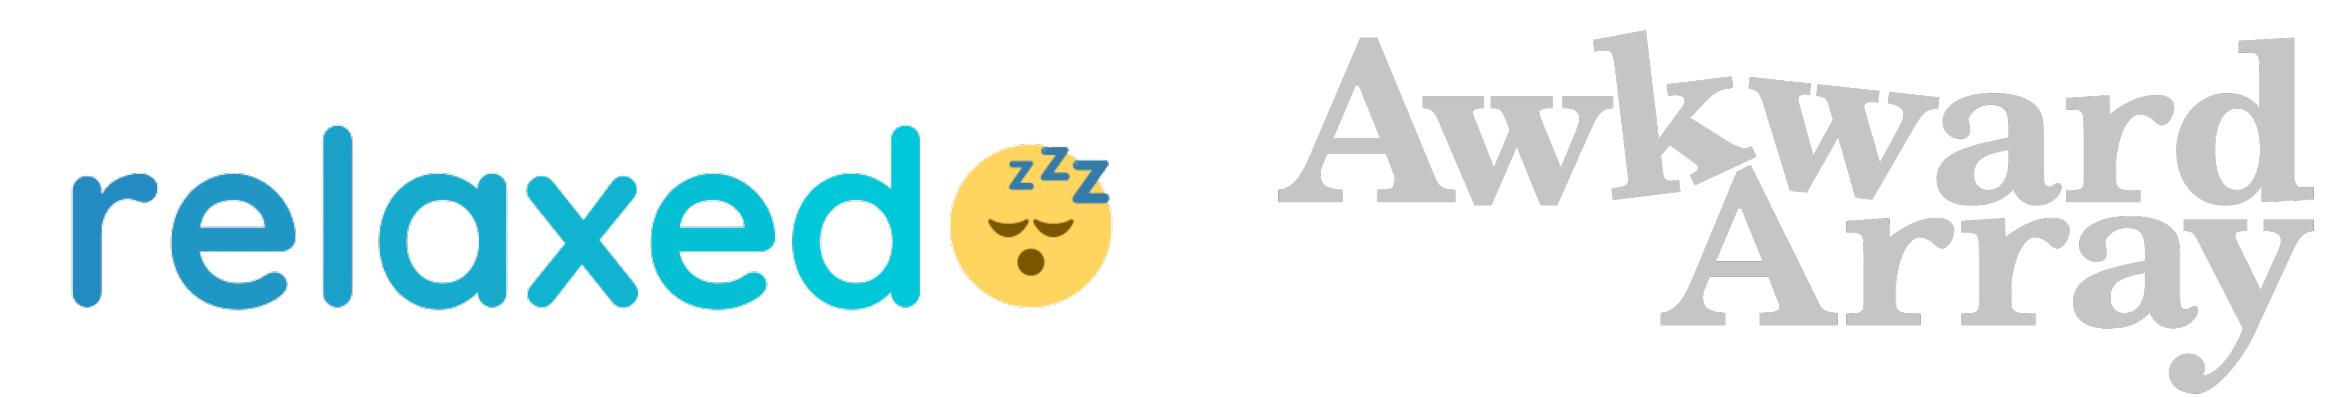
ever<sup>↑</sup> more relaxed

Evermore model is incompatible with  
what relaxed expects, and tracing  
issues arise when a model wrapper is  
attempted

# Differentiating the analysis

Tools that did not work very well together

---

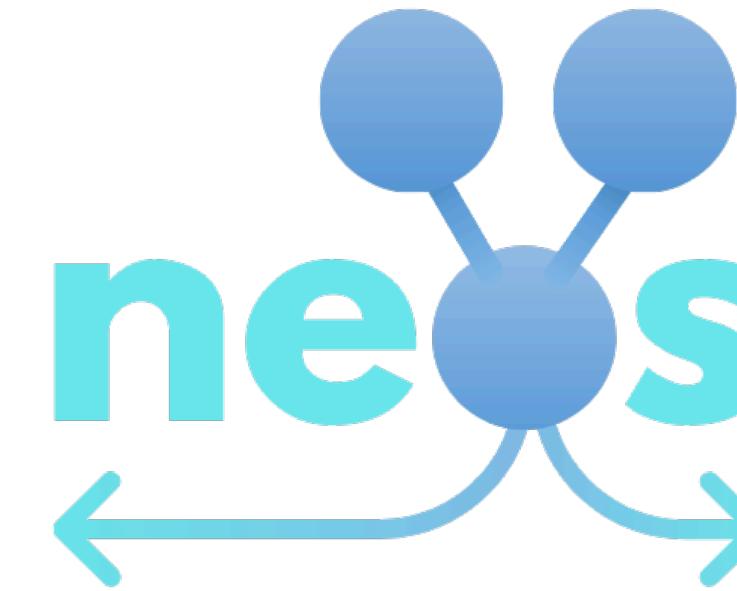


relaxed is unable to handle awkward  
arrays during data processing

# Differentiating the analysis

Tools that did not work very well together

---

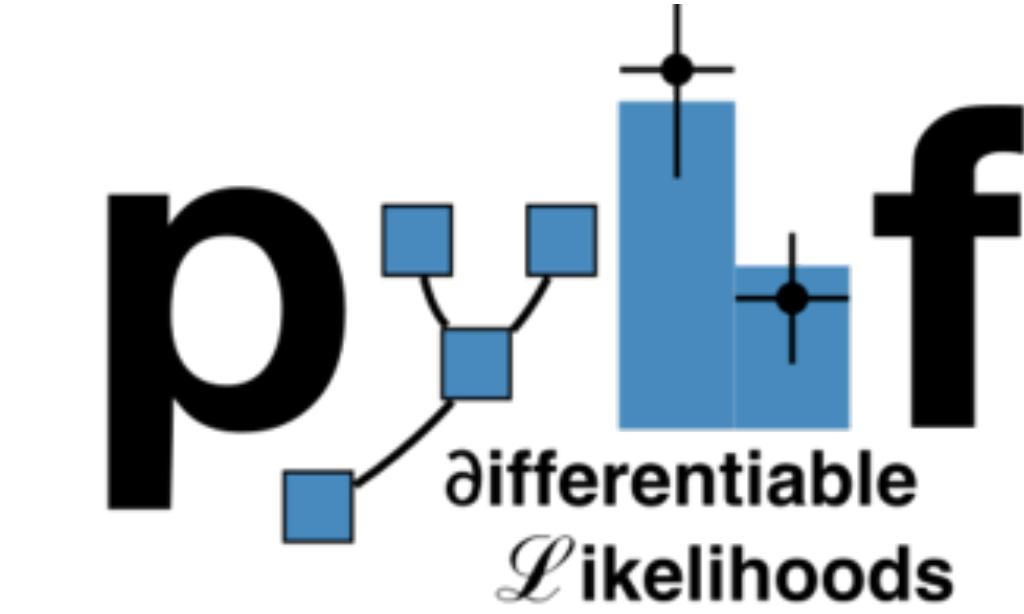
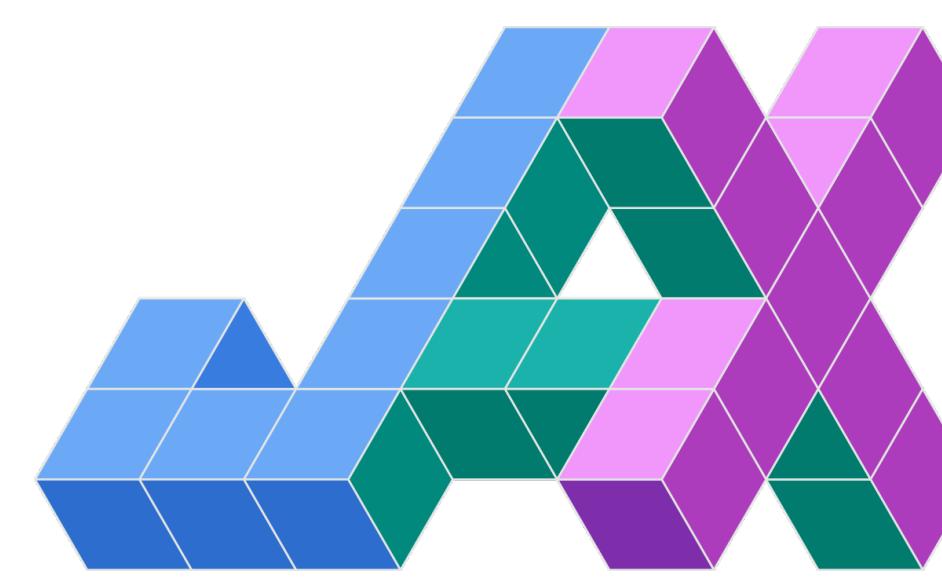


**Difficult to navigate version clashes  
between old dependencies**

# Differentiating the analysis

Tools that did not work very well together

---



JAX backend support is limited in  
PyHF — difficulties in model building

# Differentiating the analysis

Tools that did not work very well together

---



JAX backend support is limited in  
awkward array — many missing  
operations