

ORACLE

Concurrent programming: From theory to practice

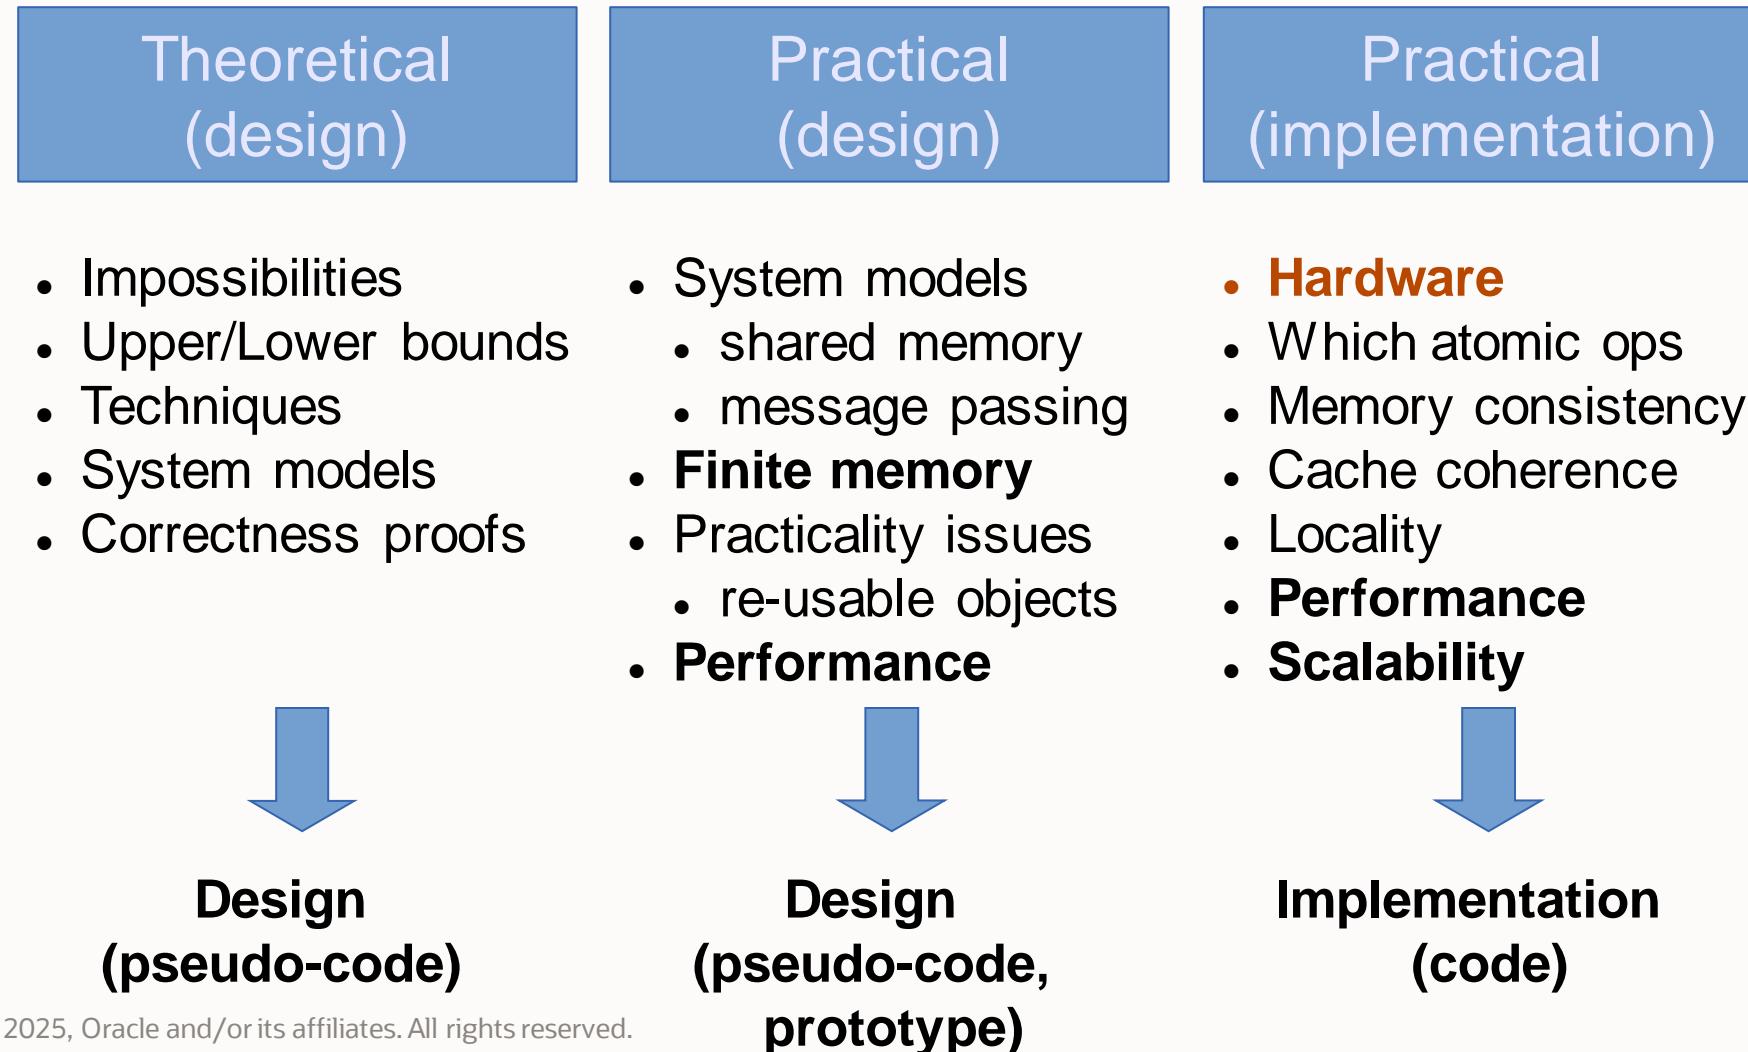
Concurrent Computing 2025

Vasileios Trigonakis

Consulting Member of Technical Staff
Oracle Zurich

08.Dec.2025

From theory to practice



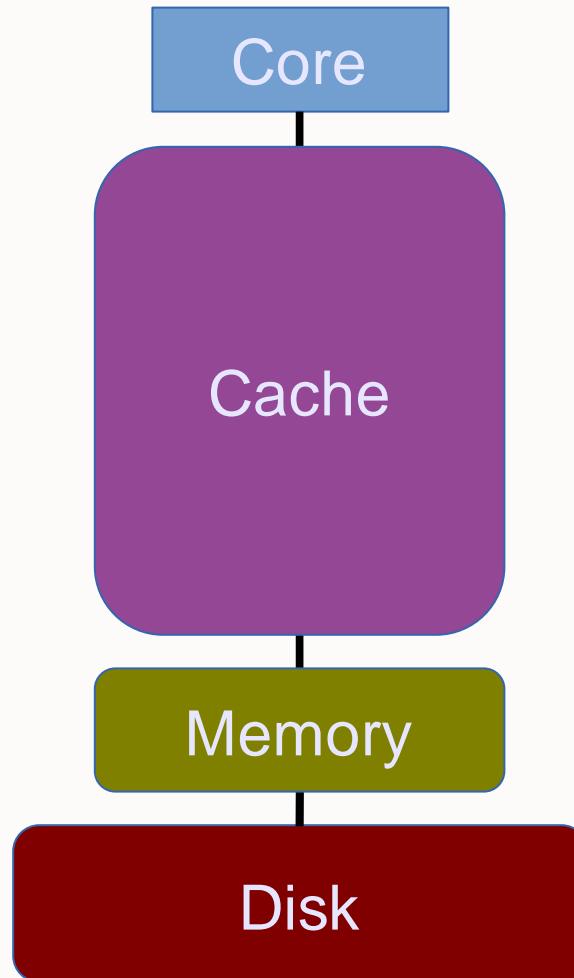
Outline

- CPU caches
- Cache coherence
- Placement of data
- Graph processing: Concurrent data structures

Outline

- **CPU caches**
- Cache coherence
- Placement of data
- Graph processing: Concurrent data structures

Why do we use caching?



Core freq: 2GHz = 0.5 ns / instr

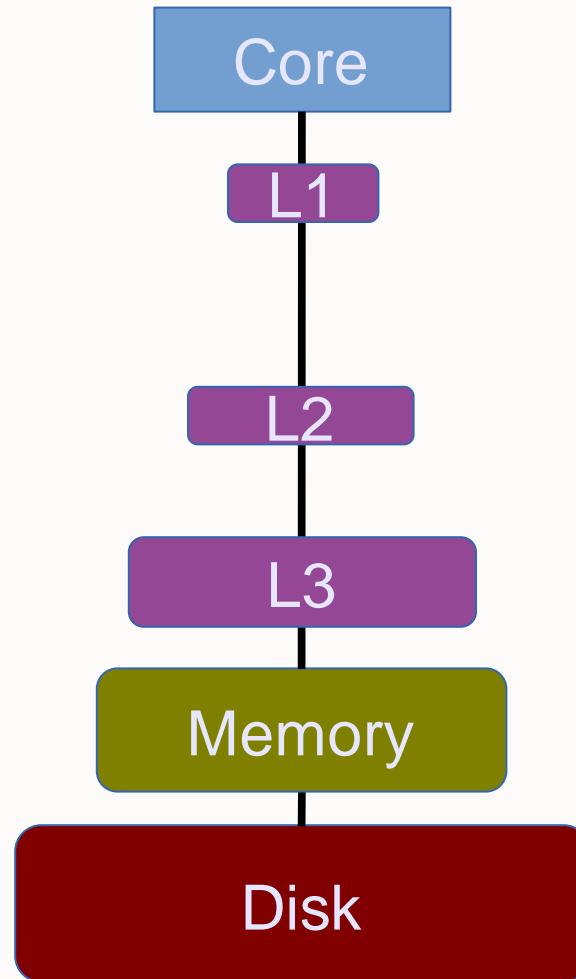
Core → Disk = ~ms

Core → Memory = ~100ns

Cache

- Large = slow
- Medium = medium
- Small = fast

Why do we use caching?



Core freq: 2GHz = 0.5 ns / instr

Core → Disk = ~ms

Core → Memory = ~100ns

Cache

- Core → L3 = ~20ns
- Core → L2 = ~7ns
- Core → L1 = ~1ns

From typical server configurations a few years back to the ERA of Gen AI

Intel® Xeon®

- 14 cores @ 2.4GHz
- L1: 32KB
- L2: 256KB
- L3: 40MB
- Memory: 512GB

AMD Opteron™

- 18 cores @ 2.4GHz
- L1: 64KB
- L2: 512KB
- L3: 20MB
- Memory: 512GB

Intel® Xeon® 6 Processors with P(erformance)-Cores

> 70 cores, > 400MB L3

&

Intel® Xeon® 6 Processors with E(nergy)-Cores

> 60 cores, > 90MB L3

<https://www.intel.com/content/www/us/en/products/details/processors/xeon.html>

AMD EPYC™ 9005 Series

Max config:
192 cores, 384MB L3

&

AMD EPYC™ 9004, 8004, 7003, 4004 Series

<https://www.amd.com/en/products/processors/server/epyc.html>

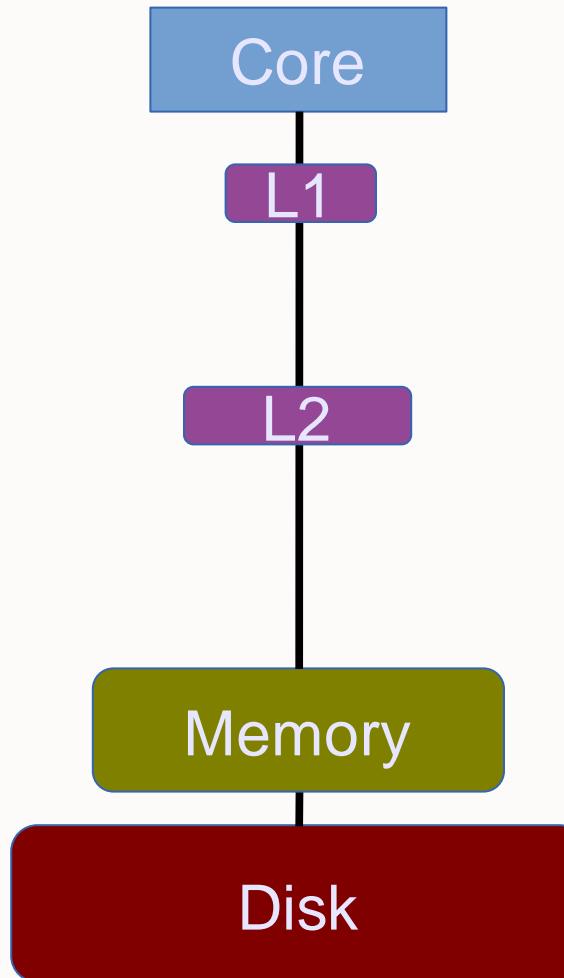
Experiment

Throughput of accessing some memory,
depending on the memory size

Outline

- CPU caches
- **Cache coherence**
- Placement of data
- Graph processing: Concurrent data structures

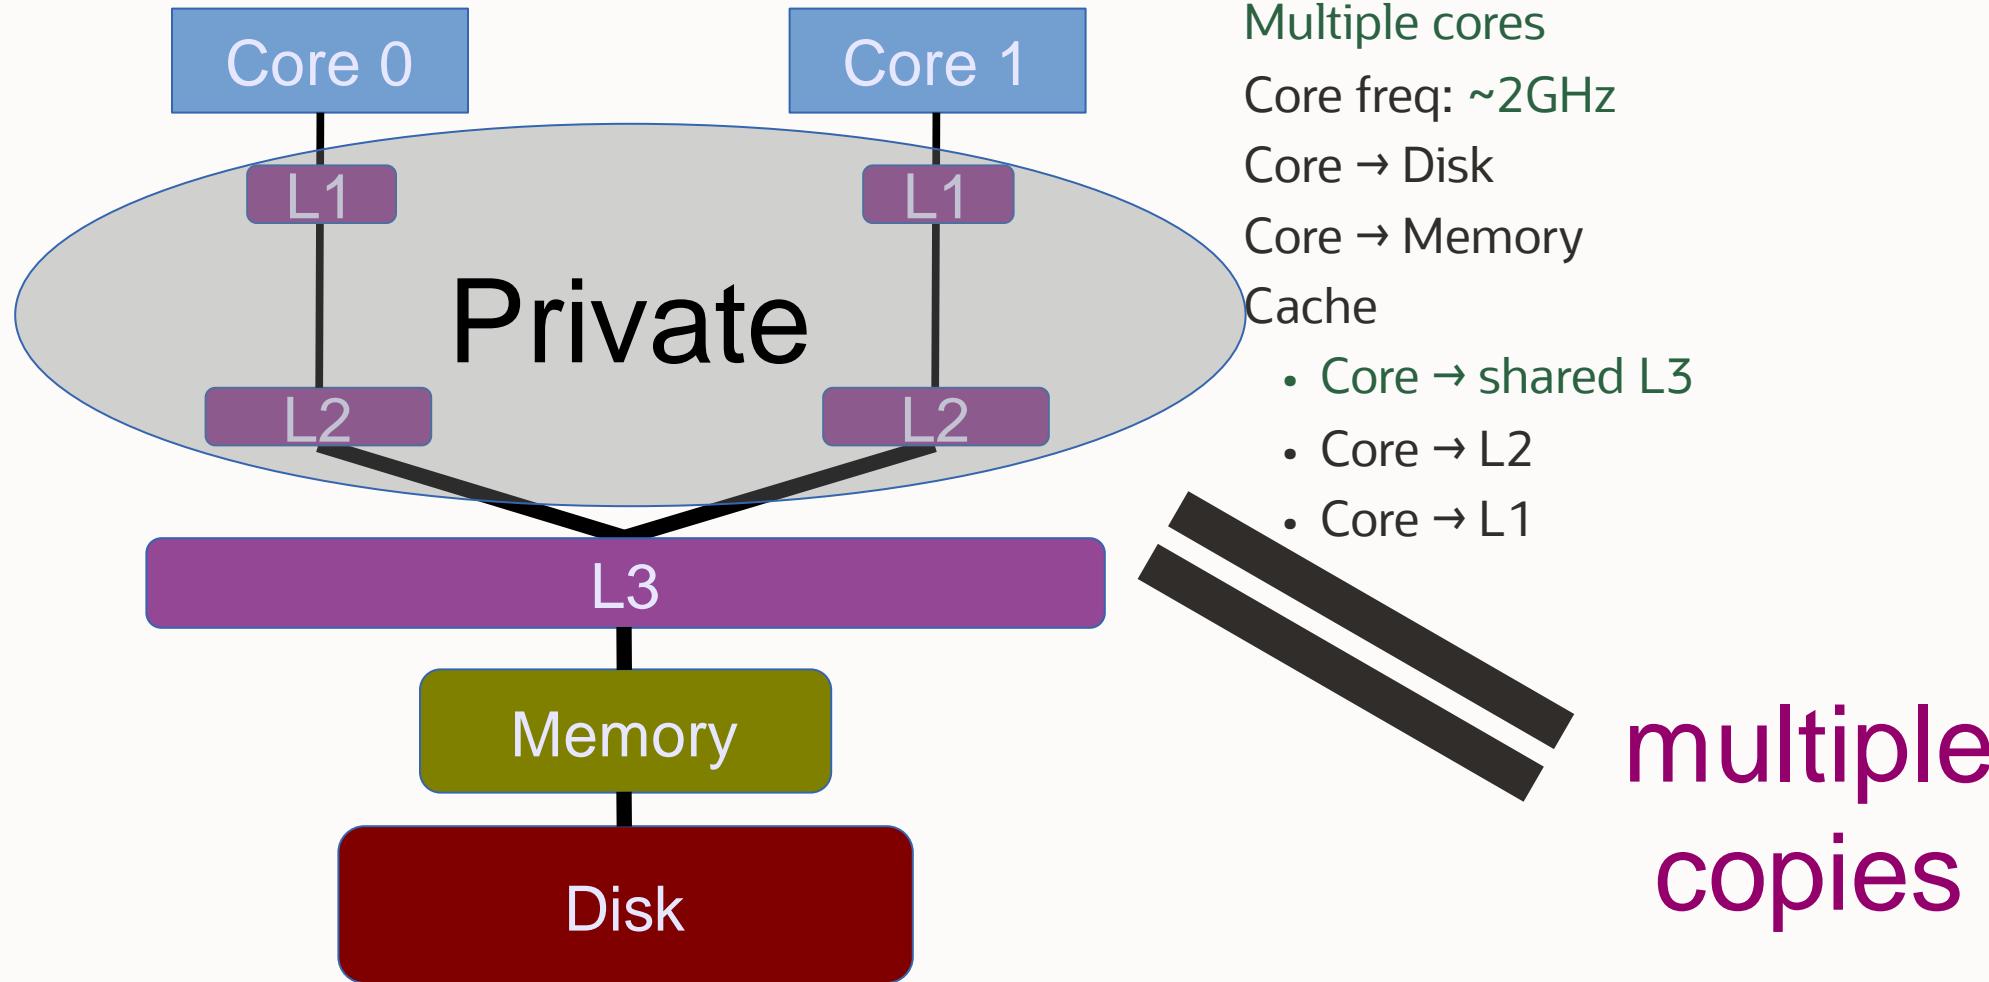
Until ~2004: single-cores



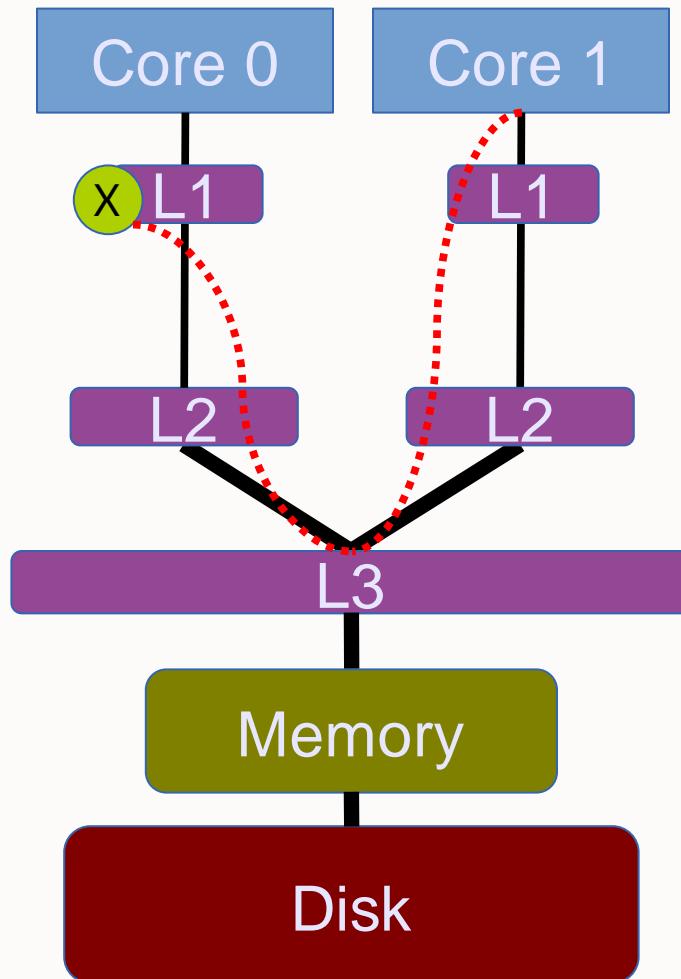
Single core
Core freq: 3+GHz
Core → Disk
Core → Memory
Cache

- Core → L2
- Core → L1

After ~2004: multi-cores



Cache coherence for consistency



Core 0 has X and Core 1

- wants to write on X
- wants to read X
- did Core 0 write or read X?

To perform a **write**

- invalidate all readers, or
- previous writer

To perform a **read**

- find the latest copy

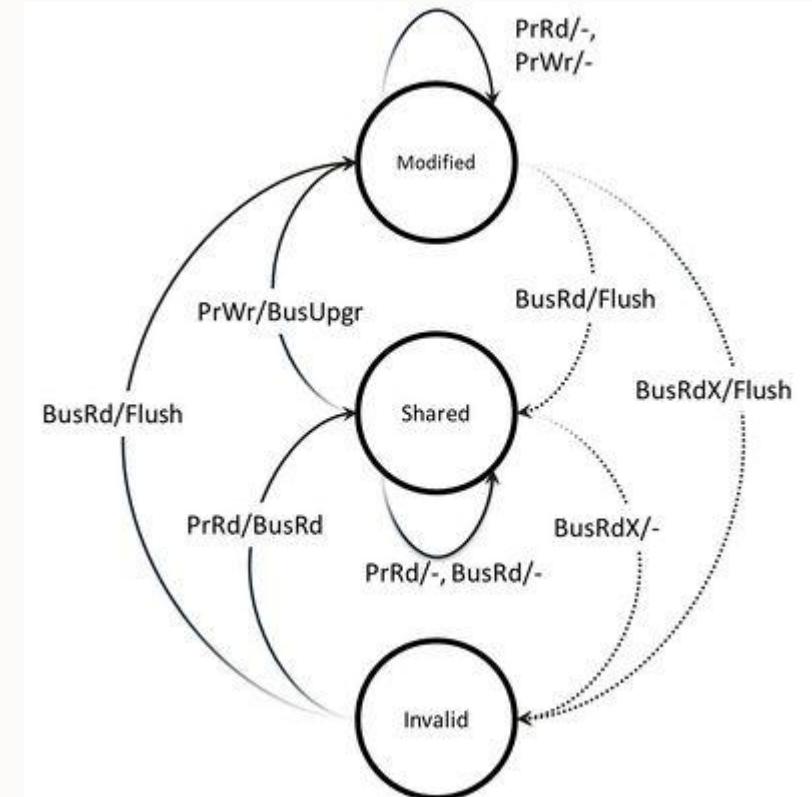
Cache coherence with MESI

A state diagram

State (per cache line)

- **Modified:** the only dirty copy
- **Exclusive:** the only clean copy
- **Shared:** a clean copy
- **Invalid:** useless data

Which state is our “favorite?”



The ultimate goal for scalability

A state diagram

State (per cache line)

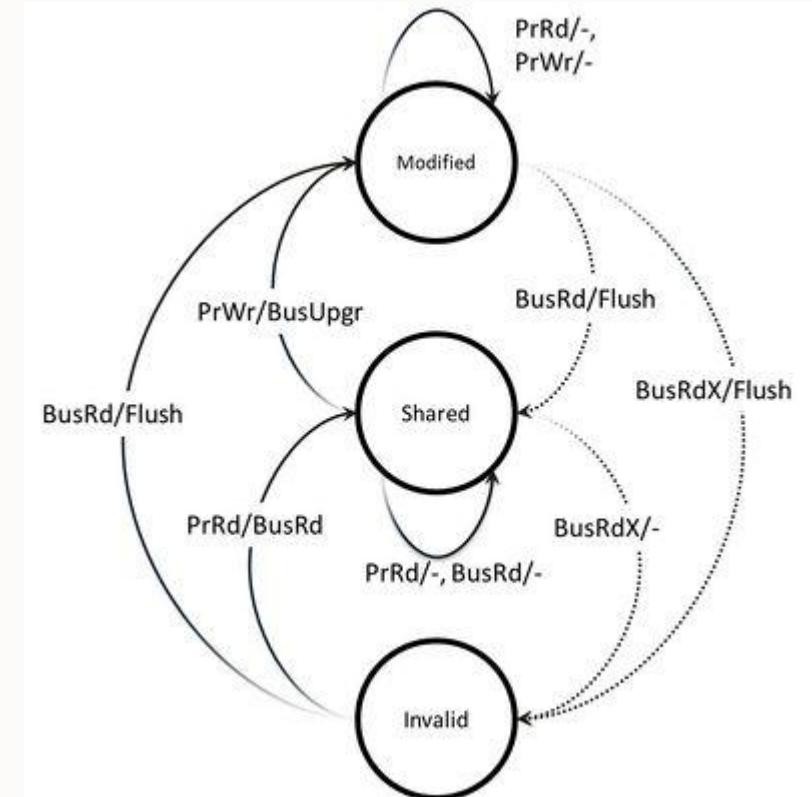
- Modified: the only dirty copy
- Exclusive: the only clean copy

• **Shared**: a clean copy

- Invalid: useless data

= threads can keep the data close (L1 cache)

= faster



—

Experiment

The effects of false sharing

Outline

- CPU caches
- Cache coherence
- **Placement of data**
- Graph processing: Concurrent data structures

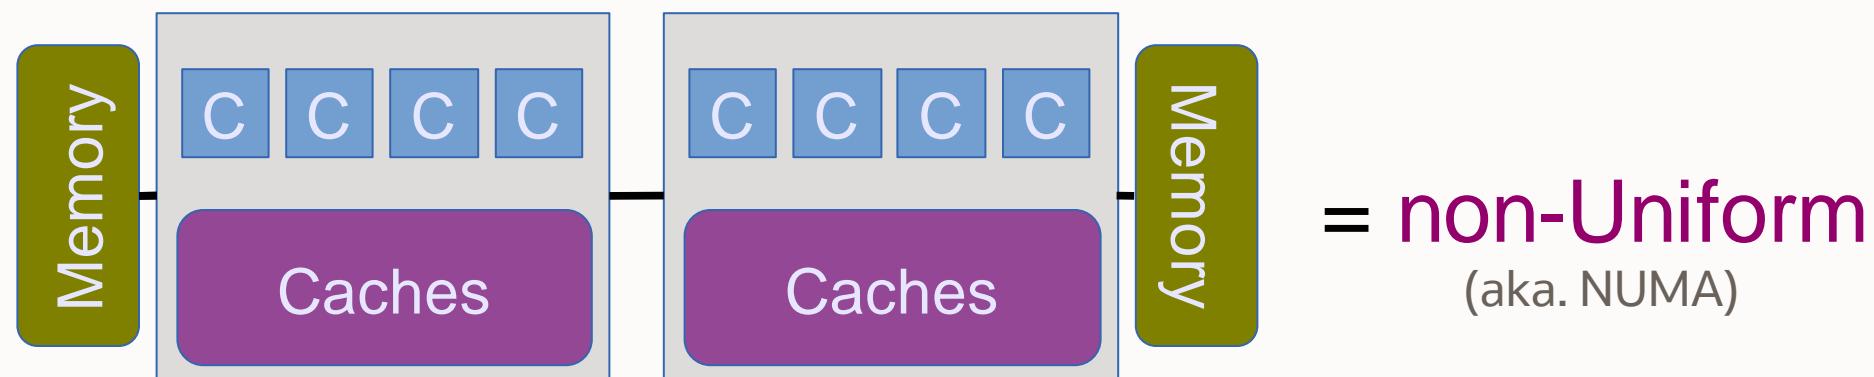
Uniformity vs. non-uniformity

Typical desktop machine



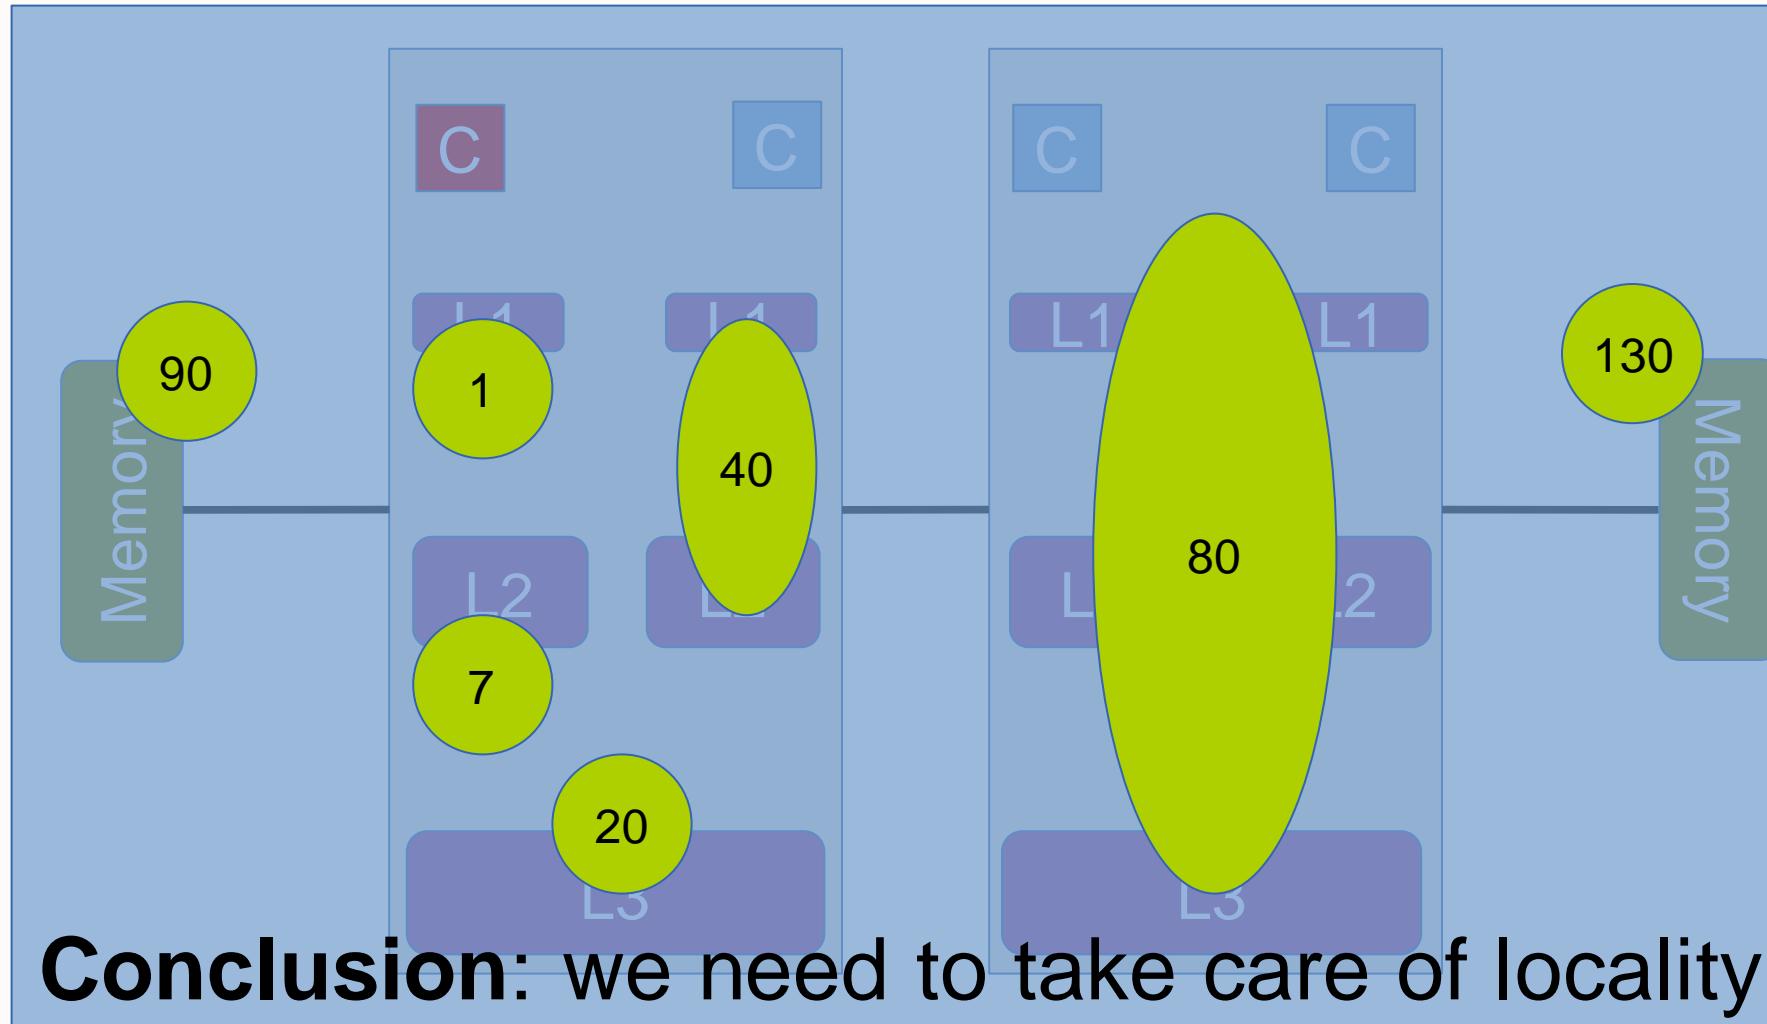
= Uniform

Typical server machine



= non-Uniform
(aka. NUMA)

Latency (ns) to access data in a NUMA multi-core server



—

Experiment

The effects of locality

Experiment

The effects of locality

```
vtrigona $ ./test_locality -x0 -y1
Size:          8 counters = 1 cache lines
Thread 0 on core : 0
Thread 1 on core : 2
Number of threads: 2
Throughput       : 104.27 Mop/s
```

Same memory node

```
vtrigona $ ./test_locality -x0 -y10
Size:          8 counters = 1 cache lines
Thread 0 on core : 0
Thread 1 on core : 10
Number of threads: 2
Throughput       : 43.16 Mop/s
```

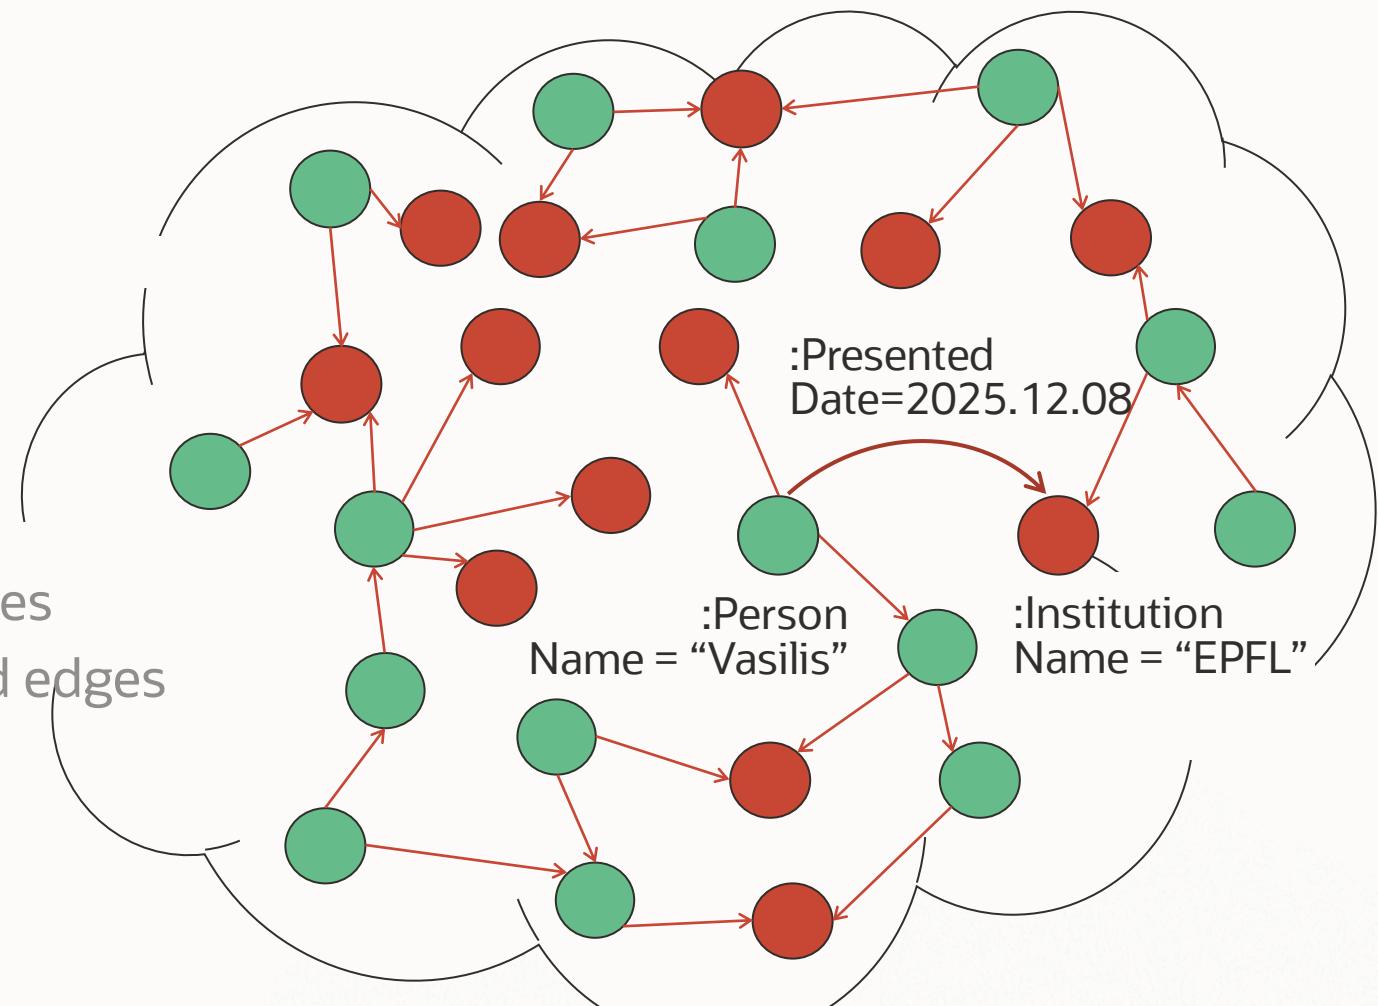
Different memory nodes

Outline

- CPU caches
- Cache coherence
- Placement of data
- **Graph processing: Concurrent data structures**

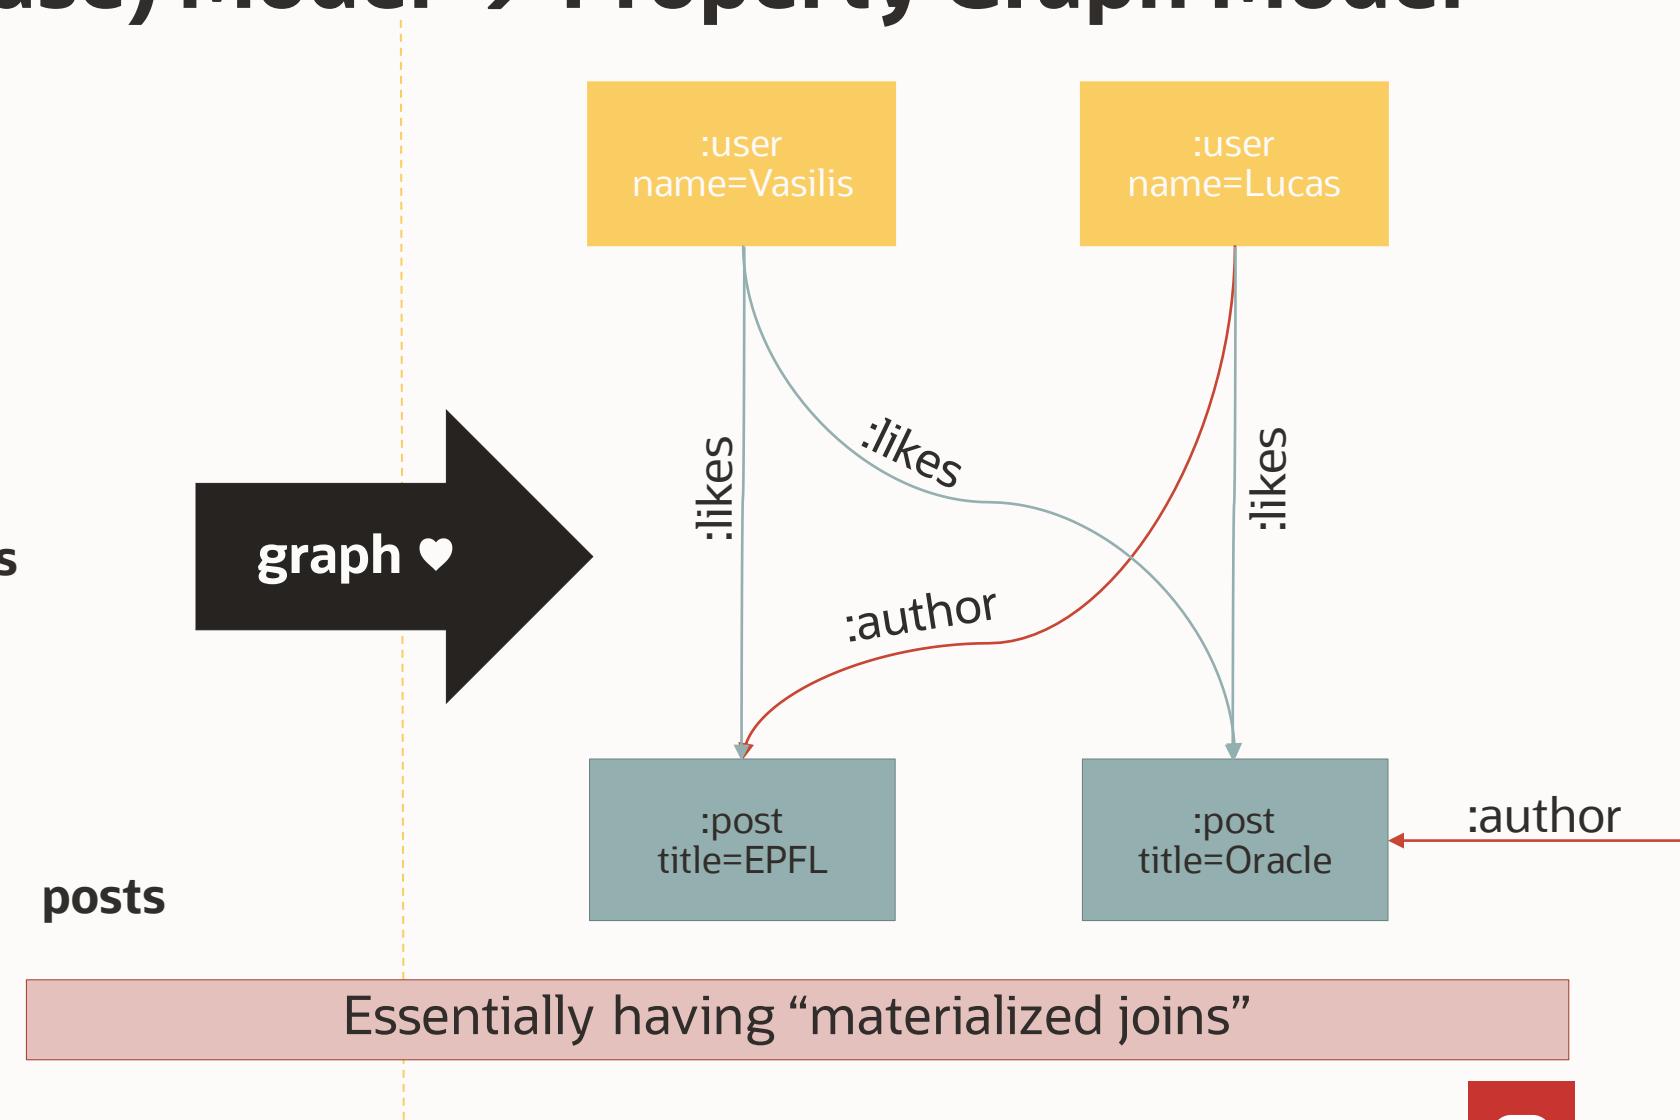
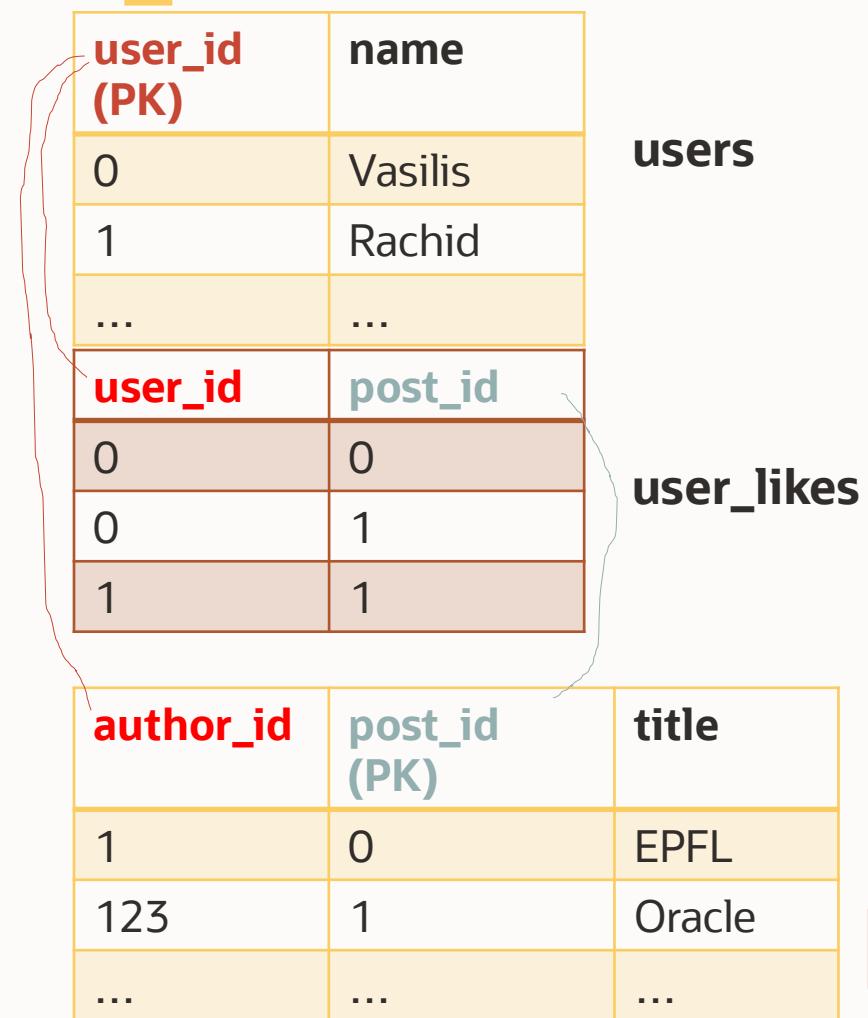
Your Data is a Graph!

- Represent it as a **property graph**
 - Entities are **vertices**
 - Relationships are **edges**
- Annotate your graph
 - **Labels** identify vertices and edges
 - **Properties** describe vertices and edges
- For the purpose of
 - Data modeling
 - Data analysis



Navigate multi-hop relationships quickly (instead of joins)

Relational (Database) Model → Property Graph Model



Main Approaches of Graph Processing

1. Computational graph analytics [ASPO'S12, VLDB'16]

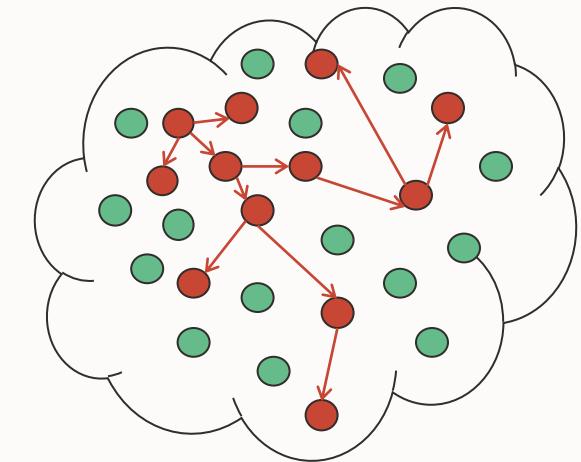
- Iterate the graph multiple times and compute mathematical properties using **Greenmarl / PGX Algorithm** (e.g., Pagerank)
- e.g., `graph.getVertices().foreach (n -> ...)`

2. Graph querying and pattern matching [GRADES'16/23, VLDB'16, Middleware Ind. 23]

- Query the graph using **PGQL** or **SQL/PGQ** to find sub-graphs that match to the given relationship pattern
- e.g., `SELECT ... MATCH (a) -[edge]-> (b) ...`

3. Graph ML

- Use the structural information latent in graphs
- e.g., graph similarity



$$PR(p_i) = \frac{1-d}{N} + d \sum_{p_j \in M(p_i)} \frac{PR(p_j)}{L(p_j)}$$

4. Vector similarity graph indices

- Hierarchical navigable small world (HNSW)

5. Graph RAG

- Retrieval-Augmented Generation (RAG)
- Enhancing RAG with knowledge graphs

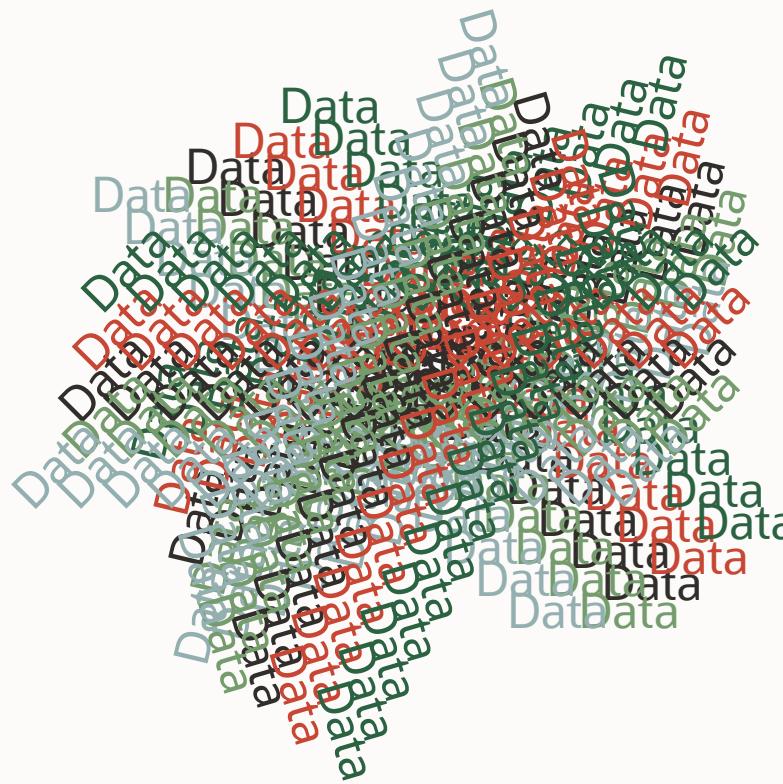
Dissecting a graph processing system

with a focus on (concurrent) data structures

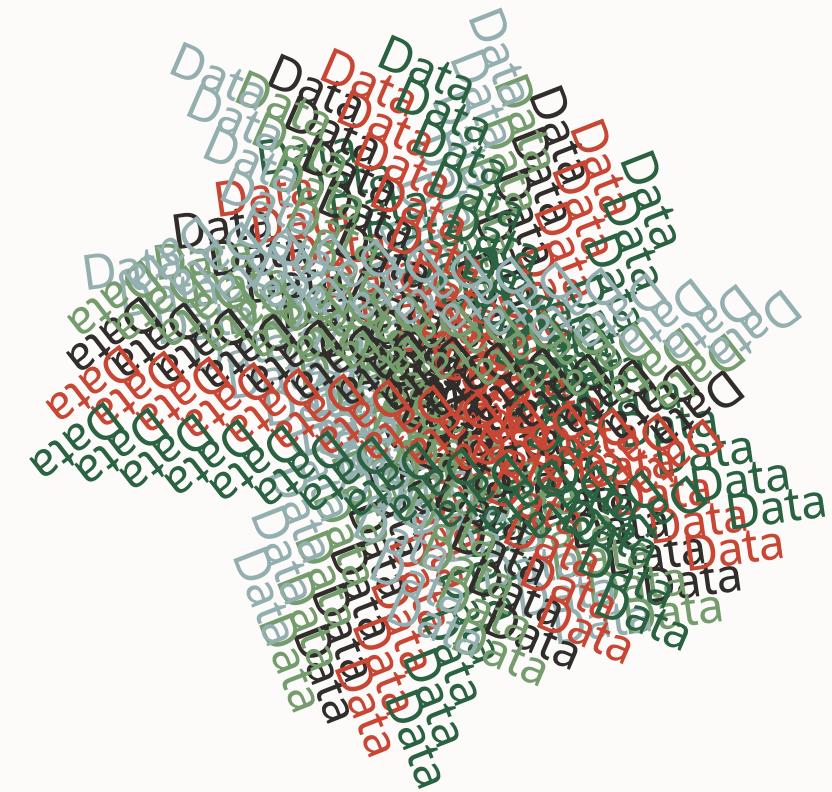
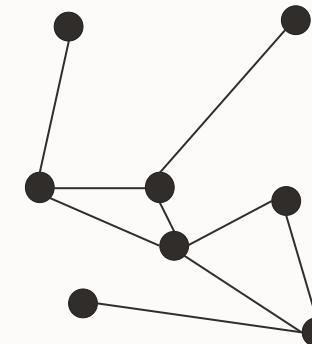
Dissecting a graph processing system and preparing for a job interview

with a focus on (concurrent) data structures

Architecture of a graph processing system



Graph



Tons of other data and metadata to store

Graph

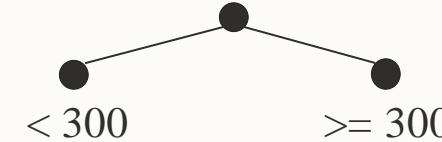
Runtime

Operations

tmp graph structure

"Vasilis", "Breaking bad", :likes
 "Rachid", "Dexter", :likes
 "Vasilis", "Dexter", :likes
 "Dexter", "Breaking bad", :similar
 "Breaking bad", "Dexter", :similar

indices / metadata



group by / join

Vasilis, Breaking bad
 Rachid, Dexter
 Vasilis, Dexter

Vasilis, 2
Rachid, 1

graph structure



buffer management



distinct

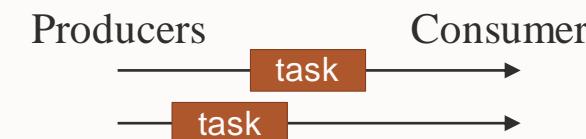
Vasilis
 Rachid
 Vasilis

Vasilis
Rachid

user-ids - internal ids

Vasilis → 0	0 → Vasilis
Rachid → 1	1 → Rachid
Breaking bad → 2	2 → Breaking bad
Dexter → 3	3 → Dexter

task / job scheduling



limit (top k)

11 12 0 9 8 13
 8 9 11 23 32 9
 1 2 3 5 7 3 2 0

32
23
13

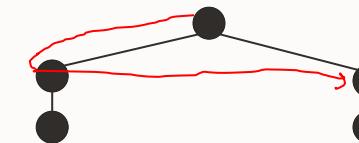
labels

:likes, :people, :similar, ...

labels

:likes, :people, :similar, :male ...

BFS



properties

"Vasilis", {people, male}, 20, Zurich
 "Rachid", {people, male}, ??, Lausanne

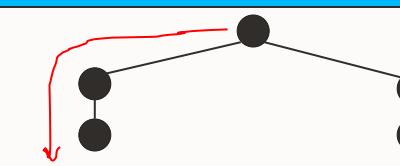
lifetime management

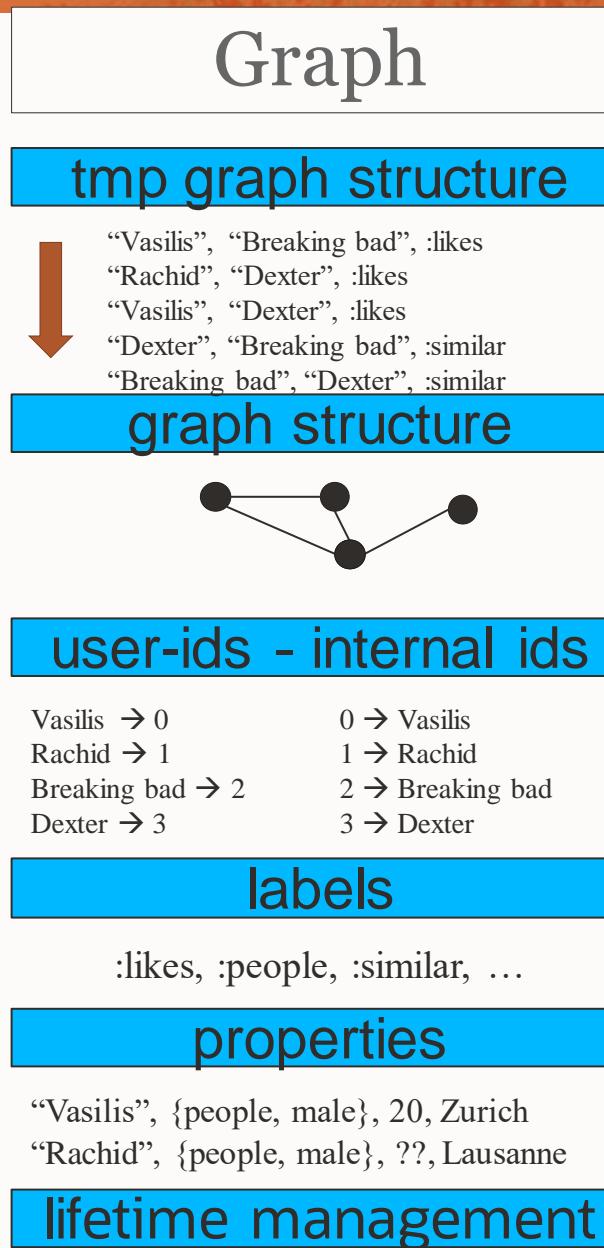
number_of_references: X

renaming (ids)



DFS





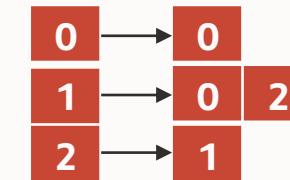
- tmp graph structure
 - append only
 - dynamic schema

→ **dataframe** = segmented buffer

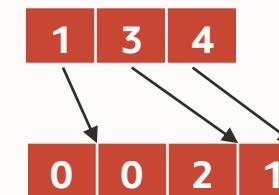
- Classic graph structures
 1. adjacency matrix

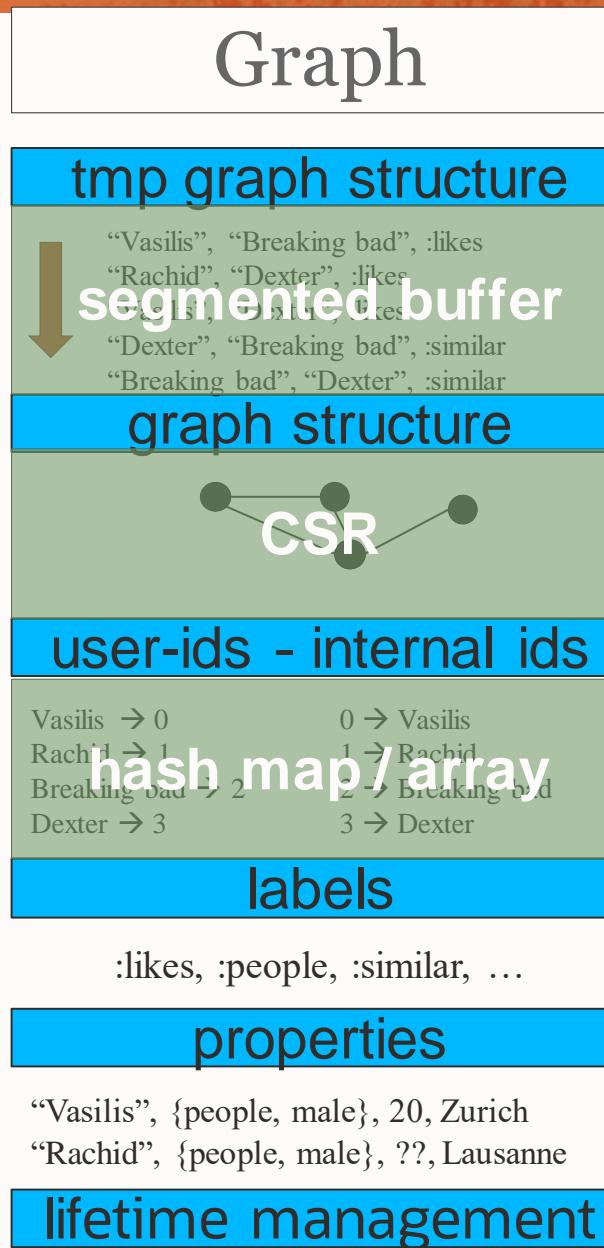
	0	1	2
0	x		
1	x		x
2		x	

2. adjacency list

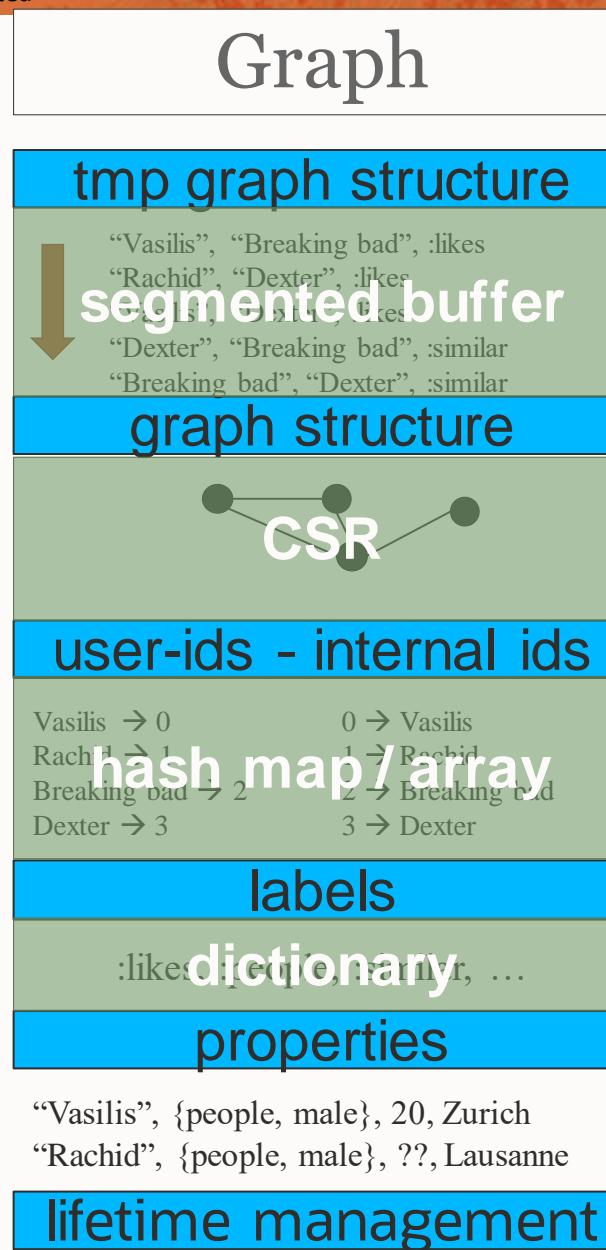


3. compressed source row (CSR)





- **Storing labels**
 - usually a small enumeration
e.g., person, female, male
 - storing strings is expensive
“person” → ~ 7 bytes
 - comparing strings is expensive
- **dictionary encoding**, e.g.,
 - person → 0
 - female → 1
 - male → 2
- Ofc, **hash map** to
 - store those
 - translate during runtime



- **Property**
 - one type per property, e.g., int
 - 1:1 mapping with vertices/edges

→ **(sequential) arrays**

- **Lifetime management (and other counters)**
 - cache coherence: atomic counters can be expensive
 - Two potential solutions
 - approximate counters**
 - stripped counters**

Thread local: counter[0] counter[1] counter[2]

```

increment(int by) { counter[my_thread_id] += by; }
int value() {
    int sum = 0;
    for (int i = 0; i < num_threads; i++) { sum += counter[i]; }
    return sum;
}

```

Graph

tmp graph structure

↓
 "Vasilis", "Breaking bad", :likes
 "Rachid", "Dexter", :likes
 "Vasilis", "Dexter", :likes
 "Dexter", "Breaking bad", :similar
 "Breaking bad", "Dexter", :similar

graph structure



user-ids - internal ids

Vasilis → 0	0 → Vasilis
Rachid → 1	1 → Rachid
Breaking bad → 2	2 → Breaking bad
Dexter → 3	3 → Dexter

hash map / array

labels

dictionary (=map)

properties

"Vasilis", {people, male}, 20, Zurich
"Rachid", {people, male}, ??, Lausanne

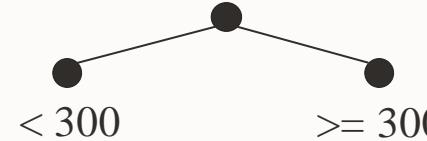
lifetime management

stripped counter

Score

Structure	# Usages
array / buffer	5
map	2

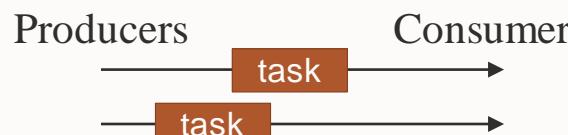
Runtime



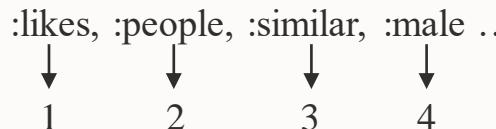
buffer management



task / job scheduling



labels



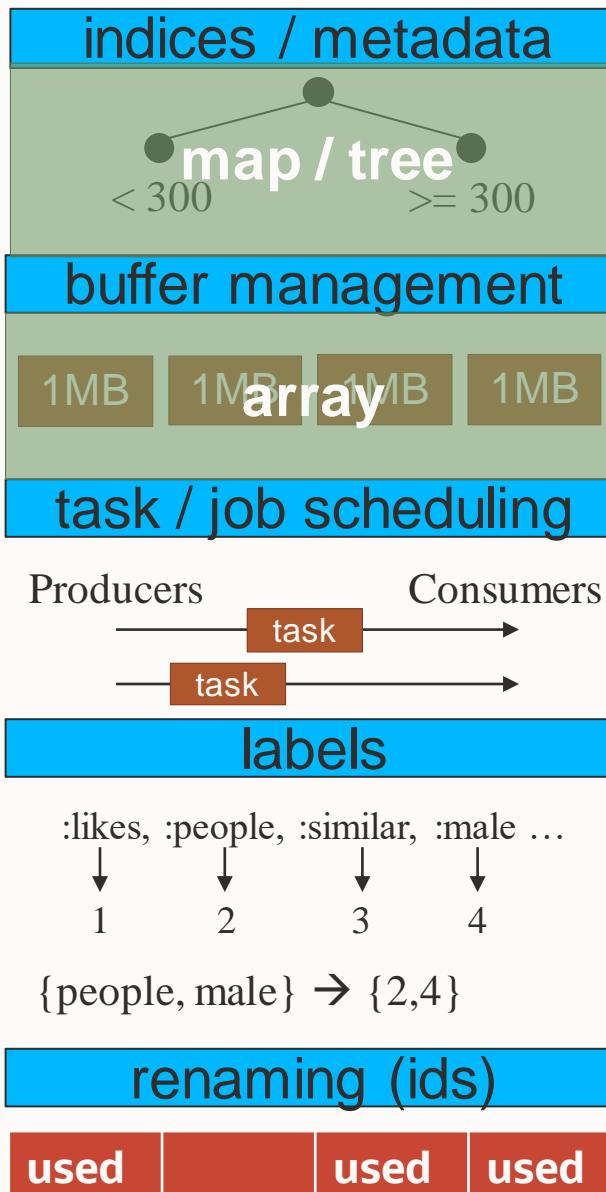
{people, male} → {2,4}

renaming (ids)



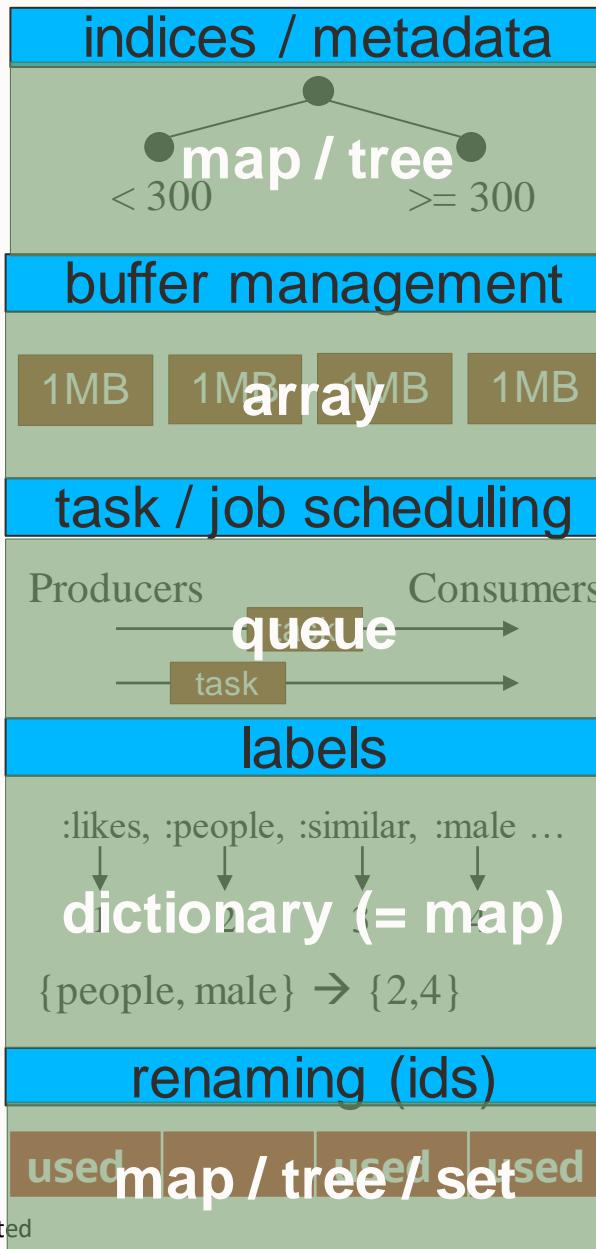
- Indices
 - Used for speeding up “queries”
 - Which vertices have label :person?
 - Which edges have value > 1000?
 - **maps, trees**
 - Buffer management
 - In “real” systems, resource management is very important
 - buffer pools
 - no order
 - insertions and deletions
 - no keys
 - Fixed num object pool: **array**
 - Otherwise: **list**
 - Variable-sized elements: **heap**

Runtime



- Task and job scheduling
 - producers create and share tasks
 - consumers get and handle tasks
 - insertions and deletions
 - usually FIFO requirements
 → **queues**
- Storing / querying sets of labels
 - set equality expensive
 - usually common groups
e.g., {person, female}, {person, male}
 → **2-level dictionary** encoding
 - {person, female} → 0
 - {person, male} → 1
- Giving unique ids (renaming)
 - **tree, map, set, counter, other?**

Runtime



Score

Structure	# Usages
array / buffer	6
map	5
tree / heap	2
set	1
queue	1

Operations

group by / join

Vasilis, Breaking bad
Rachid, Dexter
Vasilis, Dexter

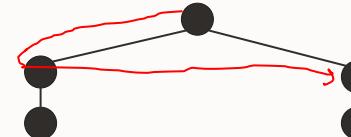
distinct

Vasilis
Rachid
Vasilis

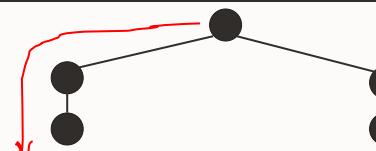
limit (top k)

11 12 0 9 8 13
8 9 11 23 32 9
1 2 3 5 7 3 2 0

BFS



DFS



- Group by
 1. Mapping from keys to values
 2. Atomic value aggregations
e.g., COUNT, SUM, MAX
 - insertion only

→ **hash map**
→ **atomic inc / sum / max, etc.**

- Join

- create a map of the small table
 - insertion phase, followed by
 - probing phase
- **hash map, lock-free probing**

Operations

group by / join

Vasilis, Breaking bad Rachid, Dexter Vasilis, Dexter

map / atomics

distinct

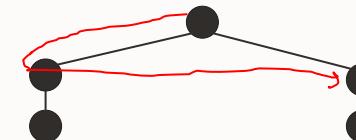
Vasilis Rachid Vasilis Rachid



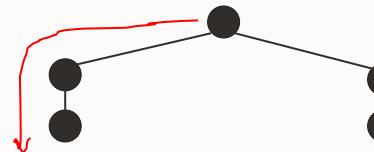
limit (top k)

$$\begin{array}{r} 11 \ 12 \ 0 \ 9 \ 8 \ 13 \\ 8 \ 9 \ 11 \ 23 \ 32 \ 9 \\ 1 \ 2 \ 3 \ 5 \ 7 \ 3 \ 2 \ 0 \end{array} \quad \rightarrow \quad \begin{array}{r} 32 \\ 23 \\ 13 \end{array}$$

BFS



DFS



- Distinct
 - can be solved with sorting, or

Operations

group by / join

Vasilis, Breaking bad
Rachid, Dexter
Vasilis, Dexter

map / atomics

Vasilis
Rachid
Vasilis



Vasilis
Rachid

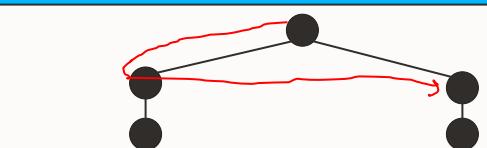
distinct

11 12 0 9 8 13
8 9 11 23 32 9
1 2 3 5 7 3 2 0

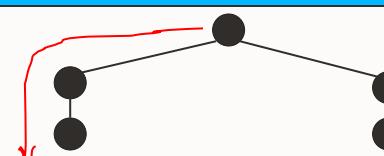


32
23
13

limit (top k)



DFS



- **Distinct**
 - can be solved with sorting, or
→ **hash set**
- **Limit (top k)**
 - can be solved with sorting, or
 - different specialized structures
→ **tree**
→ **heap**
→ ~ **list**
→ **array** (e.g., 2 elements only)
→ **register** (1 element only)

Operations

group by / join

Vasilis, Breaking bad
Rachid, Dexter
Vasilis, Dexter

map / atomics

Vasilis, 2
Rachid, 1
Vasilis, 1

distinct

Vasilis
Rachid
Vasilis

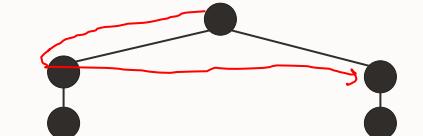
hash set

limit (top k)

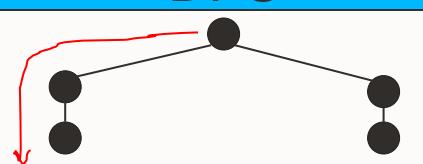
11 12 0 9 8 13
8 9 1 2 3 29
1 2 3 5 7 3 20

tree / heap / list

BFS

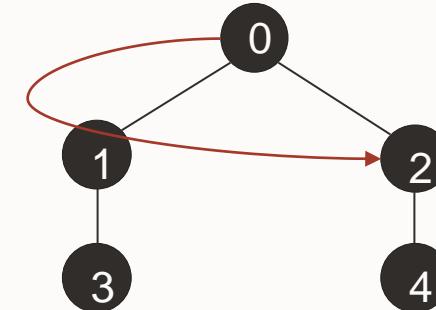


DFS



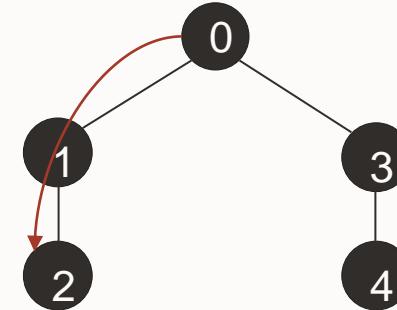
- Breadth-first search (BFS)

- FIFO order
 - track visited vertices
- **queue**
→ **set**



- Depth-first search (DFS)

- LIFO order
 - track visited vertices
- **stack**
→ **set**



Operations

group by / join

Vasilis, Breaking bad
Rachid, Dexter
Vasilis, Dexter

distinct

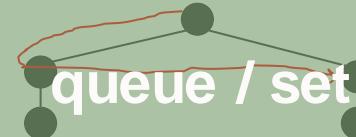
Vasilis
Rachid
Vasilis

limit (top k)

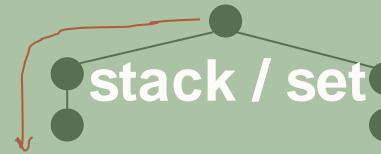
11 12 0 9 8 13
8 9 1 2 3 2 9
1 2 3 5 7 3 2 0

32
23
13

BFS



DFS



map / atomics

hash set

tree / heap / list

Score

Structure	# Usages
array / buffer	7
map	6
set	4
tree / heap	3
queue	2
stack	1
list	1

Your new cheatsheet for interviews!

Graph

tmp graph structure

↓
segmented buffer

“Vasilis”, “Breaking bad”, :likes
 “Rachid”, “Dexter”, :likes
 “Vasilis”, “Dexter”, :likes
 “Dexter”, “Breaking bad”, :similar
 “Breaking bad”, “Dexter”, :similar

graph structure



user-ids - internal ids

Vasilis → 0 0 → Vasilis
 Rachid → 1 1 → Rachid
 Breaking bad → 2 2 → Breaking bad
 Dexter → 3 3 → Dexter

labels

:likes, :people, :similar, ...

properties

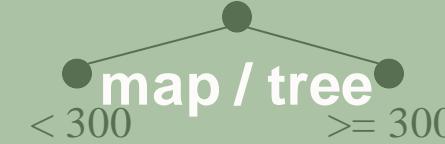
“Vasilis”, {people, male}, 20, Zurich
 “Rachid”, {people, male}, ??, Lausanne

lifetime management

number of references: X
stripped counter

Runtime

indices / metadata



buffer management

1MB 1MB **array** 1MB 1MB

task / job scheduling



labels

:likes, :people, :similar, :male ...

1 2 3 4

{people, male} → {2,4}

renaming (ids)

used **map / tree / set** **used** **used**

Operations

group by / join

Vasilis, Breaking
 bad
 Rachid, Dexter
 Vasilis, Dexter

map / atomics

distinct

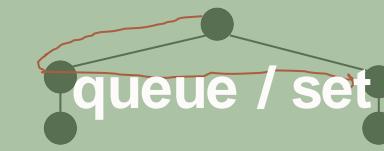
Vasilis
 Rachid
 Vasilis

limit (top k)

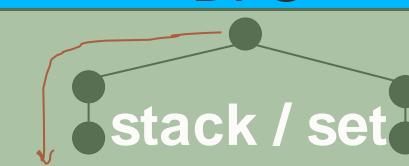
11 12 0 9 8 13
 8 9 1 2 3 9
 1 2 3 5 7 3 2 0

tree / heap / list

BFS



DFS



Conclusions

- Both **theory** and **practice** are necessary for
 - Designing, and
 - Implementing fast / scalable data structures
- **Hardware** plays a huge role on implementations
 - How and which memory access patterns to use
- **(Concurrent) Data structures**
 - The backbone of every system
 - An “open” and challenging area of research

Our mission is to help people
see data in new ways, discover insights,
unlock endless possibilities.



Internship and job opportunities

Visit our Internship Page: labs.oracle.com/pls/apex/labs/r/labs/internships
or find our topics in EPFL's portal

- AI/ML in Database
- AI/ML Technology for Enterprise Applications
- Cloud Security at Oracle
- Extending Oracle (Labs) Security and Compliance Applications
- Graal Cloud Service
- List of All Topics
- Optimizing Data Access in Microservice Applications with GraalOS-in-DB

or just send us an email at epfl-labs_ch@oracle.com or to me Vasileios.Trigonakis@oracle.com

Using the Oracle Cloud for free



Everybody

Oracle Cloud Always-Free Tier: oracle.com/cloud/free/

Universities and Schools

Oracle Academy: academy.oracle.com

Research Institutions

Oracle For Research: oracle.com/oracle-for-research/

Our mission is to help people
see data in new ways, discover insights,
unlock endless possibilities.

