


Transactional Memory

R. Guerraoui, EPFL

Locking is “history”


Lock-freedom is “difficult”

Wanted



*A synchronisation abstraction that is
simple, robust and efficient*

Transactions



Historical perspective

- ⌚ Eswaran et al (CACM'76) Databases
- ⌚ Papadimitriou (JACM'79) Theory
- ⌚ Liskov/Sheifler (TOPLAS'82) Language
- ⌚ Knight (ICFP'86) Architecture
- ⌚ Herlihy/Moss (ISCA'93) Hardware
- ⌚ Shavit/Touitou (PODC'95) Software
- ⌚ Herlihy et al (PODC'03) Software – Dynamic
- ⌚ Kapalka/Guerraoui (Morgan Claypool 2010)

Back to the sequential level

- ⌚ accessing object 1;
- ⌚ accessing object 2;

Back to the sequential level

atomic {


- ⌚ accessing object 1;
- ⌚ accessing object 2;

}

Semantics (serialisability)


Every transaction appears to execute
at an indivisible point in time
(linearizability of transactions)

Double-ended queue



```
 ↵ class Queue {  
 ↵     QNode head;  
 ↵     QNode tail;  
 ↵     public enq(Object x) {  
 ↵         atomic {  
 ↵             QNode q = new QNode(x);  
 ↵             q.next = head;  
 ↵             head = q;  
 ↵         }  
 ↵     }  
 ↵     ... }
```

Queue composition



```
 ↵ class Queue {  
 ↵     ...  
 ↵     public transfer(Queue q) {  
 ↵         atomic {  
 ↵             Qnode n = this.dequeue();  
 ↵             q.enqueue(n) }  
 ↵         }  
 ↵         ... }
```

Simple example

(consistency invariant)

$$0 < x < y$$

Simple example

(transaction)

⌚ T: x := x+1 ; y:= y+1

The illusion of a critical section

atomic {

- ⌚ accessing object 1;
- ⌚ accessing object 2;

}

“It is better for Intel to get involved in this [Transactional Memory] now so when we get to the point of having ...tons... of cores we will have the answers” **Justin Rattner, Intel Chief Technology Officer**

“...we need to explore new techniques like transactional memory that will allow us to get the full benefit of all those transistors and map that into higher and higher performance.” **Bill Gates**

“...manual synchronization is intractable...transactions are the only plausible solution....” **Tim Sweeney, Epic Games**

The TM Topic has been a VERY HOT topic

- ➊ Sun/Oracle, Intel, AMD, IBM, MSR
- ➋ Fortress (Sun); X10 (IBM); Chapel (Cray)

The TM API (a simple view)

- ➊ ***begin()*** returns *ok*

- ➋ ***read()*** returns a value or *abort*
- ➌ ***write()*** returns an *ok* or *abort*

- ➍ ***commit()*** returns *ok* or *abort*
- ➎ ***abort()*** returns *ok*

Two-phase locking

- ➊ To ***write*** or ***read*** O, T requires a ***lock*** on O;
T ***waits*** if some T' acquired a ***lock*** on O
- ➋ At the end, T ***releases*** all its locks

Two-phase locking (more details)

- Every object O , with state $s(O)$ (a *register*), is protected by a lock $l(O)$ (a **c&s**)
- Every transaction has local variables $wSet$ and $wLog$
- Initially: $l(O) = \text{unlocked}$, $wSet = wLog = \emptyset$

Two-phase locking

Upon op = *read()* or *write(v)* on object O

if O \notin wSet then

 wait until unlocked= I(O).c&s(unlocked,locked)

 wSet = wSet U O

 wLog = wLog U S(O).read()

if op = read() then return S(O).read()

 S(O).write(v)

return ok

Two-phase locking (cont'd)

Upon *commit()*

cleanup()

return ok

Upon *abort()*

rollback()

cleanup()

return ok

Two-phase locking (cont'd)

Upon *rollback()*

for all $O \in wSet$ do $S(O).write(wLog(O))$

$wLog = \emptyset$

Upon *cleanup()*


for all $O \in wSet$ do $I(O).write(unlocked)$

$wSet = \emptyset$

Why two phases? (what if?)

- ➊ To ***write*** or ***read*** O, T requires a ***lock*** on O;
T ***waits*** if some T' acquired a ***lock*** on O
- ➋ T ***releases*** the lock on O when T is done with O

Why two phases?



Two-phase locking (read-write lock)

- ➊ To ***write*** O, T requires a ***write-lock*** on O;
T ***waits*** if some T' acquired a ***lock*** on O
- ➋ To ***read*** O, T requires a ***read-lock*** on O;
T ***waits*** if some T' acquired a ***write-lock*** on O
- ➌ Before committing, T ***releases*** all its locks

Two-phase locking

- better dead than wait -

- ☛ To **write** O, T requires a **write-lock on O**; T **aborts** if some T' acquired a **lock** on O
- ☛ To **read** O, T requires a **read-lock** on O; T **aborts** if some T' acquired a **write-lock** on O
- ☛ Before committing, T releases all its locks
- ☛ A transaction that aborts restarts again

Two-phase locking

- better kill than wait -

- ☛ To **write** O, T requires a **write-lock on O**; T **aborts T'** if some T' acquired a **lock** on O
- ☛ To **read** O, T requires a **read-lock** on O; T **aborts T'** if some T' acquired a **write-lock** on O
- ☛ Before committing, T releases all its locks
- ☛ A transaction that is aborted restarts again

Two-phase locking

- better kill than wait -

- ➊ To **write** O, T requires a **write-lock on** O;
T **aborts T'** if some T' acquired a **lock** on O
- ➋ To **read** O, T requires a **read-lock** on O;
T **waits** if some T' acquired a **write-lock** on O
- ➌ Before committing, T releases all its locks
- ➍ A transaction that is aborted restarts again

Visible Read

(SXM, RSTM, TLRW)

- ☞ ***Write is mega killer.*** to write an object, a transaction aborts any live one which has read or written the object
- ☞ ***Visible but not so careful read:*** when a transaction reads an object, it says so

Visible Read

- ➊ A visible read invalidates cache lines
- ➋ For read-dominated workloads, this means a lot of traffic on the bus between processors
 - This reduces the throughput
 - Not a big deal with single-CPU, but with many core machines

Two-phase locking with invisible reads

- ➊ To ***write*** O, T requires a ***write-lock on O***; T **waits** if some T' acquired a ***write-lock*** on O
- ➋ To ***read*** O, T checks if ***all objects read remain valid*** - else T **aborts**
- ➌ Before committing, T checks if all objects read remain valid and releases all its locks

Invisible reads (more details)

- Every object O , with state $s(O)$ (register), is protected by a lock $I(O)$ (c&s)
- Every transaction maintains, besides wSet and wLog:
- A local variable $rset(O)$ for every object

Invisible reads

Upon ***write(v)*** on object O

if $O \notin wSet$ then

 wait until $unlocked = l(O).c \& s(unlocked, locked)$

$wSet = wSet \cup O$

$wLog = wLog \cup S(O).read()$

$(*, ts) = S(O).read()$

$S(O).write(v, ts)$

 return ok

Invisible reads

Upon *read()* on object O

$(v, ts) = S(O).read()$

if $O \in wSet$ then return v

if $I(O) = \text{locked}$ or $\text{not validate}()$ then $\text{abort}()$

if $rset(O) = 0$ then $rset(O) = ts$

return v

Invisible reads

Upon ***validate()***

for all O s.t $rset(O) > 0$ do

$(v, ts) = S(O).read()$

if $ts \neq rset(O)$ or

$(O \notin wset \text{ and } I(O) = \text{locked})$

then return false

else return true

Invisible reads

Upon *commit()*

if not validate() then abort()

for all $O \in wset$ do

$(v, ts) = S(O).read()$

$S(O).write(v, ts+1)$

 cleanup()

Invisible reads

Upon *rollback()*

for all $O \in wSet$ do $S(O).write(wLog(O))$

$wLog = \emptyset$

Upon *cleanup()*

for all $O \in wset$ do $I(O).write(unlocked)$

$wset = \emptyset$

$rset(O) = 0$ for all O

DSTM (SUN)

- ☛ To **write** O, T requires a ***write-lock on*** O;
T aborts T' if some T' acquired a ***write-lock*** on O
- ☛ To **read** O, T checks if all objects read remain valid – else T **abort**
- ☛ Before committing, T releases all its locks

DSTM

- ☛ ***Killer write*** (ownership)
- ☛ ***Careful read*** (validation)

More efficient algorithm?

Apologizing versus asking permission

⌚ ***Killer write***

⌚ ***Optimistic read***

- validity check only at commit time

Example

Invariant: $0 < x < y$

Initially: $x := 1; y := 2$

Division by zero

☛ T1: $x := x + 1 ; y := y + 1$

☛ T2: $z := 1 / (y - x)$

Infinite loop

- ☛ T1: $x := 3; y := 6$
- ☛ T2: $a := y; b := x;$
repeat $b := b + 1$ until $a = b$


Opacity

- ⌚ Serializability
- ⌚ Consistent memory view

Trade-off

The read is either
visible or *careful*

Intuition



Read invisibility

- ➊ The fact that the read is invisible means T1 cannot inform T2, which would in turn abort T1 if it accessed similar objects (SXM, RSTM)
- ➋ NB. Another way out is the use of multiversions: T2 would not have written “on” T1

Conditional progress - obstruction-freedom -

- ➊ A correct transaction that eventually does not encounter ***contention*** eventually commits
- ➋ ***Obstruction-freedom*** seems reasonable and is indeed possible

DSTM

- ☞ To **write** O, T requires a **write-lock on** O (use C&S);
T aborts T' if some T' acquired a **write-lock** on O (use C&S)
- ☞ To **read** O, T checks if all objects read remain valid -
else abort (use C&S)
- ☞ Before committing, T releases all its locks (use C&S)

Progress

- ☞ If a transaction T wants to write an object O owned by another transaction T', T calls a ***contention manager***
- ☞ The contention manager can decide to wait, retry or abort T'

Contention managers

- ☛ **Aggressive:** always aborts the victim
- ☛ **Backoff:** wait for some time (exponential backoff) and then abort the victim
- ☛ **Karma:** priority = cumulative number of shared objects accessed – work estimate. Abort the victim when number of retries exceeds difference in priorities.
- ☛ **Polka:** Karma + backoff waiting

Greedy contention manager

☛ State

- Priority (based on start time)
- Waiting flag (set while waiting)


☛ **Wait** if other has

- Higher priority AND not waiting

☛ **Abort** other if

- Lower priority OR waiting

Aborting is a fatality



Concluding remarks

TM does not always replace locks:
it hides them

Memory transactions look like db
transactions but are different

The garbage-collection analogy

- ➊ In the early times, the programmers had to take care of allocating and de-allocating memory
- ➋ Garbage collectors do it for you: they are now incorporated in Java and other languages
- ➌ Hardware support was initially expected, but now software solutions are very effective



MORGAN & CLAYPOOL PUBLISHERS

Principles of Transactional Memory

Rachid Guerraoui
Michał Kapałka

*SYNTHESIS LECTURES ON
DISTRIBUTED COMPUTING THEORY*

Nancy Lynch, Series Editor