

RDMA & Concurrent Algorithms

1 Dec 2025

Igor Zablotchi

 **MystenLabs**

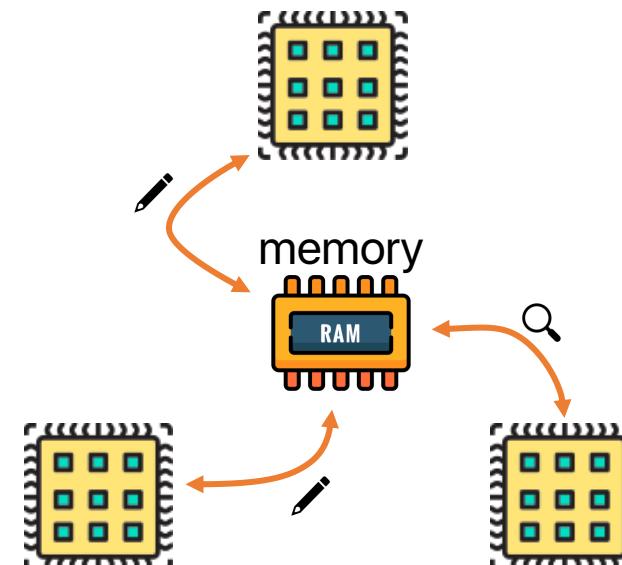
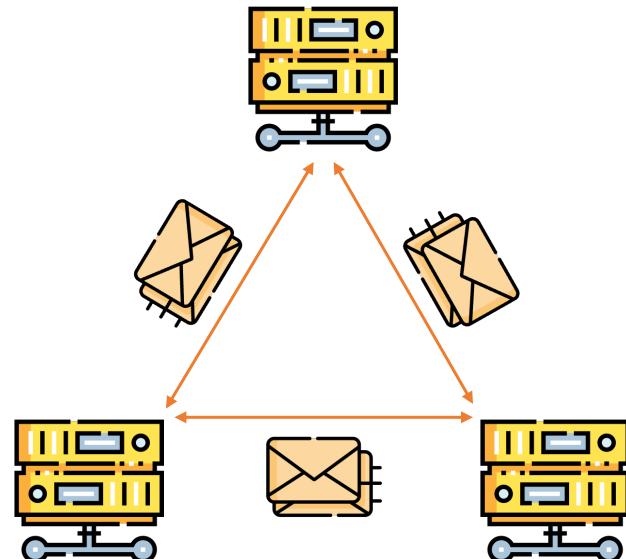
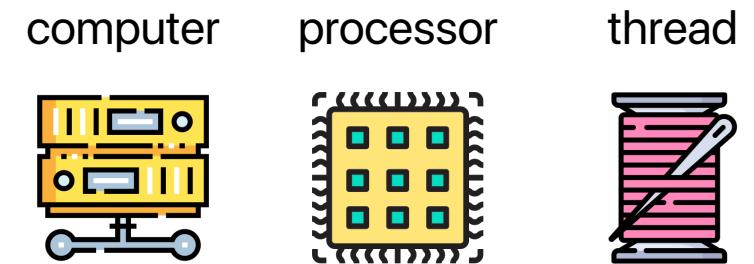
Based on joint work with, and slides from:

Marcos Aguilera, Naama Ben-David, Clément Burgelin, Rachid Guerraoui,
Virendra Marathe, Antoine Murat, Dalia Papuc, Athanasios Xygkis



A Tale of Two Models

- processes
- collaborate on some common task
- improve performance or robustness



Equal But Not Quite

The two models are equivalent [Attiya, Bar-Noy, Dolev 1995]
=
One can simulate the other

but, e.g., for solving consensus:

n = num processes f = num failures	Crash	
	Fault Tolerance	Common-case Complexity
Message Passing	$f < n/2$	2 [Lamport'98]
Shared Memory	$f < n$	4 [GL'02]

Models Reflect Technology

The two standard models reflect existing technology

BUT

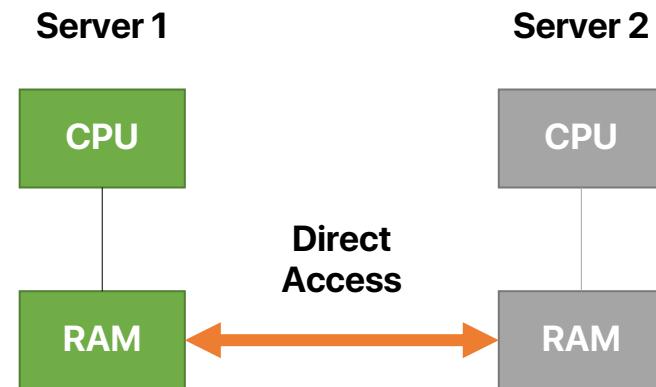
Technology evolves, new technologies emerge

SO

We need new models

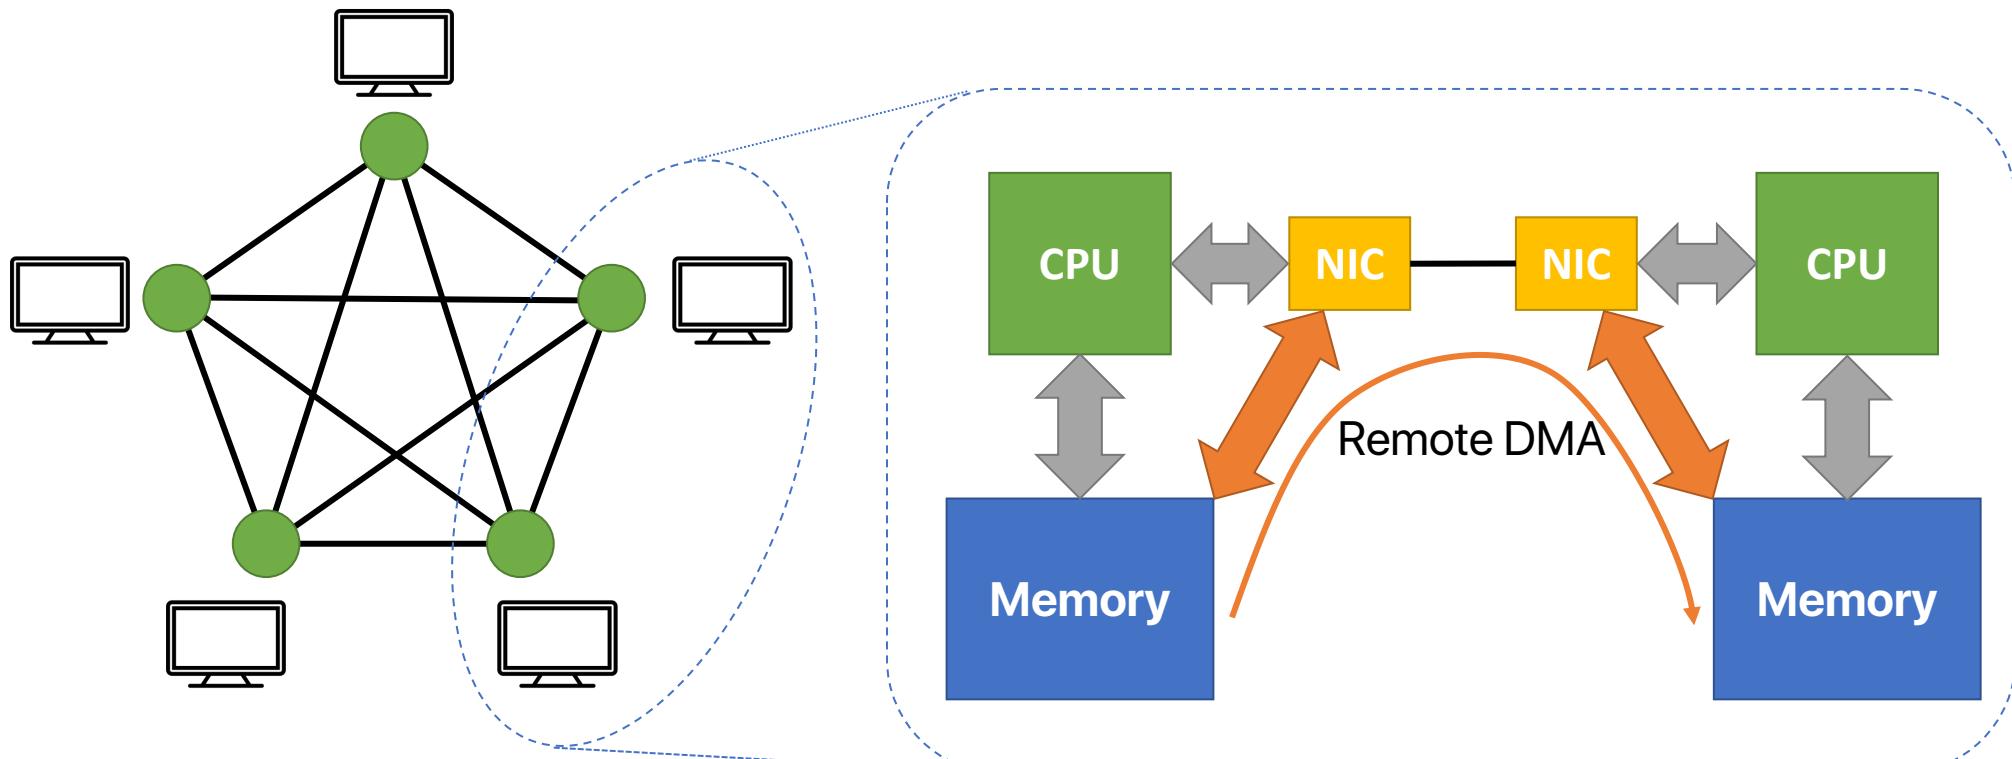
RDMA: Overview

- Networking hardware feature
- Direct access to remote memory
 - No CPU at remote side
 - No OS at either side
- Good performance
 - ~1 us latency
 - ~100-800 Gbps bandwidth
- Configurable access permissions

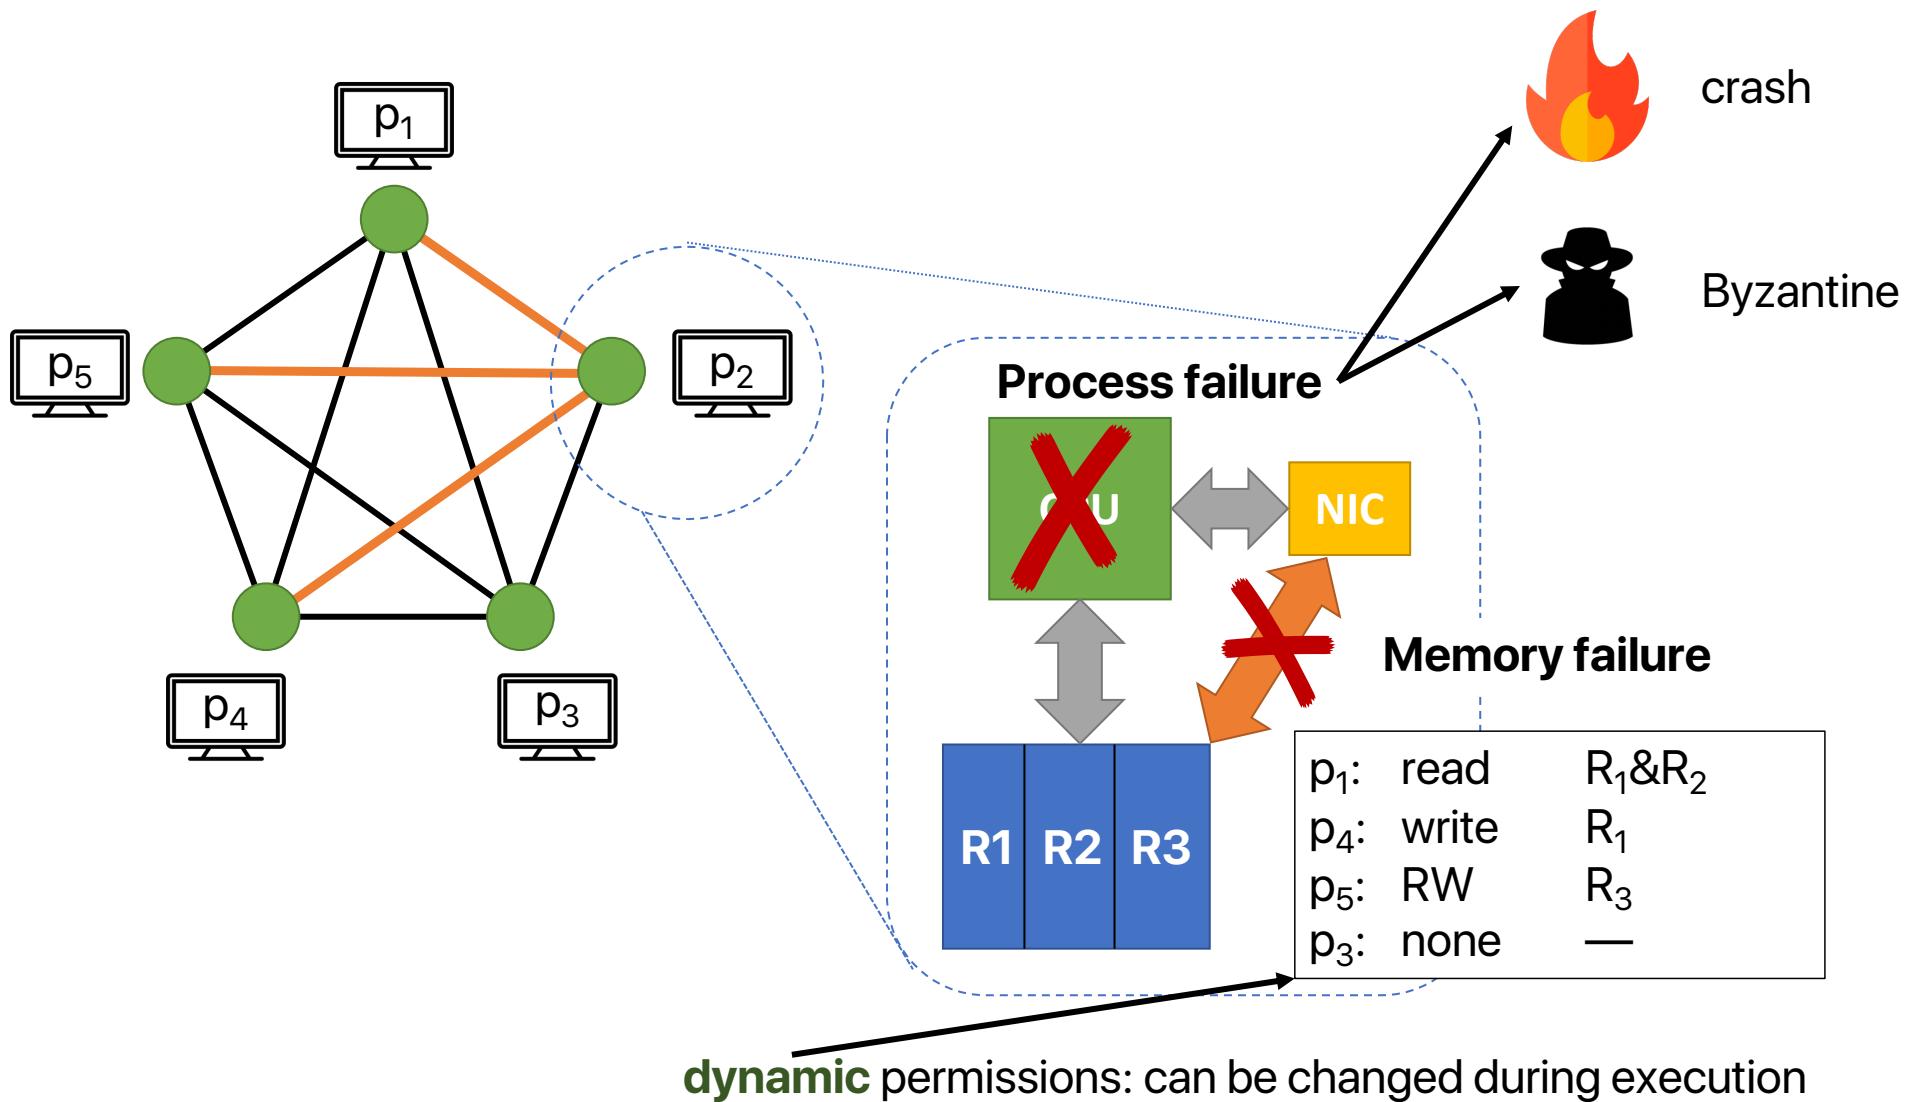


RDMA

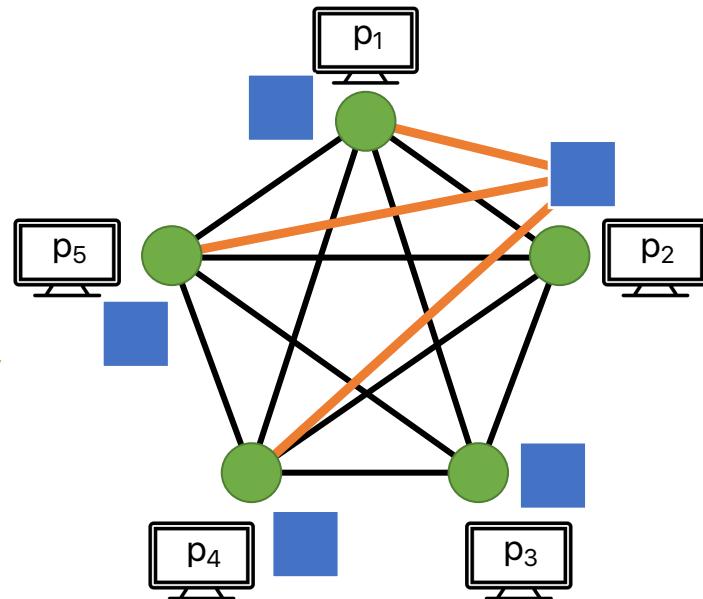
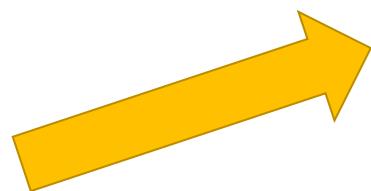
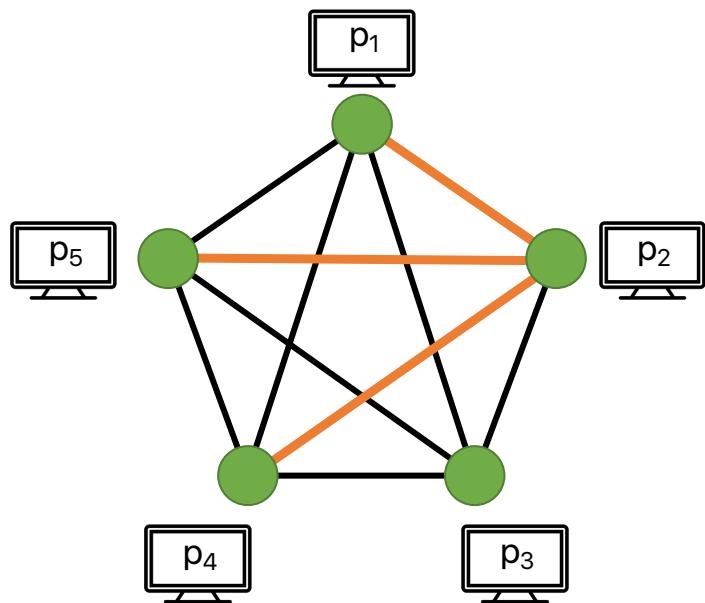
Remote Direct Memory Access (RDMA)



RDMA: Permissions and Failures



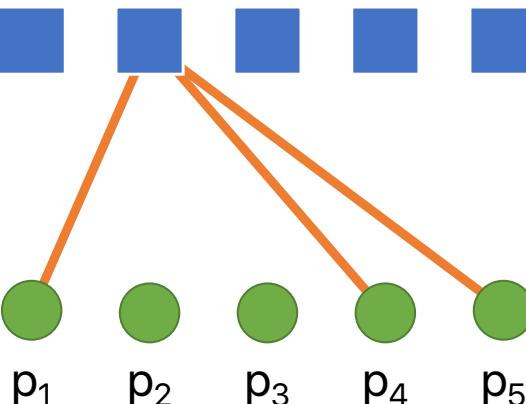
Modelling RDMA



Also called
"disaggregated memory"

memories

processes



minority of memories
can fail

Outline

- Introduction
- 3 remarkable results with RDMA:
 - Consensus with crash faults
 - Broadcast with Byzantine faults
 - Fast memory replication

Best of Both Worlds

n = num processes f = num failures	Crash	
	Fault Tolerance	Common-case Performance
Message Passing	$n > 2f$	2 [Lamport'98]
Shared Memory	$n > f$	4 [GL'02]
RDMA	$n > f$	2

Refresher: O-Consensus

Paxos in Shared Memory

```
propose(v):  
    while(true)  
        Reg[i].T.write(ts); } announce my timestamp  
        val := Reg[1,..,n].highestTspValue(); }  
        if val = ⊥ then val := v;  
        Reg[i].V.write(val,ts); } announce my value, ts  
        if ts = Reg[1,..,n].highestTsp() then }  
            return(val)  
        ts := ts + n
```

adopt value with highest ts (or mine if none)

if my timestamp is the highest, decide

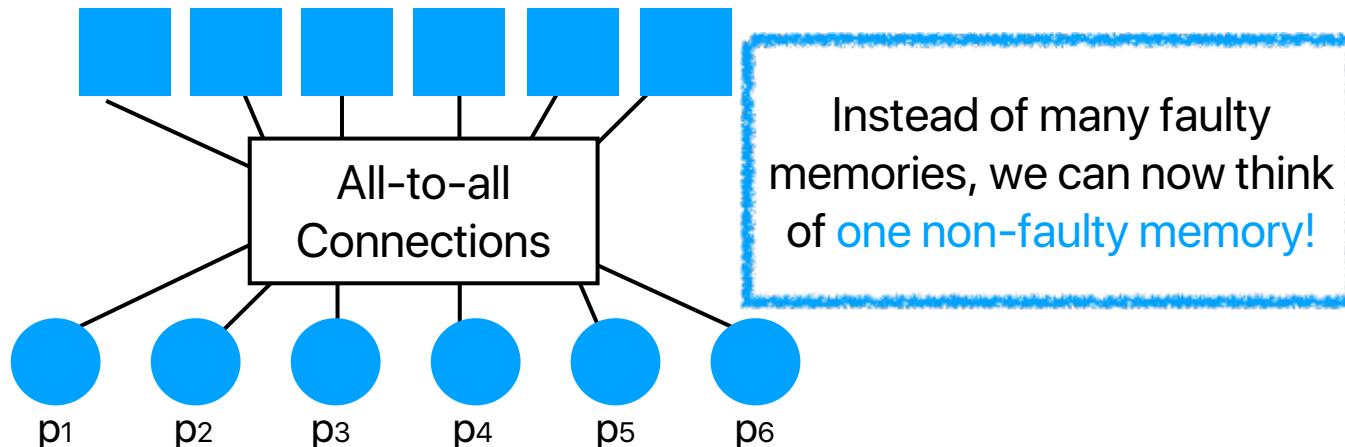
This assumes that shared memory never fails.

🤔 What if memory can fail? 🤔

Handling Memory Failures

Replication: Treat all memories the same

Send all write/read requests to all memories, wait to hear acknowledgement from majority



O-Consensus w Memory Failures

Disk Paxos [GafniLamport2002]

```
propose(v):  
    while(true)  
        for every memory m in parallel:  
            Reg[m][i].T.write(ts);  
            temp[m][1..n] = Reg[m][1..n].read();  
        until completed for majority of memories  
        val := temp[1..m][1..n].highestTspValue();  
        if val = ⊥ then val := v;  
        for every memory m in parallel:  
            Reg[m][i].V.write(val,ts);  
            temp[m][1..n] = Reg[m][1..n].read();  
        until completed for majority of memories  
        if ts = temp[1..m][1..n].highestTsp() then  
            return(val)  
        ts := ts + n
```

announce my timestamp
adopt value with highest ts (or mine if none)

announce my value, ts

if my timestamp is the highest, decide

O-Consensus w Memory Failures

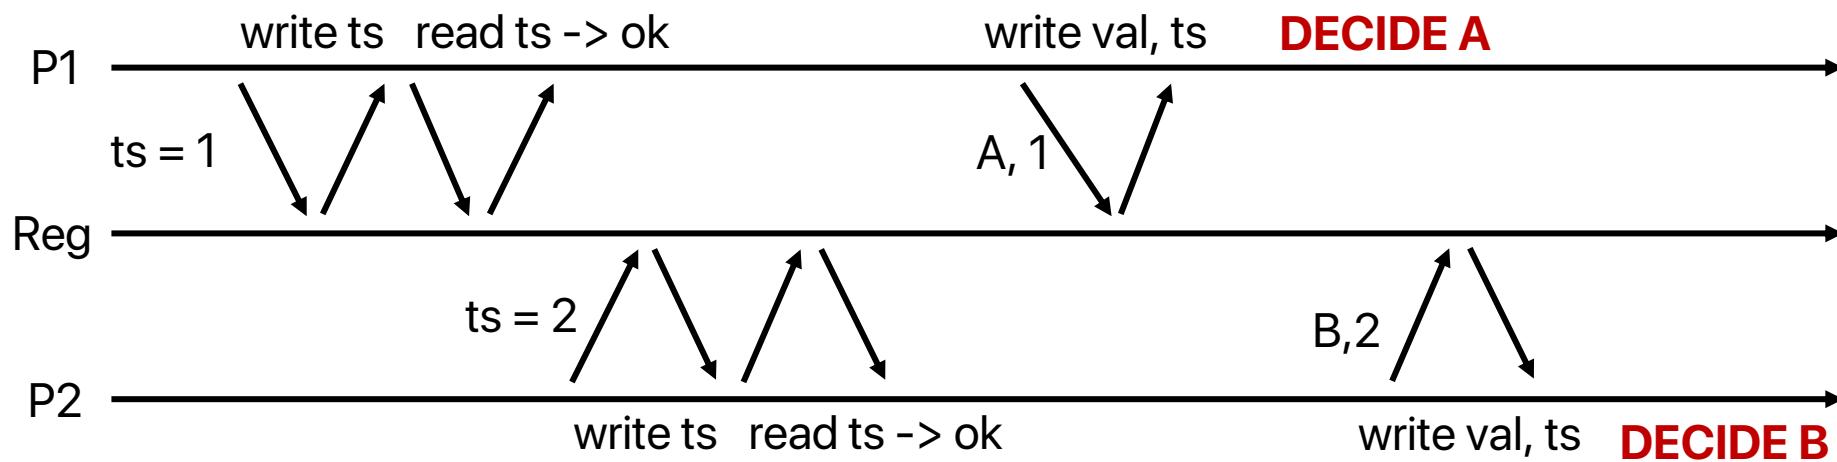
```
propose(v):  
    while(true)  
        for every memory m in parallel:  
            Reg[m][i].T.write(ts);  
            temp[m][1..n] = Reg[m][1..n].read();  
        until completed for majority of memories  
        val := temp[1..m][1..n].highestTspValue();  
        if val = ⊥ then val := v;  
        for every memory m in parallel:  
            Reg[m][i].V.write(val,ts);  
            temp[m][1..n] = Reg[m][1..n].read();  
        until completed for majority of memories  
        if ts = temp[1..m][1..n].highestTsp() then  
            return(val)  
        ts := ts + n
```

Why read again here?



👉 Need to check if I ran alone!

What If We Didn't Read?

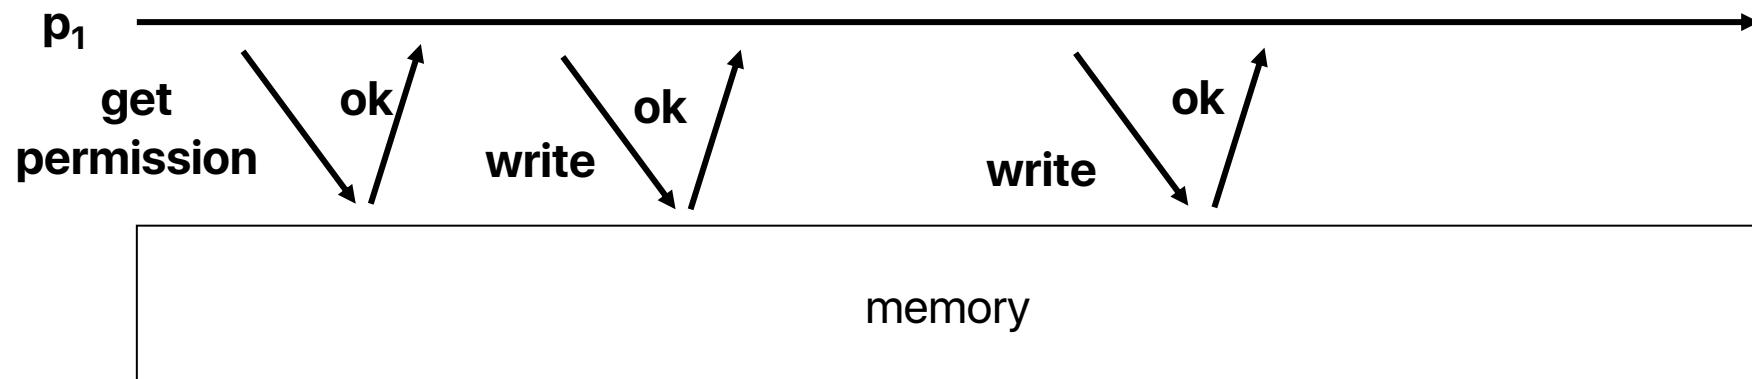


O-Consensus w Memory Failures

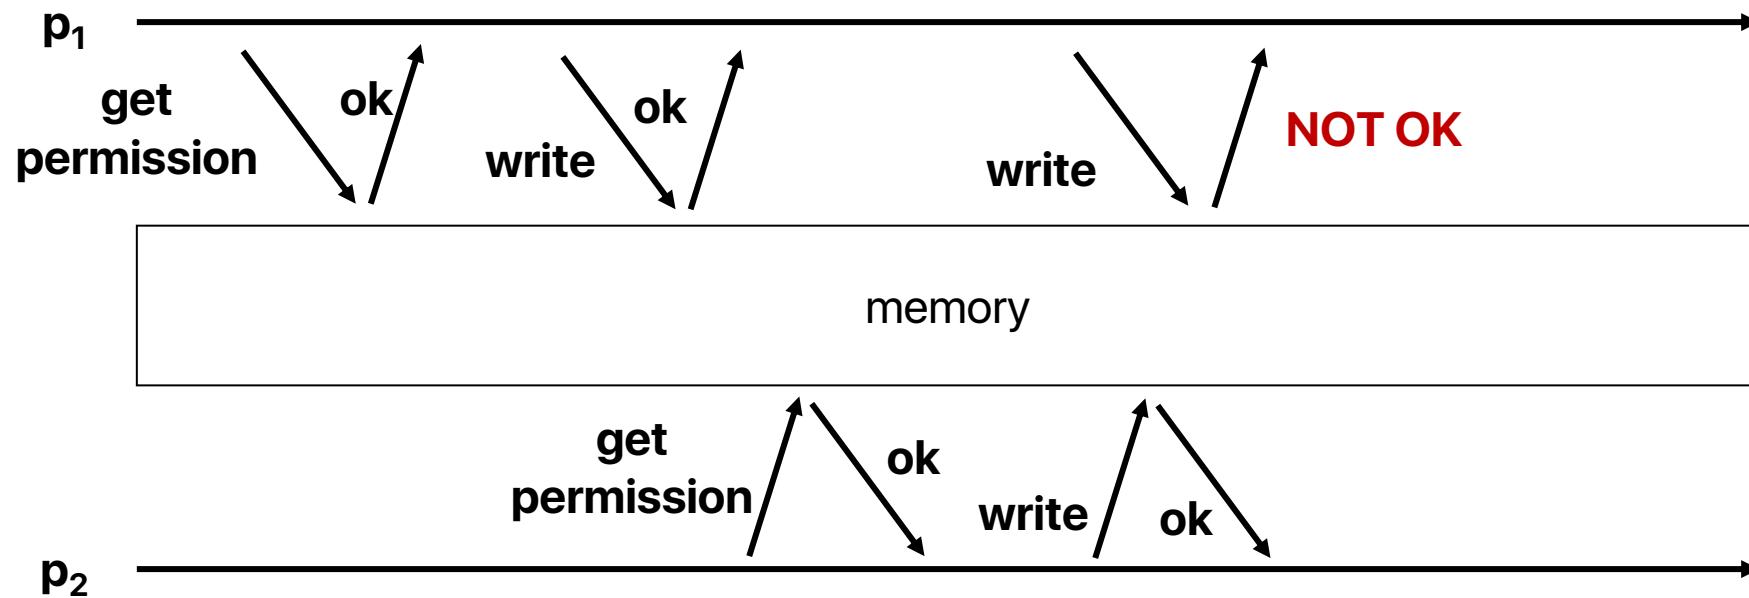
- If we don't read again, we might miss a concurrent process's timestamp
- This could lead to violation of agreement
- What if there was another way to determine if there was a concurrent process?
- We wouldn't need the last read!
→ better complexity

Solo Detection w/ Permissions

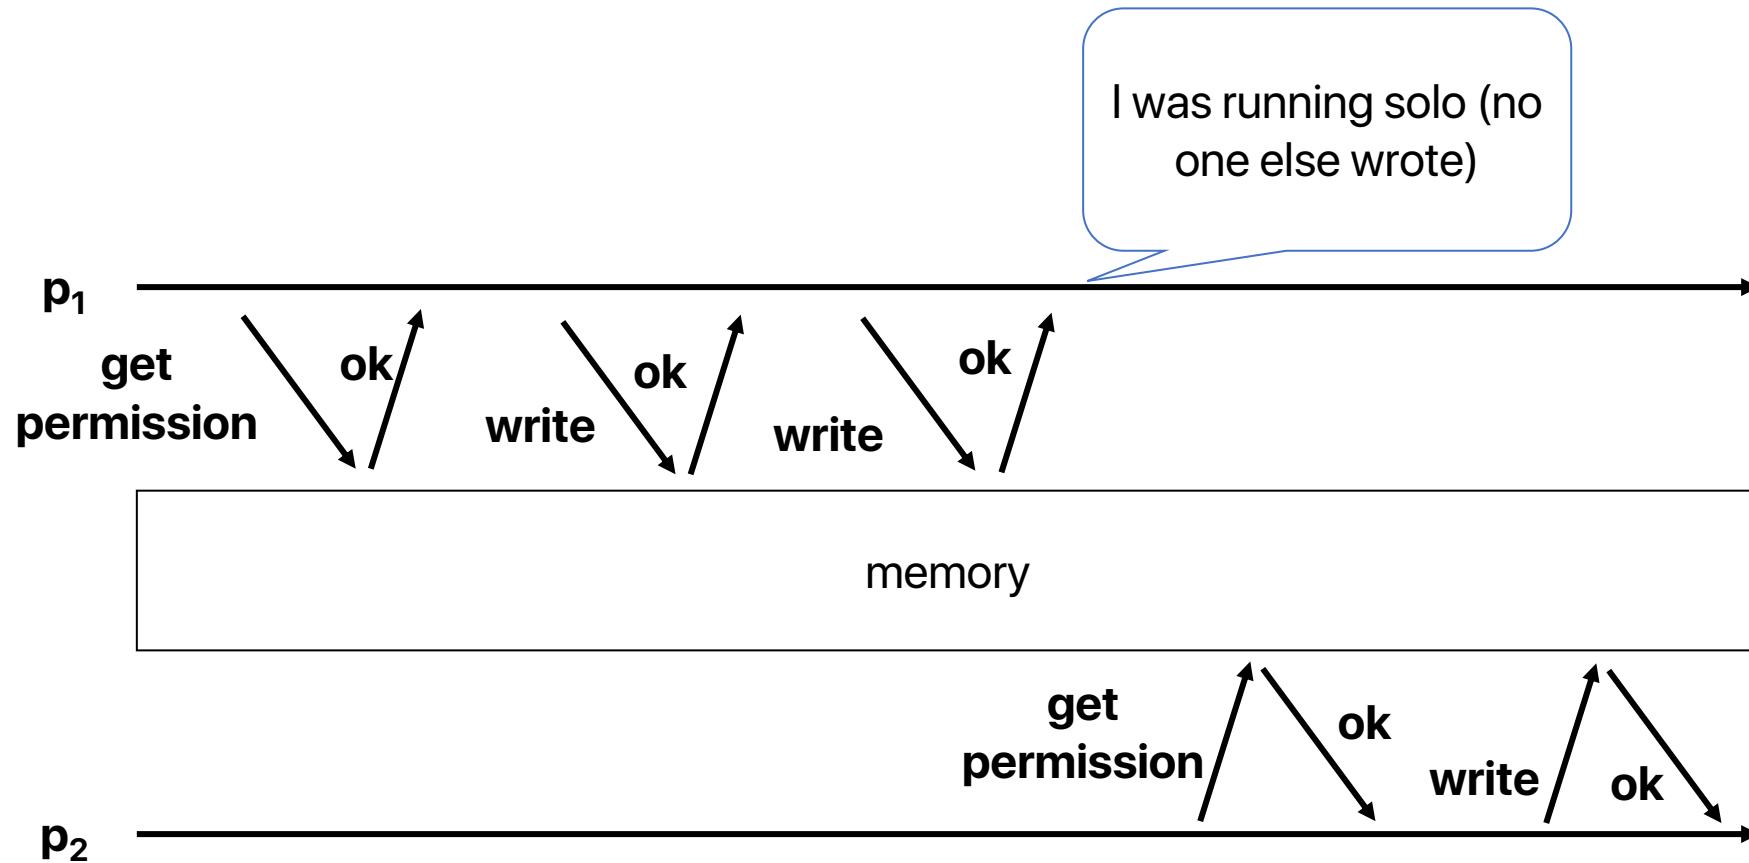
Idea: Memory gives write permission to the last process that requested it.
→ Only one process has write permission on a memory at any time.



Solo Detection w/ Permissions



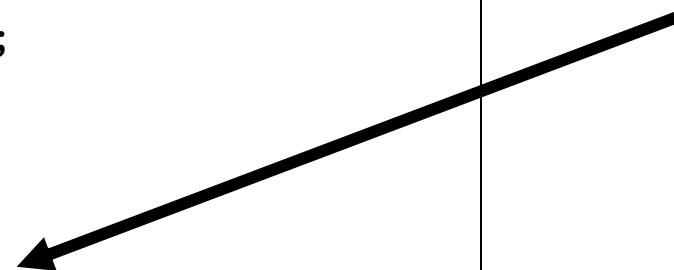
Solo Detection w/ Permissions



O-Consensus with Memory Failures and Permissions

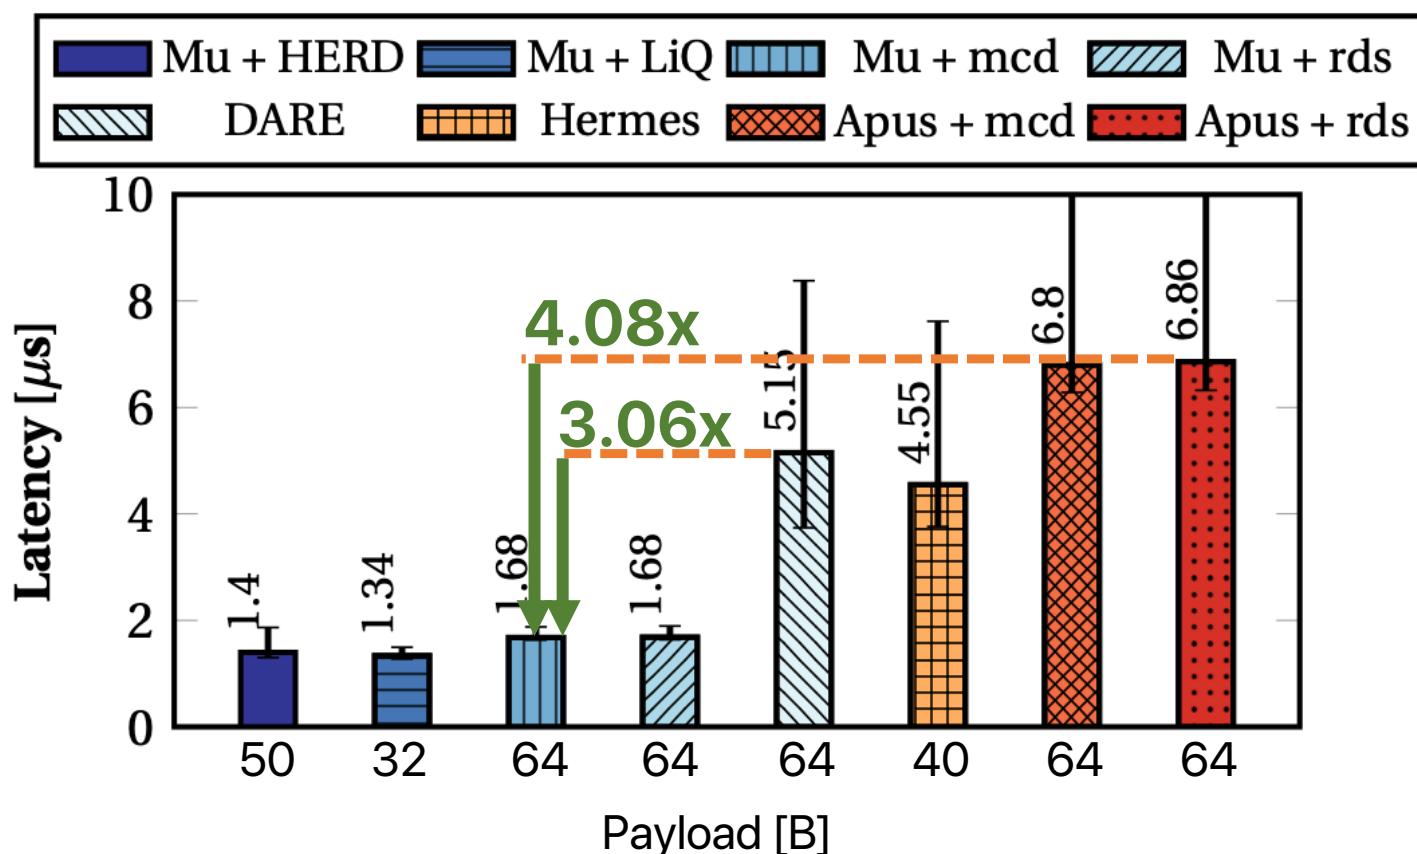
```
propose(v):  
    while(true)  
        ts := ts + n  
        for every memory m in parallel:  
            m.getPermission();  
            Reg[m][i].T.write(ts);  
            temp[m][1..n] = Reg[m][1..n].read();  
        until completed for majority of memories  
        if ts < temp[1..m][1..n].highestTsp() then continue;  
        val := temp[1..m][1..n].highestTspValue();  
        if val = ⊥ then val := v;  
        for every memory m in parallel:  
            Reg[m][i].V.write(val,ts);  
            temp[m][1..n] = Reg[m][1..n].read();  
        until completed for majority of memories  
        if writes succeeded at majority of memories then  
            return(val)
```

No need to read again!



Quick Look: Replication Latency

[3x replication, 100Gbps Infiniband]



3-4x faster than state-of-the art

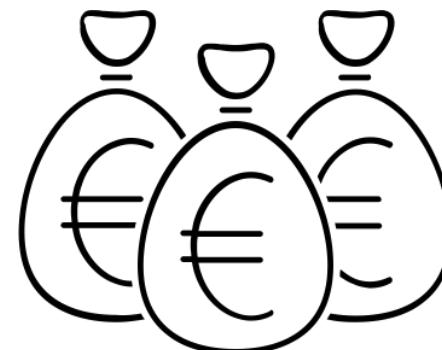
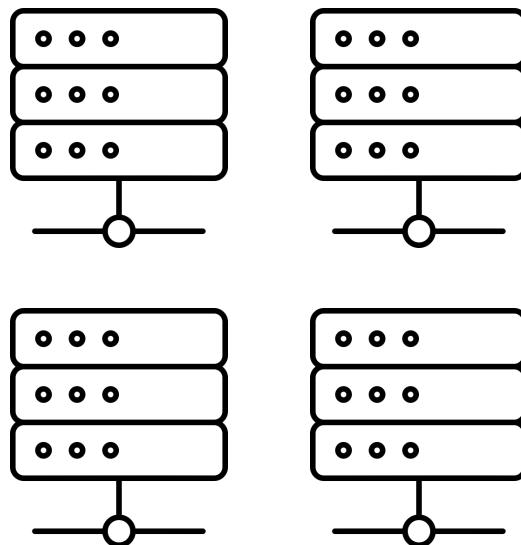
Outline

- Introduction
- 3 remarkable results with RDMA:
 - Consensus with crash faults
 - Broadcast with Byzantine faults
 - Fast memory replication

On Frugality

Number of replicas in the system

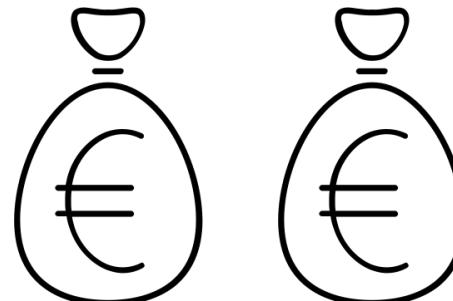
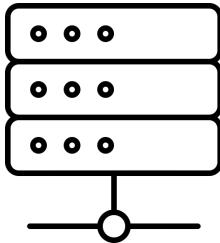
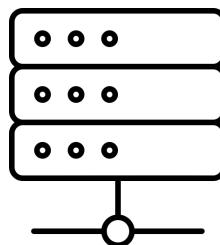
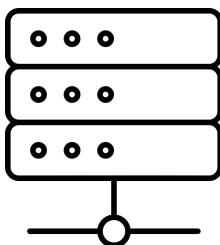
A system with $n = 3f + 1$ replicas has 33–50% more hardware than a system with $n = 2f + 1$, where f is the number of Byzantine replicas



On Frugality

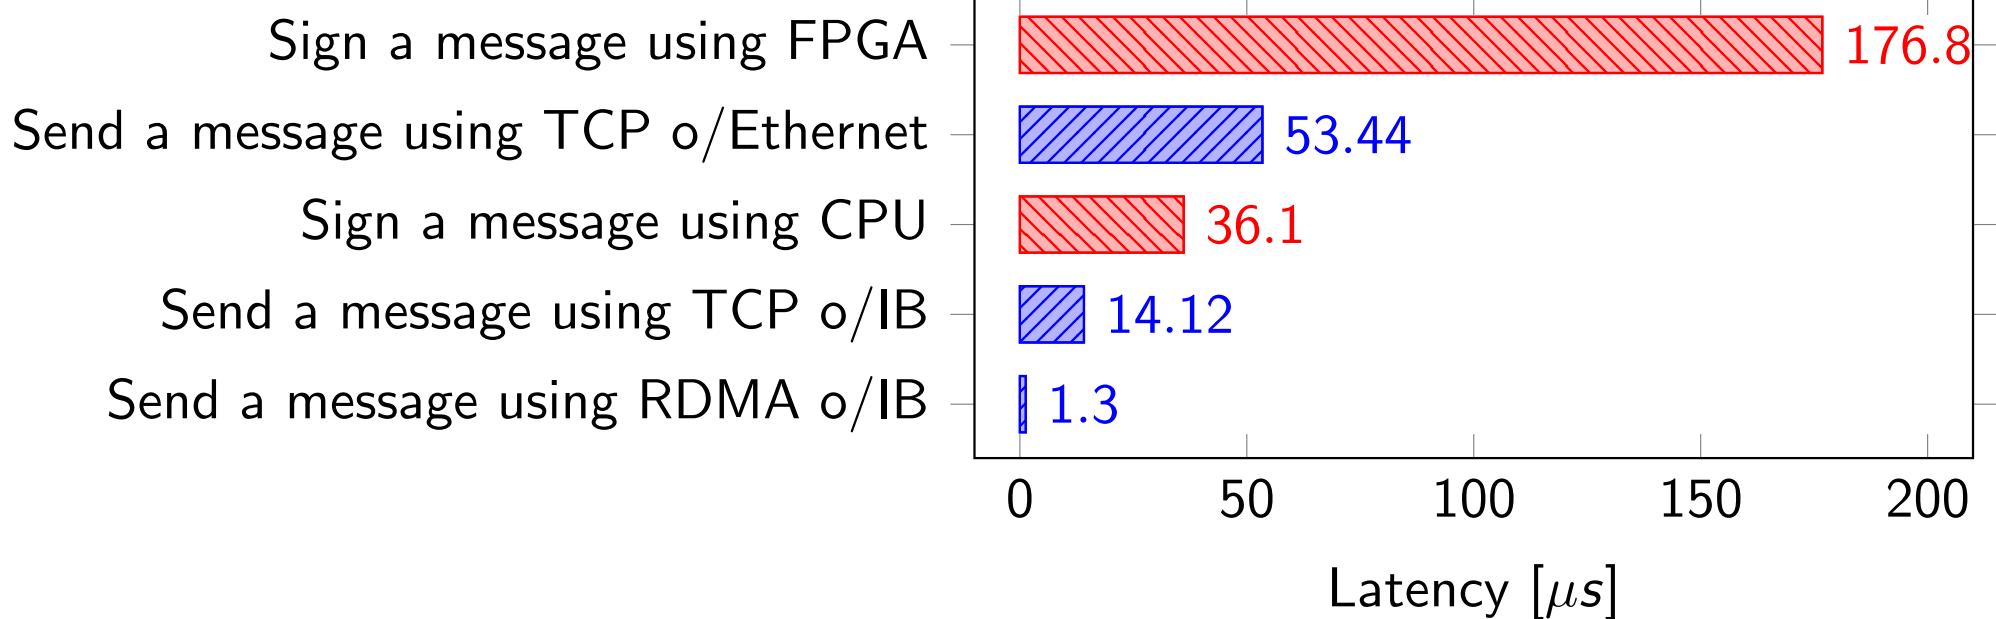
Number of replicas in the system

A system with $n = 3f + 1$ replicas has 33–50% more hardware than a system with $n = 2f + 1$, where f is the number of Byzantine replicas



On Frugality

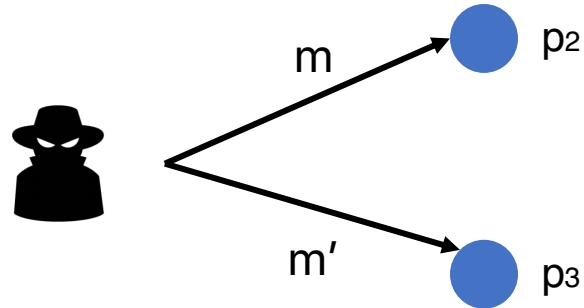
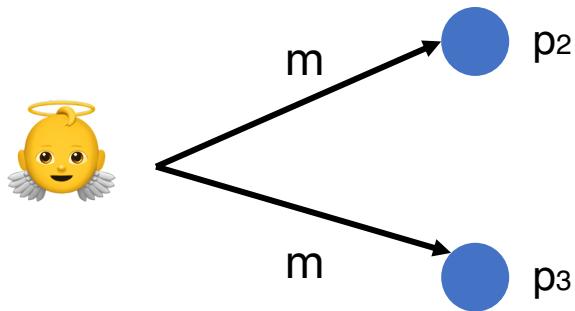
Number of digital signatures



Goal

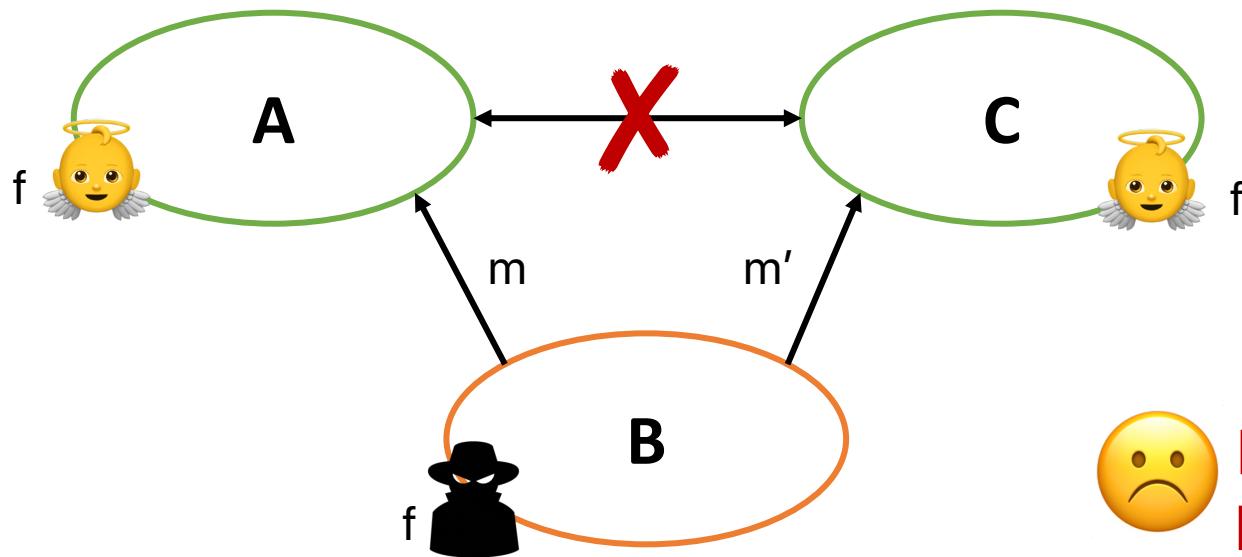
Address traditional distributed computing problems
subject to Byzantine failures with few processes,
 $n = 2f + 1$, and few signatures

Equivocation



Preventing Equivocations in Message Passing

- Requires $n=3f+1$, where n is the total number of processes and up to f processes can be Byzantine
- Intuition:



Adversary can prevent correct processes from communicating

Byzantine fault-tolerance

Non-equivocation and digital signatures improve the fault-tolerance from $3f + 1$ to $2f + 1$ for reaching agreement

Byzantine fault-tolerance

Non-equivocation and digital signatures improve the fault-tolerance from $3f + 1$ to $2f + 1$ for reaching agreement

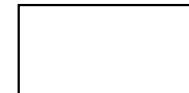
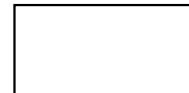
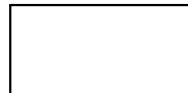
Shared memory provides non-equivocation capabilities:

Byzantine fault-tolerance

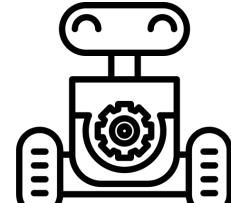
Non-equivocation and digital signatures improve the fault-tolerance from $3f + 1$ to $2f + 1$ for reaching agreement

Shared memory provides non-equivocation capabilities:

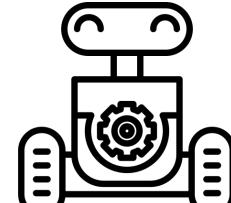
Shared memory M :



Process p_0



Process p_1

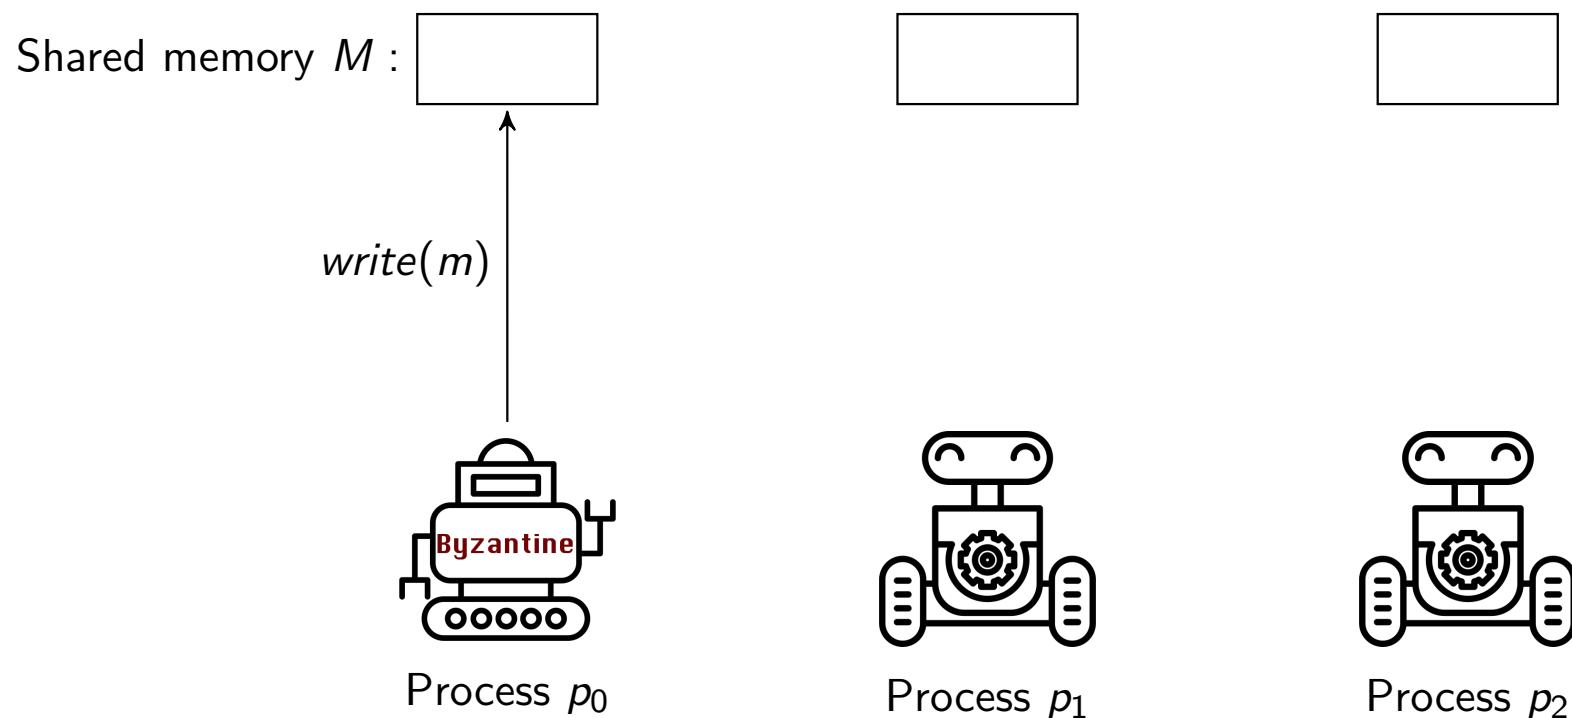


Process p_2

Byzantine fault-tolerance

Non-equivocation and digital signatures improve the fault-tolerance from $3f + 1$ to $2f + 1$ for reaching agreement

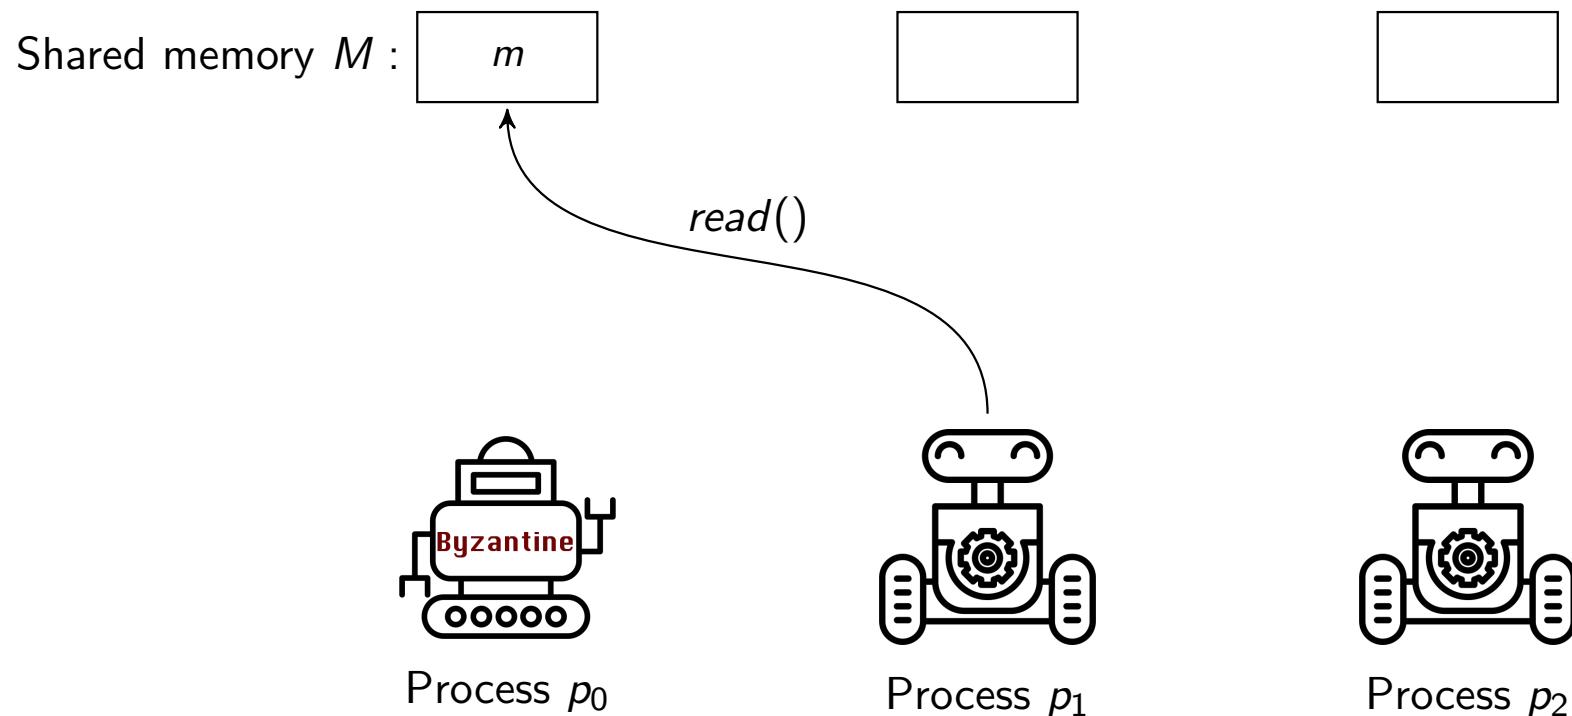
Shared memory provides non-equivocation capabilities:



Byzantine fault-tolerance

Non-equivocation and digital signatures improve the fault-tolerance from $3f + 1$ to $2f + 1$ for reaching agreement

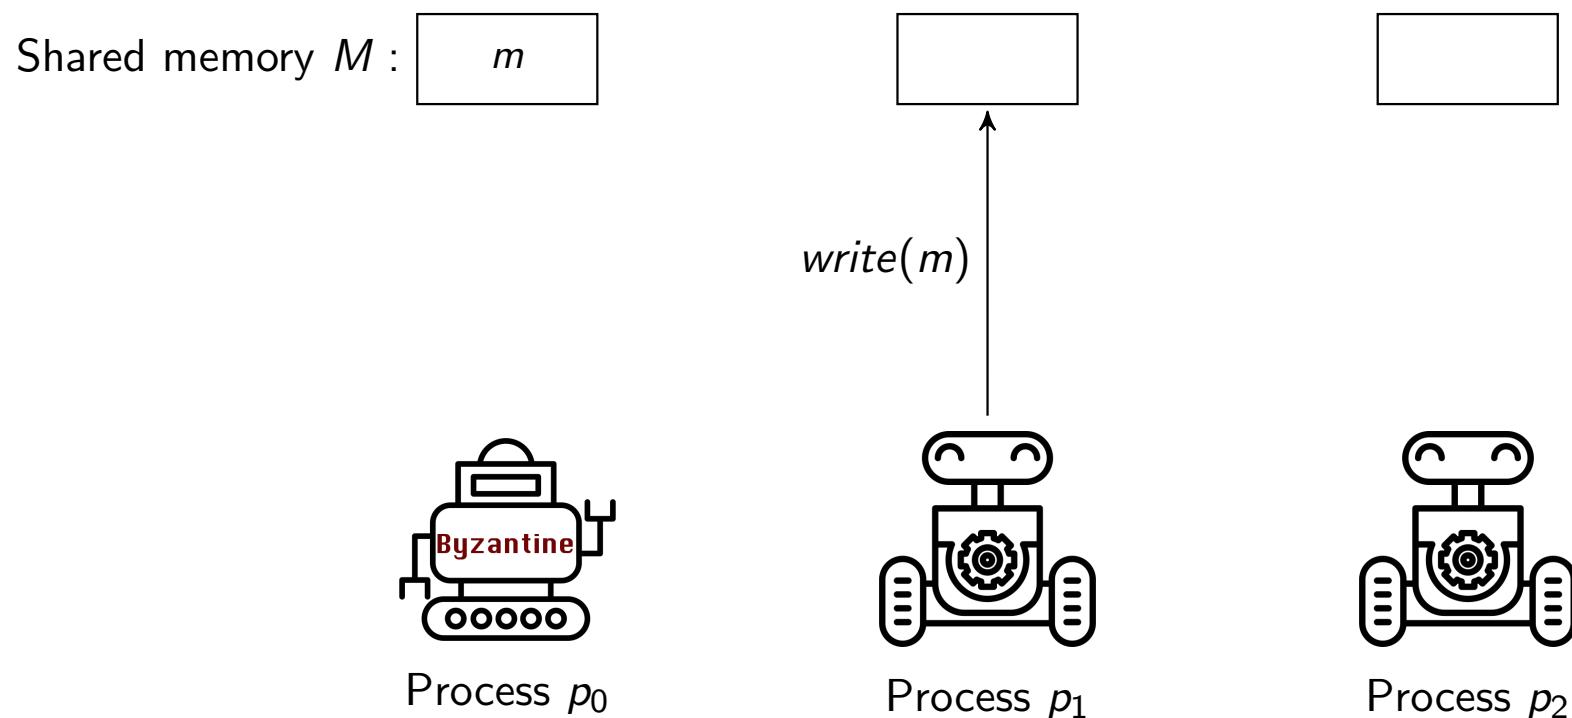
Shared memory provides non-equivocation capabilities:



Byzantine fault-tolerance

Non-equivocation and digital signatures improve the fault-tolerance from $3f + 1$ to $2f + 1$ for reaching agreement

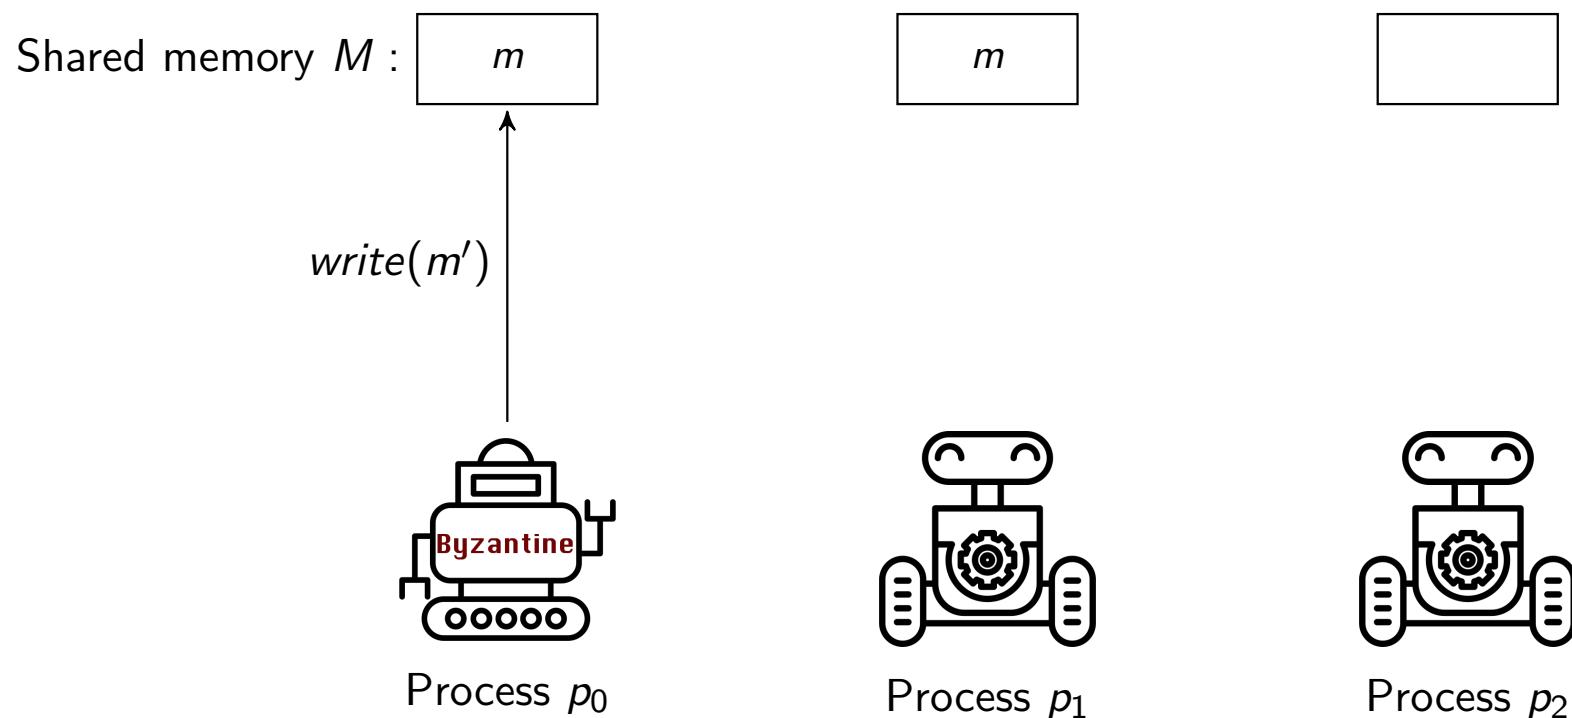
Shared memory provides non-equivocation capabilities:



Byzantine fault-tolerance

Non-equivocation and digital signatures improve the fault-tolerance from $3f + 1$ to $2f + 1$ for reaching agreement

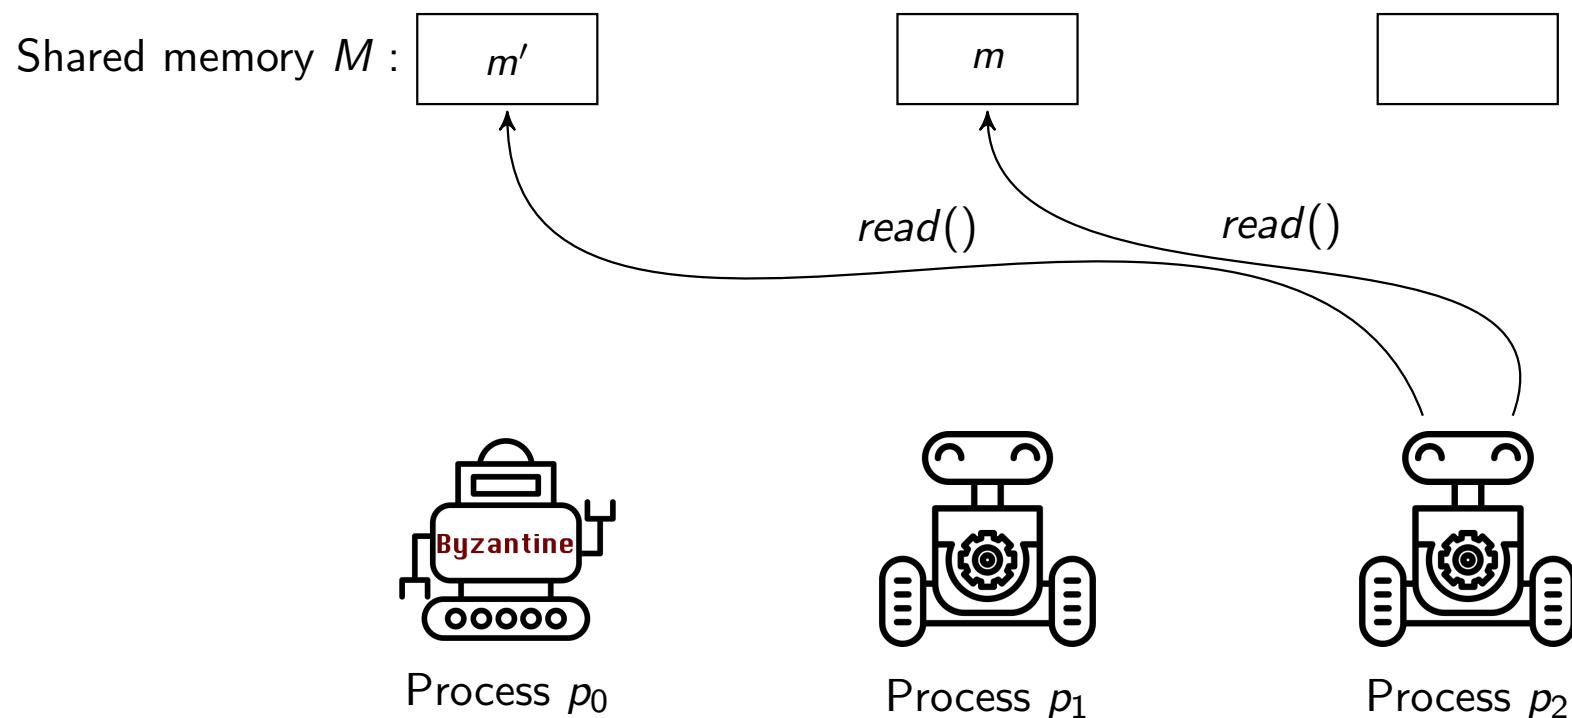
Shared memory provides non-equivocation capabilities:



Byzantine fault-tolerance

Non-equivocation and digital signatures improve the fault-tolerance from $3f + 1$ to $2f + 1$ for reaching agreement

Shared memory provides non-equivocation capabilities:



Byzantine fault-tolerance

Non-equivocation and digital signatures improve the fault-tolerance from $3f + 1$ to $2f + 1$ for reaching agreement

Shared memory provides non-equivocation capabilities:

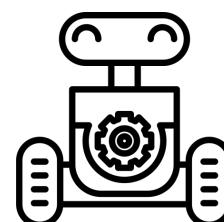
Shared memory M :

m'

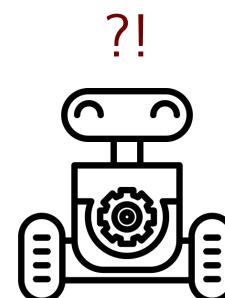
m



Process p_0



Process p_1



Process p_2

Model

Message-and-memory (M&M) [ABCGPT18] - allows processes to both pass messages and share memory M :

- Single-Writer Multi-Reader (SWMR) atomic registers
- individual memory may only fail by crashing

Model

Message-and-memory (M&M) [ABCGPT18] - allows processes to both pass messages and share memory M :

- Single-Writer Multi-Reader (SWMR) atomic registers
- individual memory may only fail by crashing

Signatures - each process has access to the primitives *sign* and *verify*

Model

Message-and-memory (M&M) [ABCGPT18] - allows processes to both pass messages and share memory M :

- Single-Writer Multi-Reader (SWMR) atomic registers
- individual memory may only fail by crashing

Signatures - each process has access to the primitives *sign* and *verify*

Up to f Byzantine processes, where $n = 2f + 1$

- cannot write on a register that is not its own
- cannot forge the signature of a correct process

Outline

- ① Algorithms for Consistent and Reliable Broadcast
 - ▶ Signature-free in well-behaved executions

Outline

① Algorithms for Consistent and Reliable Broadcast

- ▶ Signature-free in well-behaved executions

② Lower bounds for Consistent and Reliable Broadcast

Consistent Broadcast	Reliable Broadcast
1	$O(n)$

Table: Total number of signatures created by correct processes

Outline

① Algorithms for Consistent and Reliable Broadcast

- ▶ Signature-free in well-behaved executions

② Lower bounds for Consistent and Reliable Broadcast

Consistent Broadcast	Reliable Broadcast
1	$O(n)$

Table: Total number of signatures created by correct processes

③ Consensus protocol using Consistent Broadcast

Process roles

Primitives: $broadcast(m)$ and $deliver(m)$

Process roles

Primitives: $broadcast(m)$ and $deliver(m)$

Sender s - the process that invokes $broadcast(m)$

Process roles

Primitives: $broadcast(m)$ and $deliver(m)$

Sender s - the process that invokes $broadcast(m)$

Replicator r - the process that ensures broadcast properties are satisfied
(e.g., replicates messages)

Process roles

Primitives: $broadcast(m)$ and $deliver(m)$

Sender s - the process that invokes $broadcast(m)$

Replicator r - the process that ensures broadcast properties are satisfied
(e.g., replicates messages)

Receiver p - the process that invokes $deliver(m)$

Process roles

Primitives: $broadcast(m)$ and $deliver(m)$

Sender s - the process that invokes $broadcast(m)$

Replicator r - the process that ensures broadcast properties are satisfied
(e.g., replicates messages)

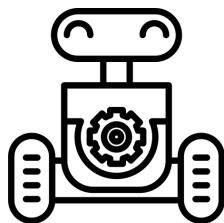
Receiver p - the process that invokes $deliver(m)$

n and f refer to the replicators

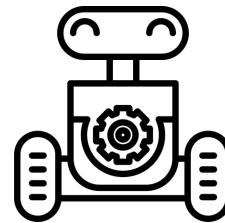
Consistent Broadcast

Validity

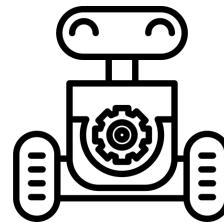
If a correct process s broadcasts m , then every correct process eventually delivers m



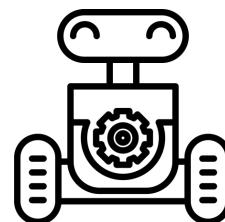
Sender s



Receiver p_0



Receiver p_1



Receiver p_2



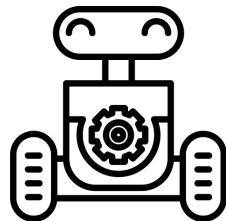
Receiver p_3

Consistent Broadcast

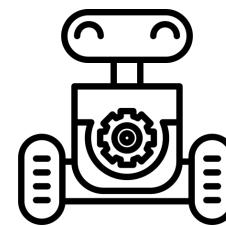
Validity

If a correct process s broadcasts m , then every correct process eventually delivers m

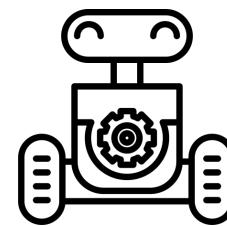
$broadcast(m)$



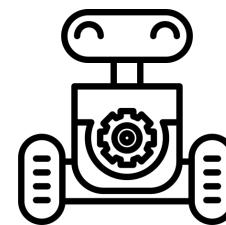
Sender s



Receiver p_0



Receiver p_1



Receiver p_2

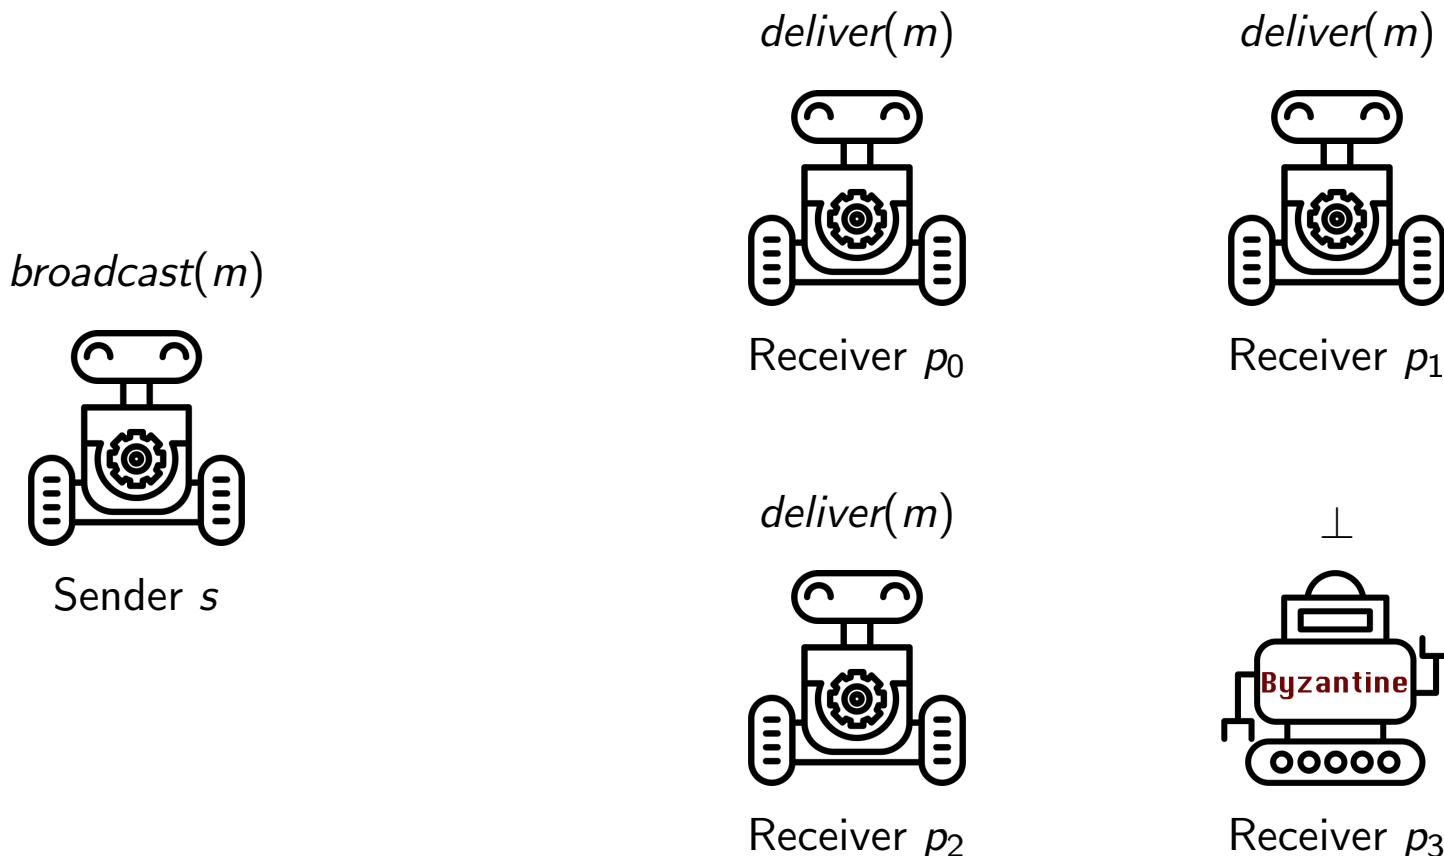


Receiver p_3

Consistent Broadcast

Validity

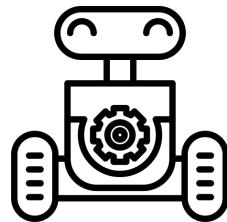
If a correct process s broadcasts m , then every correct process eventually delivers m



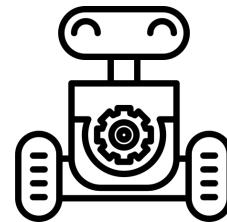
Consistent Broadcast

Consistency

If p and p' are correct processes, p delivers m , and p' delivers m' , then $m=m'$



Receiver p



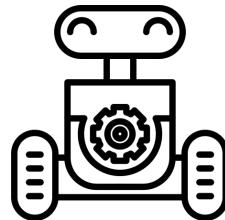
Receiver p'

Consistent Broadcast

Consistency

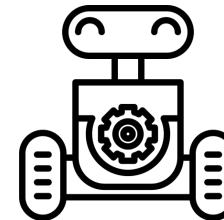
If p and p' are correct processes, p delivers m , and p' delivers m' , then $m=m'$

$deliver(m)$



Receiver p

$deliver(m')$



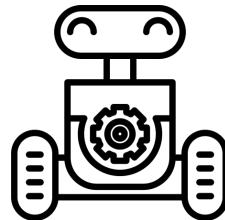
Receiver p'

Consistent Broadcast

Consistency

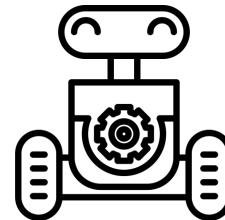
If p and p' are correct processes, p delivers m , and p' delivers m' , then $m=m'$

$deliver(m)$



Receiver p

$deliver(m)$

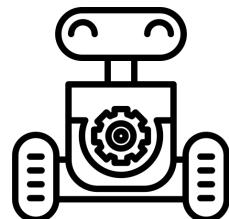


Receiver p'

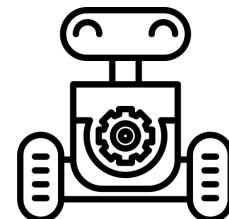
Consistent Broadcast

Integrity

If some correct process delivers m and s is correct, then s previously broadcast m



Sender s

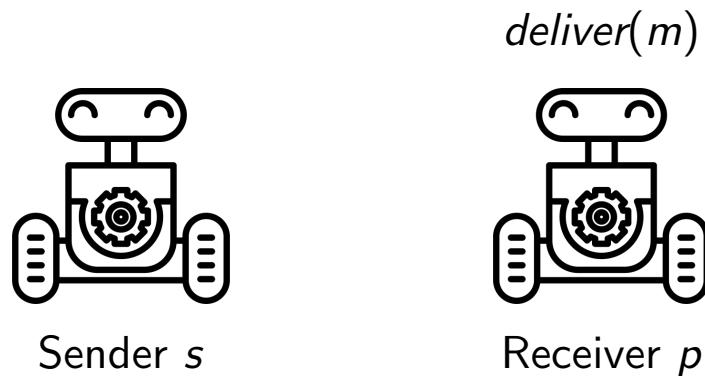


Receiver p

Consistent Broadcast

Integrity

If some correct process delivers m and s is correct, then s previously broadcast m

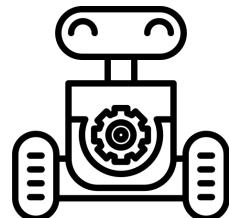


Consistent Broadcast

Integrity

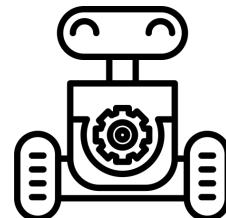
If some correct process delivers m and s is correct, then s previously broadcast m

$broadcast(m)$



Sender s

$deliver(m)$



Receiver p

Consistent Broadcast

Validity - If a correct process s broadcasts m , then every correct process eventually delivers m

Consistency - If p and p' are correct processes, p delivers m , and p' delivers m' , then $m=m'$

Integrity - If some correct process delivers m and s is correct, then s previously broadcast m

Consistent Broadcast

Algorithm sketch, $f = 1$. Fast path

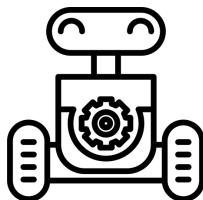
$broadcast(m)$



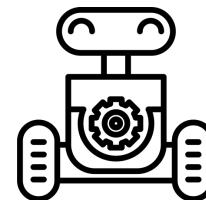
Sender s



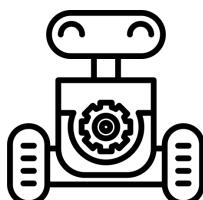
Replicator r_0



Replicator r_1



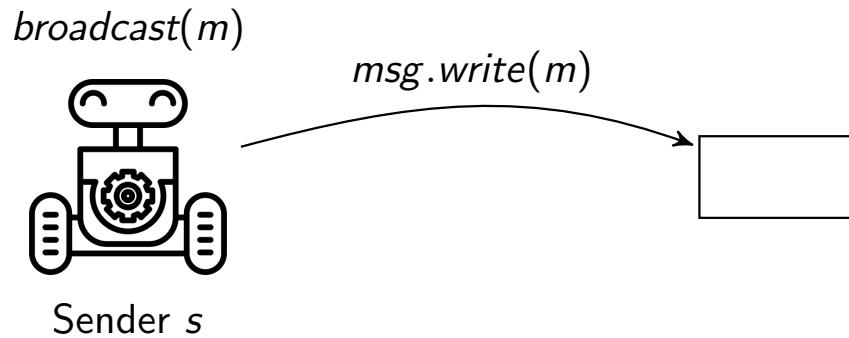
Replicator r_2



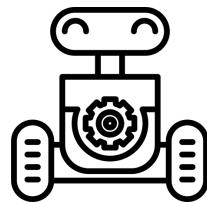
Receiver p

Consistent Broadcast

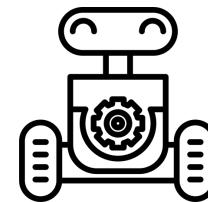
Algorithm sketch, $f = 1$. Fast path



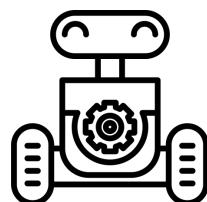
Replicator r_0



Replicator r_1



Replicator r_2

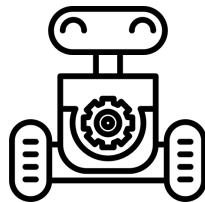


Receiver p

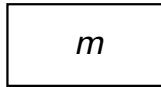
Consistent Broadcast

Algorithm sketch, $f = 1$. Fast path

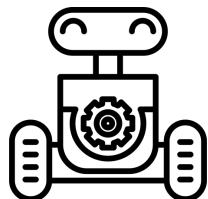
$broadcast(m)$



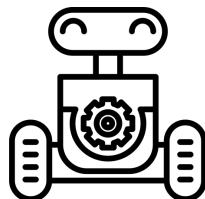
Sender s



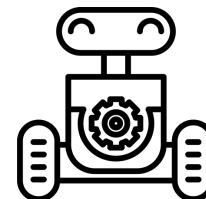
m



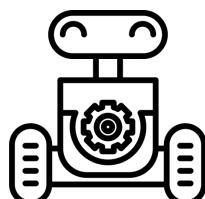
Replicator r_0



Replicator r_1



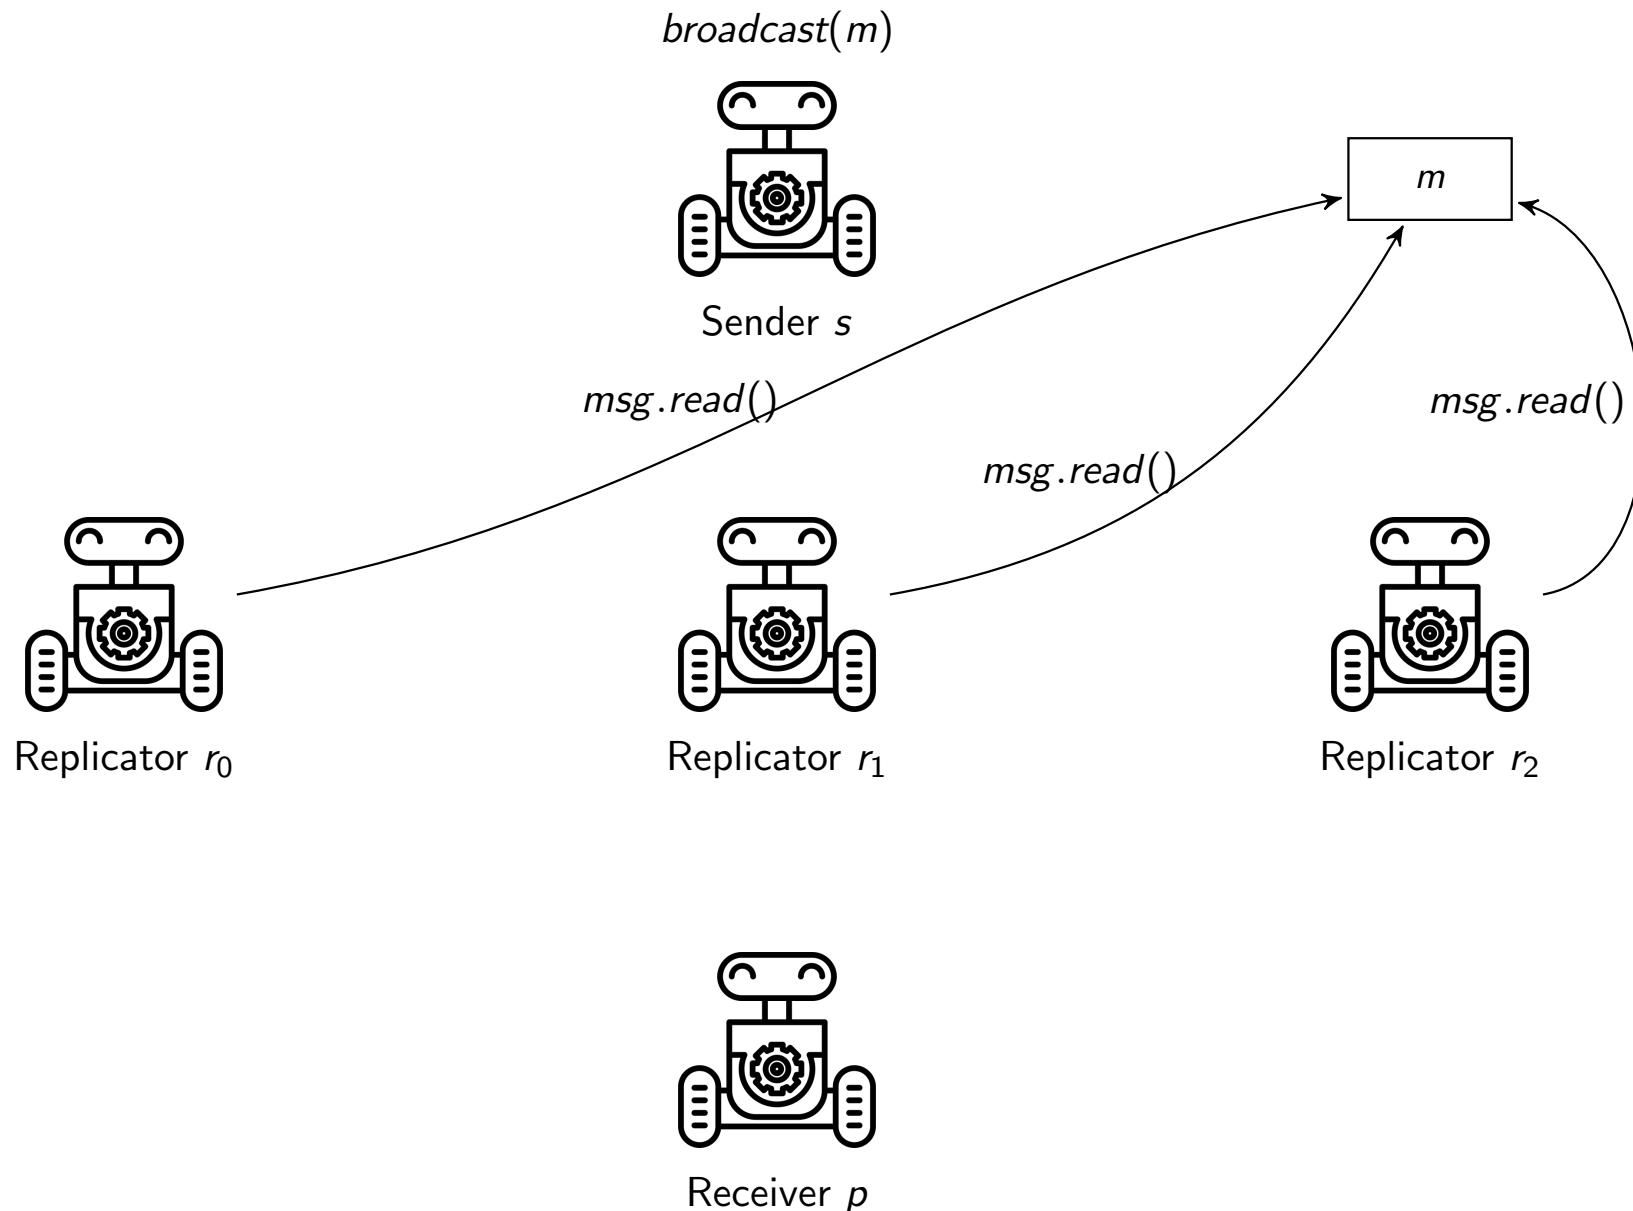
Replicator r_2



Receiver p

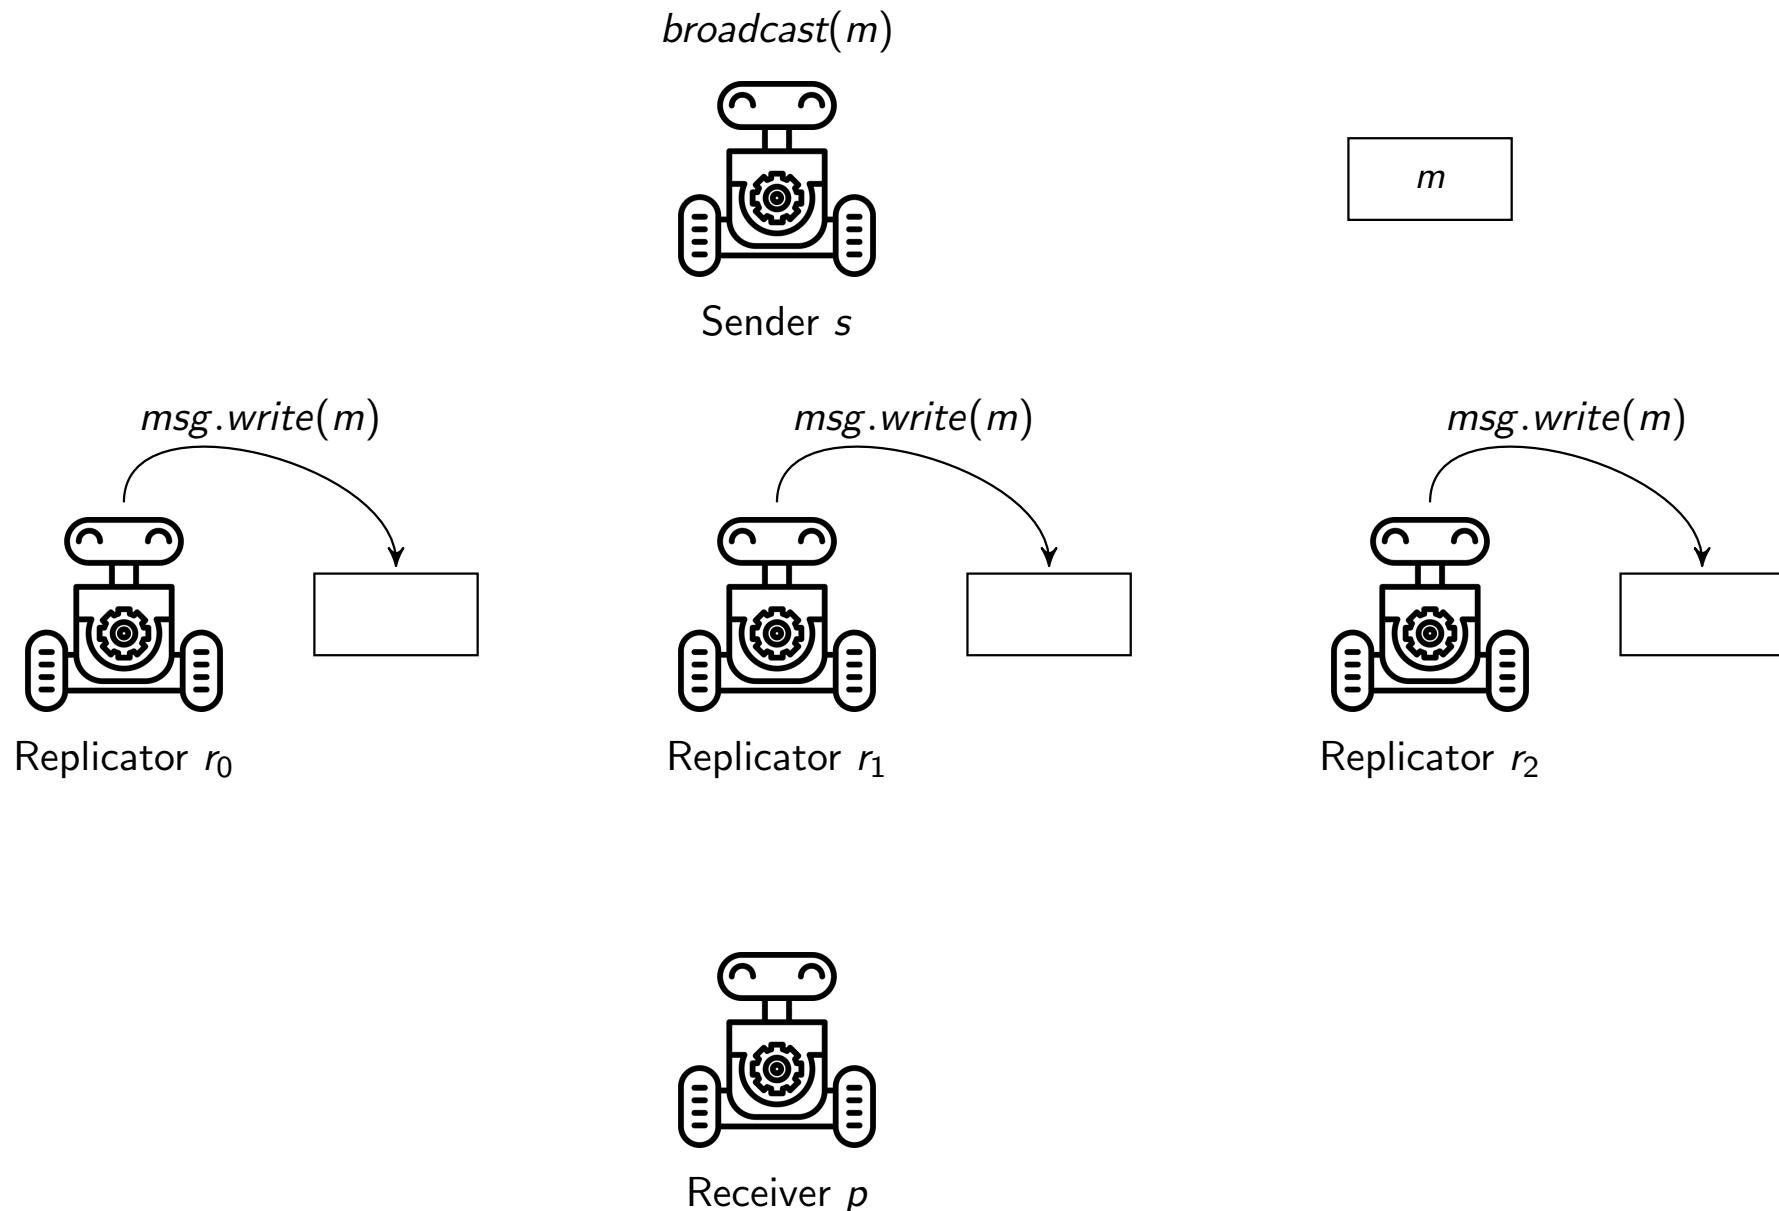
Consistent Broadcast

Algorithm sketch, $f = 1$. Fast path



Consistent Broadcast

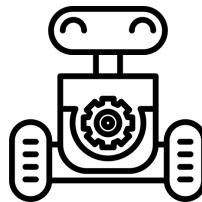
Algorithm sketch, $f = 1$. Fast path



Consistent Broadcast

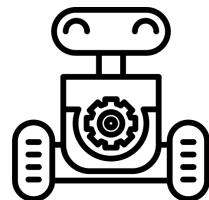
Algorithm sketch, $f = 1$. Fast path

$broadcast(m)$



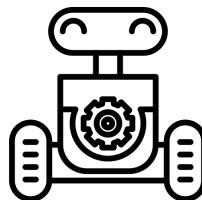
Sender s

m



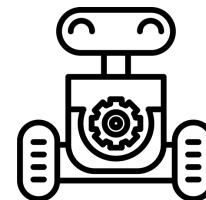
Replicator r_0

m

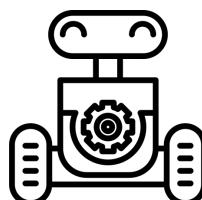


Replicator r_1

m



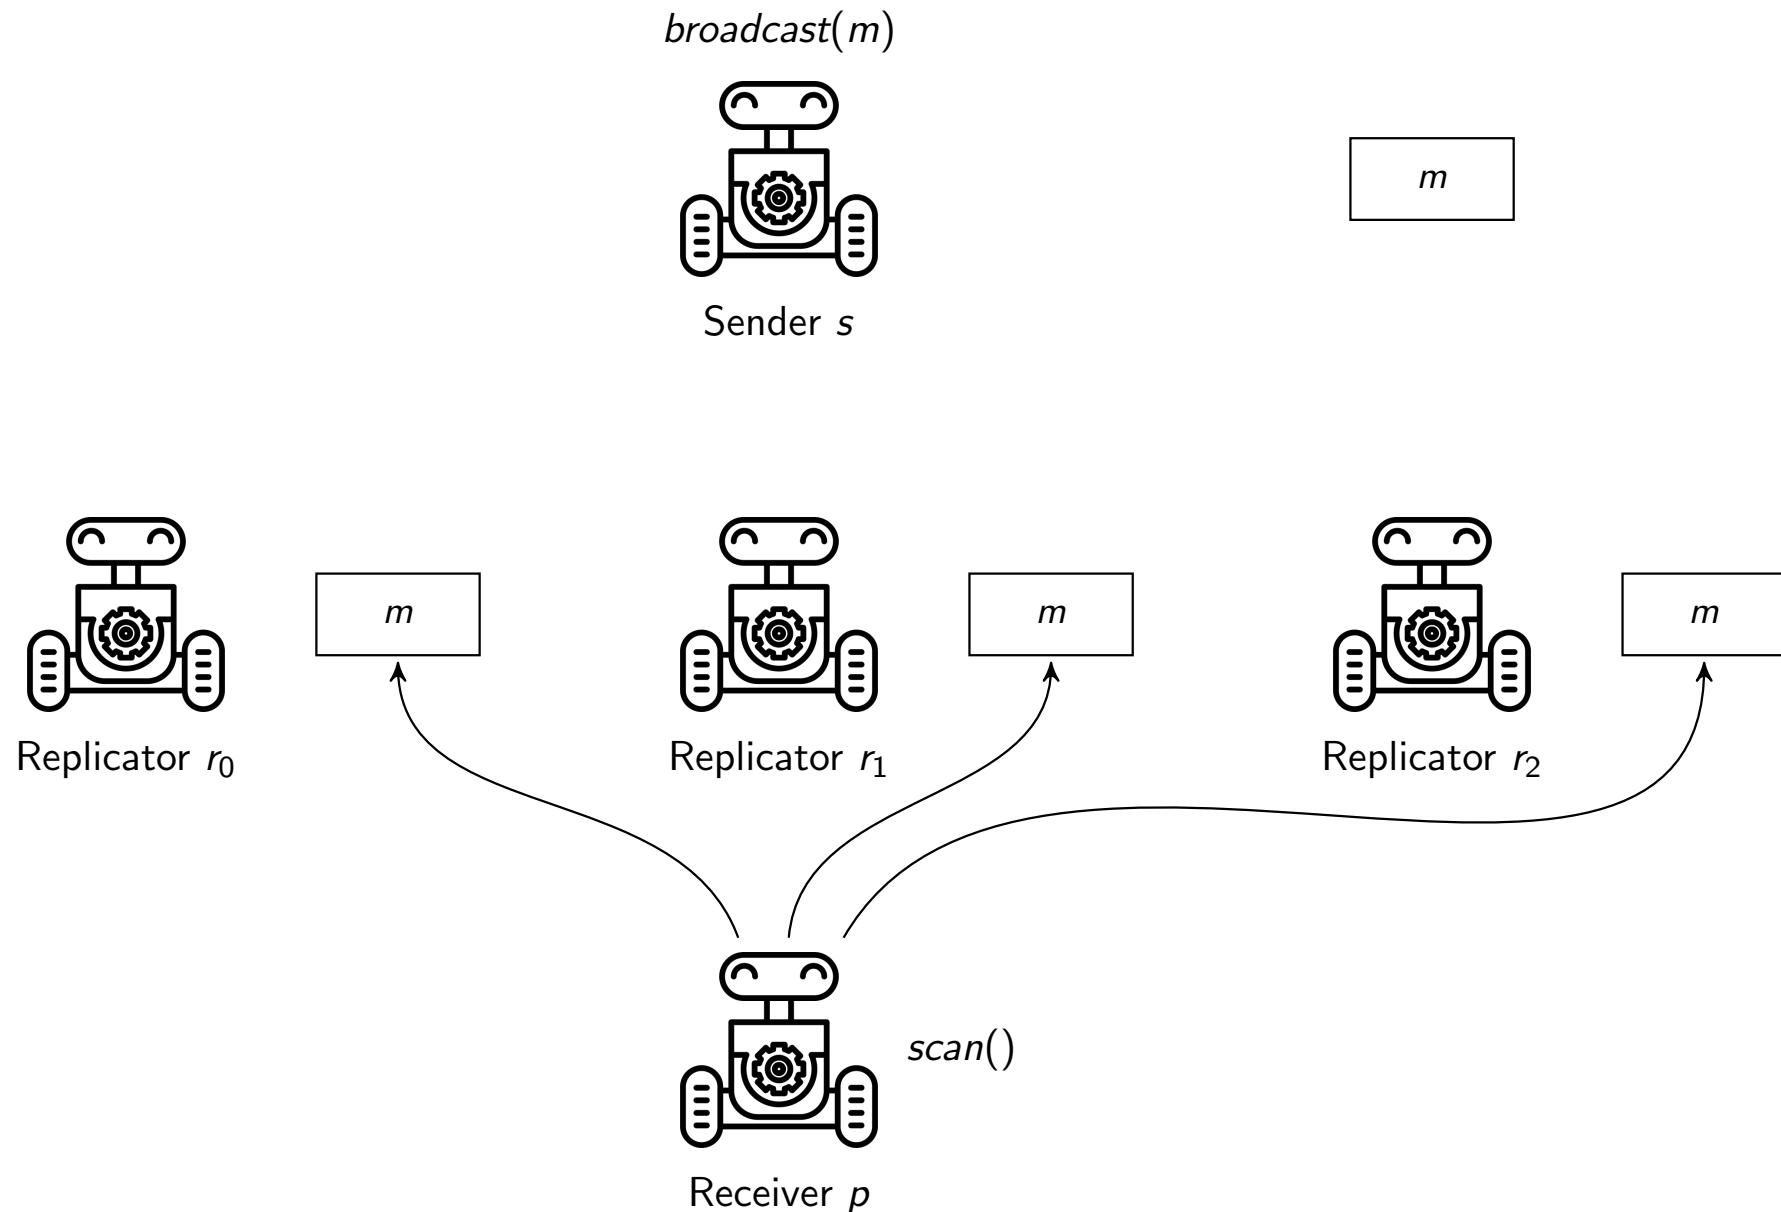
Replicator r_2



Receiver p

Consistent Broadcast

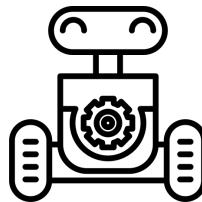
Algorithm sketch, $f = 1$. Fast path



Consistent Broadcast

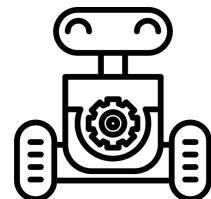
Algorithm sketch, $f = 1$. Fast path

$broadcast(m)$



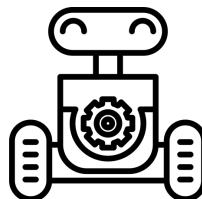
Sender s

m



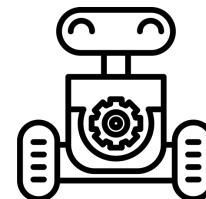
Replicator r_0

m

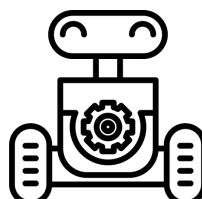


Replicator r_1

m



Replicator r_2



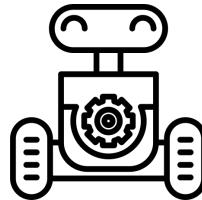
Receiver p

unanimity \implies deliver m via *fast path*

Consistent Broadcast

Algorithm sketch, $f = 1$

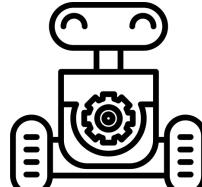
$broadcast(m)$



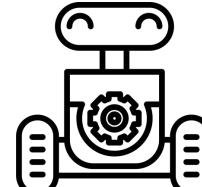
Sender s



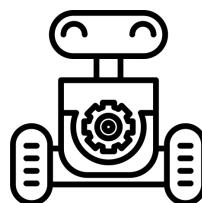
Replicator r_0



Replicator r_1



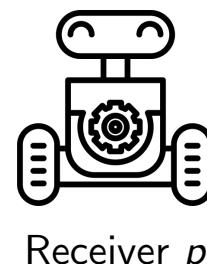
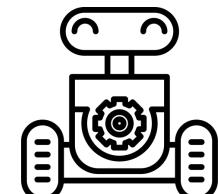
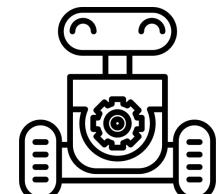
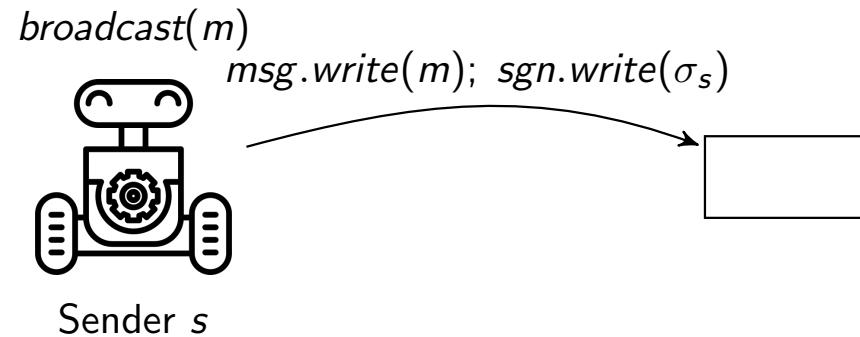
Replicator r_2



Receiver p

Consistent Broadcast

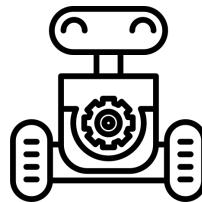
Algorithm sketch, $f = 1$



Consistent Broadcast

Algorithm sketch, $f = 1$

$broadcast(m)$

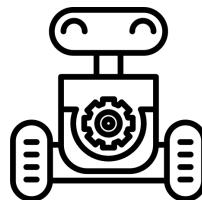


Sender s

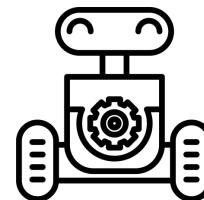
$m \mid \sigma_s$



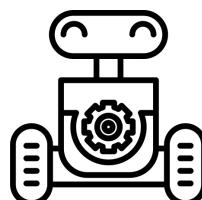
Replicator r_0



Replicator r_1



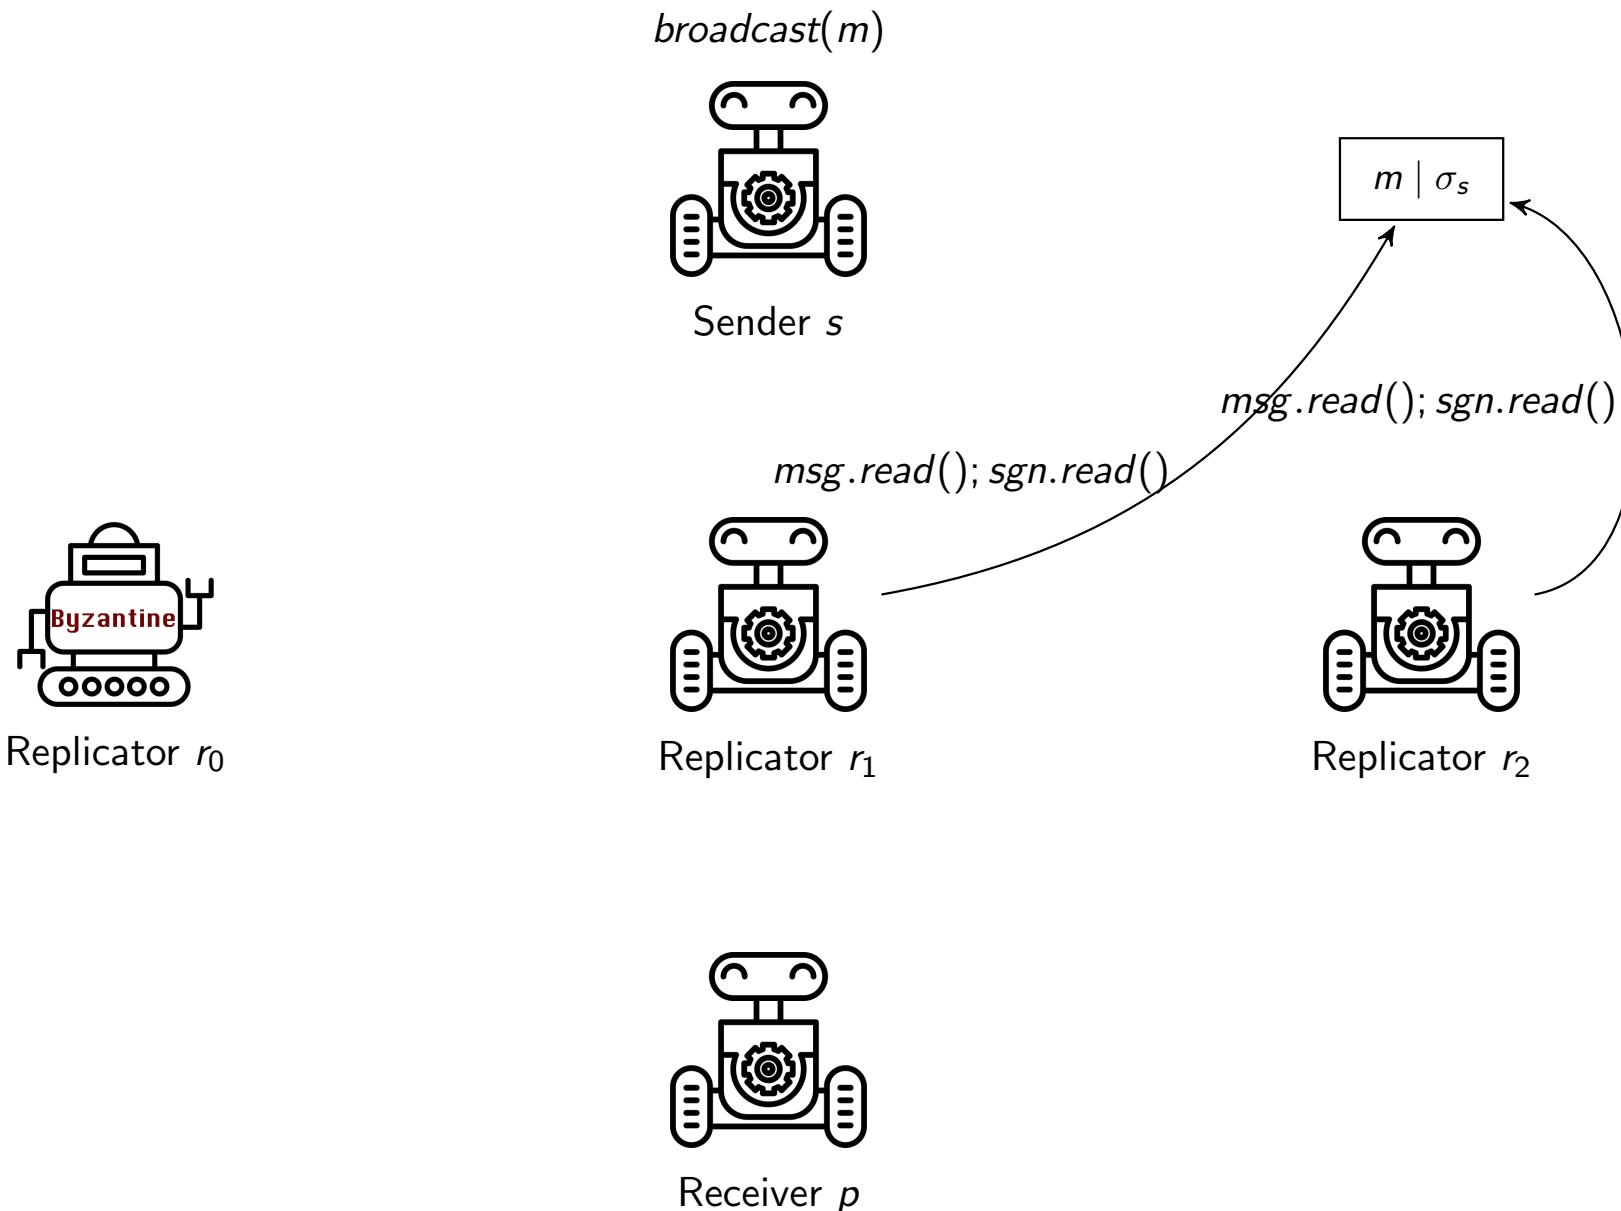
Replicator r_2



Receiver p

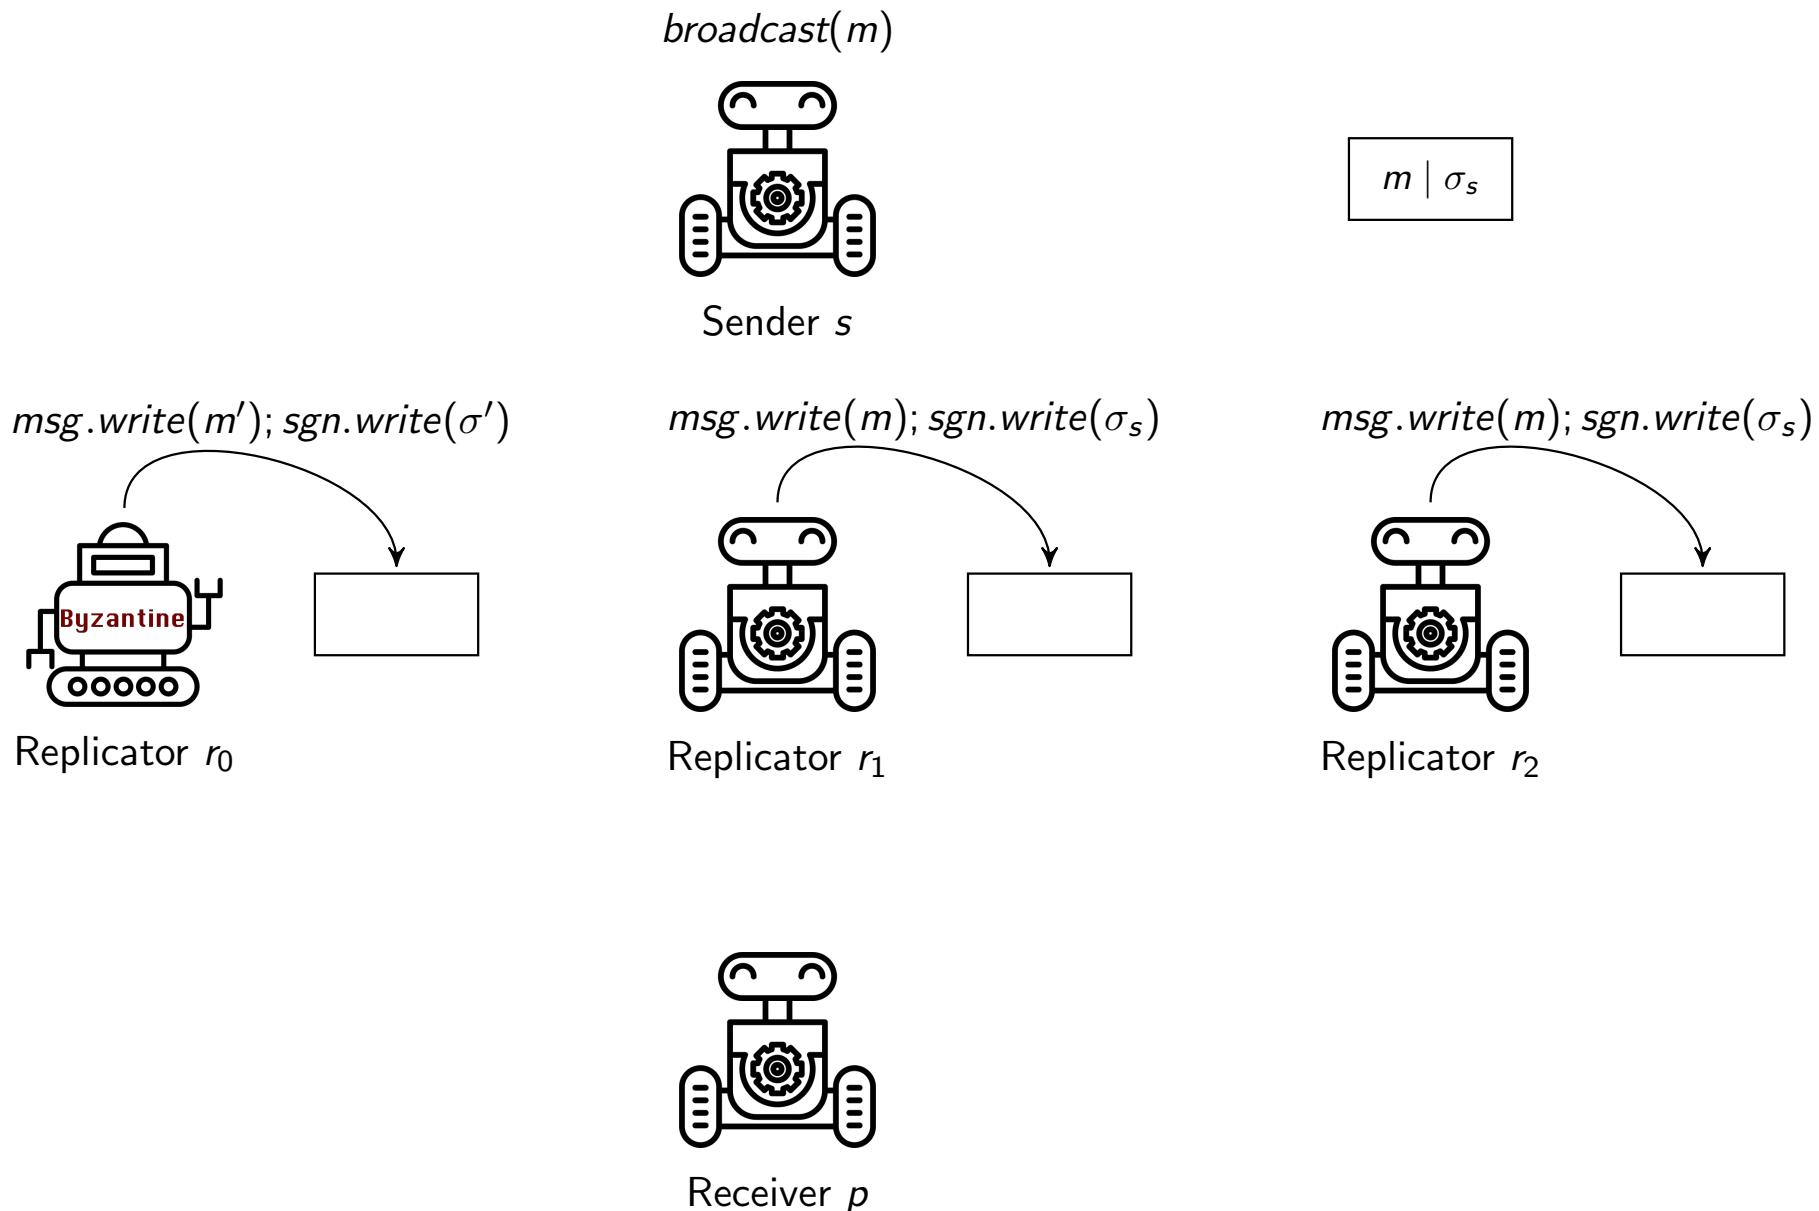
Consistent Broadcast

Algorithm sketch, $f = 1$



Consistent Broadcast

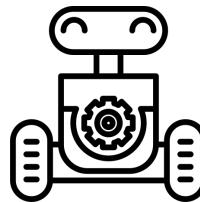
Algorithm sketch, $f = 1$



Consistent Broadcast

Algorithm sketch, $f = 1$

$broadcast(m)$



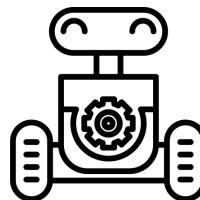
Sender s

$m \mid \sigma_s$



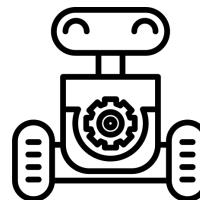
Replicator r_0

$m' \mid \sigma'$



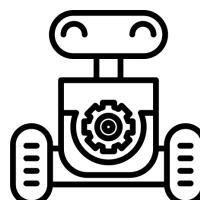
Replicator r_1

$m \mid \sigma_s$



Replicator r_2

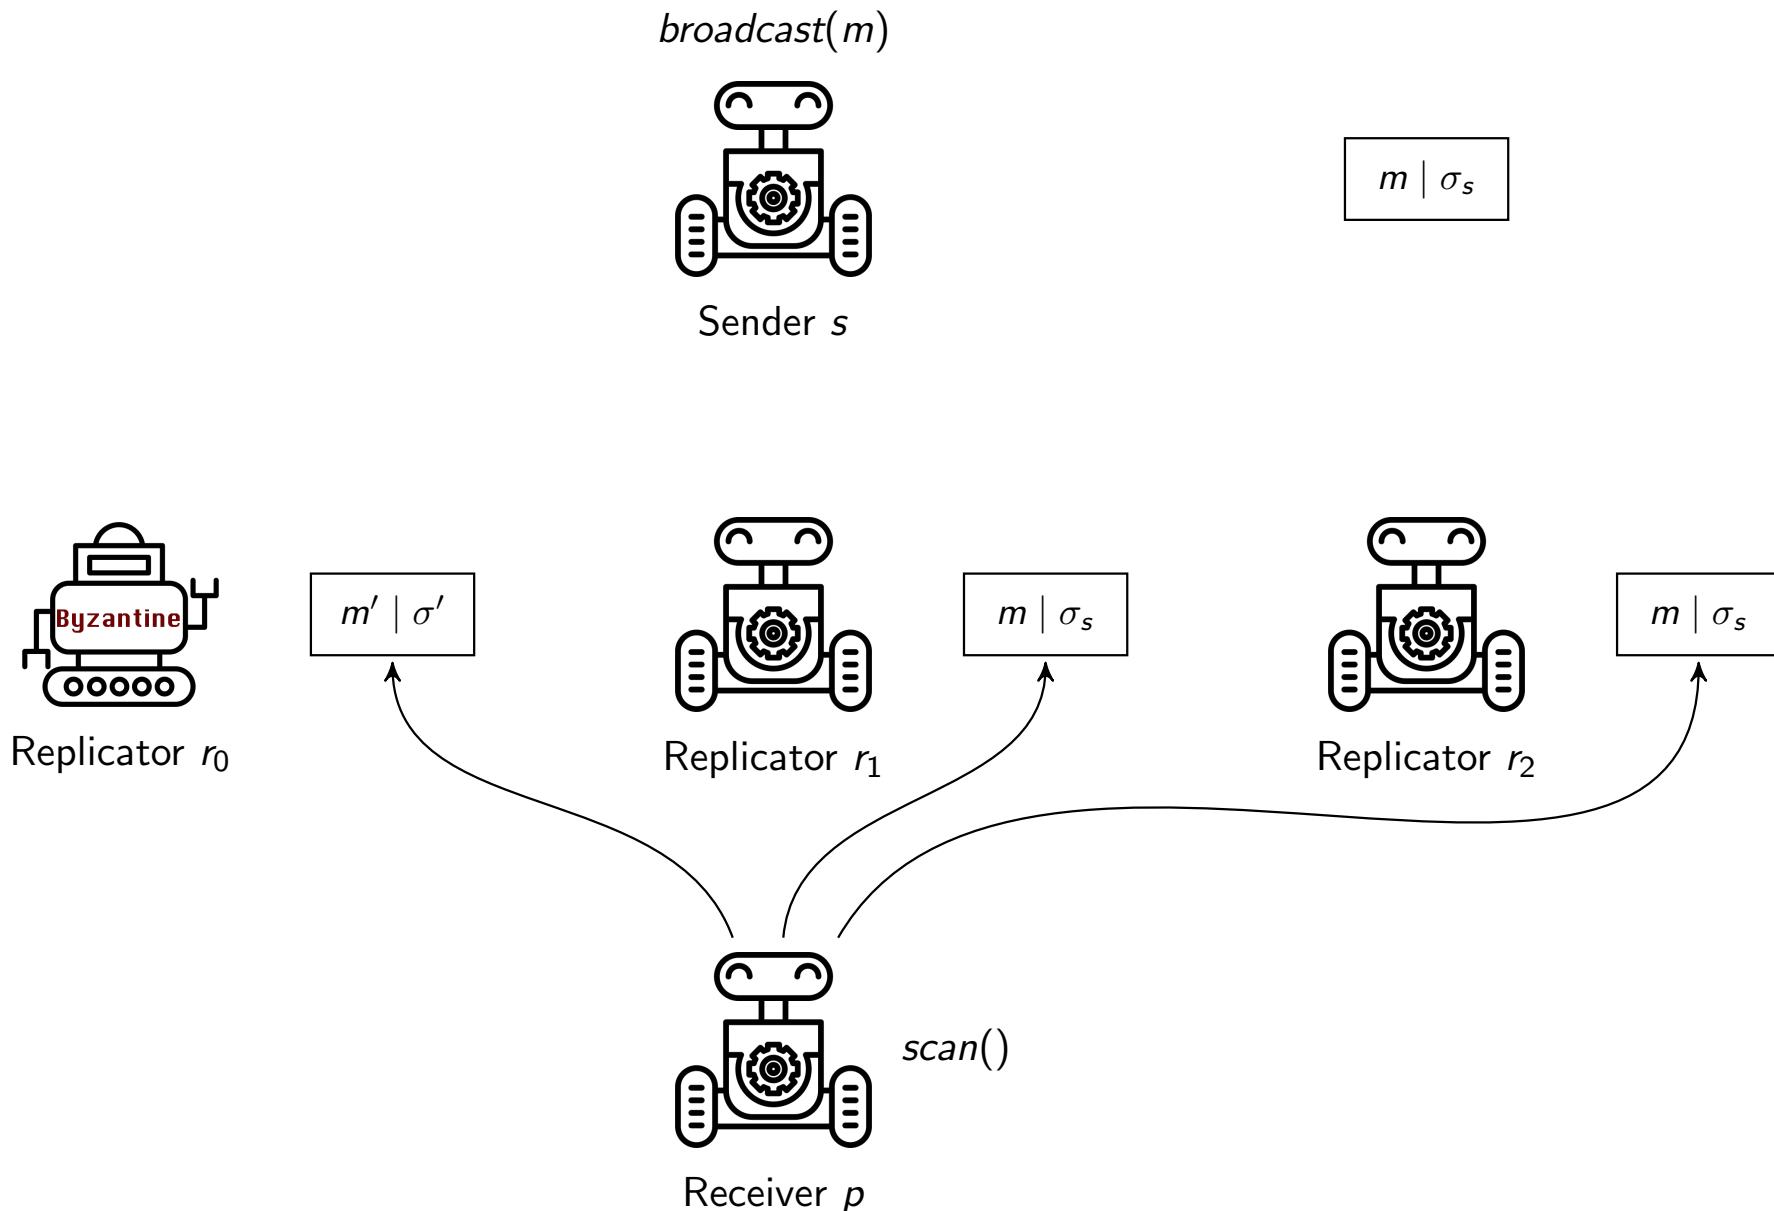
$m \mid \sigma_s$



Receiver p

Consistent Broadcast

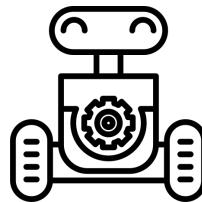
Algorithm sketch, $f = 1$



Consistent Broadcast

Algorithm sketch, $f = 1$

$broadcast(m)$



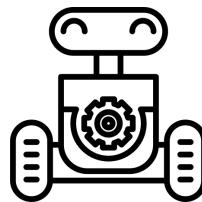
Sender s

$m \mid \sigma_s$



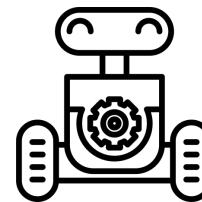
Replicator r_0

$m' \mid \sigma'$



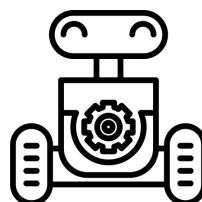
Replicator r_1

$m \mid \sigma_s$



Replicator r_2

$m \mid \sigma_s$



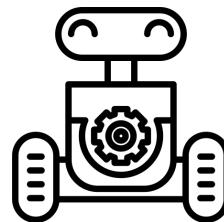
Receiver p

$n - f$ signed copies of m and
no $m' \neq m$ validly signed \implies deliver m

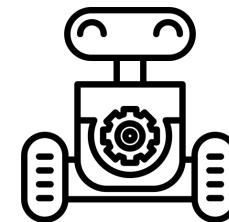
Reliable Broadcast

Same properties as Consistent Broadcast + Totality

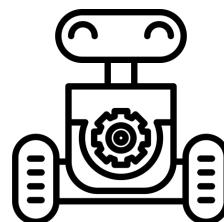
If some correct process delivers m , then every correct process eventually delivers a message



Receiver p_0



Receiver p_1



Receiver p_2

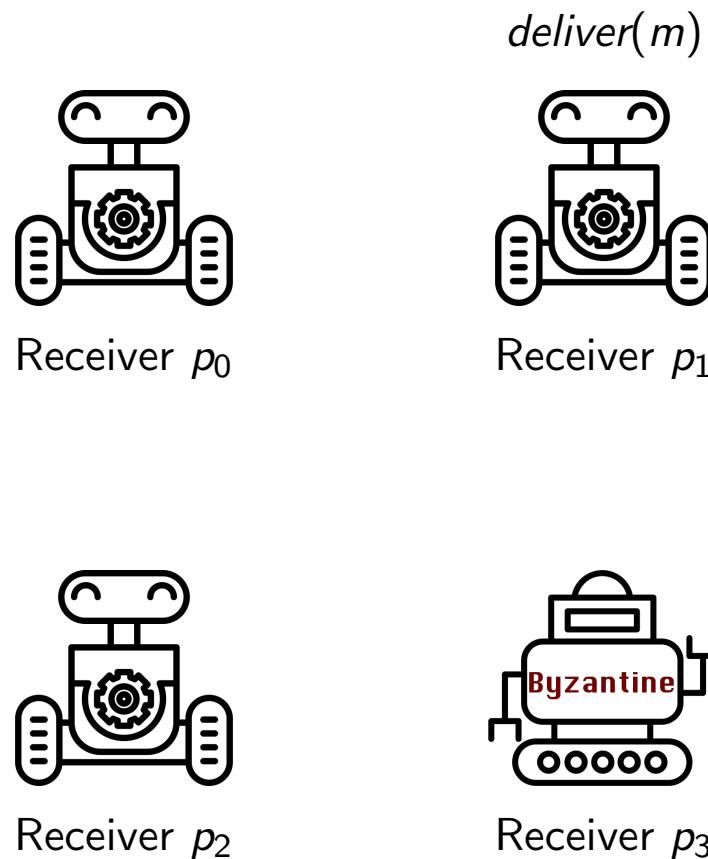


Receiver p_3

Reliable Broadcast

Same properties as Consistent Broadcast + Totality

If some correct process delivers m , then every correct process eventually delivers a message

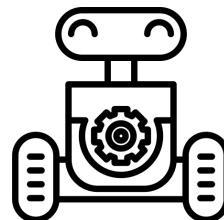


Reliable Broadcast

Same properties as Consistent Broadcast + Totality

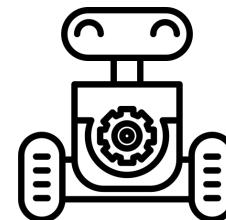
If some correct process delivers m , then every correct process eventually delivers a message

$deliver(m)$



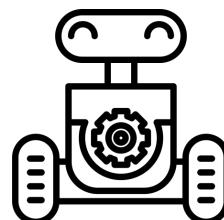
Receiver p_0

$deliver(m)$



Receiver p_1

$deliver(m)$



Receiver p_2

\perp



Receiver p_3

Reliable Broadcast

Validity - If a correct process s broadcasts m , then every correct process eventually delivers m

Consistency - If p and p' are correct processes, p delivers m , and p' delivers m' , then $m=m'$

Integrity - If some correct process delivers m and s is correct, then s previously broadcast m

Totality - If some correct process delivers m , then every correct process eventually delivers a message

Consistent Broadcast vs Reliable Broadcast

Consistent and Reliable Broadcast behave the same way when the sender s is **correct** (recall the *Validity* property)

Consistent Broadcast vs Reliable Broadcast

Consistent and Reliable Broadcast behave the same way when the sender *s* is **correct** (recall the *Validity* property)

Yet when the sender is **faulty** ...

Consistent Broadcast vs Reliable Broadcast

Consistent and Reliable Broadcast behave the same way when the sender *s* is **correct** (recall the *Validity* property)

Yet when the sender is **faulty** ...

- Consistent Broadcast has no delivery guarantees: some correct processes may deliver a message, others may not

Consistent Broadcast vs Reliable Broadcast

Consistent and Reliable Broadcast behave the same way when the sender **is correct** (recall the *Validity* property)

Yet when the sender is **faulty** ...

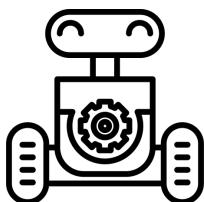
- Consistent Broadcast has no delivery guarantees: some correct processes may deliver a message, others may not
- while Reliable Broadcast guarantees every correct process eventually delivers a message as soon as one correct process delivered

Consistent Broadcast algorithm = \neg sufficient

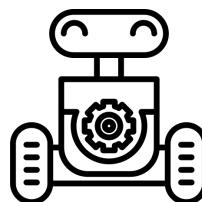
$broadcast(m)$



Sender s



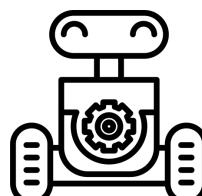
Replicator r_0



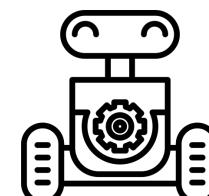
Replicator r_1



Replicator r_2

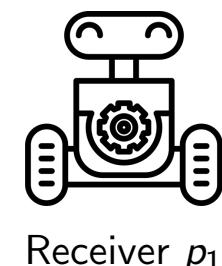
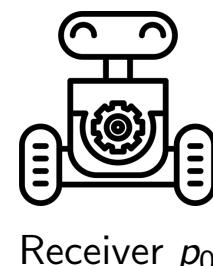
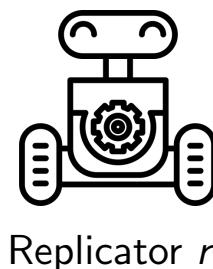
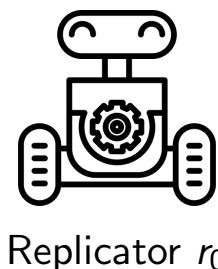
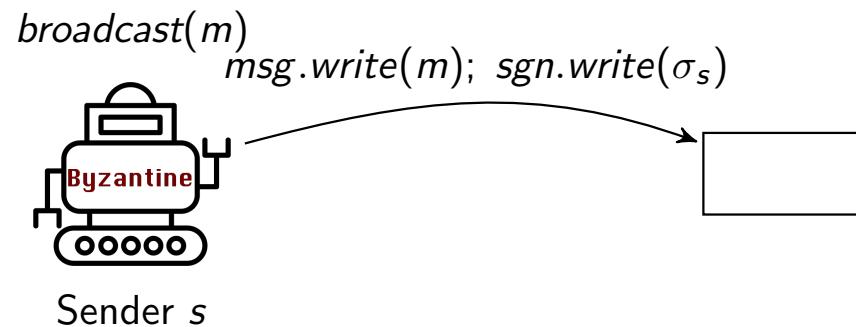


Receiver p_0



Receiver p_1

Consistent Broadcast algorithm = \neg sufficient



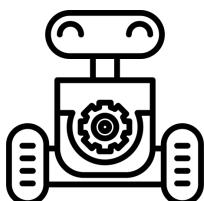
Consistent Broadcast algorithm = \neg sufficient

$broadcast(m)$

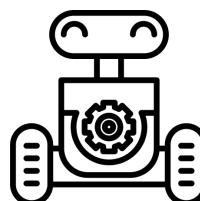


Sender s

$m \mid \sigma_s$



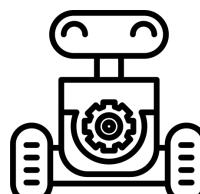
Replicator r_0



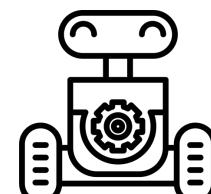
Replicator r_1



Replicator r_2

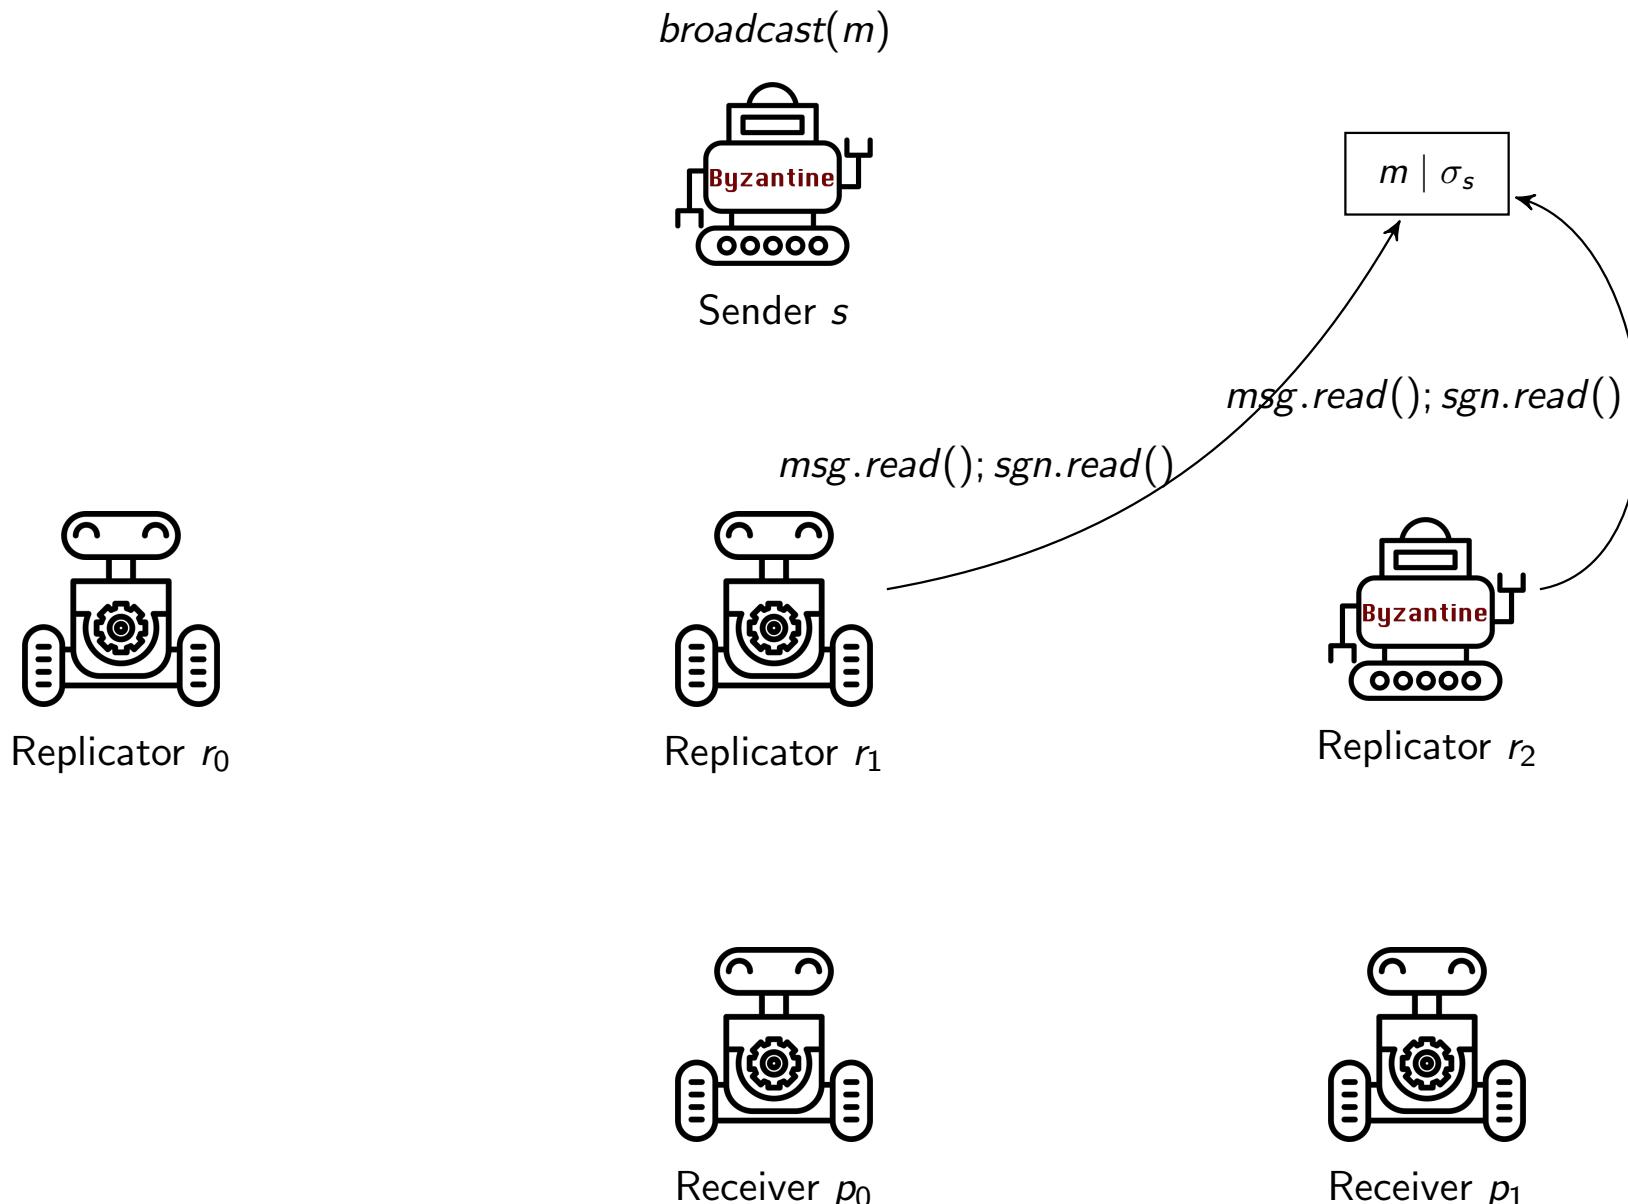


Receiver p_0

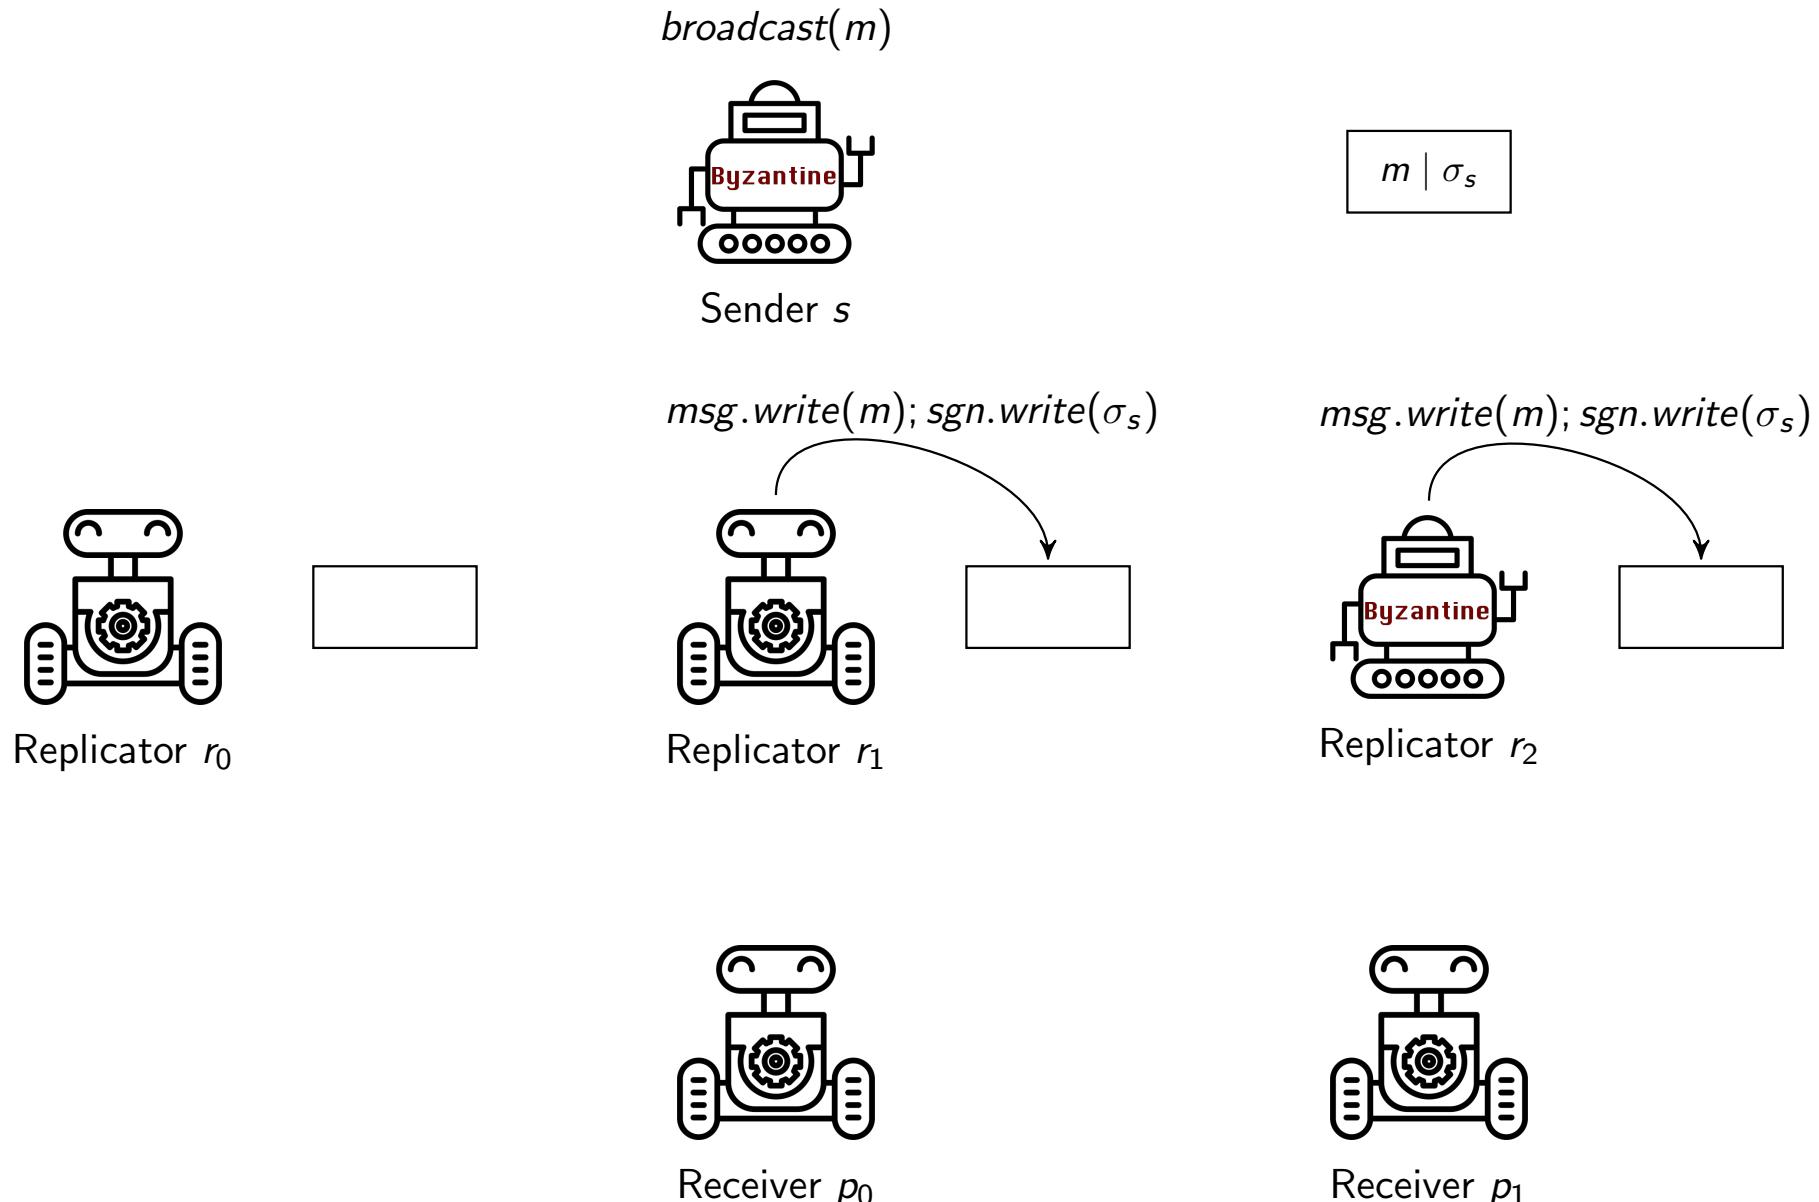


Receiver p_1

Consistent Broadcast algorithm = \neg sufficient



Consistent Broadcast algorithm = \neg sufficient



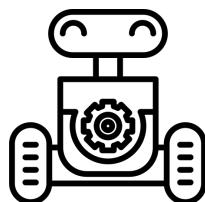
Consistent Broadcast algorithm = \neg sufficient

$broadcast(m)$

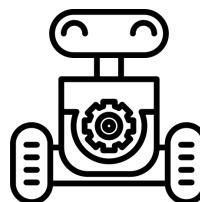


Sender s

$m \mid \sigma_s$



Replicator r_0



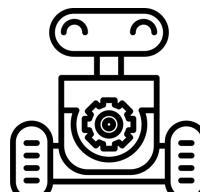
Replicator r_1

$m \mid \sigma_s$

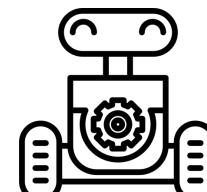


Replicator r_2

$m \mid \sigma_s$

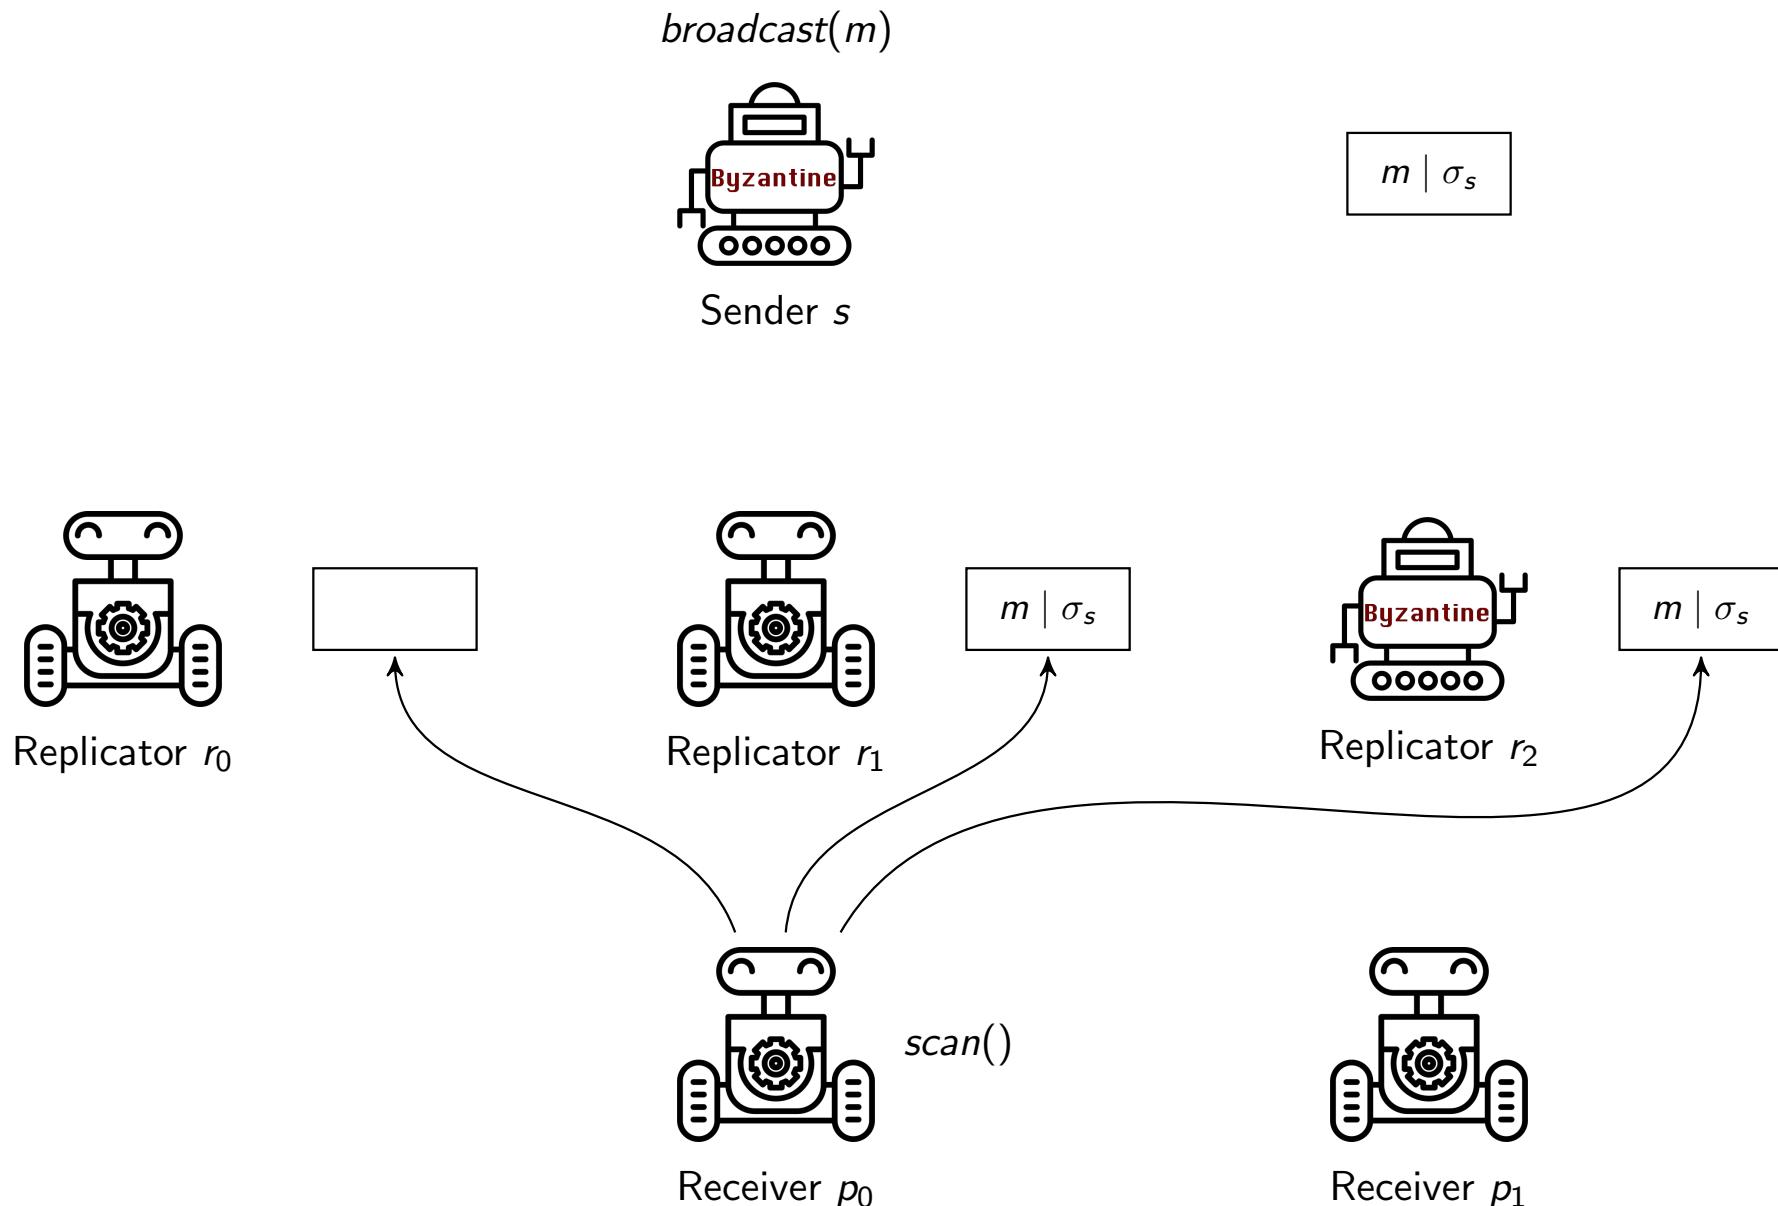


Receiver p_0



Receiver p_1

Consistent Broadcast algorithm = \neg sufficient



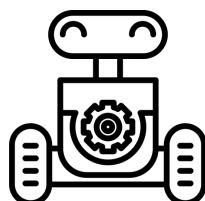
Consistent Broadcast algorithm = \neg sufficient

$broadcast(m)$

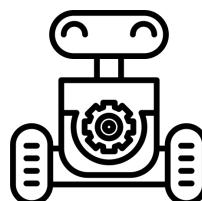


Sender s

$m \mid \sigma_s$



Replicator r_0



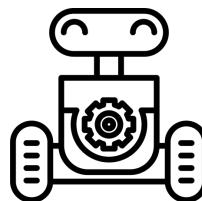
Replicator r_1

$m \mid \sigma_s$



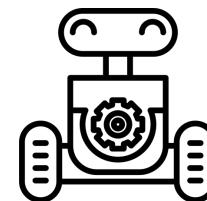
Replicator r_2

$m \mid \sigma_s$



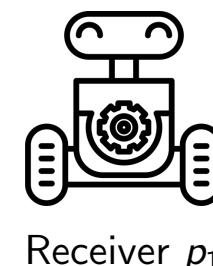
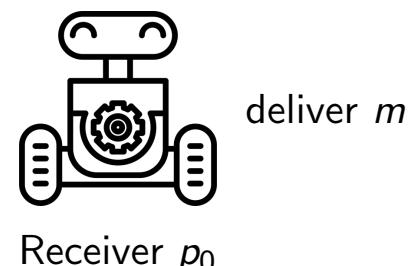
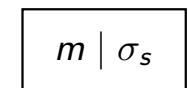
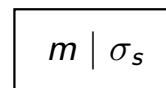
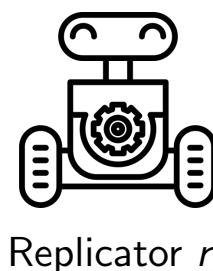
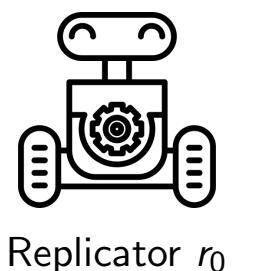
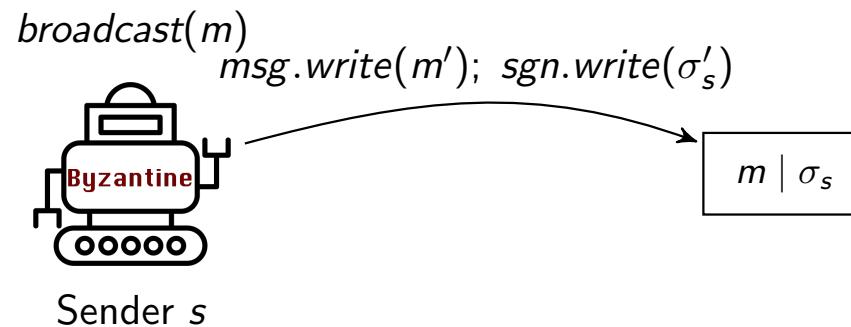
deliver m

Receiver p_0



Receiver p_1

Consistent Broadcast algorithm = \neg sufficient



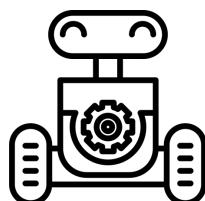
Consistent Broadcast algorithm = \neg sufficient

$broadcast(m)$

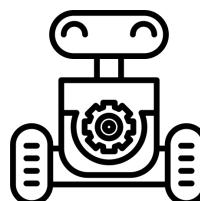


Sender s

$m' \mid \sigma'_s$



Replicator r_0



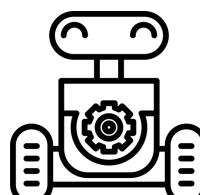
Replicator r_1

$m \mid \sigma_s$



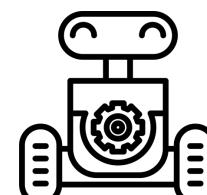
Replicator r_2

$m \mid \sigma_s$



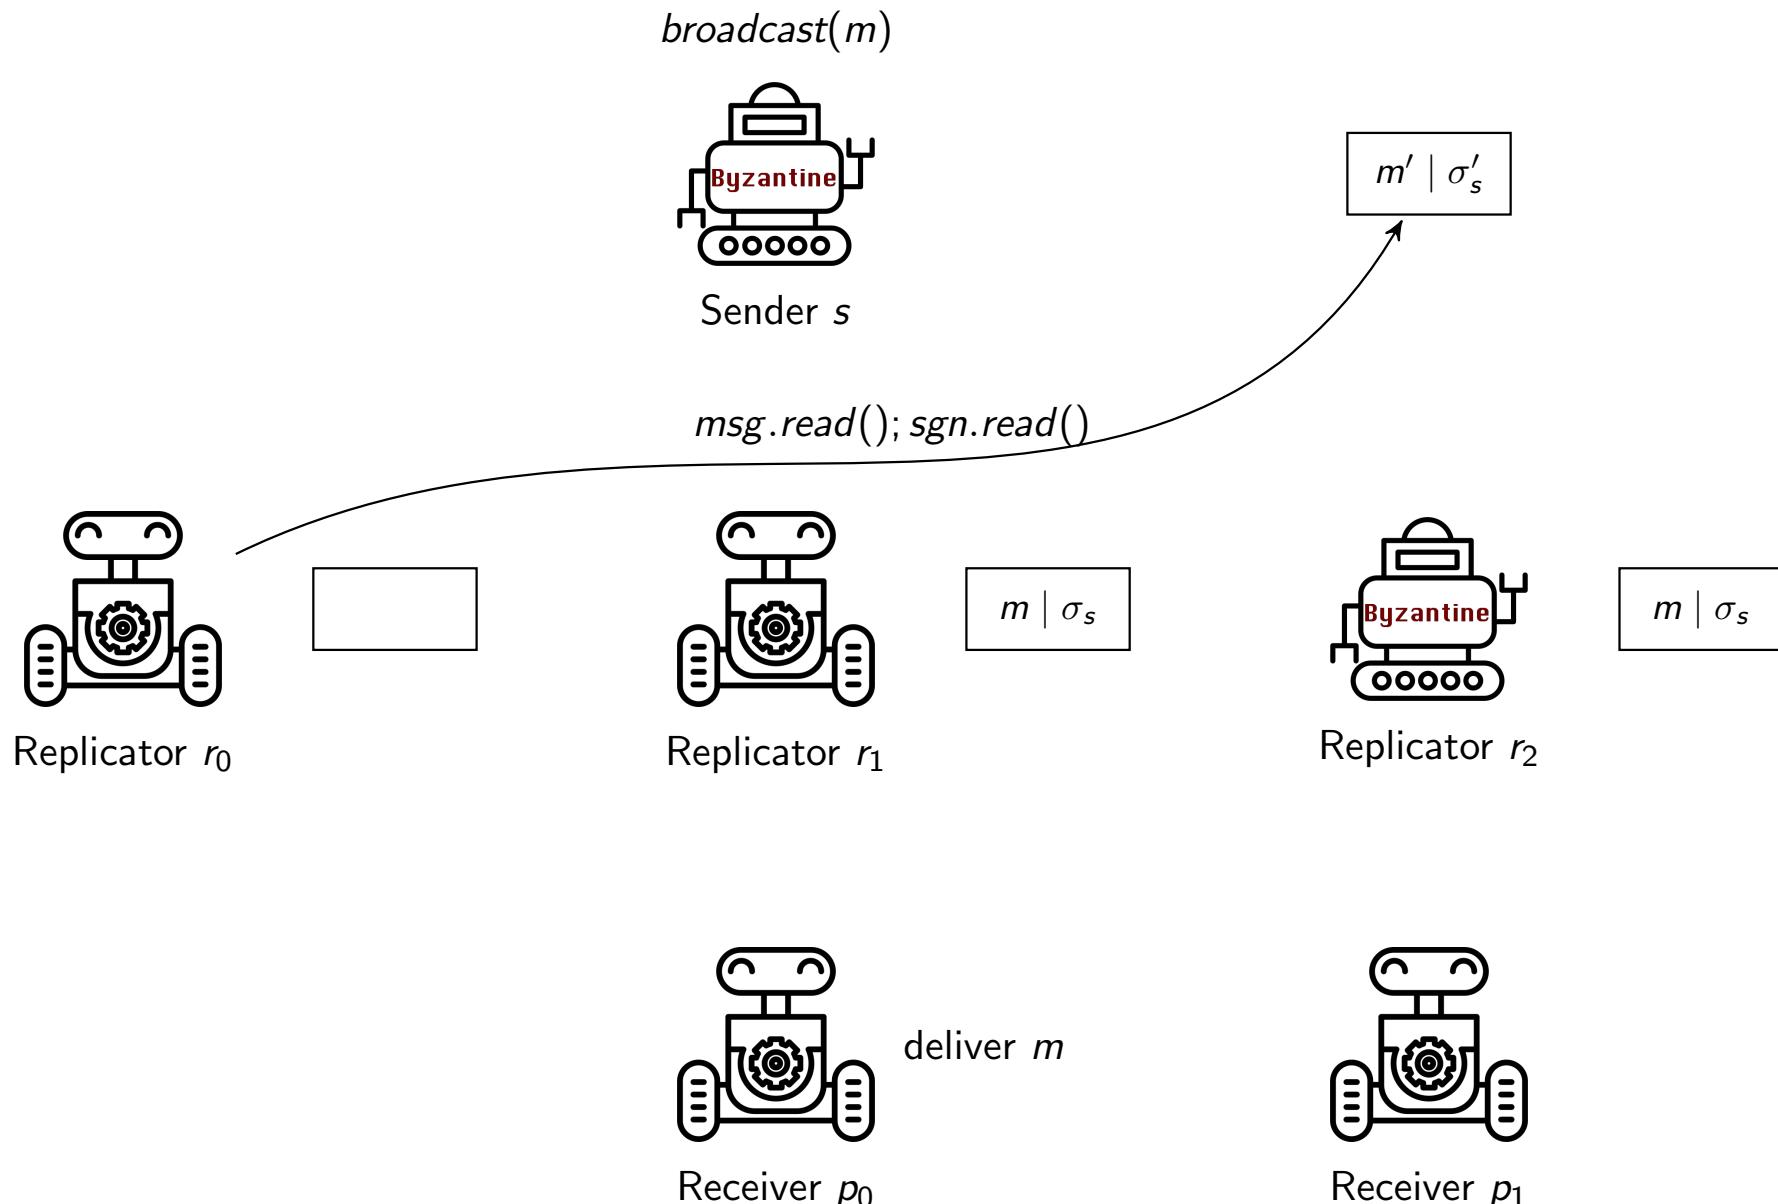
deliver m

Receiver p_0

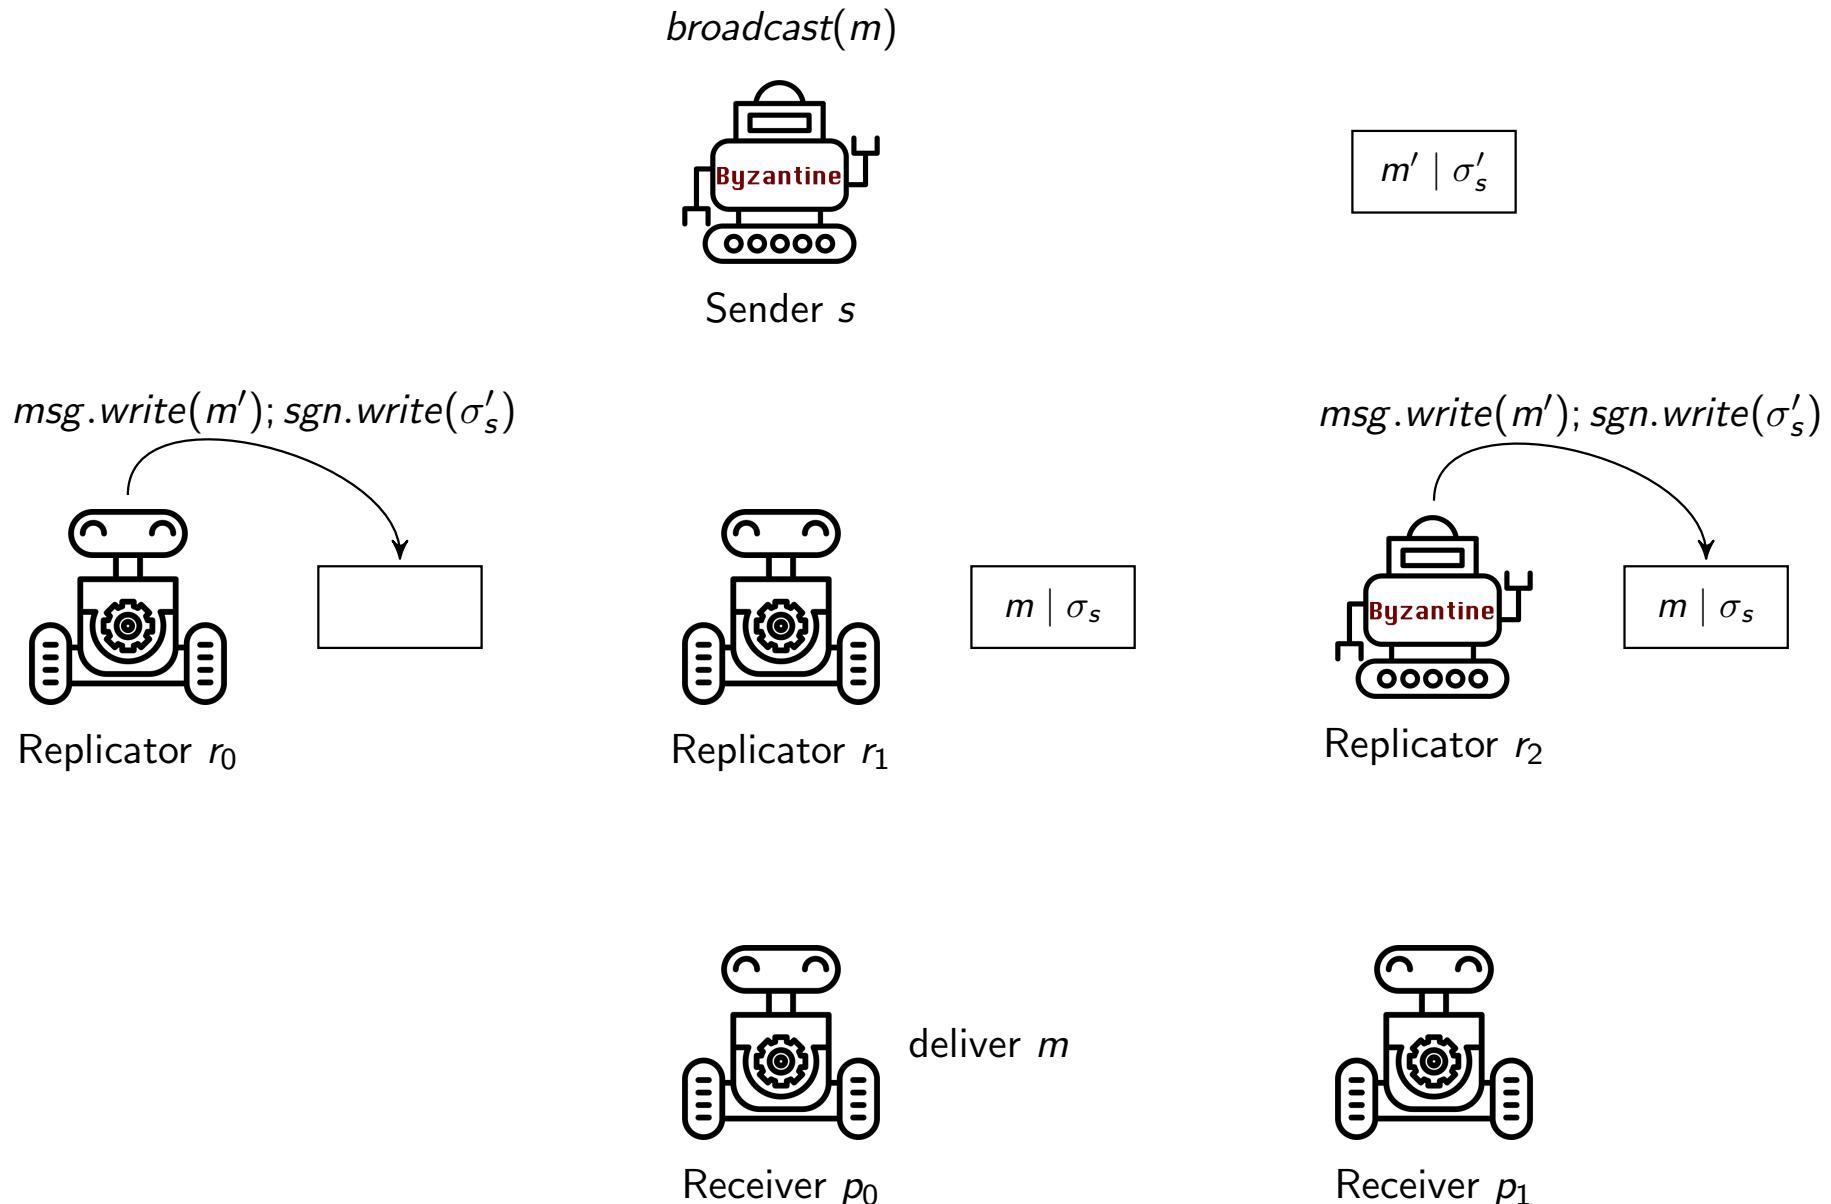


Receiver p_1

Consistent Broadcast algorithm = \neg sufficient



Consistent Broadcast algorithm = \neg sufficient



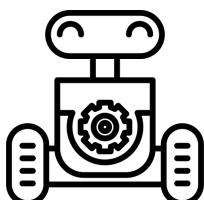
Consistent Broadcast algorithm = \neg sufficient

$broadcast(m)$



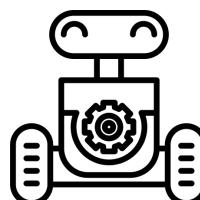
Sender s

$m' \mid \sigma'_s$



Replicator r_0

$m' \mid \sigma'_s$



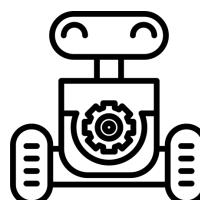
Replicator r_1

$m \mid \sigma_s$



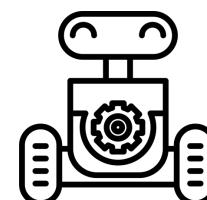
Replicator r_2

$m' \mid \sigma'_s$



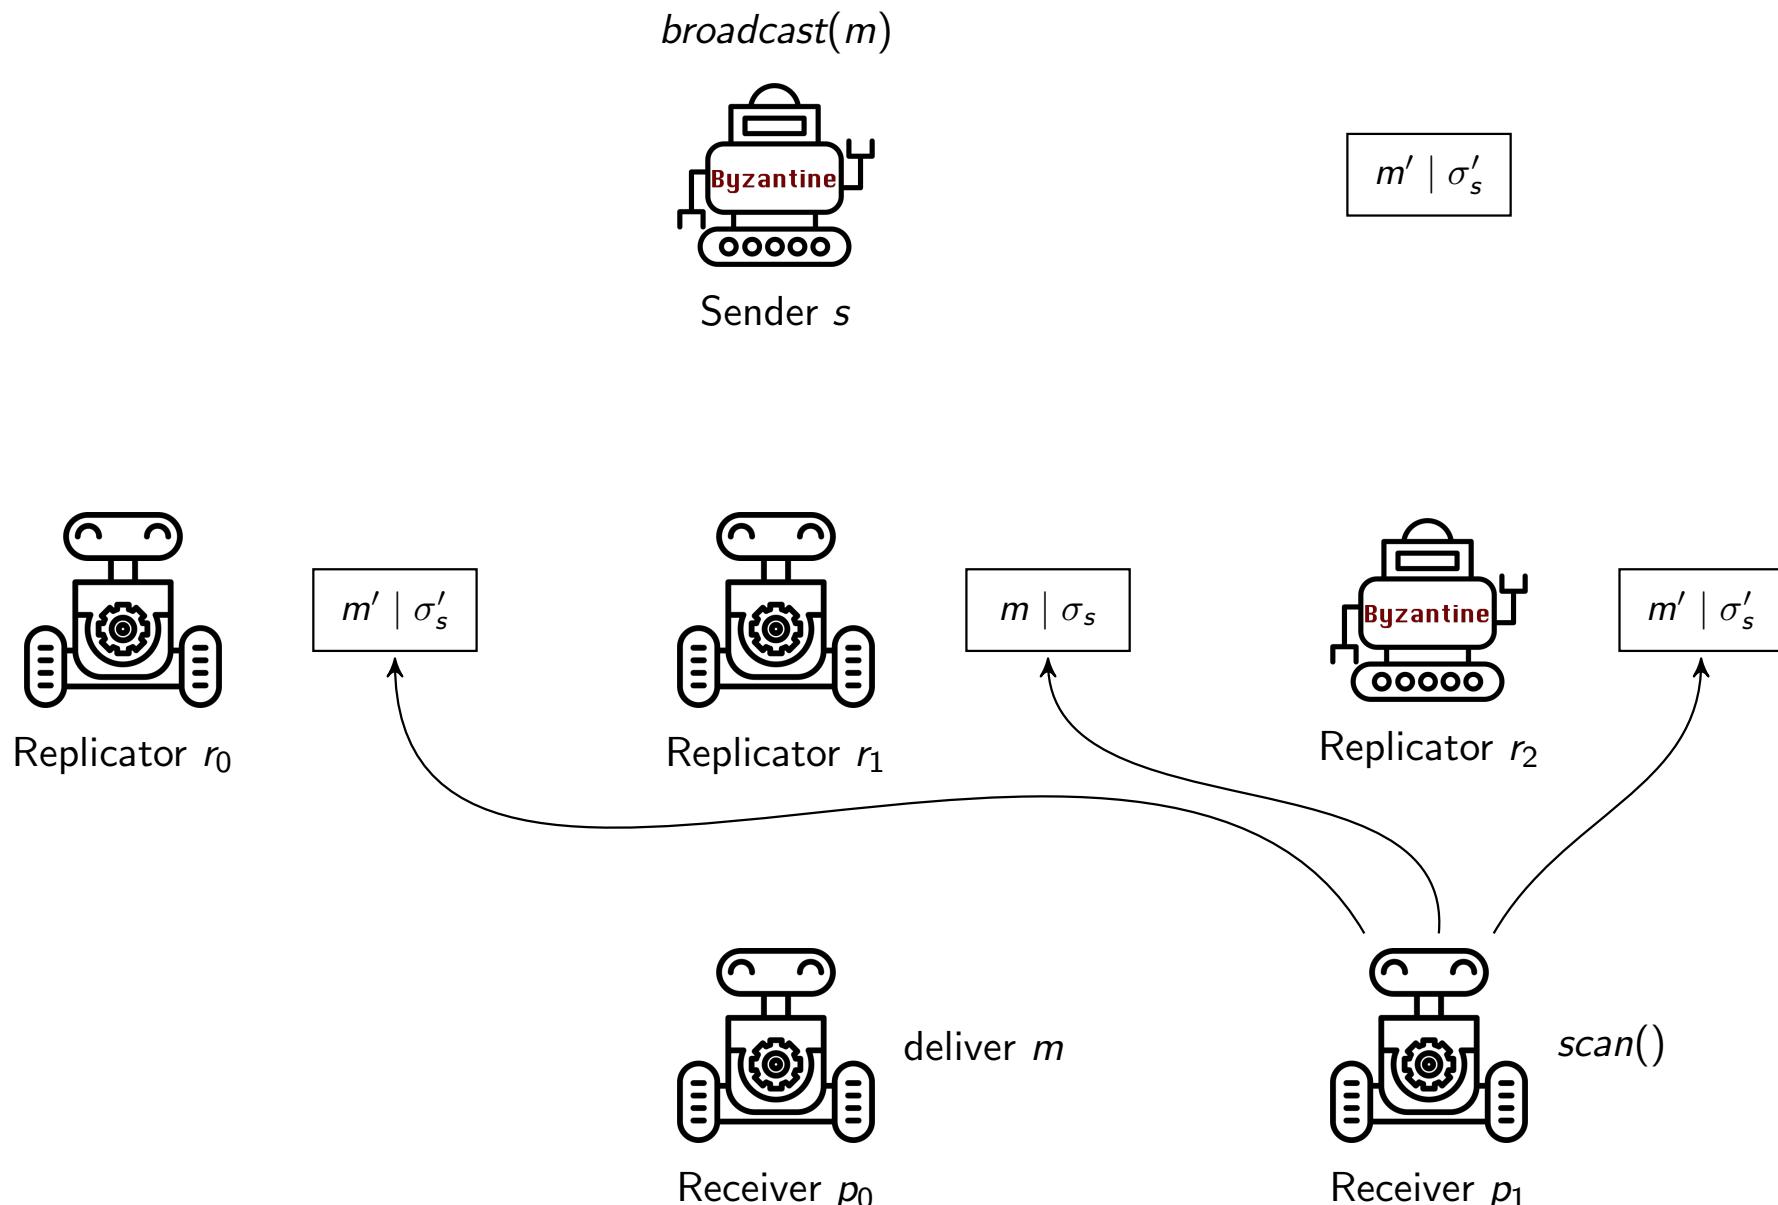
deliver m

Receiver p_0



Receiver p_1

Consistent Broadcast algorithm = \neg sufficient



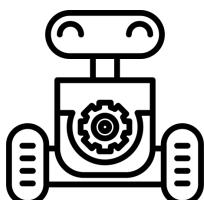
Consistent Broadcast algorithm = \neg sufficient

$broadcast(m)$



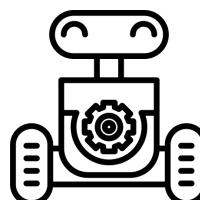
Sender s

$m' \mid \sigma'_s$



Replicator r_0

$m' \mid \sigma'_s$



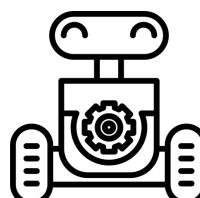
Replicator r_1

$m \mid \sigma_s$



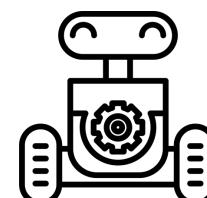
Replicator r_2

$m' \mid \sigma'_s$



deliver m

Receiver p_0



no delivery

Receiver p_1

Reliable Broadcast

Algorithm details

Init - Echo - Ready mechanism

Reliable Broadcast

Algorithm details

Init - Echo - Ready mechanism

Uses Consistent Broadcast

Reliable Broadcast

Algorithm details

Init - Echo - Ready mechanism

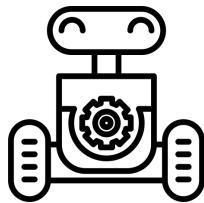
Uses Consistent Broadcast

Similar delivery strategy to Consistent Broadcast: **fast path**, i.e., when there is unanimity and otherwise when $\exists n - f$ valid proof sets for m

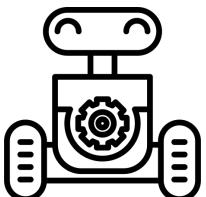
Reliable Broadcast

Algorithm sketch, $f = 1$. Fast path

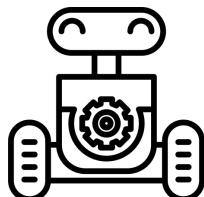
$broadcast(m)$



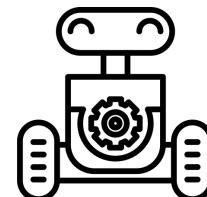
Sender s



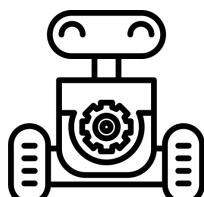
Replicator r_0



Replicator r_1



Replicator r_2

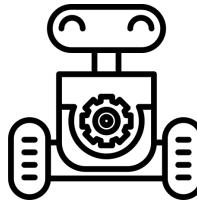


Receiver p

Reliable Broadcast

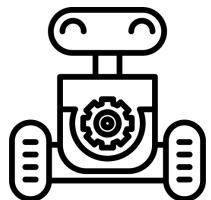
Algorithm sketch, $f = 1$. Fast path

$broadcast(m)$

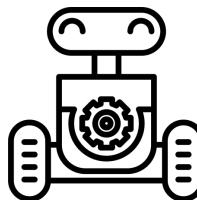


Sender s

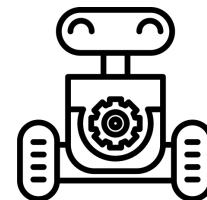
$cb\text{-}broadcast(\langle Init, m \rangle)$



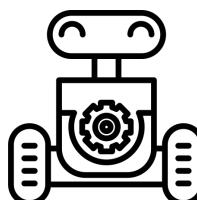
Replicator r_0



Replicator r_1



Replicator r_2

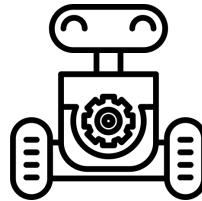


Receiver p

Reliable Broadcast

Algorithm sketch, $f = 1$. Fast path

$broadcast(m)$

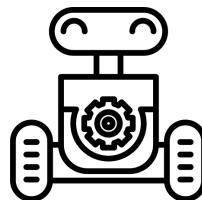


Sender s



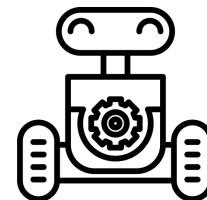
Replicator r_0

$cb\text{-}deliver(\langle Init, m \rangle)$



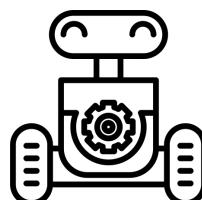
Replicator r_1

$cb\text{-}deliver(\langle Init, m \rangle)$



Replicator r_2

$cb\text{-}deliver(\langle Init, m \rangle)$

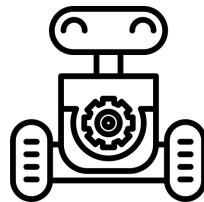


Receiver p

Reliable Broadcast

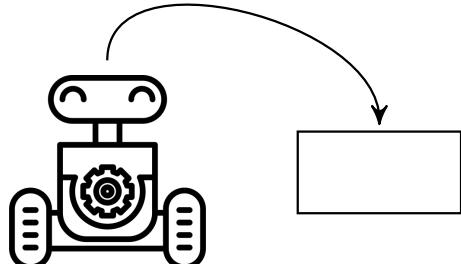
Algorithm sketch, $f = 1$. Fast path

$broadcast(m)$



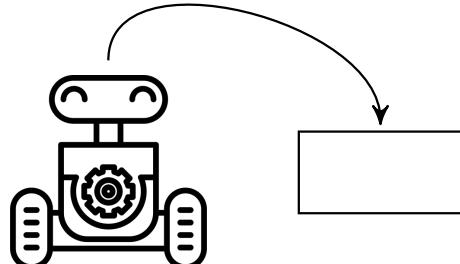
Sender s

$Echo.msg.write(m)$



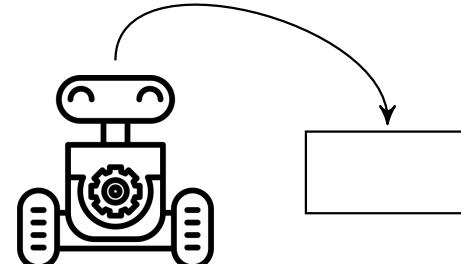
Replicator r_0

$Echo.msg.write(m)$

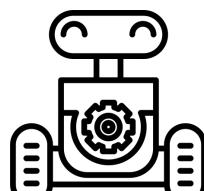


Replicator r_1

$Echo.msg.write(m)$



Replicator r_2

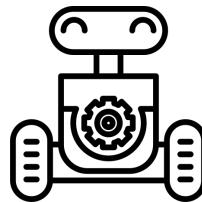


Receiver p

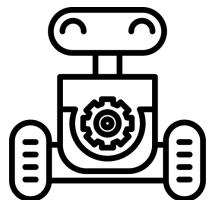
Reliable Broadcast

Algorithm sketch, $f = 1$. Fast path

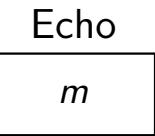
$broadcast(m)$



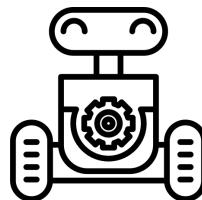
Sender s



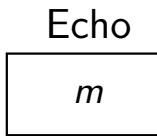
Replicator r_0



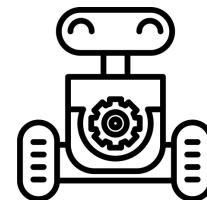
Echo
 m



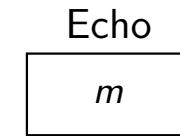
Replicator r_1



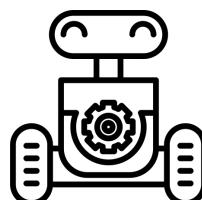
Echo
 m



Replicator r_2



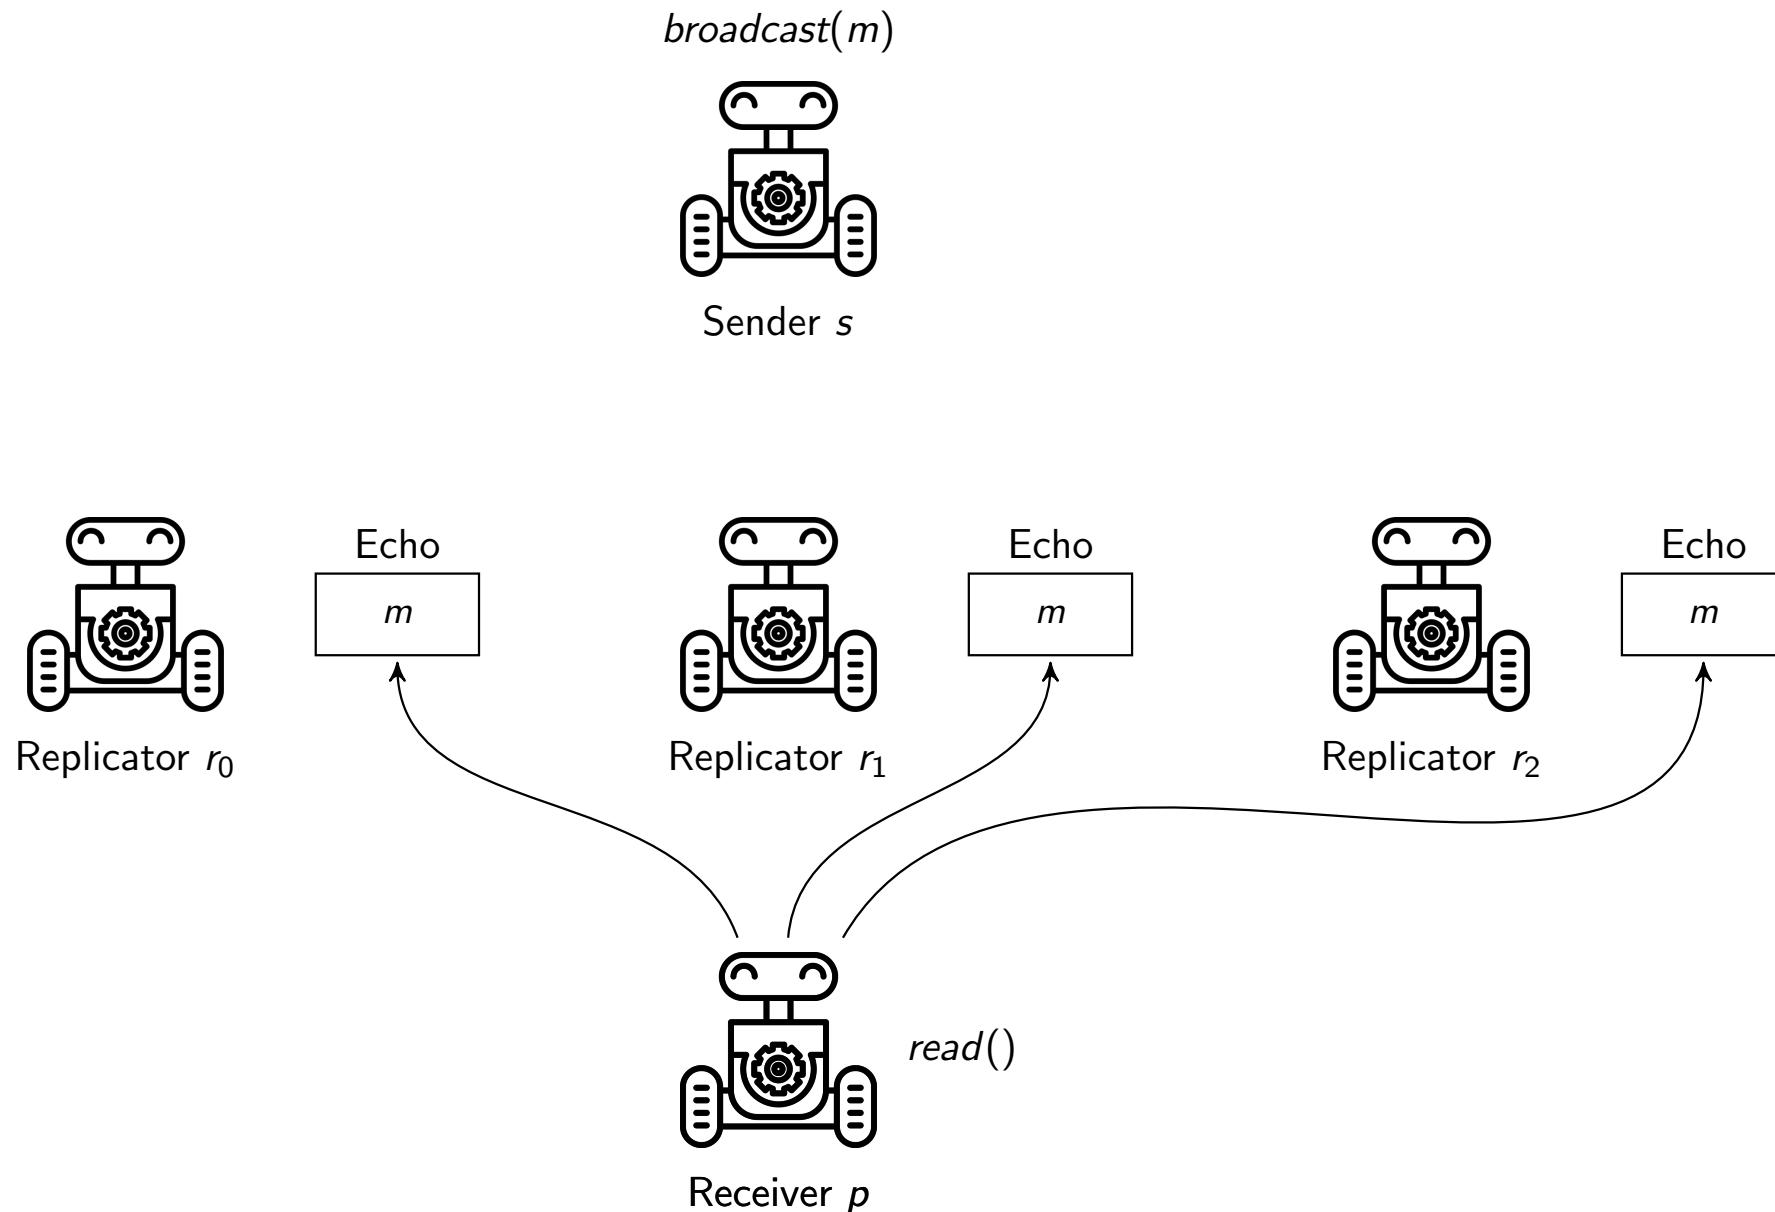
Echo
 m



Receiver p

Reliable Broadcast

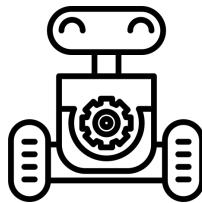
Algorithm sketch, $f = 1$. Fast path



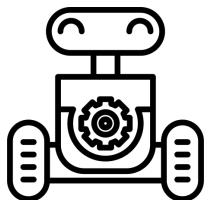
Reliable Broadcast

Algorithm sketch, $f = 1$. Fast path

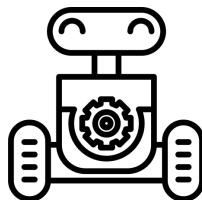
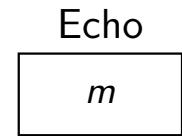
broadcast(m)



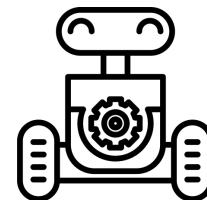
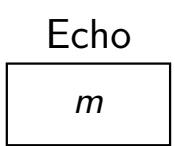
Sender s



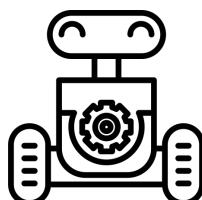
Replicator r_0



Replicator r_1



Replicator r_2



Receiver p

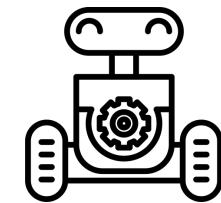
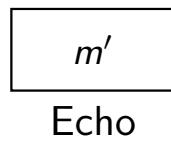
unanimity \implies deliver m via *fast path*

Reliable Broadcast

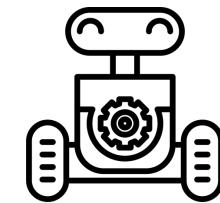
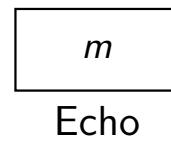
Algorithm sketch, $f = 1$. Construction of *ReadySet*



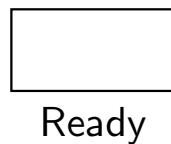
Replicator r_0



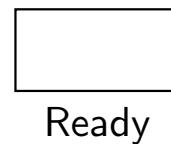
Replicator r_1



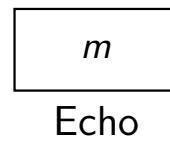
Replicator r_2



Ready



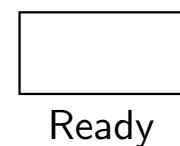
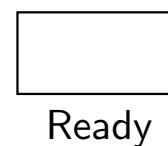
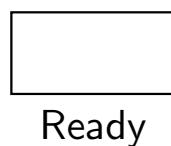
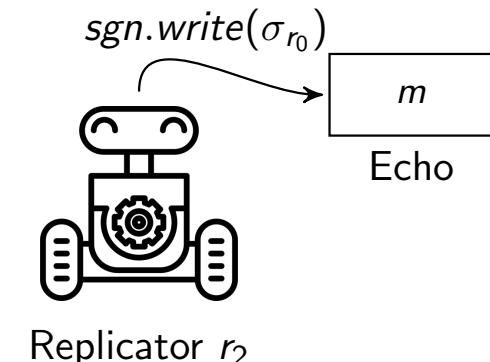
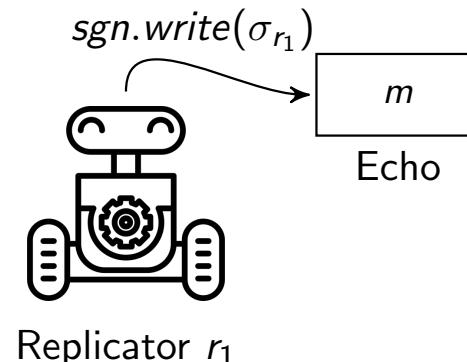
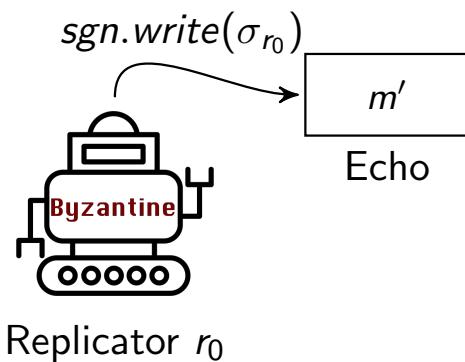
Ready



Ready

Reliable Broadcast

Algorithm sketch, $f = 1$. Construction of *ReadySet*



Reliable Broadcast

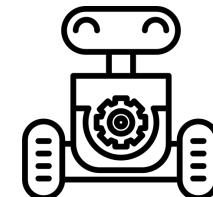
Algorithm sketch, $f = 1$. Construction of *ReadySet*



Replicator r_0

$m' \mid \sigma_{r_0}$

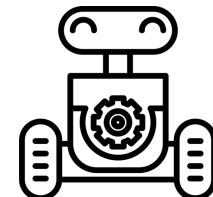
Echo



Replicator r_1

$m \mid \sigma_{r_1}$

Echo



Replicator r_2

$m \mid \sigma_{r_2}$

Echo



Ready



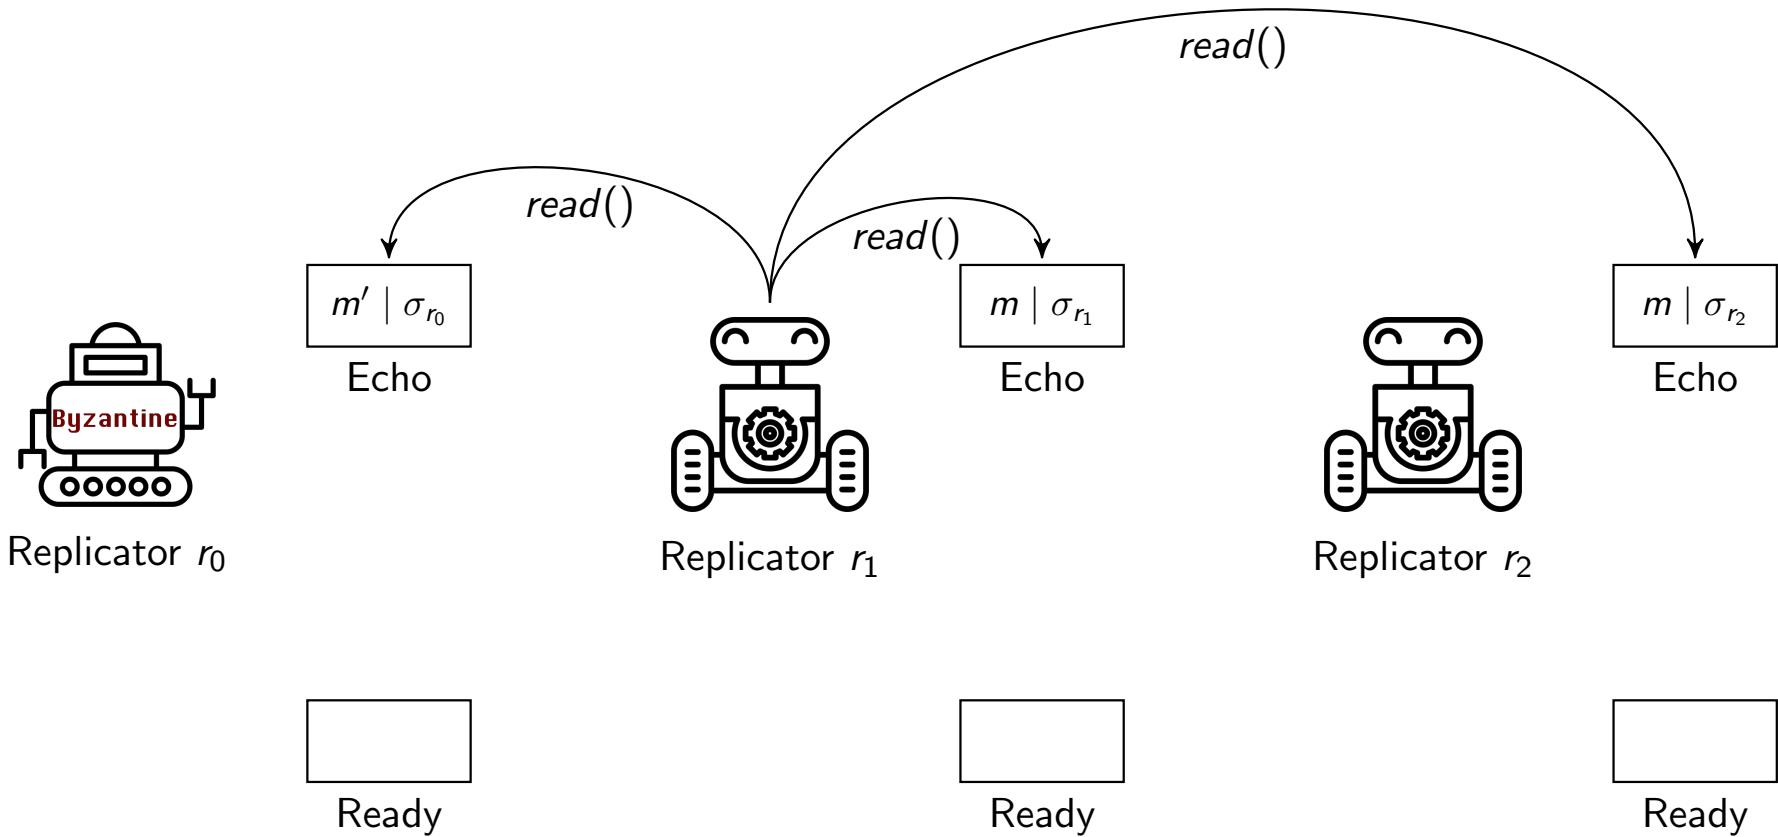
Ready



Ready

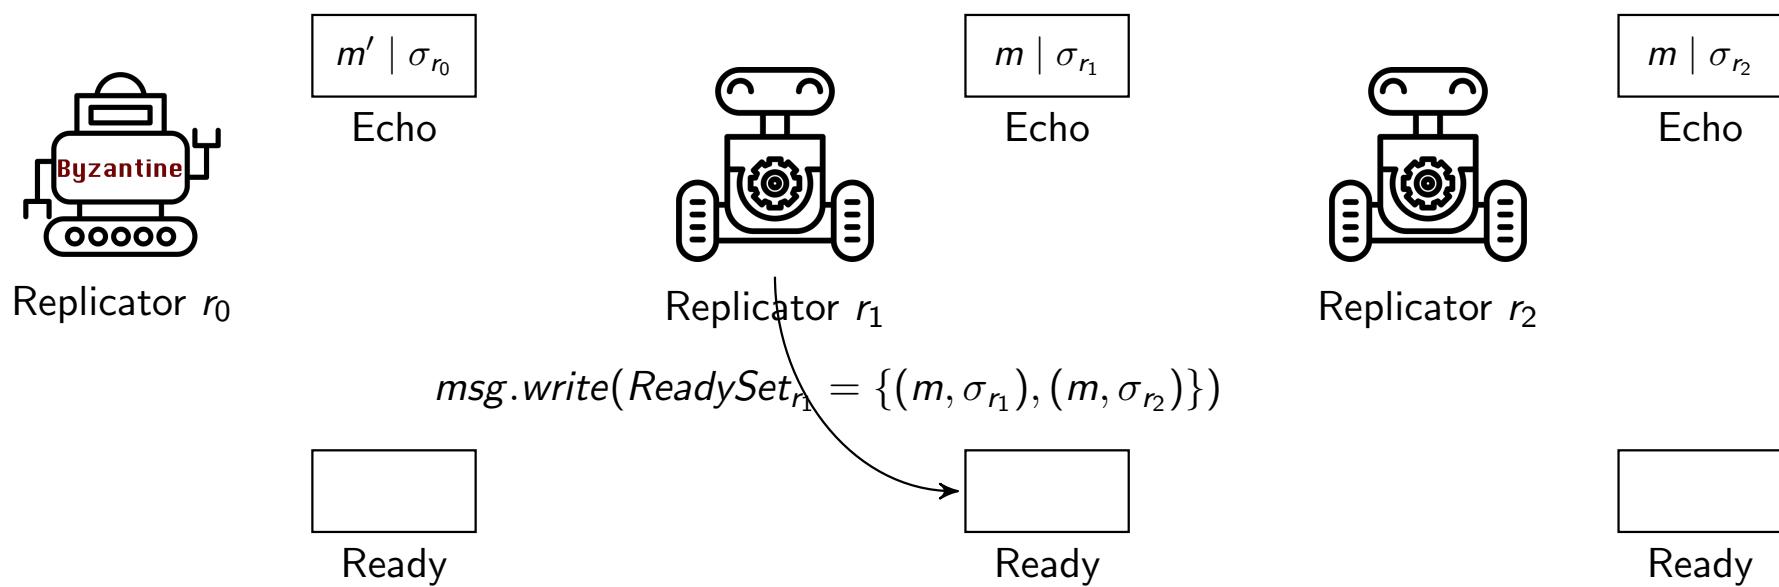
Reliable Broadcast

Algorithm sketch, $f = 1$. Construction of *ReadySet*



Reliable Broadcast

Algorithm sketch, $f = 1$. Construction of *ReadySet*



Reliable Broadcast

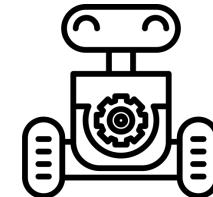
Algorithm sketch, $f = 1$. Construction of *ReadySet*



Replicator r_0

$m' \mid \sigma_{r_0}$

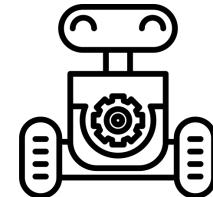
Echo



Replicator r_1

$m \mid \sigma_{r_1}$

Echo



Replicator r_2

$m \mid \sigma_{r_2}$

Echo

Ready

$ReadySet_{r_1}$

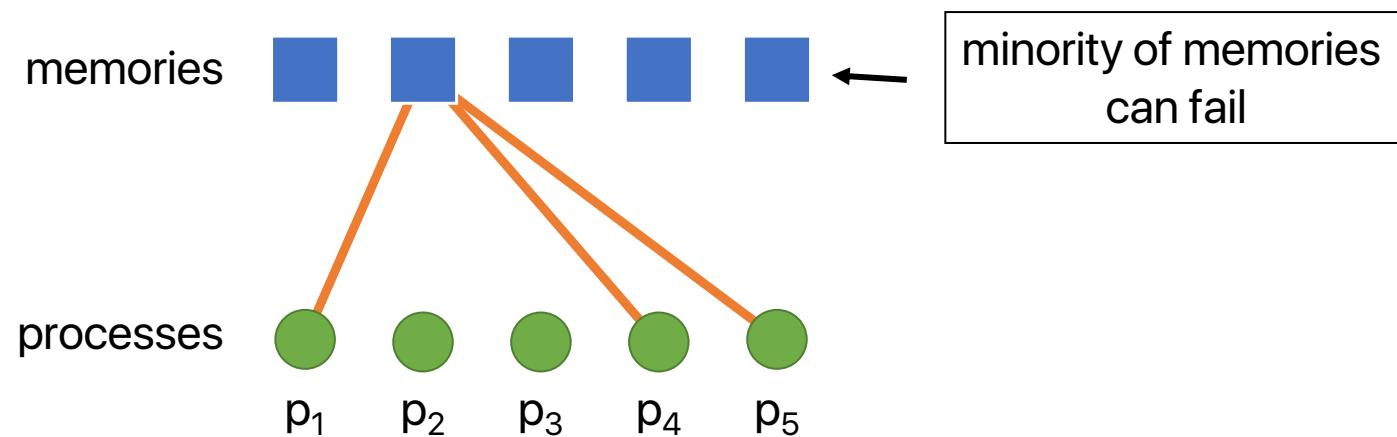
Ready

Ready

Outline

- Introduction
- 3 remarkable results with RDMA:
 - Consensus with crash faults
 - Broadcast with Byzantine faults
 - Fast memory replication

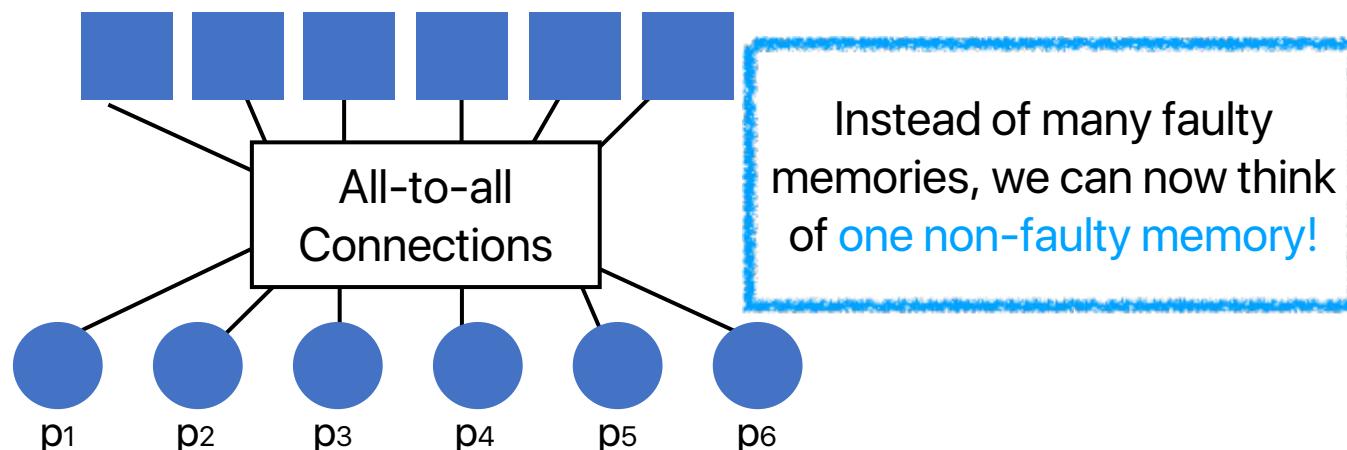
Recall: Disaggregated Memory



Handling Memory Failures

Replication: Treat all memories the same

Send all write/read requests to all memories, wait to hear acknowledgement from majority



Show that this implements a regular register, but not an atomic register!

Reliable MRMW Atomic Register

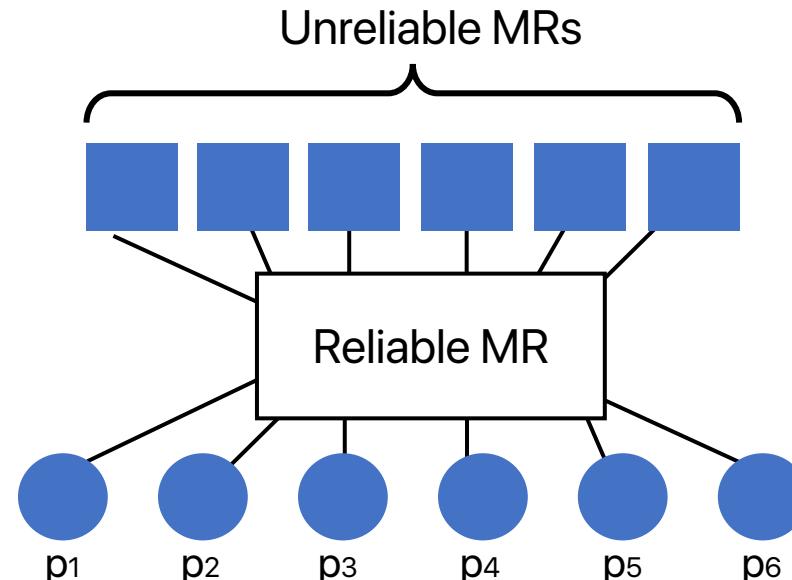
- We want to implement an atomic MRMW register on a set of unreliable (fault-prone) memories
- We want to minimize the number of round trips (RTTs) per operation.
- Proven: cannot be solved s.t. each operation always takes 1 RTT.
- But can it be done s.t. operations take 1 RTT most of the time?
- To simplify the problem, we assume each memory has plenty of **max registers**.

Max Registers

- Two operations: read and write
- Intuitively: read returns highest value written so far
- Formally:
 - **Validity:** If read R returns ν , then either (a) $\nu = \perp$, or (b) some operation $\text{write}(\nu)$ was invoked before R returns.
 - **Read-read monotonicity:** If a read returns value y and a preceding read returns value x , then $x \leq y$.
 - **Write-read monotonicity:** If a read returns value y and a preceding write writes value x , then $x \leq y$.
 - **Liveness (wait-freedom):** Every invoked operation eventually returns.

Step 1: Reliable Max Register

- Implement a reliable max register from a set of unreliable max registers
- Writes should complete in 0-1 RTTs, reads should complete in 1-2 RTTs.
- Common case: both operations should take 1 RTT.
- Hint: use caching.



Step 2: Atomic MRMW Register

Classic Algorithm

M = Reliable max register. Each value is a tuple (timestamp, id, value).
Lexicographic ordering.

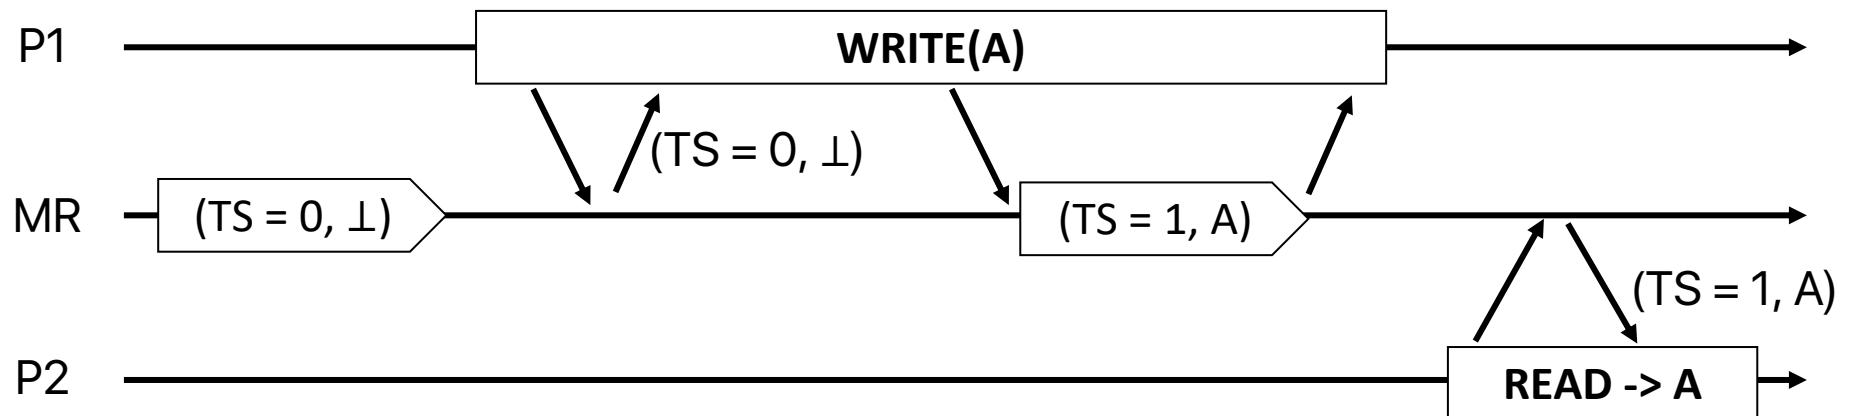
```
def WRITE(v): // this block is not atomic
    fresh_ts = (M.READ().ts.i + 1, tid) ← 2 RTTs
    M.WRITE((fresh_ts, v)) ←

def READ():
    return M.READ().v ← 1 RTT
```

Why 2 RTTs for WRITE? Can we do better?

Example

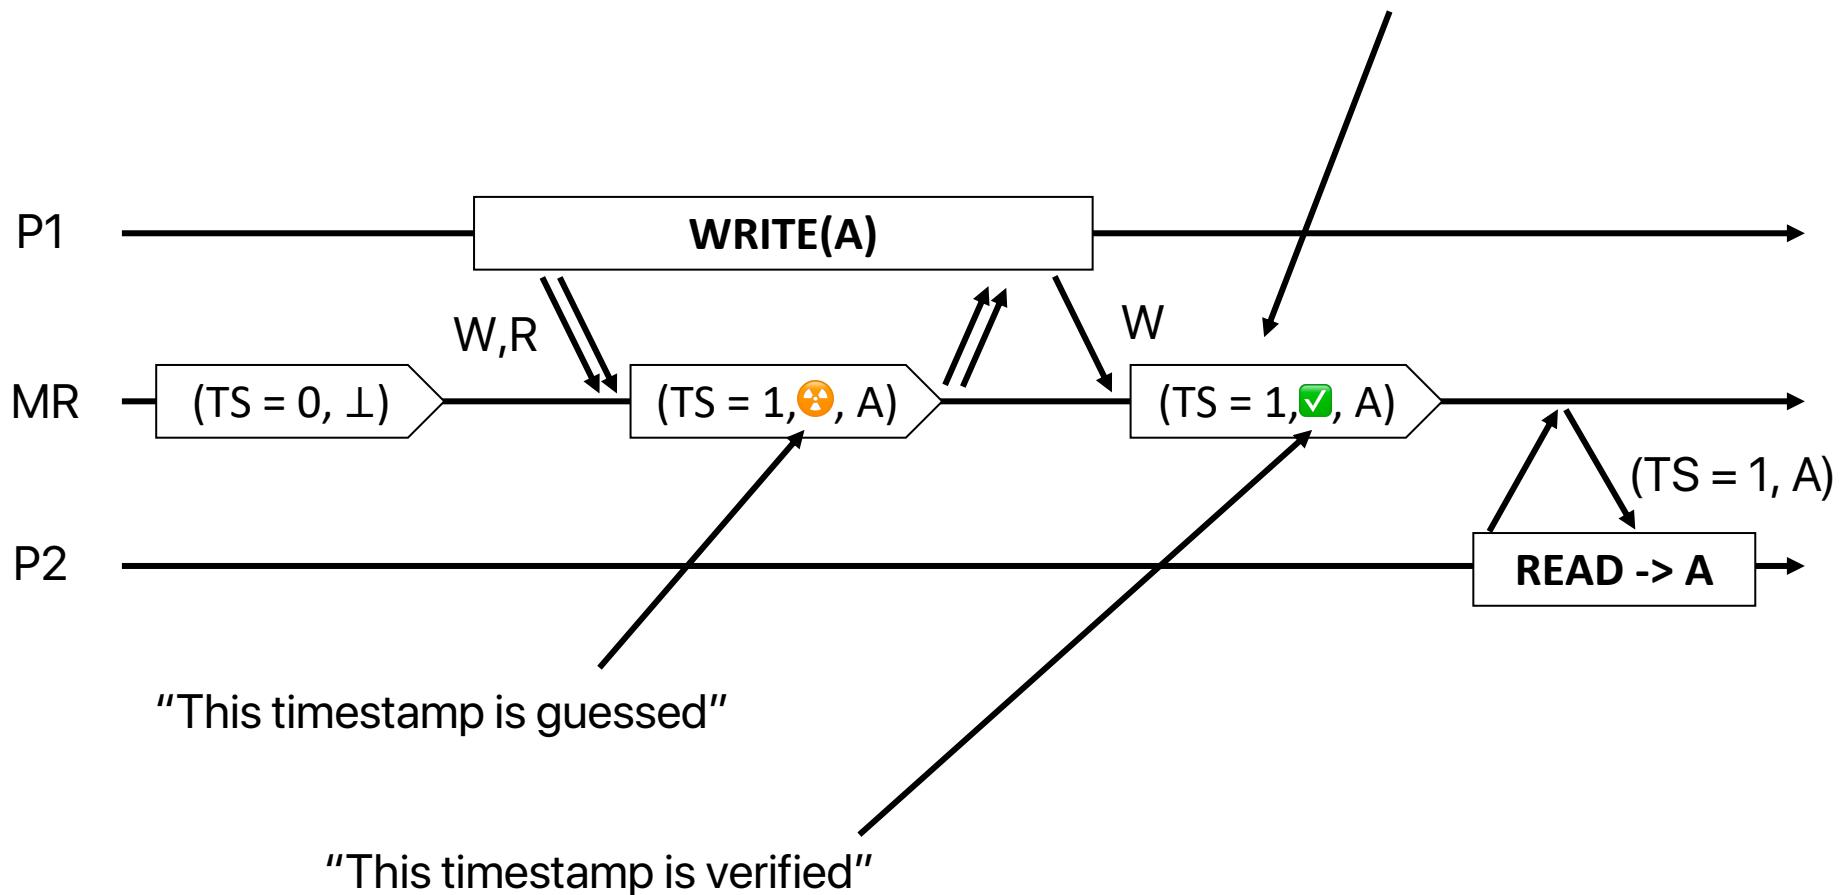
- Each write needs to use a fresh timestamp, i.e., higher than all preceding (why?)
- Finding a fresh timestamp takes 1 RTT.



2^{nd} WRITE RTT is unavoidable.
 1^{st} WRITE RTT: Could we guess the timestamp?

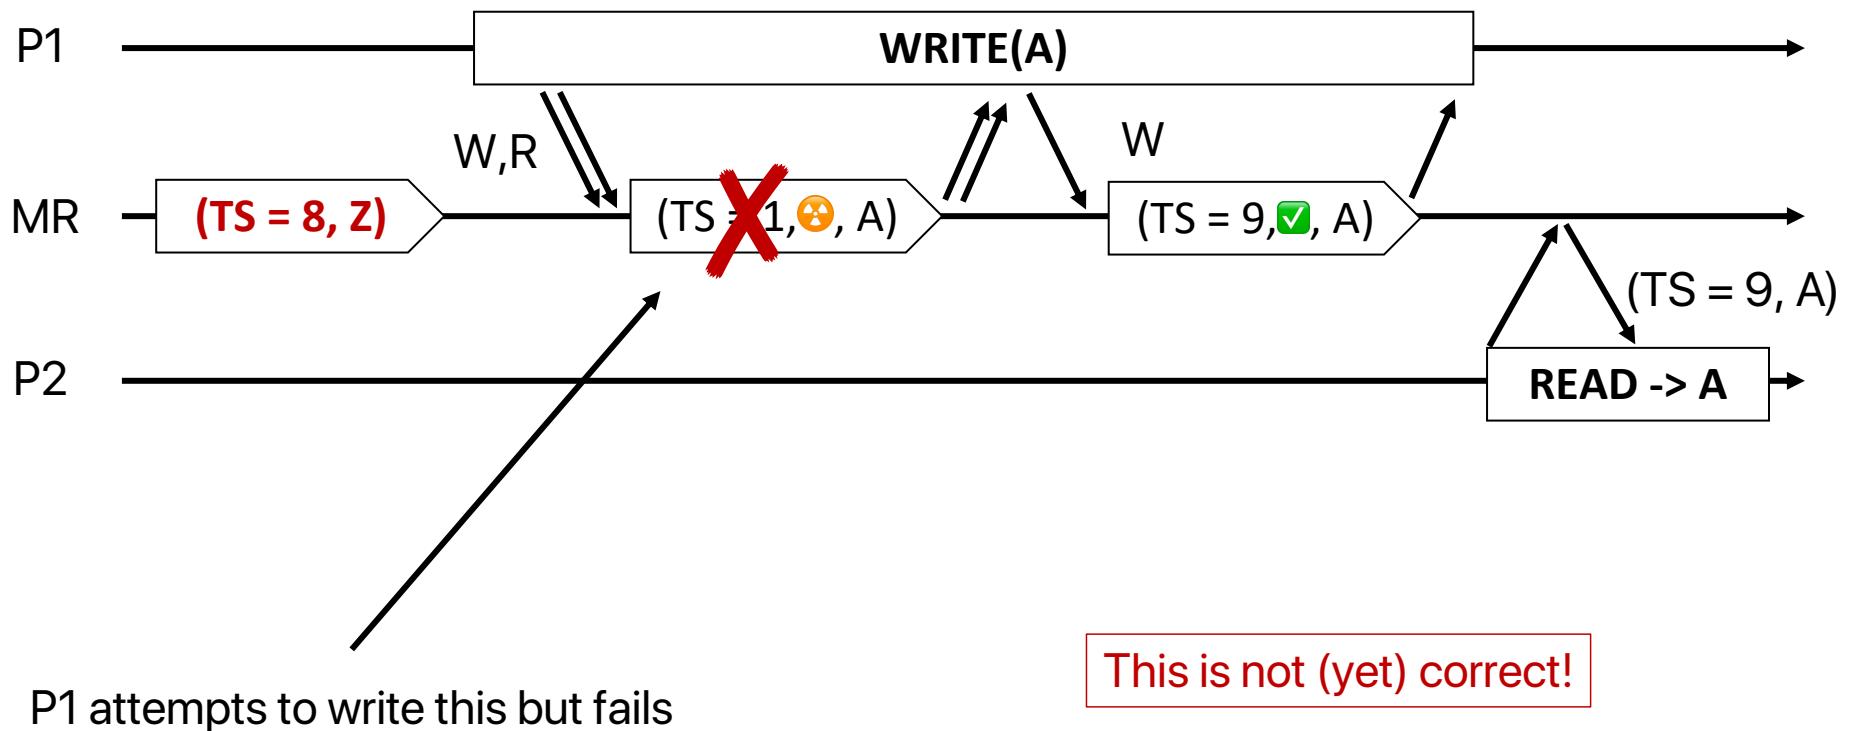
Guessing Timestamps

Done in background, can return before it completes
-> does not count as RTT on critical path

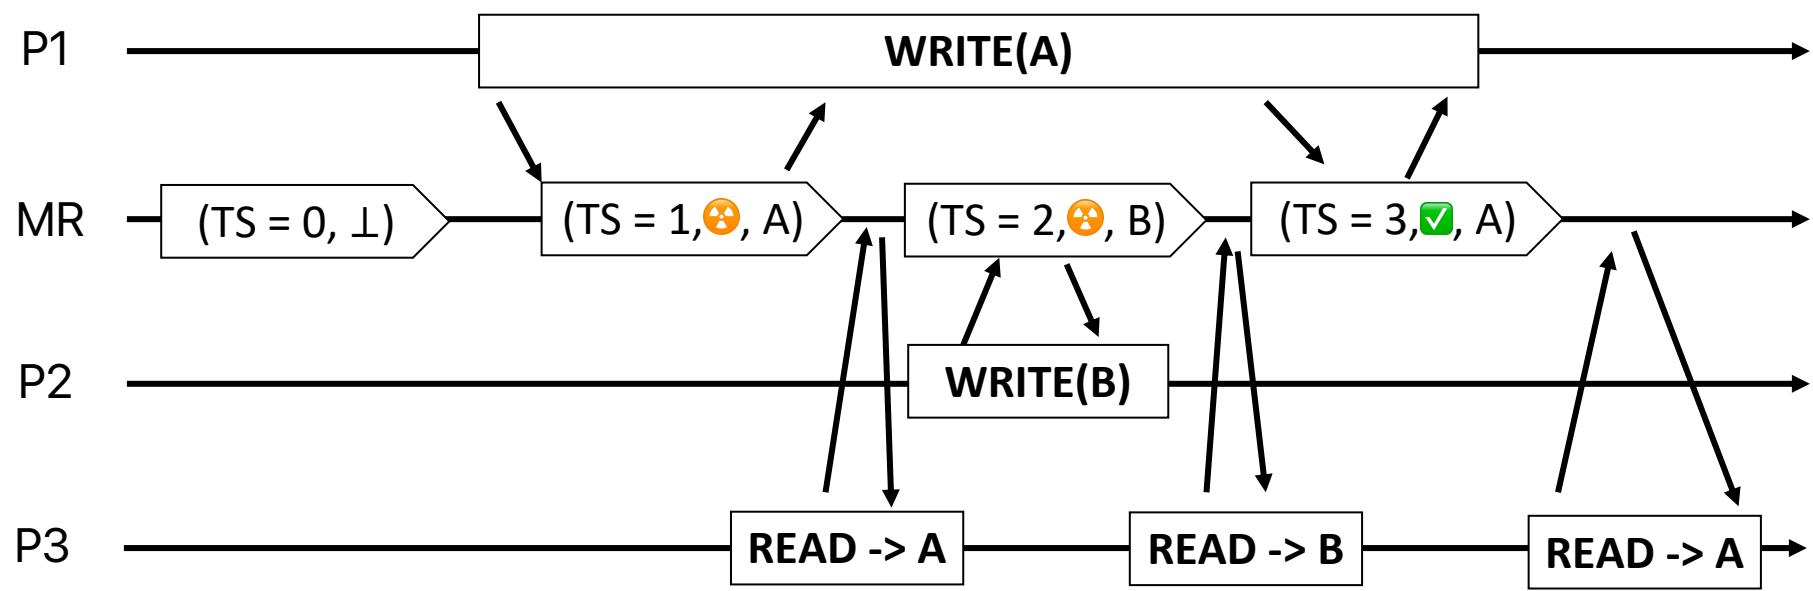


Guessing Timestamps

What if guessed timestamp is wrong?

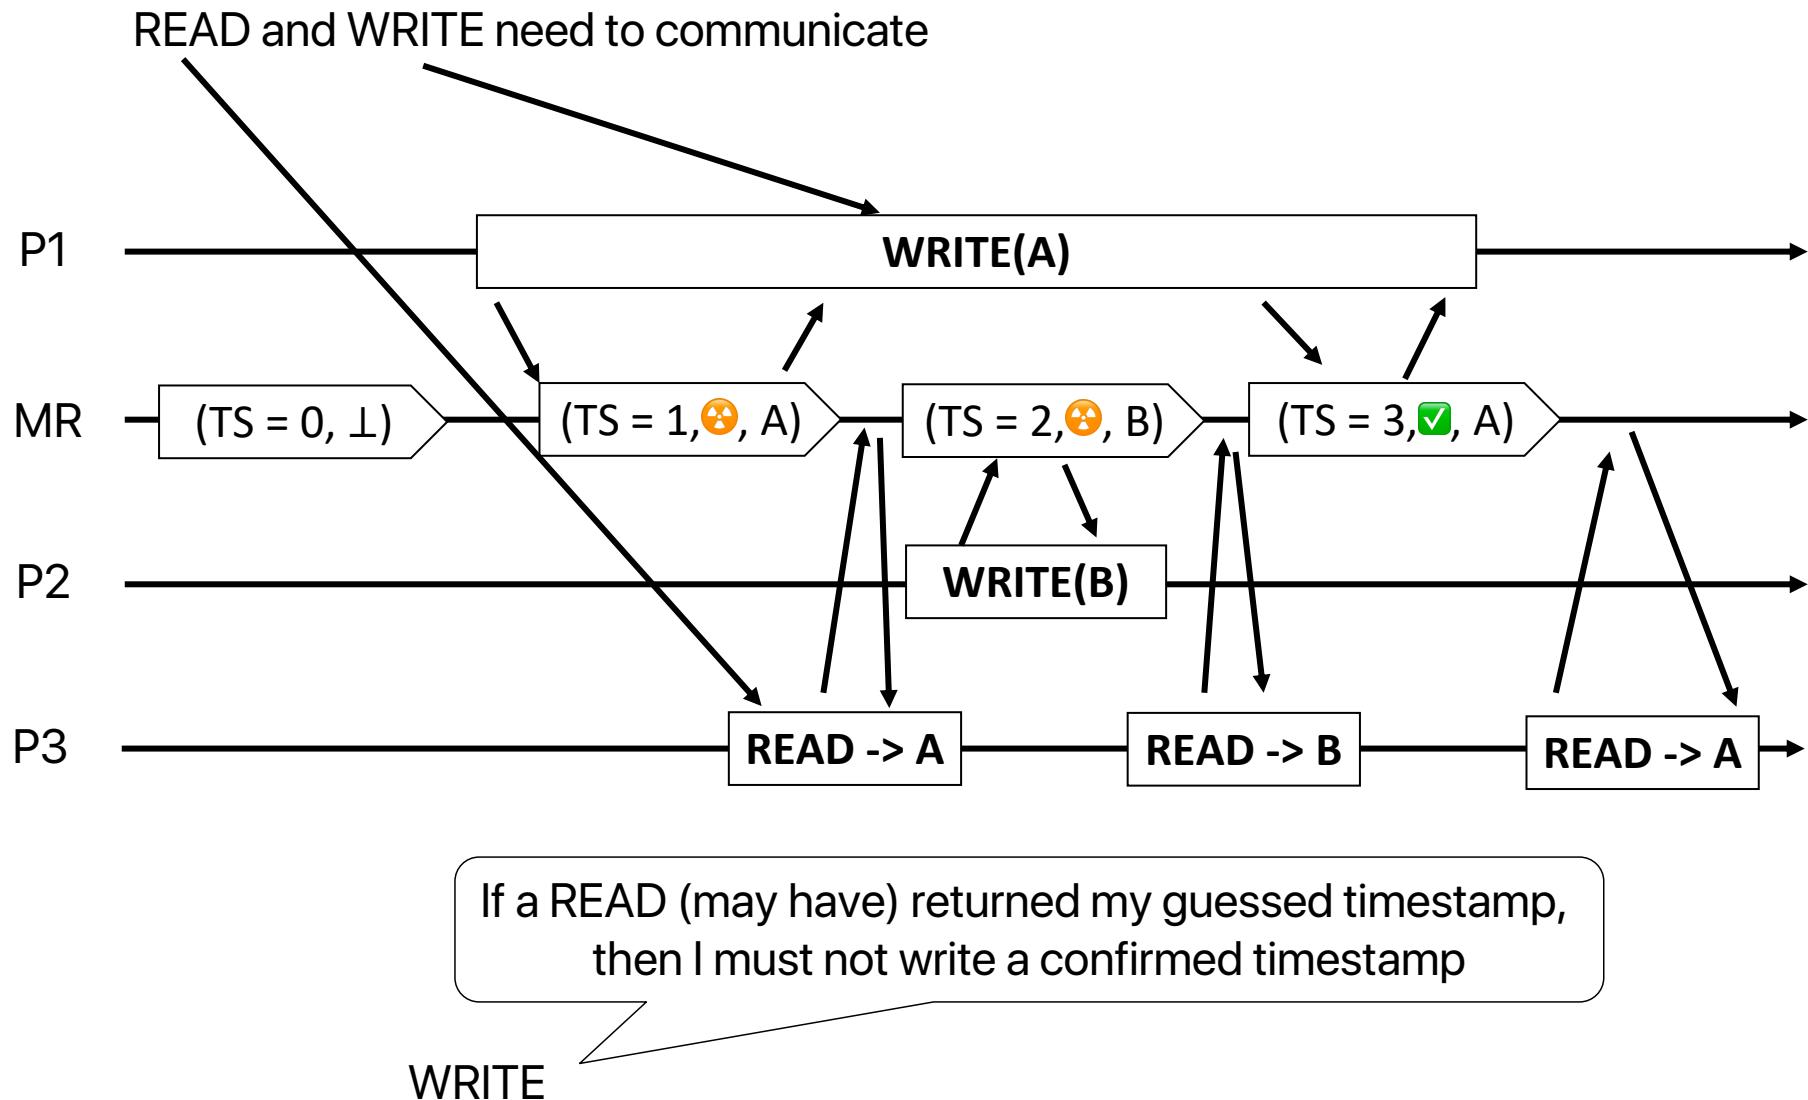


Guessing Timestamps



Not atomic/linearizable ☹

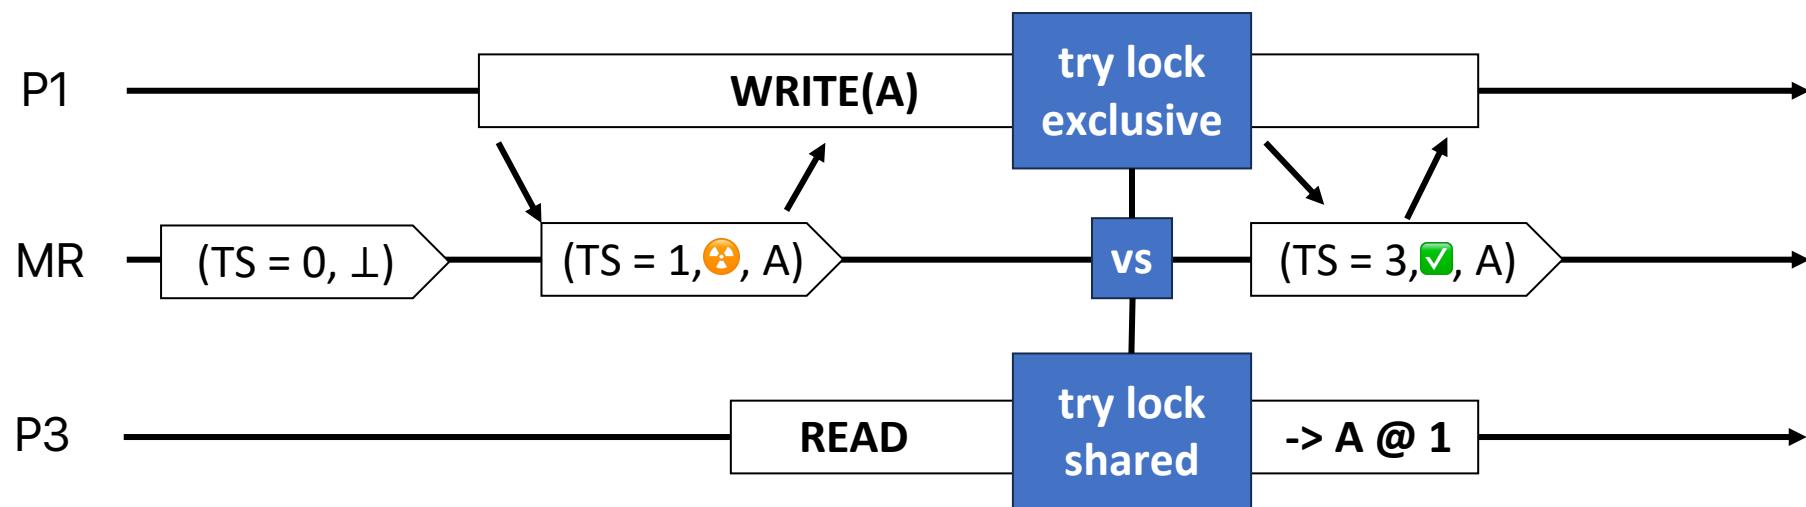
Solution:



Solution

If a READ (may have) returned my guessed timestamp,
then I must not write a confirmed timestamp

WRITE



Putting It All Together

Write algorithm

```
M = ((0, ⊥), VERIFIED, ⊥) // Max Register
TSL[tid] = {} // Timestamp Lock

def WRITE(v):
    w = (guessTs(), GUESSED, v)
    in parallel {m = M.READ(), M.WRITE(w)}
    if m <= w: // Fast path (fresh timestamp)
        in bg: M.WRITE(w with VERIFIED) // Spdup reads
    else: // Slow path (potentially stale timestamp)
        if TSL[tid].TRYLOCK(w.ts, WRITE):
            M.WRITE((m.ts.i+1, tid), VERIFIED, v))
```

guess a timestamp
write guessed ts + read current ts
if guessed ts is fresh:
write verified ts in bg
if guessed ts is stale:
try to take exclusive lock
if successful, write fresh ts

Common case: 1 RTT!

Putting It All Together

Read algorithm

```
def READ():
    seen: dict<ThreadId, MValue> = {}
    while True:
        m = M.READ()
        if m is VERIFIED: return m.v // Fast path
        if m in seen.values: // Fresh timestamp
            if TSL[m.ts.tid].TRYLOCK(m.ts, READ):
                in bg: M.WRITE(m with VERIFIED) // Spdup rds
                return m.v
        elif m.ts.tid in seen.keys: // Wait-free path
            return seen[m.ts.tid].v
        seen[m.ts.tid] = m
```

read from MR
if ts is verified, return it

try to take shared lock
if successful, help reads & return

Common case: 1 RTT!

Putting It All Together

Read algorithm

```
def READ():
    seen: dict<ThreadId, MValue> = {}
    while True:
        m = M.READ()
        if m is VERIFIED: return m.v // Fast path
        if m in seen.values: // Fresh timestamp
            if TSL[m.ts.tid].TRYLOCK(m.ts, READ):
                in bg: M.WRITE(m with VERIFIED) // Spdup rds
                return m.v
        elif m.ts.tid in seen.keys: // Wait-free path
            return seen[m.ts.tid].v
        seen[m.ts.tid] = m
```

What about all this
other stuff?

It's for wait-freedom. Check out the paper for the full explanation:

“SWARM: Replicating Shared Disaggregated-Memory Data in No Time”

Further Reading

1. ABGMZ. *The Impact of RDMA on Agreement*. PODC 2019.
2. ABGMXZ. *Microsecond Consensus for Microsecond Applications*. OSDI 2020.
3. ABGPXZ. *Frugal Byzantine Computing*. DISC 2021.
4. ABGMXZ. *uBFT: Microsecond-Scale BFT using Disaggregated Memory*. ASPLOS 2023.
5. MBXZAG. *SWARM: Replicating Shared Disaggregated-Memory Data in No Time*. SOSP 2024