Evaluation of Novel Fast Machine Learning Algorithms for Knowledge-Distillation-Based Anomaly Detection at CMS

Lino Gerlach
Princeton University
lg0508@princeton.edu

Elliott Kauffman
Princeton University
ek8842@princeton.edu

Abhishikth Mallampalli
University of Wisconsin-Madison
amallampalli@wisc.edu

Abstract

The CICADA (Calorimeter Image Convolutional Anomaly Detection Algorithm) project aims to detect anomalous physics signatures without bias from theoretical models in proton–proton collisions at the Compact Muon Solenoid (CMS) experiment at the Large Hadron Collider. CICADA identifies anomalies in low-level calorimeter trigger data using a convolutional autoencoder, whose behavior is transferred to compact student models via knowledge distillation. Careful model design and quantization ensure sub-200 ns inference times on FPGAs. We investigate novel student model architectures that employ differentiable relaxations to enable extremely fast inference at the cost of slower training—a welcome tradeoff in the knowledge distillation context. Evaluated on CMS open data and under emulated FPGA conditions, these models achieve comparable anomaly detection performance to classically quantized baselines with significantly reduced resource usage. The savings in resource usage enable the possibility to look at a richer input granularity.

1 Introduction

At CERN's Large Hadron Collider (LHC), tens of millions of proton-proton collisions are produced per second. The Compact Muon Solenoid (CMS) detector at the LHC captures complex high-dimensional data from each collision. With such large data volume and rate, it is imperative to have an effective data reduction strategy. The Level-1 Trigger (L1T) in the CMS experiment is the first data reduction step, bringing the event rate from 40 MHz to 100 kHz–a 99.75% reduction within 3.8 microseconds [1][2]. The L1T is comprised of real-time algorithms built on Field Programmable Gate Arrays (FPGAs). Traditional physics-process oriented triggers are highly efficient for specific, known signal signatures but are inherently inflexible, rendering them ineffective for broad searches for unknown phenomena. Conversely, more general physics object-oriented triggers capture a wider variety of signals but must impose high momentum thresholds to control data rates, severely limiting their acceptance for new physics involving low-energy particles. To broaden the search for new physics to more possibilities, CMS has implemented Machine Learning (ML)-based implemented anomaly detection algorithms in the L1T. These new triggers aim to provide an optimal tradeoff by having a low data rate while maintaining a wider reach, all within the demanding, low-latency environment of the L1T, Fig. 1.

One such trigger is the Calorimeter Image Convolutional Anomaly Detection Algorithm

Preprint.

(CICADA) [3]. CICADA is trained using a convolutional autoencoder that is trained on CMS Zero Bias data, which is a random trigger that selects events with no bias towards any physics signature.

1.1 Our contribution

This paper explores three different quantization schemes that aim to reduce the resources required to run the CICADA algorithm without sacrificing performance. We show that two novel quantization libraries (HGQ, LGN) offer a vastly superior trade-off between physics performance and resource efficiency compared to the established QKeras baseline. This resource saving directly benefits the LHC physics program by enabling the deployment of more complex real-time algorithms — for example, analyzing calorimeter energy deposits at finer spatial granularity than previously possible — while meeting CMS's stringent real-time data selection constraints. We also utilize CMS Open Data to ensure the project's reach to the broader community [4].

2 Background

The input to the autoencoder is an $18\phi \times 14\eta$ unrolled cylindrical grid of calorimeter energy deposits, as shown in Fig. 2. Here, ϕ refers to the azimuthal angle with respect to the proton beamline and η refers to the a pseudorapidity (effecitvely the polar angle). The reconstruction loss is defined as the mean squared error (MSE) of the transverse energy deposits (E_T) in the input and output $\phi \times \eta$ grid:

$$\text{loss} = \frac{1}{252} \sum_{i=1}^{252} \left(E_{T,i}^{\text{original}} - E_{T,i}^{\text{reconstructed}} \right)^2 \tag{1}$$

This value is used to define the anomaly score, based on the principle that the autoencoder will reconstruct rare physics processes less accurately than the bulk of Zero Bias data it was trained on.

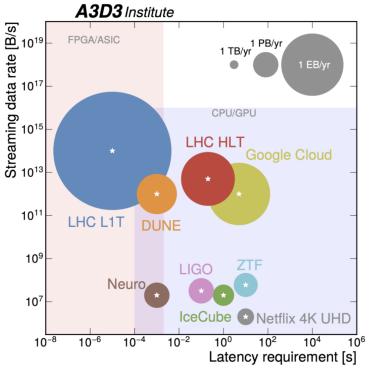

Figure 1: Streaming data rate versus latency requirements across various scientific and industrial applications.

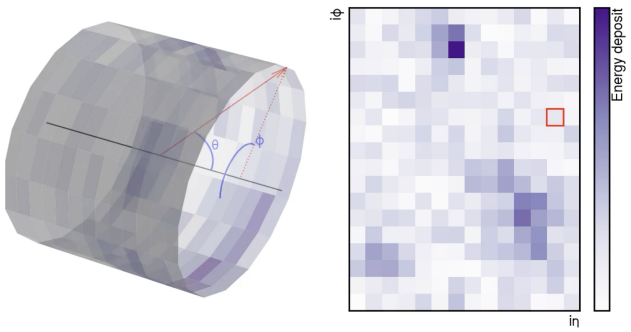

Figure 2: Schematic depiction of the topology of CICADA's inputs.

This autoencoder model consists of around 300k parameters and therefore cannot be deployed within the FPGA-based L1T. To solve this problem, CICADA employs a knowledge distillation strategy, in which the autoencoder is treated as a "teacher" model and a much smaller, ∼10k parameter convolutional neural network acts as a "student" [5]. The student model is trained to learn a soft target from the teacher model (the anomaly score) from the original $18\phi \times 14\eta$ input, using the Mean Absolute Error (MAE) loss function. The soft target is defined from the teacher reconstruction loss as:

$$Q_{log} := \text{quantize}\left(\log\left(32 \cdot \text{loss}\right)\right) \tag{2}$$

This transformation of the loss reduces the range of the targets to assist in training and makes use of the fixed bit-width requirement of the trigger system. Since the teacher model is trained on Zero Bias data and thus will result in only low anomaly scores, the student model is exposed to events with a known high anomaly score so that it can sufficiently learn the characteristics of a high score.

This process is known as *outlier exposure*. To optimize for low-precision inference on hardware, the student model was trained with quantization-aware training (QAT) using QKeras [6]. The trained QKeras student model was then translated into synthesizable C++ code with hls4ml. hls4ml inserts pragmas into the code to optimize the high-level synthesis [7]. The C++ code was further fine-tuned in order to meet the resource and latency requirements imposed by the trigger system.

3 Methods and Experimental Setup

3.1 Quantization Schemes

HGQ The High Granularity Quantization (HGQ) library [8] offers a more fine-grained approach to model compression than QKeras, which operates at the layer level. HGQ achieves this sub-layer granularity by treating the bitwidths of individual weights and biases as trainable parameters. To optimize these bitwidths, HGQ employs a two-pronged strategy within the loss function. First, it intentionally minimizes a proxy for hardware cost by introducing a term called Effective Bit Operations (EBOPs), which is proportional to the number of bitwise operations in the network. Second, it adds an L1 regularization penalty directly to the bitwidth parameters, which explicitly pushes them towards lower values. The EBOPs term is weighted by a hyperparameter, β, controlling the trade-off between model accuracy and resource efficiency. By minimizing this combined loss, the training process simultaneously learns the optimal value for each weight and its ideal bitwidth, automatically pruning connections by learning a bitwidth of zero.

LGN Another strategy for quantizing the CICADA student model is by changing the architecture to a logic gate network (LGN). LGNs are comprised of binary logic gates accepting only two inputs, so they can execute very quickly. To enable gradient-based training, differentiable functions are used as substitutes for the logic gates. Upon inference, each neuron is collapsed to the logic gate with the highest probability [9]. Prior to these quantization efforts, the performance of the CICADA student model increased when using a convolutional-based architecture.

3.2 Experimental design

Dataset Our study is conducted entirely using public CMS Open Data to ensure reproducibility and enable further community-driven research. The model is trained on the 2017 Zero Bias proton-proton collision dataset, which serves as the background distribution. For outlier exposure, we use top quark pair production ($t\bar{t}$), which has three possible final states characterized by number of leptons (l), neutrinos (ν), and quarks (q). The $t\bar{t} \to 2l + 2\nu$ sample is split into three sets: two are used for signal exposure during training and validation, while the third is reserved for the final evaluation. The other two samples, $t\bar{t} \to l + \nu + 2q$ and $t\bar{t} \to 4q$, are used exclusively for evaluation. These processes were chosen as they are the primary relevant physics signals available in the open data collection. Additionally, $t\bar{t}$ processes have a much lower cross-section than the majority of processes present in the Zero Bias dataset, meaning that they happen at a comparatively negligible rate, so that they should be anomalous with respect to the majority of events.

Experiments To explore the trade-off between performance and resource usage, we performed a hyperparameter scan for both HGQ and LGN. For HGQ, multiple student models were trained by varying the β parameter in the loss function. For LGN, multiple student models with different architectures as well as different bit representations for the input data were trained.

Performance was evaluated using a number of metrics. First, the teacher and student model scores were compared using the Earth Mover's Distance (EMD) to assess how well the student model learned the target teacher anomaly scores. Then, the physics performance was quantified using the Receiver Operating Characteristic Area Under the Curve (ROC AUC) to measure the trade-off between signal efficiency and background rejection.

Hardware Cost was estimated by synthesizing each model onto the target FPGA platform, an AMD Virtex-7 (XC7VX690T) FPGA, using Vitis HLS [10, 11]. The synthesis reposts provide key implementation metrics: resource utilization, measured in Look-Up Tables (LUTs), Flip-Flops (FFs), and Digital Signal Processors (DSPs), as well as the processing latency, expressed in the number of FPGA clock cycles required for a single inference. For the HGQ model, the hls4ml framework was used to generate HLS-compatible code. For the LGN, the structure was directly translared into C code for HLS synthesis.

4 Results

Fig. 3 (left) shows the EMD as a function of EBOPs (EBOPs $\approx$ #LUTs + 55 $*$ #DSPs) for different CICADA student models, with the QKeras implementation serving as a baseline. EBOPs is used because the HGQ library documentation claims that the final resource count, after the place-and-route stage follows this formula. EMD is reported for three datasets: Zero Bias, $t\bar{t} \to 4q$, and $t\bar{t} \to l + \nu + 2q$. The HGQ student models generally achieve comparable or sometimes lower EMD than the QKeras model, demonstrating strong fidelity in learning the teacher outputs. The LGN student learns the Zero Bias sample reasonably well but struggles with the signal datasets, likely because the limited number of high-value outliers in the training set is insufficient given the binary input representation required, compounded by the constraints of limited open data. Despite this, LGN models exhibit very low resource utilization, suggesting strong potential if outlier training can be improved. HGQ models also maintain slightly lower resource usage while achieving similar physics performance, as illustrated in the ROC AUC versus EBOPs plot (right). Table 1 shows the resource utilization estimates in detail.

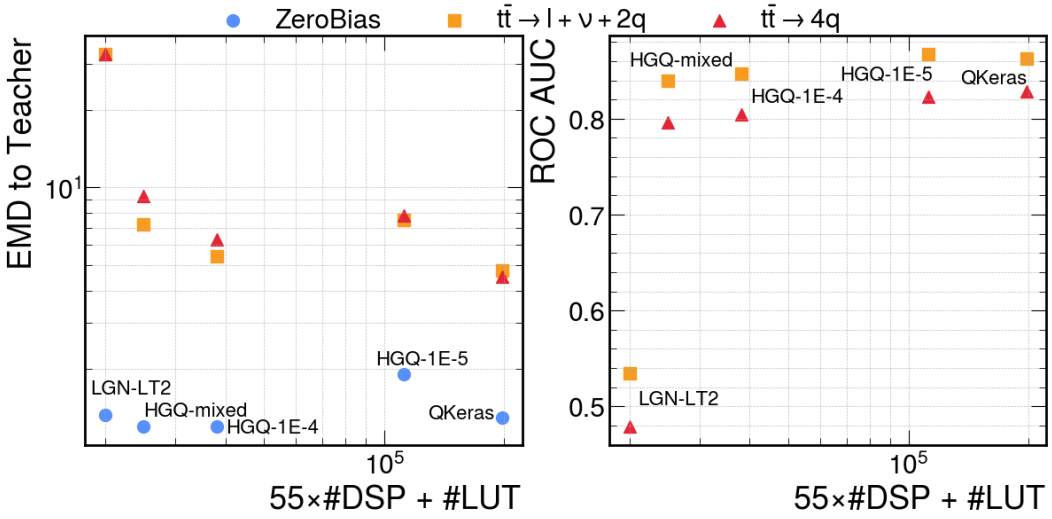

Figure 3: Performance results for the different CICADA student models given as a function of resource utilization calculated in EPOPs. Any given model lies along a vertical line. On the left, the EMD between the targets from the teacher and the students' prediction is shown. On the right, the ROC-AUC for two simulated signals is shown. The teacher achieved ROC-AUCs of 0.81 and 0.85 on the $t\bar{t} \to l + \nu + 2q$ and $t\bar{t} \to 4q$ signals, respectively.

Table 1: Pre-place-and-route hardware cost comparison for select student models. Latency is reported in clock cycles (cc), where one cycle is 6.25 ns.

Model Label	Quantization Library	Latency (in cc)	DSPs	FFs	LUTs	HGQ EBOPs
QKeras	QKeras	16	697	50368	159447	-
HGQ-1E-5	HGQ	17	4	27776	111848	39170
HGQ-1E-4	HGQ	11	1	6229	38111	7570
HGQ-mixed	HGQ	8	0	3019	24947	3301
LGN-LT2	LGN	3	0	856	19977	-

5 Limitations and Future work

The results presented here serve as a proof of concept, and we anticipate that performance can be further enhanced through a detailed exploration of synthesis parameters. We are working on further improving the performance of the students using these novel quantization methods. Although resource utilization estimates are provided in this paper, performing place-and-route is required to

determine the real hardware cost of the student models. Future work also includes extending the LGN to a convolutional version (C-LGN), which uses kernels composed of trees of logic gates instead of individual logic gates [12].

6 Broader Impact

By demonstrating a significant reduction in hardware requirements without compromising physics performance, this work now allows us to consider more advanced algorithms to be deployed within the LHC's stringent constraints. For example, it now becomes possible to consider using extra information bits which encode information about the finer energy distribution pattern upstream and also potential real-time monitoring of any potential input drifts. We believe these findings offer the potential to expand the search for new physics and may also be of interest to the wider community involved in resource-constrained real-time edge ML.

References

[1] *CMS TriDAS project: Technical Design Report, Volume 1: The Trigger Systems*. Technical design report. CMS. URL: https://cds.cern.ch/record/706847.

[2] A Tapper and Darin Acosta. *CMS Technical Design Report for the Level-1 Trigger Upgrade*. Tech. rep. Additional contacts: Jeffrey Spalding, Fermilab, Jeffrey.Spalding@cern.ch Didier Contardo, Universite Claude Bernard-Lyon I, didier.claude.contardo@cern.ch. 2013. URL: https://cds.cern.ch/record/1556311.

[3] CMS Collaboration. *Model-Independent Real-Time Anomaly Detection at the CMS Level-1 Calorimeter Trigger with CICADA*. CMS Detector Performance Summary CMS-DP-2024-121. Available at: https://cds.cern.ch/record/2917884. CERN, Nov. 2024.

[4] CMS Data preservation and open access group. *CERN Open Data Portal*. https://opendata.cern.ch/. Accessed: 2025-08-30. 2025.

[5] Adrian Alan Pol et al. *Knowledge Distillation for Anomaly Detection*. 2023. arXiv: 2310.06047 [cs.LG]. URL: https://arxiv.org/abs/2310.06047.

[6] Claudionor N. Coelho et al. "Automatic heterogeneous quantization of deep neural networks for low-latency inference on the edge for particle detectors". In: *Nature Machine Intelligence* 3.8 (2021), pp. 675–686.

[7] Javier Duarte et al. "Fast inference of deep neural networks in FPGAs for particle physics". In: *JINST* 13.07 (2018), P07027. DOI: 10.1088/1748-0221/13/07/P07027. arXiv: 1804.06913 [physics.ins-det].

[8] Sun Chang et al. *Gradient-based Automatic Per-Weight Mixed Precision Quantization for Neural Networks On-Chip*. en. 2024. DOI: 10.7907/HQ8JD-RHG30. URL: https://authors.library.caltech.edu/doi/10.7907/hq8jd-rhg30.

[9] Felix Petersen et al. *Deep Differentiable Logic Gate Networks*. 2022. arXiv: 2210.08277 [cs.LG]. URL: https://arxiv.org/abs/2210.08277.

[10] Xilinx Inc. *Xilinx Vivado Design Suite User Guide: Release Notes, Installation, and Licensing*. Version 2021.1. San Jose, CA, June 2021. URL: https://www.xilinx.com/support/documentation/sw_manuals/xilinx2021_1/ug973-vivado-release-notes-install-license.pdf.

[11] Xilinx Inc. *Xilinx Vitis High-Level Synthesis User Guide*. Version 2021.1. San Jose, CA, June 2021. URL: https://www.xilinx.com/support/documentation/sw_manuals/xilinx2021_1/ug1399-vitis-hls.pdf.

[12] Felix Petersen et al. "Convolutional differentiable logic gate networks". In: *Proceedings of the 38th International Conference on Neural Information Processing Systems*. NIPS '24. Vancouver, BC, Canada: Curran Associates Inc., 2025. ISBN: 9798331314385.

A Existing CICADA Architectures

The teacher model for CICADA is a convolutional autoencoder that follows the architecture outlined in Fig. 4. The encoding step comprises two convolutional layers followed by a dense layer, which compresses the input into the low-dimensional latent space. The decoding step follows these steps in reverse. The currently implemented student model on which this paper attempts to improve is a much smaller convolutional autoencoder, quantized using QKeras. The student model follows the architecture outlined in Fig. 5.

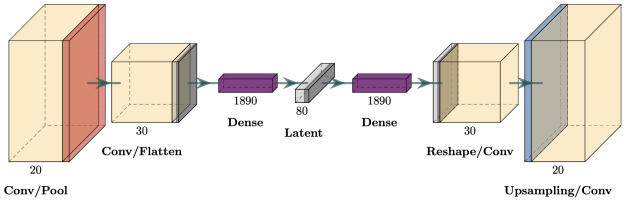

Figure 4: CICADA teacher model architecture.

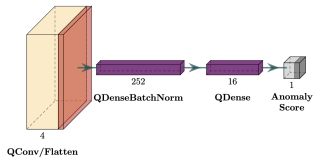

Figure 5: CICADA legacy student model architecture.

Figure 6: The numbers at the convolutional layers represent the number of filters used in each layer (e.g., 20, 30), while the numbers at the dense layers indicate the number of nodes in each fully connected layer (e.g., 1890, 80).

B Data Representation for LGN Models

During training, the LGN architecture can handle float inputs, since each logic gate is represented as a continuous function. However, the compiled model uses discrete logic gates in place of the continuous functions and can therefore only accept binary inputs. The schema for representing the $18\phi \times 14\eta$ integer input grid in binary form during training and for the compiled model has a large effect on the quality of anomaly score reconstruction. This allows values between 0 and 1 while training, such that when the inputs are collapsed to binary values upon compiling the model, important information is retained.

The bit representation lt2 follow a log transformation function. Let x be the input and $\{t_i\}_{i=1}^{2}$ be a list of thresholds. Then the transformation is calculated as:

$$y_i = \log\left(1 + \frac{\max(x - t_i, 0)}{2}\right), \quad i = 1, \ldots, 2 \tag{3}$$

Different thresholds and number of thresholds are selected for different student models. The untransformed inputs range from 0 to 256, but high values are extremely rare in the input dataset. The threshold values are therefore chosen as [0, 2] to reflect this.

C Source Code

Our source code can be found here https://anonymous.4open.science/r/ml4physics-paper-plots-7C16/.