

# Rapid HEP ML inference with logic gate neural nets

ACAT 2025

09.09.25

*This presentation contains several animations and is best viewed as .key or .pptx, or you can view it online [here](#)*

Lino Gerlach, Elliott Kauffman, [Liv Våge](#)



# Contents

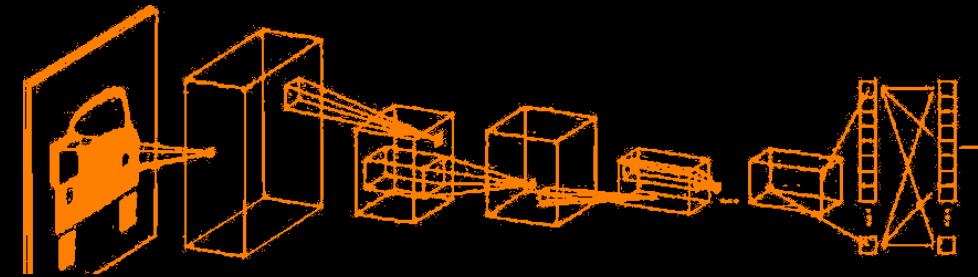
Differentiable logic  
gate networks



CICADA



Differentiable logic gates  
for CICADA



FPGA  
implementation

**AMD Vitis™ HLS**

## Foundational work



SOTA inference times on MNIST  
- 10-20x faster than others

[Original difflogic paper](#) presented at NEURIPS 2022 [1]

[Code repo](#) for the original paper

[Follow up paper](#) on difflogic CNN from NEURIPS 2024

## TLDR

Good physics performance and  
rapid inference for CICADA student model  
with LGCN

# Logic gate neural networks

Construct neural network from logic gates  
instead of nodes, we have logic gates

Convert the input to a binary representation  
different representations can produce different results

Each node receives two inputs  
The connections are randomly initialised

Outputs are aggregated so we can classify or regress

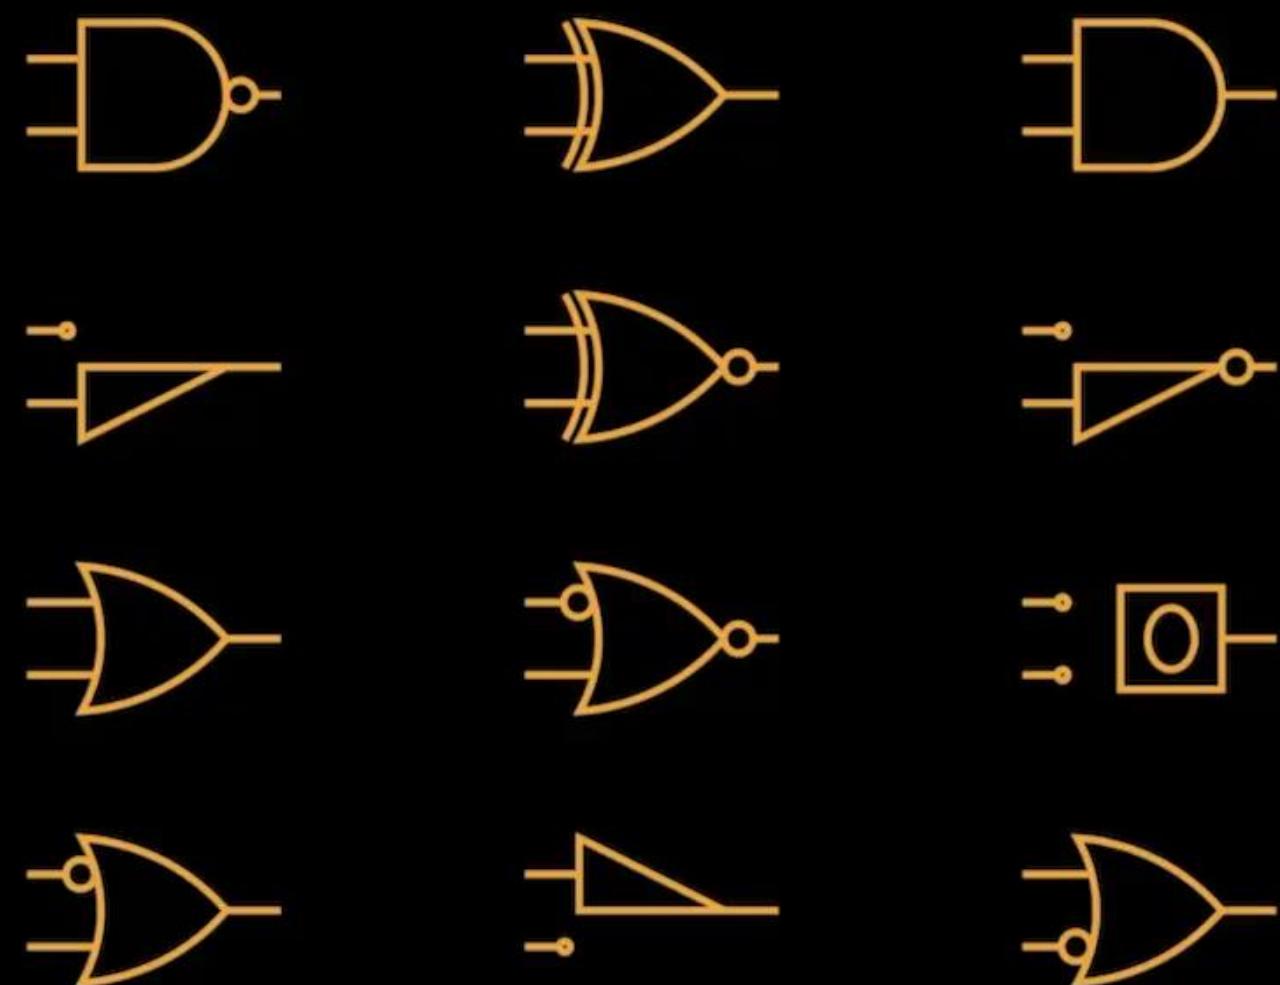
$$\hat{y}_i = \text{Final output} = \frac{(i+1) \cdot n/k}{\sum_{j=i \cdot n/k + 1}^{(i+1) \cdot n/k} a_j / \tau + \beta} \text{Neuron output}$$

Number of output neurons

Output dimension

Normalisation temperature

Optional offset



## Why it is fast

At inference each 16 gate block is replaced by most probable gate

Binary computations are fast

Compiler can optimise the binary logic

No matrix multiplications!

Video adapted  
from [2], made by Felix  
Petersen et al.

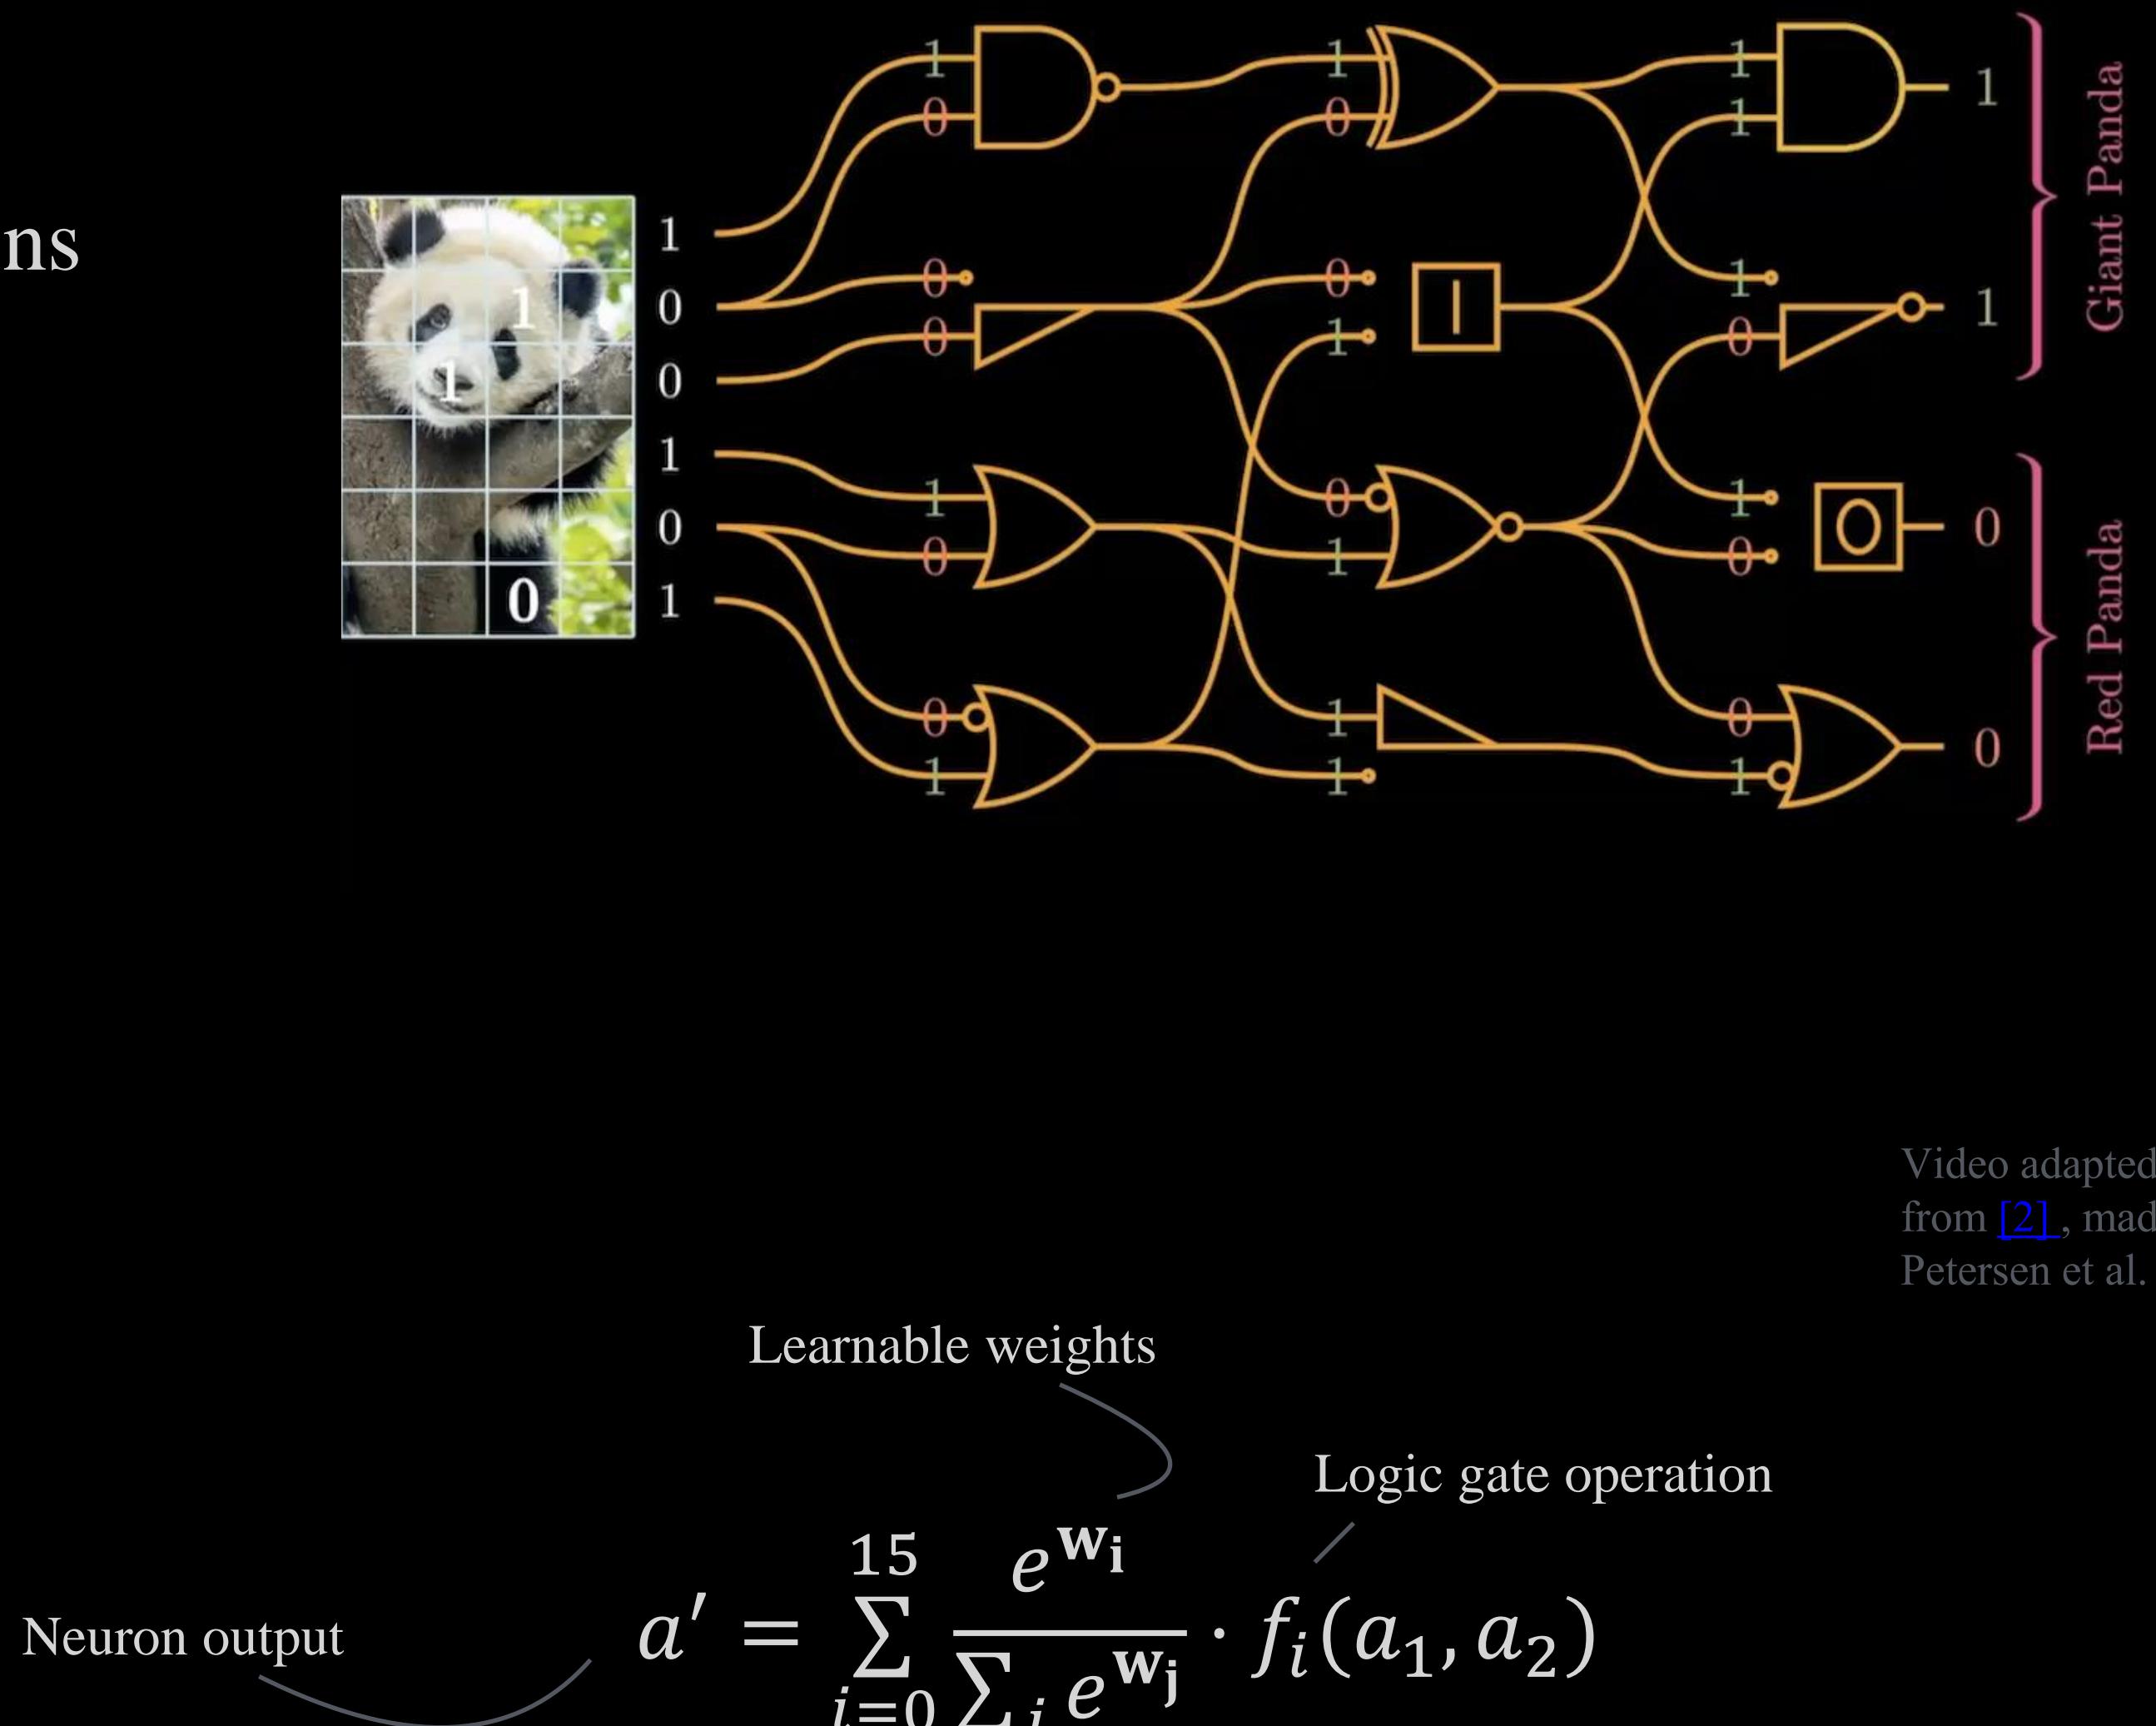
# Making it differentiable

Relax gate operations to continuous approximations  
binary inputs are usually not differentiable

Softmax of the 16 gate blocks  
we need learn which gates to use

During training evaluate 16 gate blocks  
slow training, but quick inference as we replace each block of 16  
with one gate

ID	Operator	real-valued	00	01	10	11
0	False	0	0	0	0	0
1	$A \wedge B$	$A \cdot B$	0	0	0	1
2	$\neg(A \Rightarrow B)$	$A - AB$	0	0	1	0
3	$A$	$A$	0	0	1	1
4	$\neg(A \Leftarrow B)$	$B - AB$	0	1	0	0
5	$B$	$B$	0	1	0	1
6	$A \oplus B$	$A + B - 2AB$	0	1	1	0
7	$A \vee B$	$A + B - AB$	0	1	1	1
8	$\neg(A \vee B)$	$1 - (A + B - AB)$	1	0	0	0
9	$\neg(A \oplus B)$	$1 - (A + B - 2AB)$	1	0	0	1
10	$\neg B$	$1 - B$	1	0	1	0
11	$A \Leftarrow B$	$1 - B + AB$	1	0	1	1
12	$\neg A$	$1 - A$	1	1	0	0
13	$A \Rightarrow B$	$1 - A + AB$	1	1	0	1
14	$\neg(A \wedge B)$	$1 - AB$	1	1	1	0
15	True	1	1	1	1	1



Video adapted  
from [2], made by Felix  
Petersen et al.

# Convolutional differentiable logic gate neural networks

0	1	1	0
1	1	1	0
0	0	.	

Linear layers is not enough  
empirically struggled to train over 6 layers

Replace normal kernel by binary tree  
aggregates information while keeping expressivity

One channel for each input bit  
We learn the significance of each bit, analogous to colour

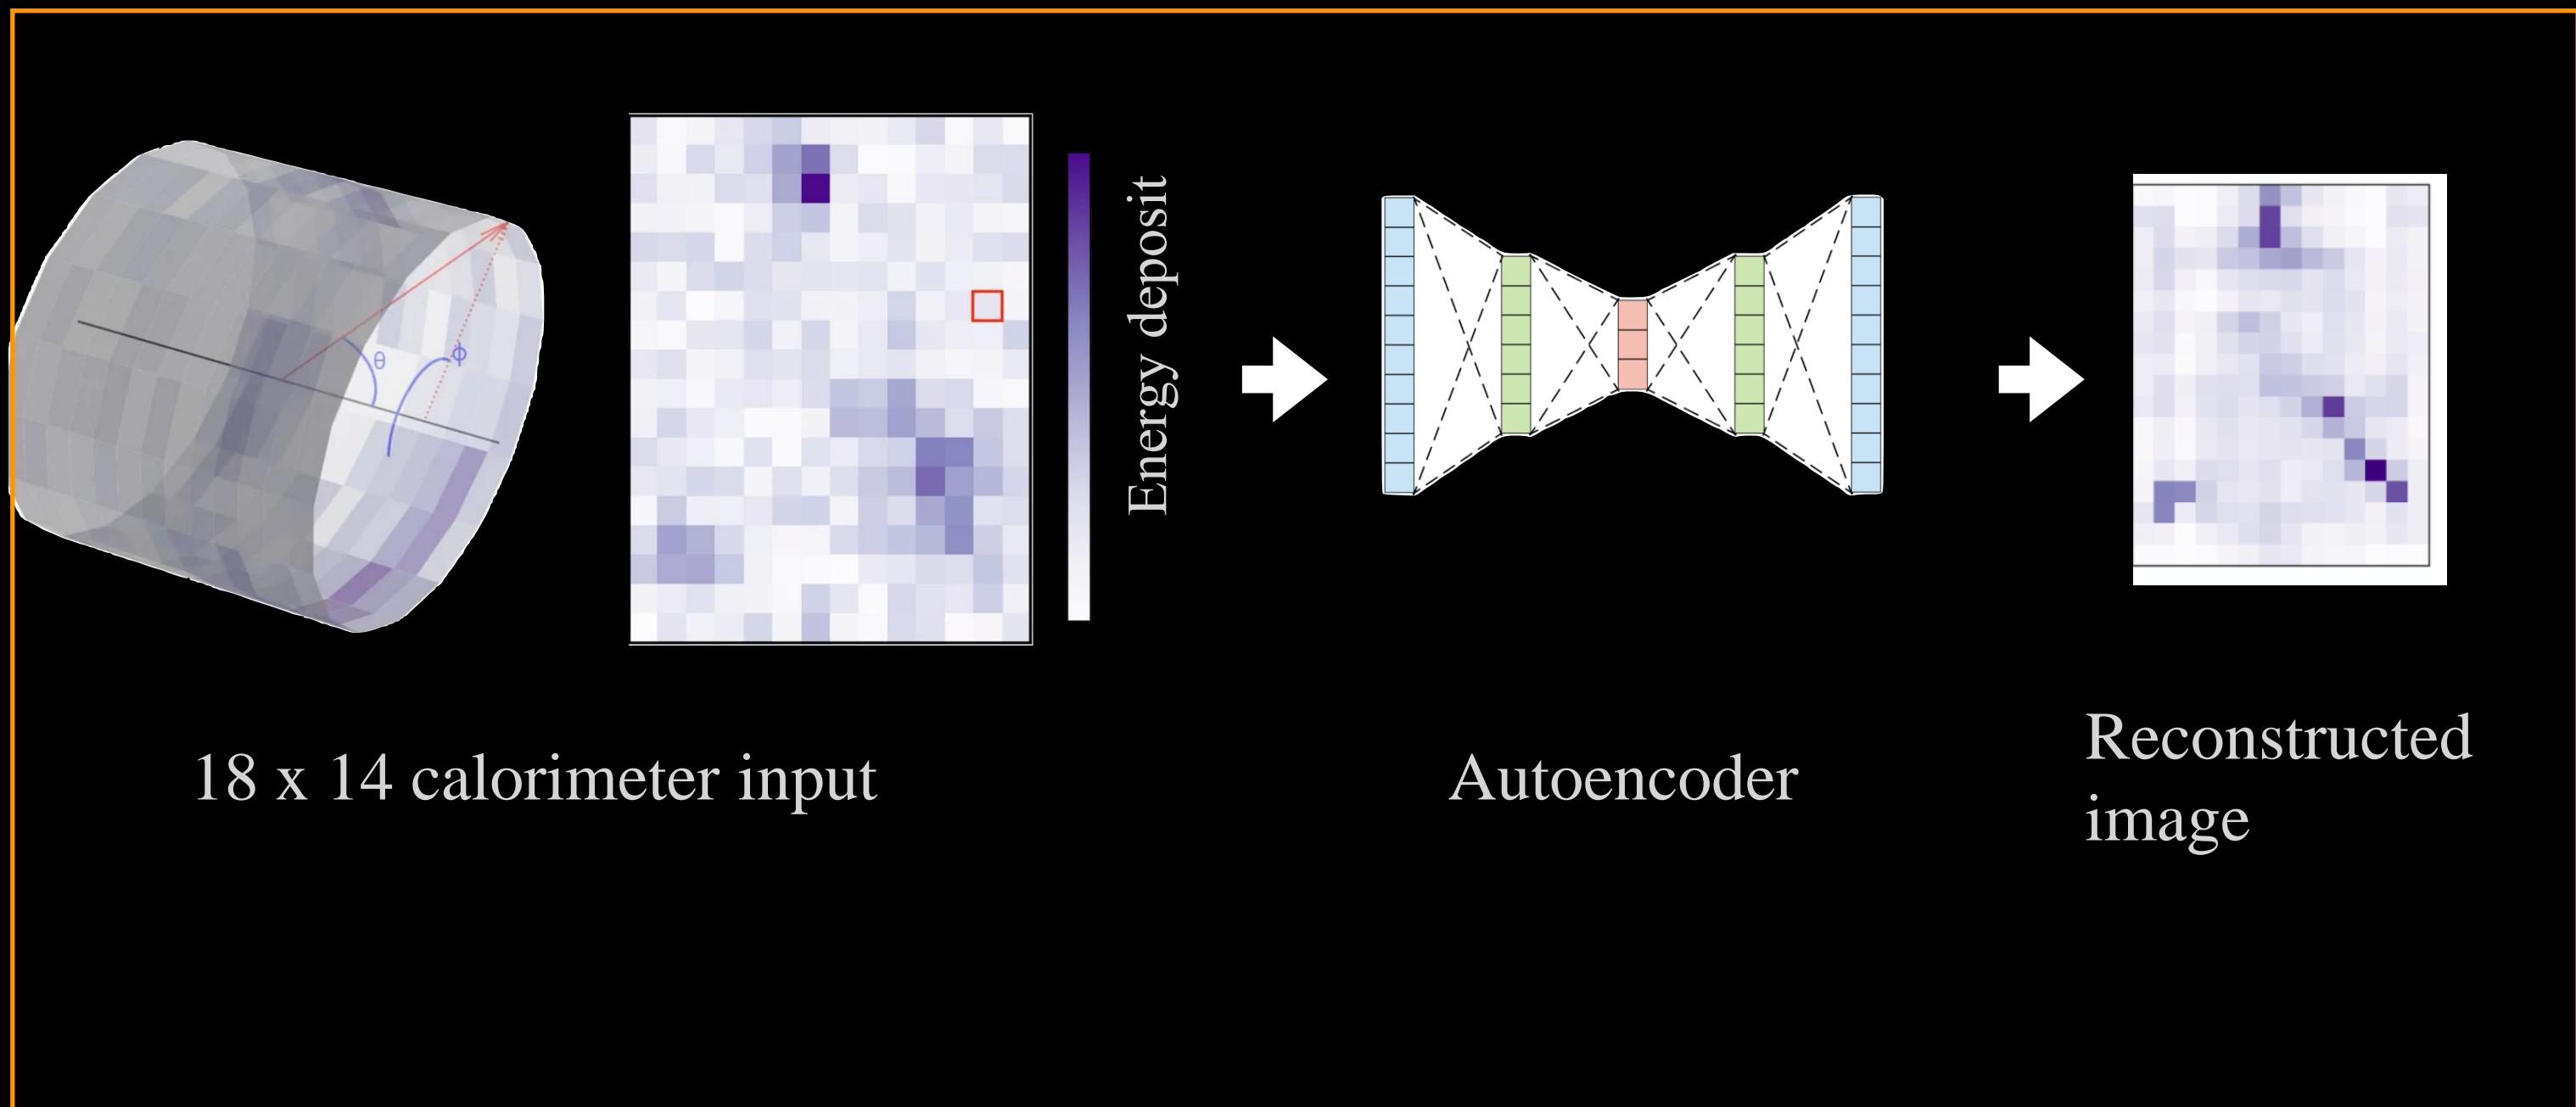
Create special *Or* pooling layers  
fast and only need to propagate through the maximum activations

$$\perp_{max} (a, b) = \max(a, b)$$

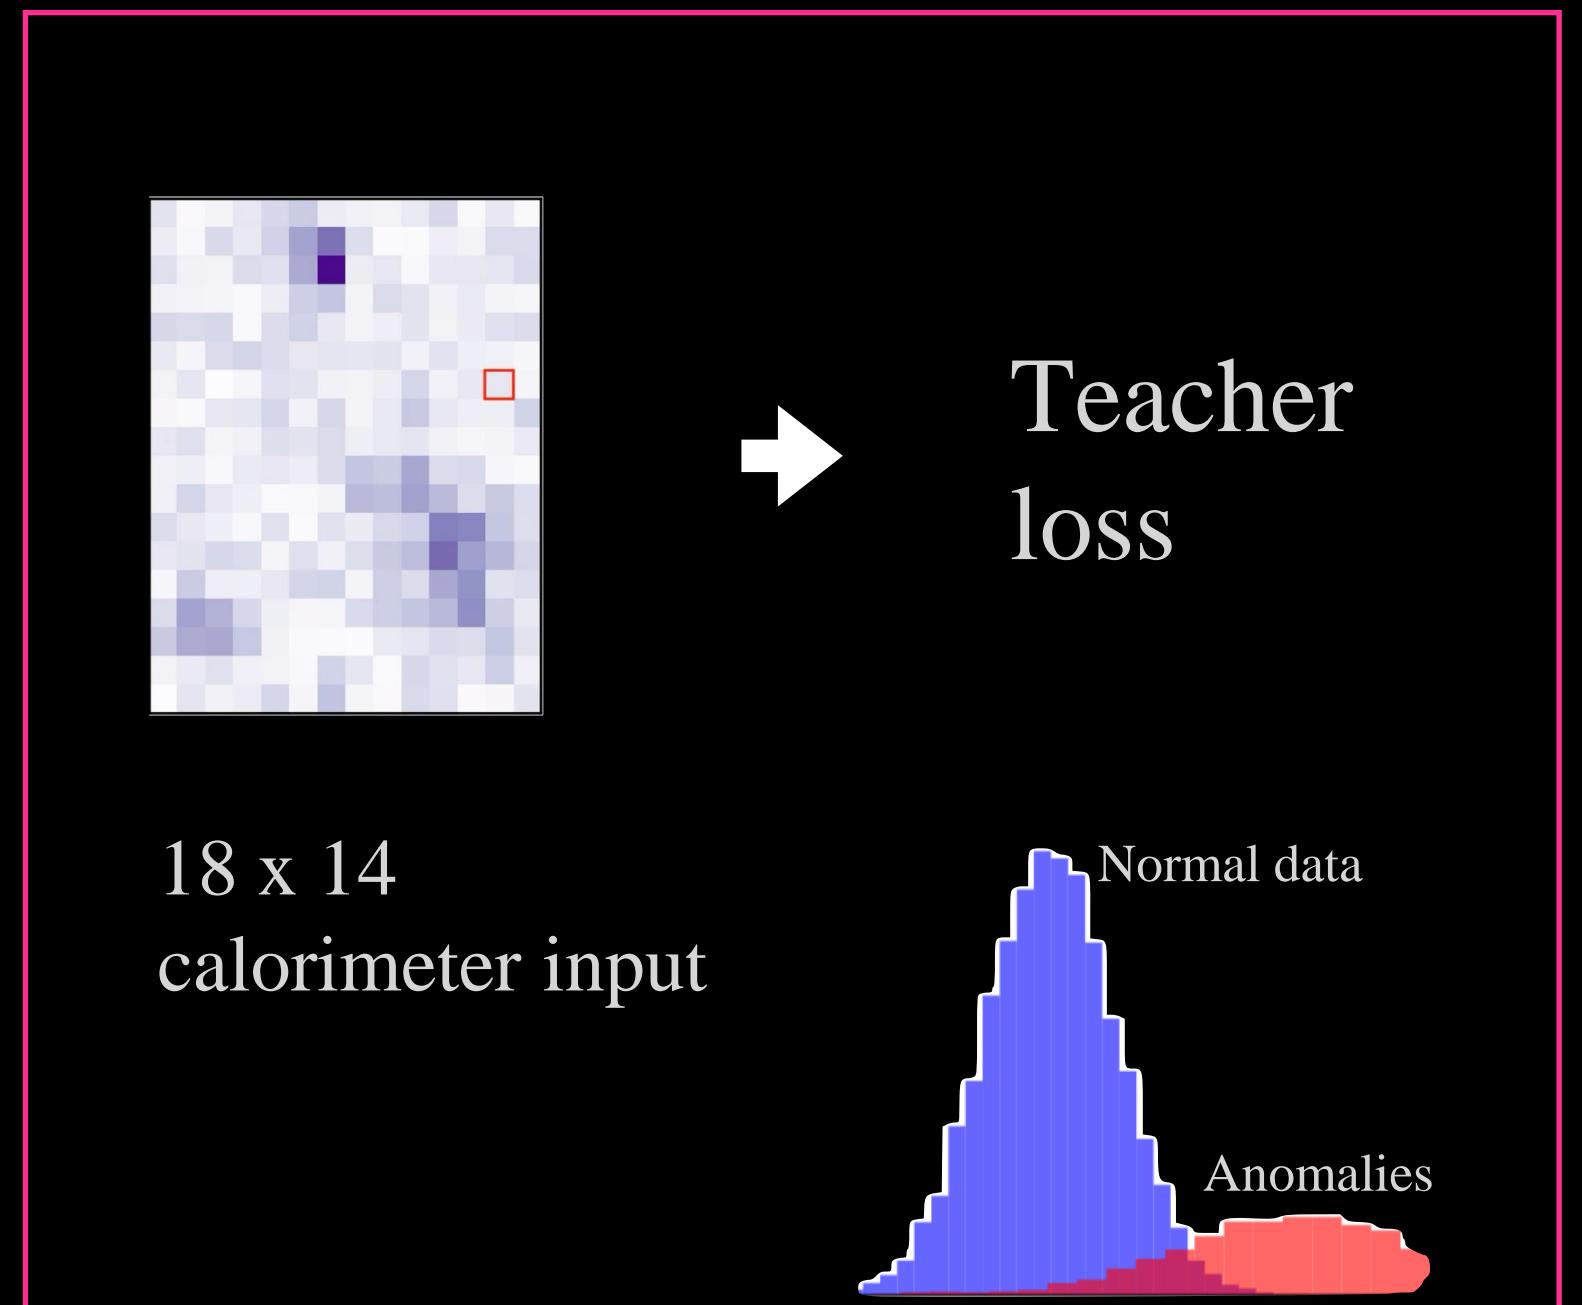


# CICADA anomaly detection at the CMS Level-1 Trigger

## Teacher



## Student



Student is implemented on FPGA in the L1T which has output rate of 100 kHz

Can we achieve better physics performance or latency?

# Results

Using 2017 CMS Open data  
will be openly available very soon

Background is Zero Bias,  $\bar{t}t$  is outlier

There are three possible outlier final states, two are reserved for validation and testing

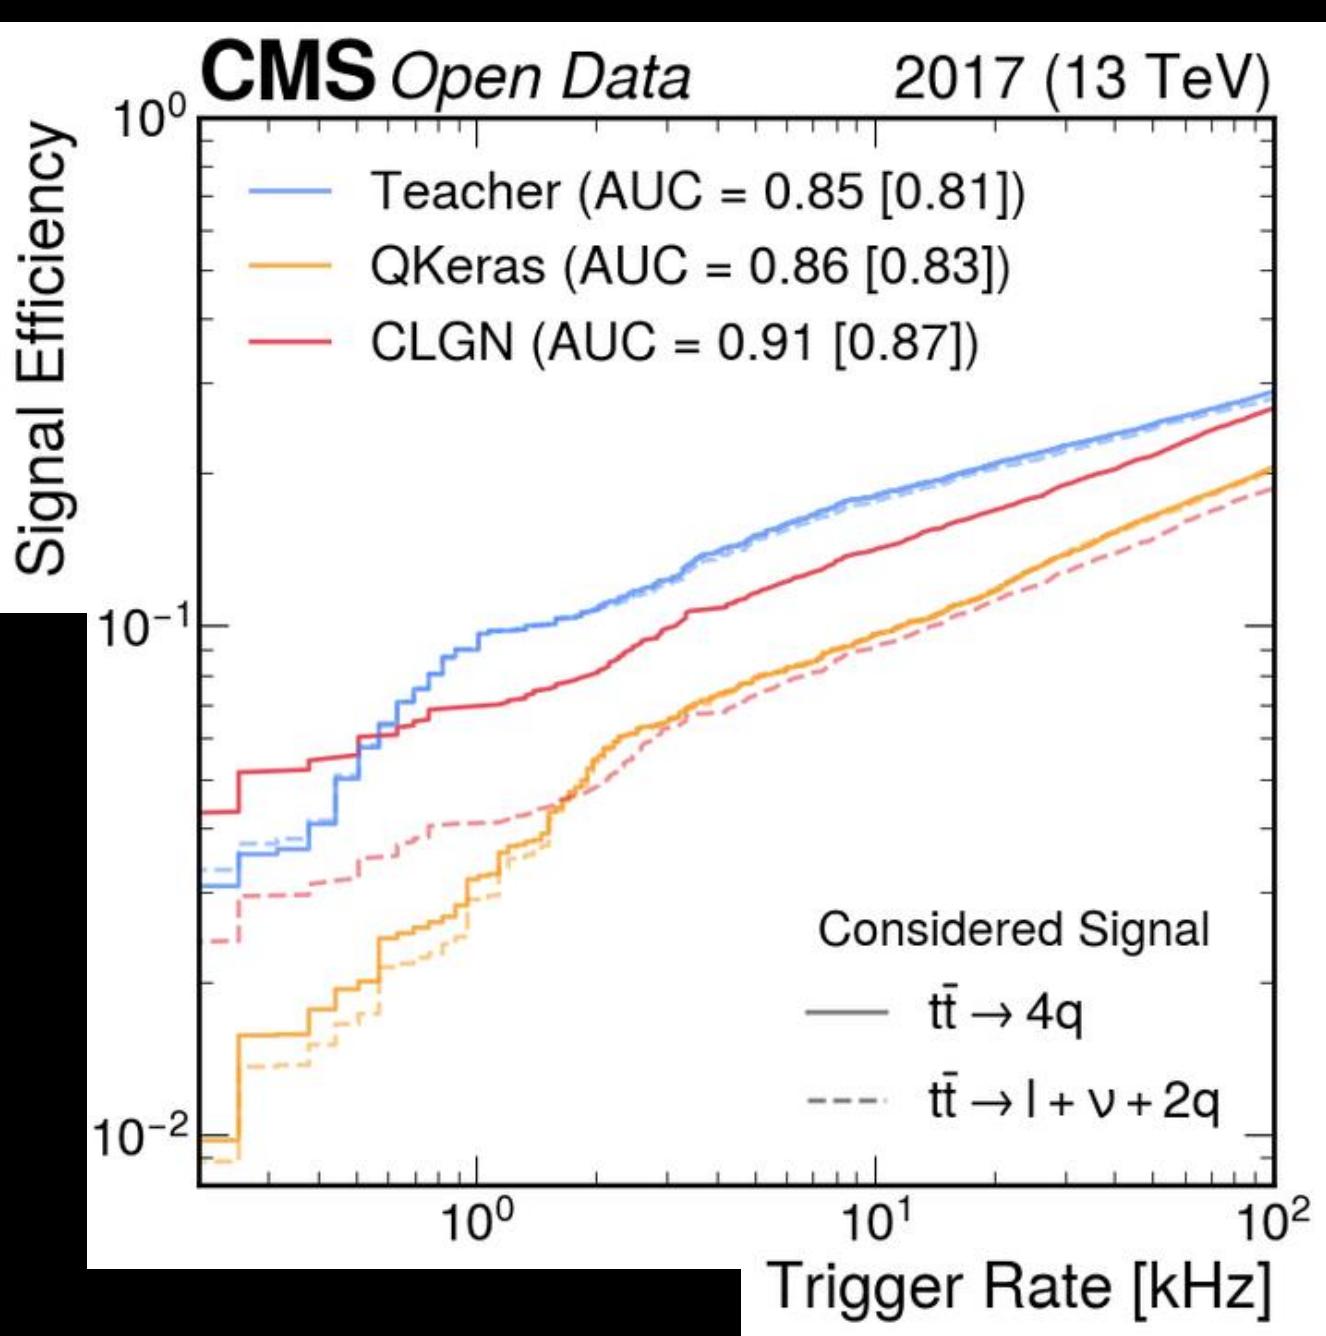
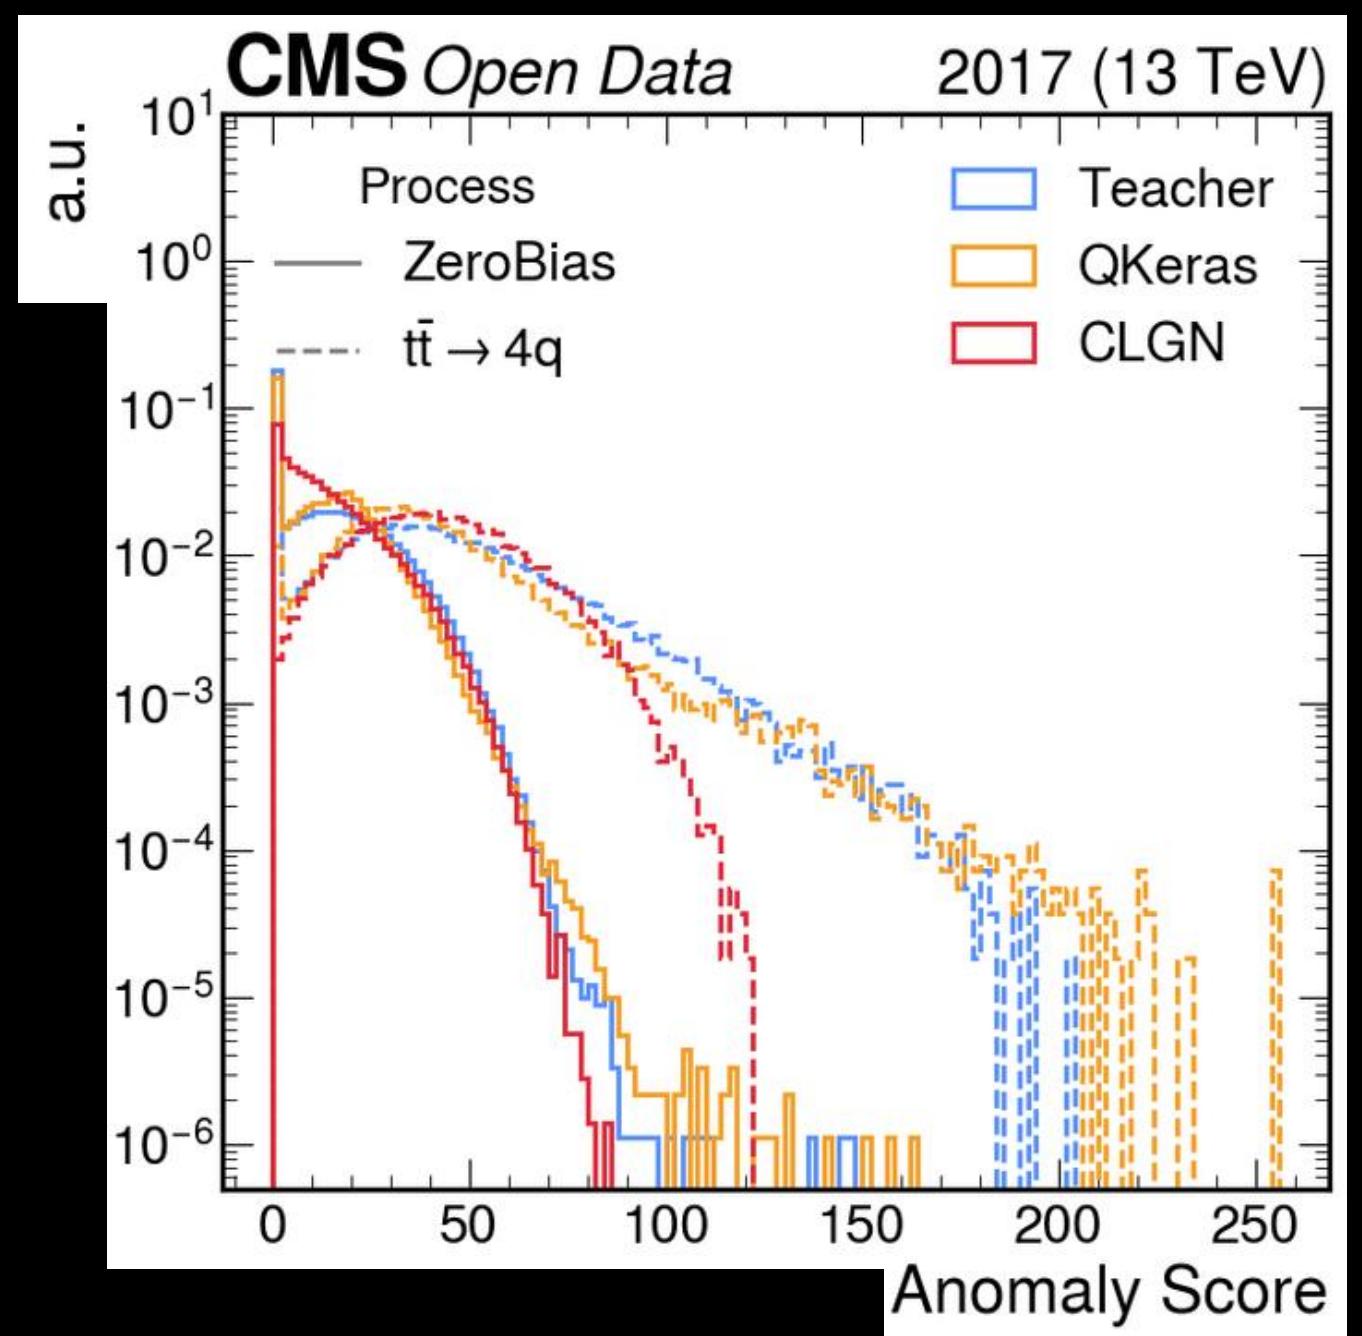
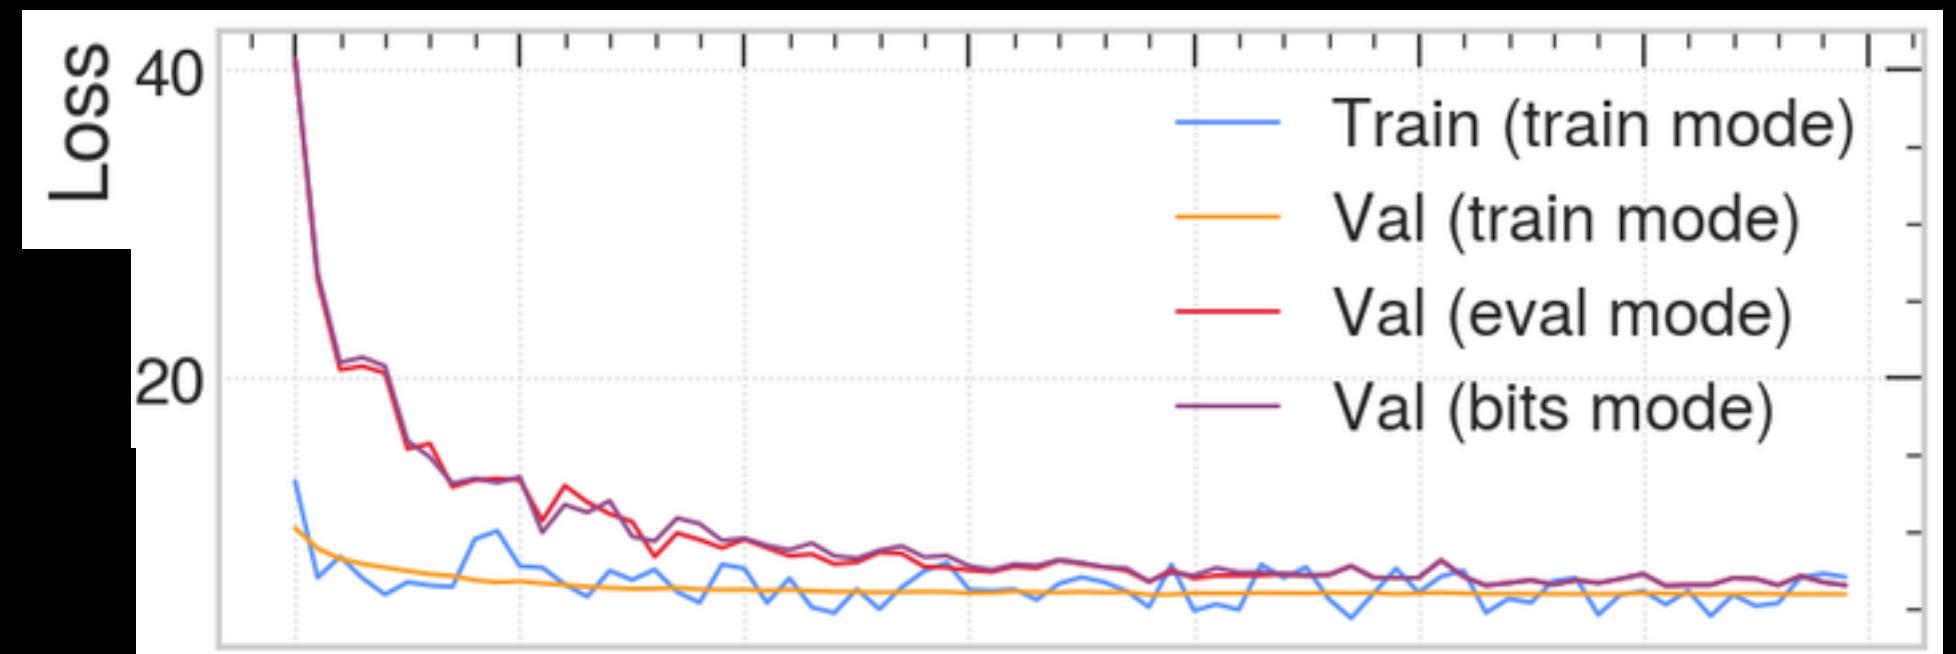
To best convert inputs to binary implement thermometer encoding  
learns how to threshold binary numbers during training,  
as in [4]

Very high anomaly scores lower since  
binarised inputs

This is okay, we can set the trigger threshold

Convolutional LGN outperforms QKeras  
implementation

QKeras is especially worse on low trigger rates



# FPGA synthesis

Synthesised for AMD Virtex-7  
using AMD Vitis HLS

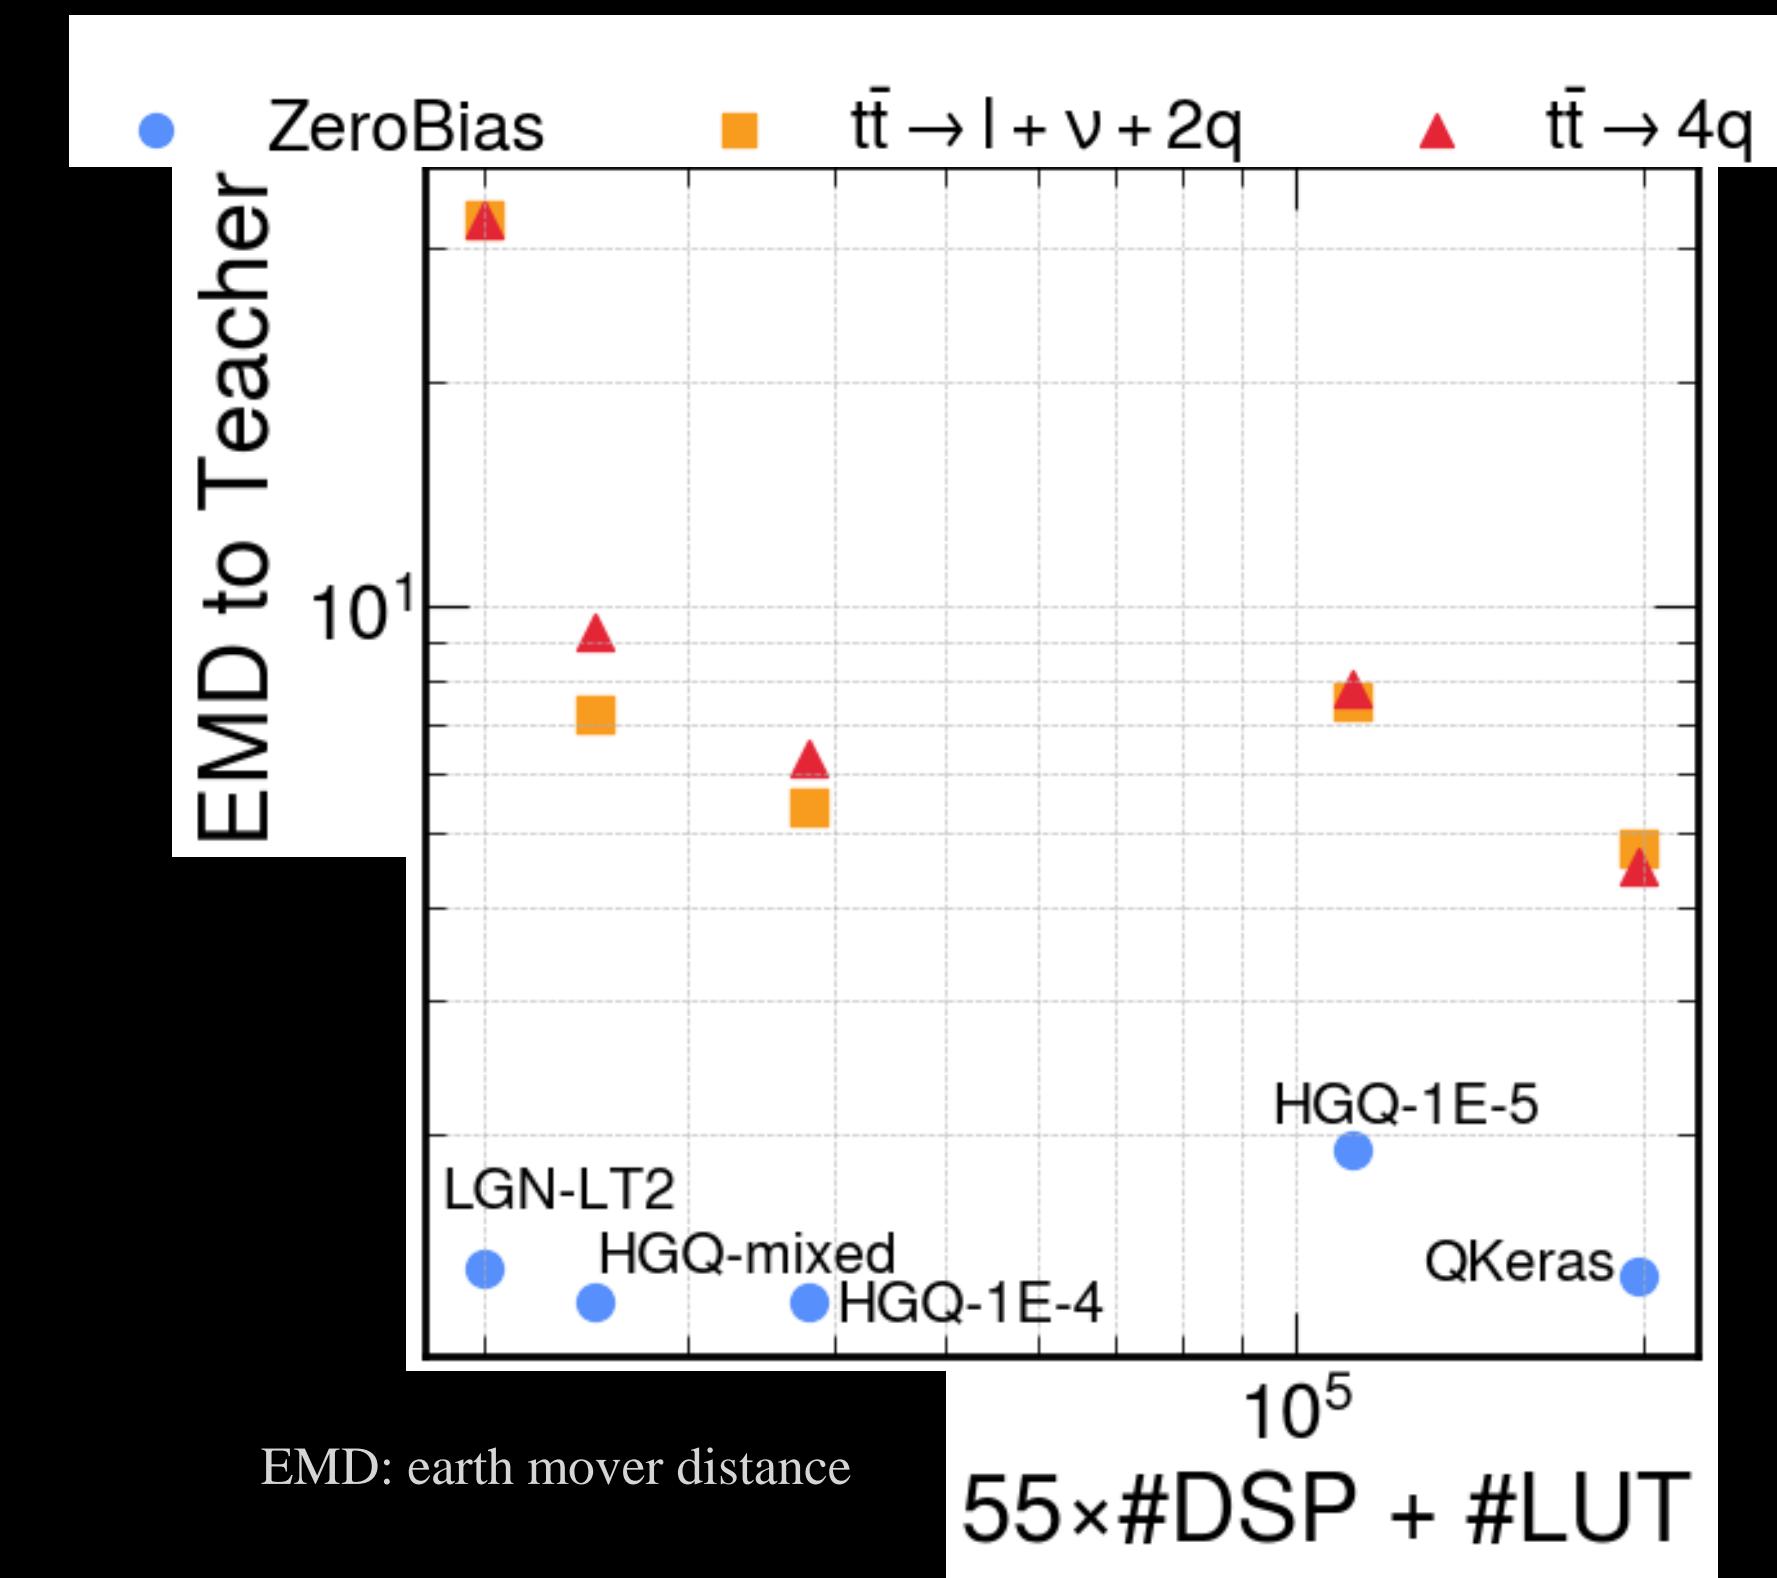
This had ~half the trainable parameters as  
compared to previous slide  
the FPGA performance is therefore likely an underestimation  
results are good on background, not on signal

Still we see very promising FPGA  
performance

no DSP usage since we don't have matrix multiplications

Latency outperforms QKeras and HGQ  
implementations

details of those can be found in [5]



Vitis synthesis estimates

## Conclusion

Implemented difflog convolutions with:  
CPU & GPU  
FPGA synthesis

### Results:

Good physics performance for CICADA student  
Promising FPGA performance with 0 DSP use

## Further work

Almost done with a custom CUDA implementation

There have been [developments in the field](#) to consider

## Code implementation

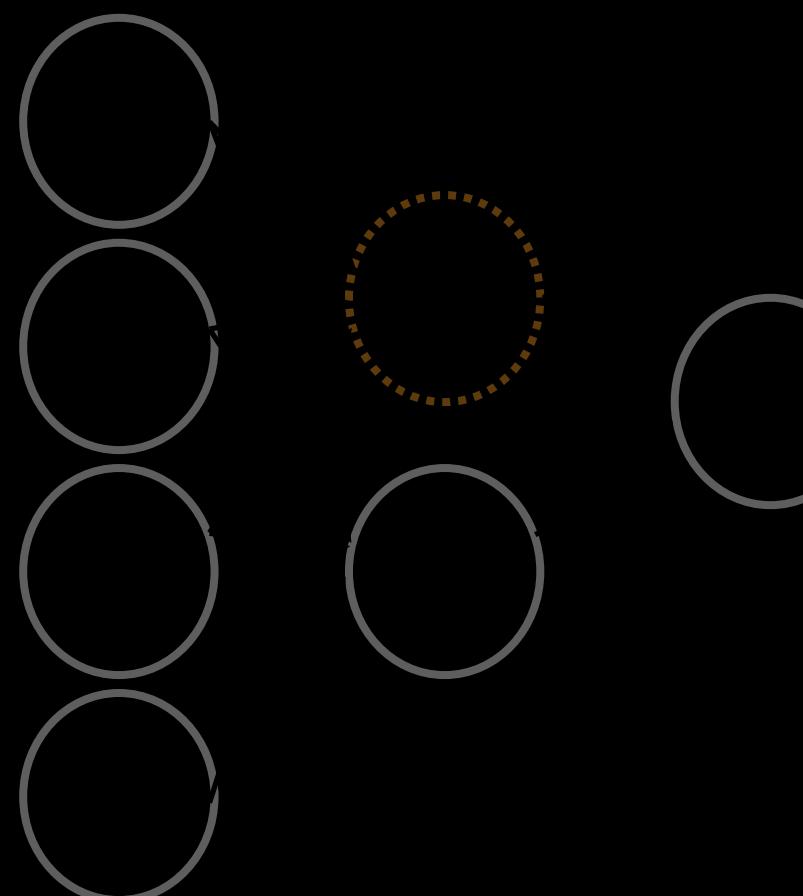
Will be included in proceedings



# Backup

# MLP vs difflog MLP

Normal neural net



Fully connected

Matrix multiplication in nodes

Same in training as inference

Logic gate neural net



Randomly connected; each node has two inputs

Nodes evaluate logic gates

*At inference, inputs are binary and only use most probable logic gate for each node*

No matrix multiplication, binary logic can be simplified by compiler, binary computations faster

ID	Operator	real-valued	00	01	10	11
0	False	0	0	0	0	0
1	$A \wedge B$	$A \cdot B$	0	0	0	1
2	$\neg(A \Rightarrow B)$	$A - AB$	0	0	1	0
3	$A$	$A$	0	0	1	1
4	$\neg(A \Leftarrow B)$	$B - AB$	0	1	0	0
5	$B$	$B$	0	1	0	1
6	$A \oplus B$	$A + B - 2AB$	0	1	1	0
7	$A \vee B$	$A + B - AB$	0	1	1	1
8	$\neg(A \vee B)$	$1 - (A + B - AB)$	1	0	0	0
9	$\neg(A \oplus B)$	$1 - (A + B - 2AB)$	1	0	0	1
10	$\neg B$	$1 - B$	1	0	1	0
11	$A \Leftarrow B$	$1 - B + AB$	1	0	1	1
12	$\neg A$	$1 - A$	1	1	0	0
13	$A \Rightarrow B$	$1 - A + AB$	1	1	0	1
14	$\neg(A \wedge B)$	$1 - AB$	1	1	1	0
15	True	1	1	1	1	1

# Thermometer thresholding

We need to convert continuous numbers to binary

We do this based on thresholds, as implemented in [4]

We learn the bin edges as we train  
each bin checks whether the number  $\geq$  threshold  
1 if it is, 0 otherwise

Our representation matters  
the best edges will depend on the problem

In our case, we let it change, then freeze it  
we use a 6 bit representation  
one of the bins settles on a high threshold - likely to capture  
high anomaly scores

